

# Deep learning for image processing

STAT 4710

November 22, 2022

# Where we are

-  **Unit 1:** R for data mining
-  **Unit 2:** Prediction fundamentals
-  **Unit 3:** Regression-based methods
-  **Unit 4:** Tree-based methods
- Unit 5:** Deep learning

**Lecture 1:** Deep learning preliminaries

**Lecture 2:** Neural networks

**Lecture 3:** Deep learning for images

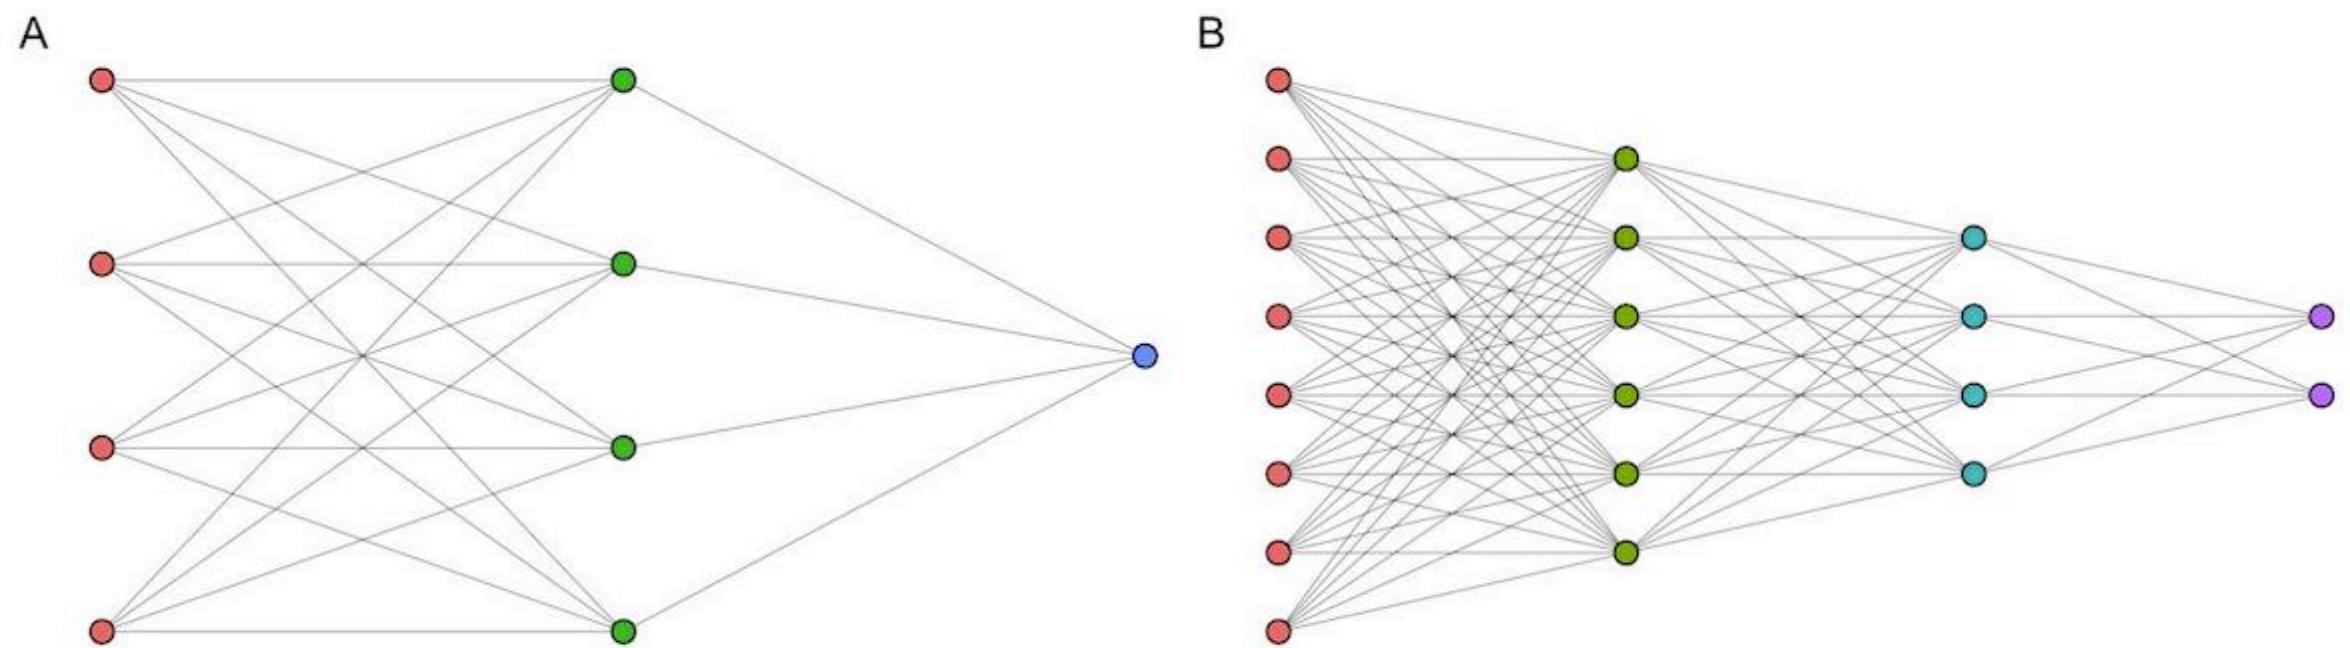
**Lecture 4:** Deep learning for text

**Lecture 5:** Unit review and quiz in class

# Network architectures

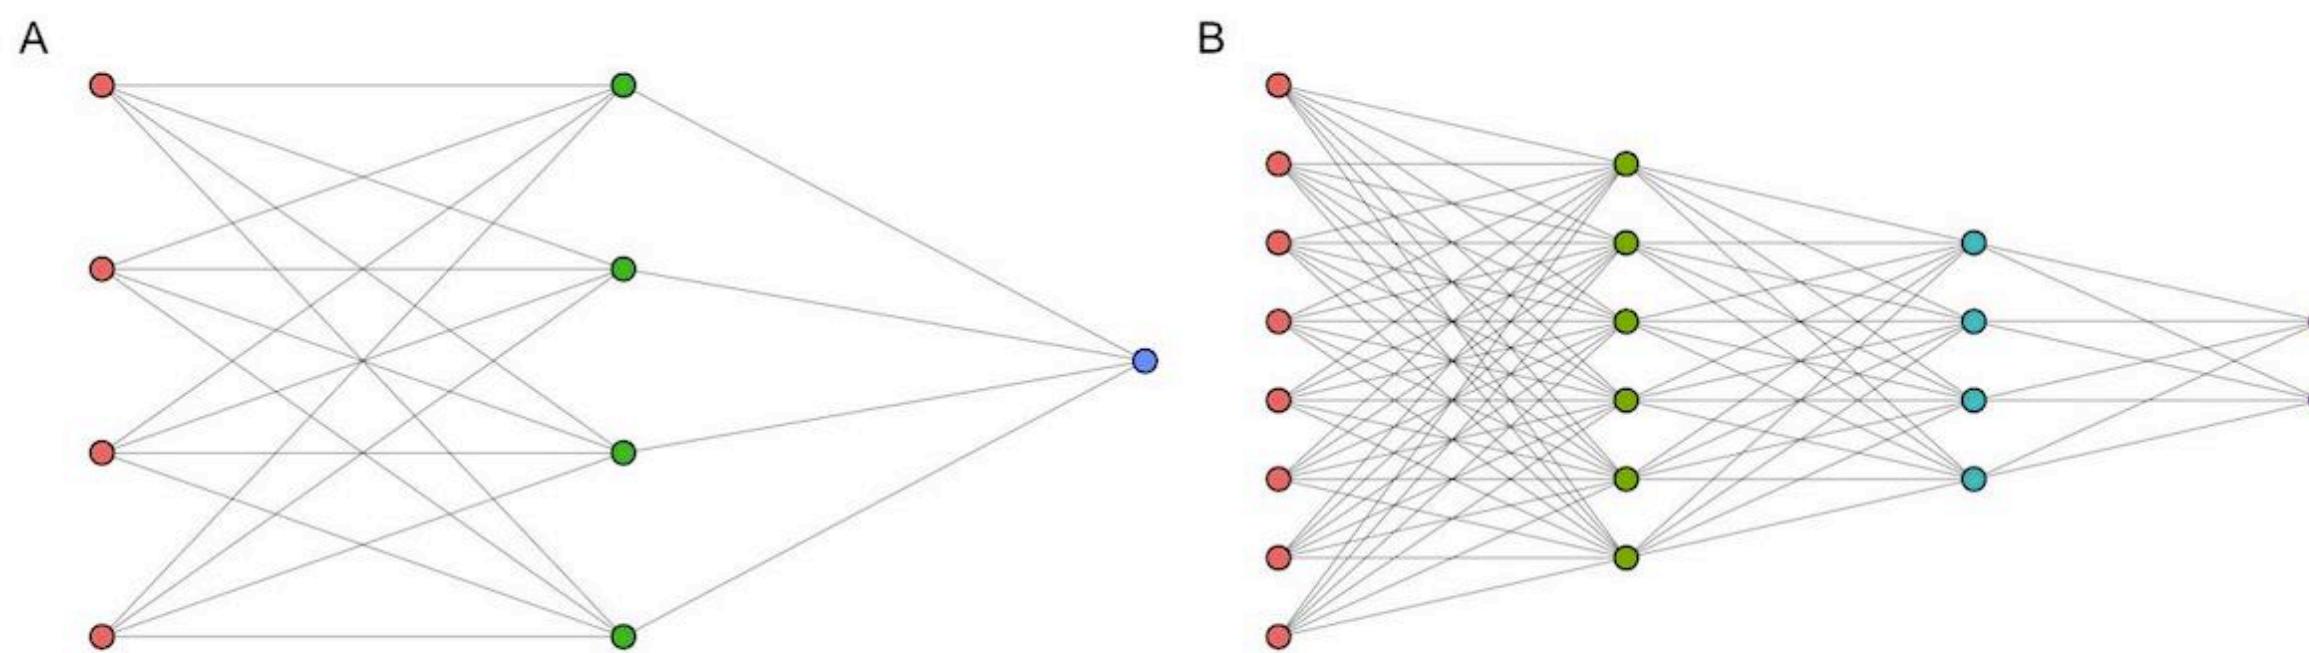
# Network architectures

## Fully connected architectures



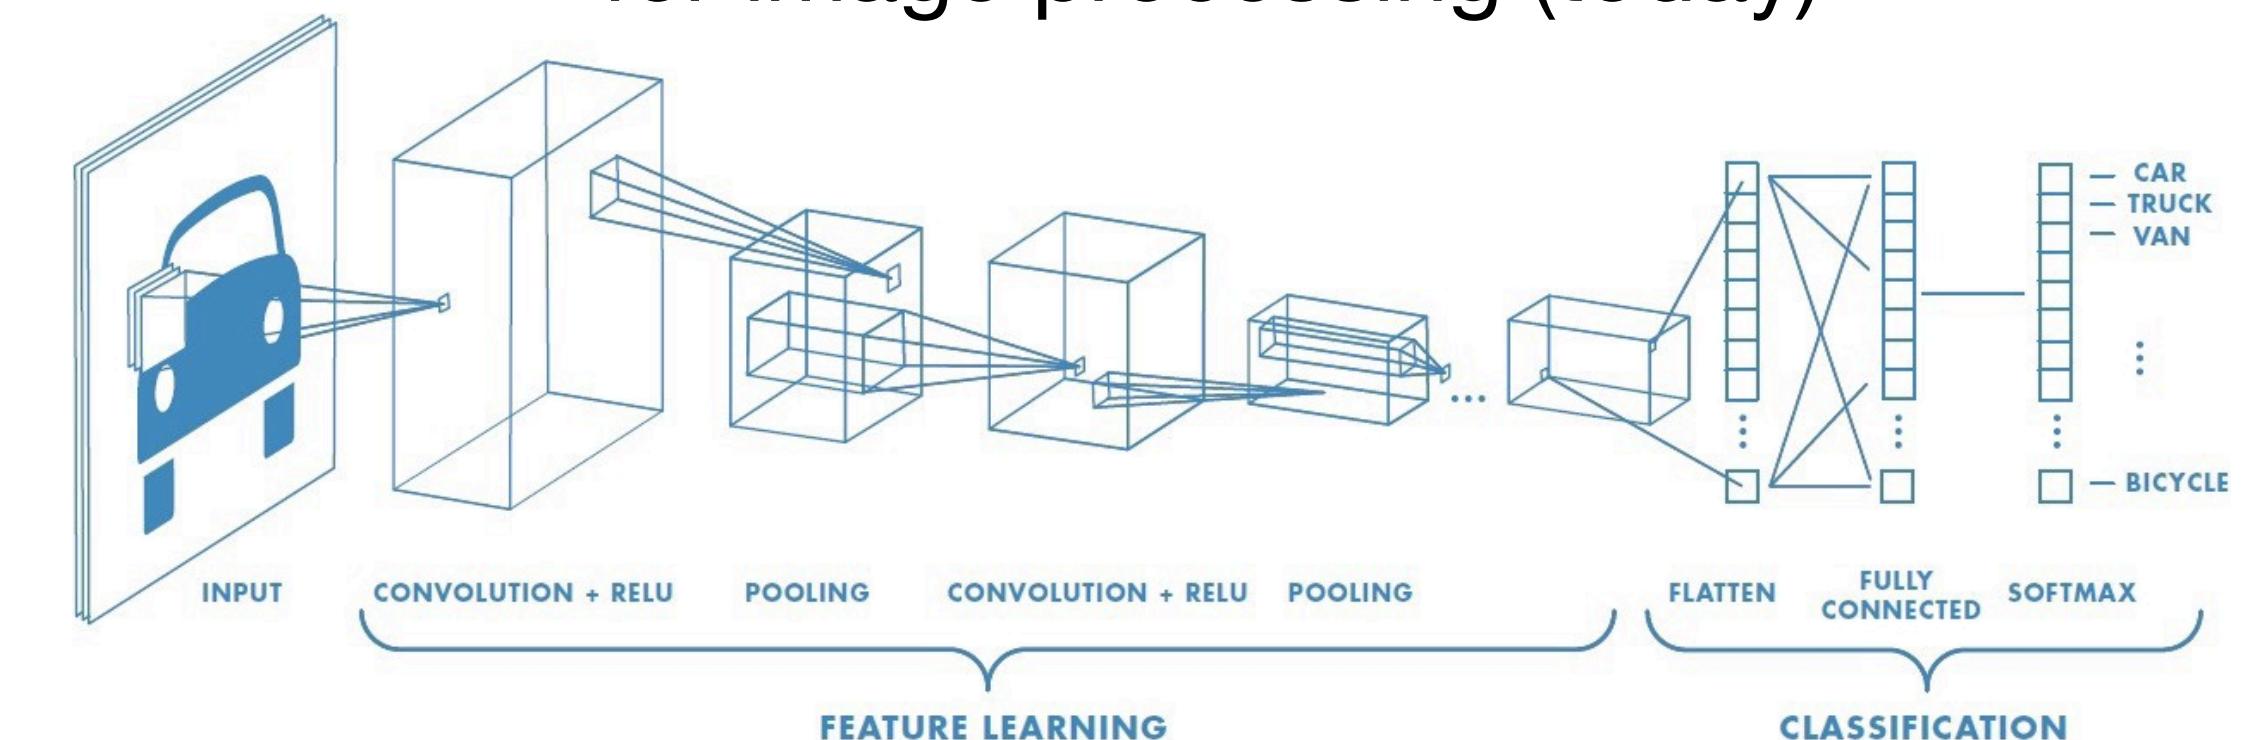
# Network architectures

## Fully connected architectures



<https://community.rstudio.com/t/visualising-neural-network-architectures/41723>

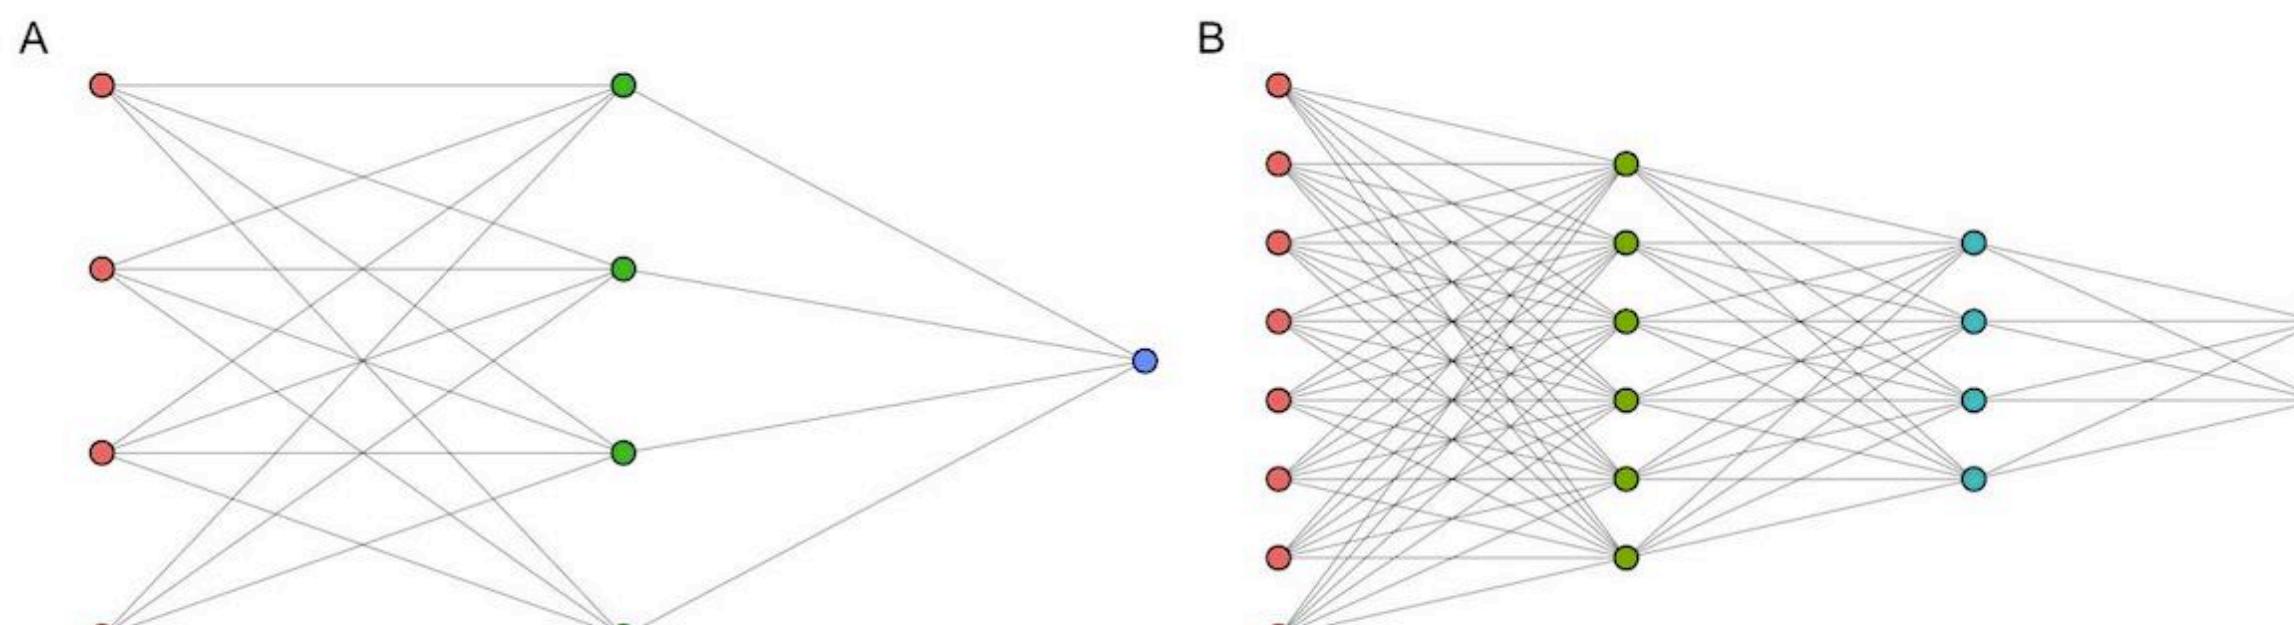
## Convolutional neural network (CNN) architectures for image processing (today)



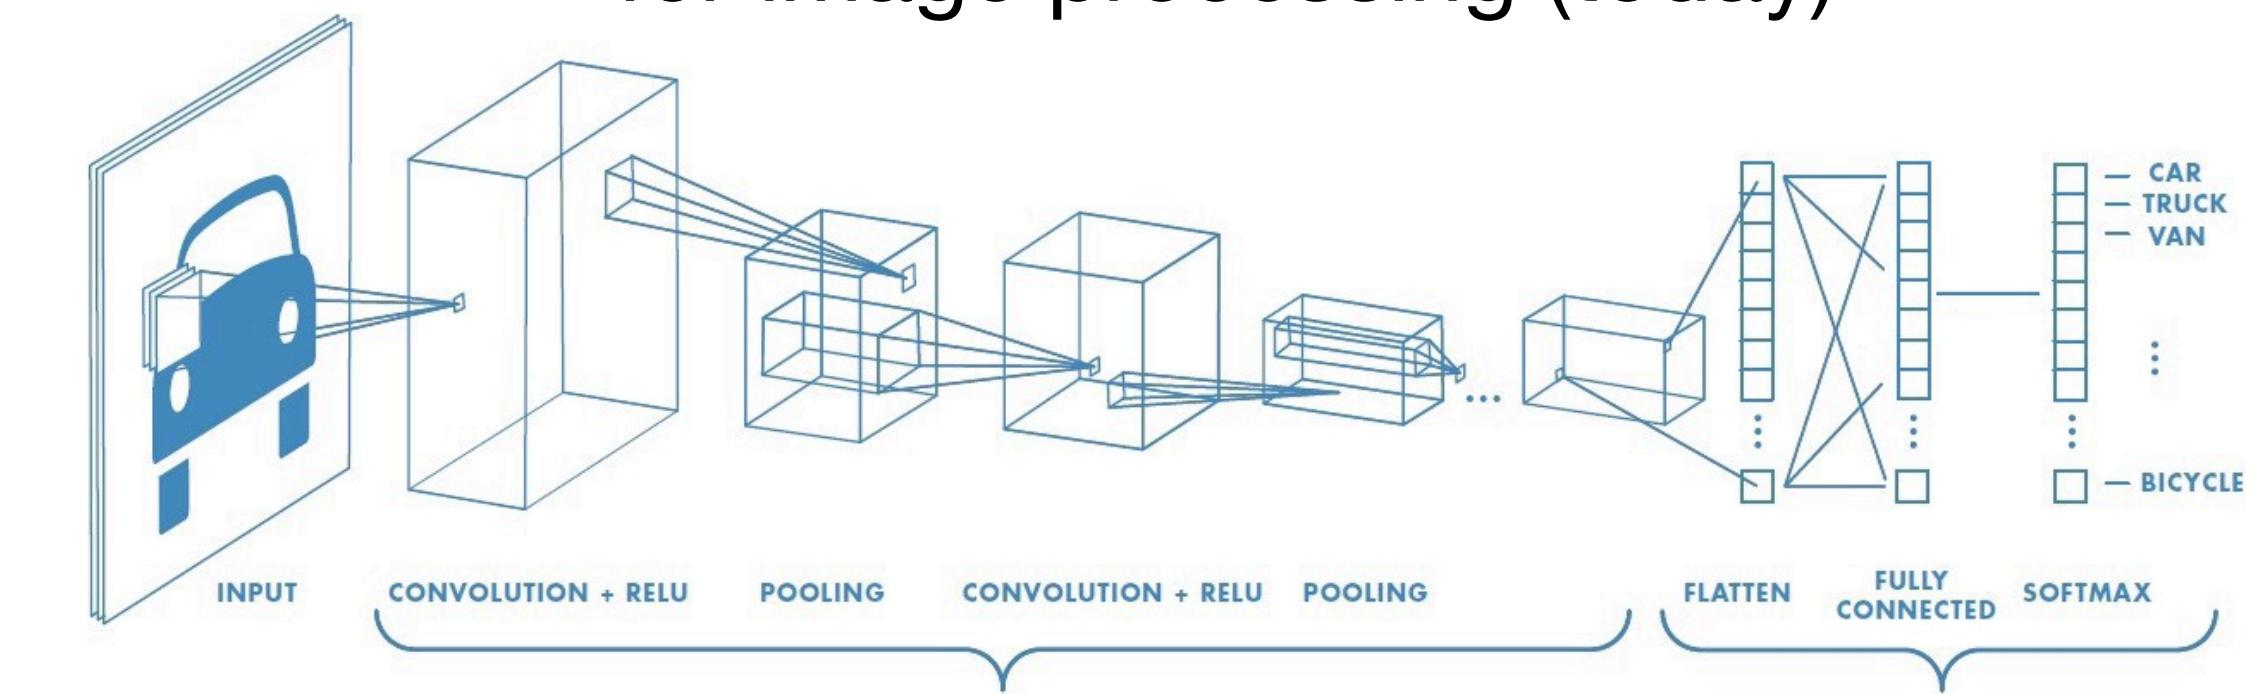
<https://towardsdatascience.com/a-comprehensive-guide-to-convolutional-neural-networks-the-eli5-way-3bd2b1164a53>

# Network architectures

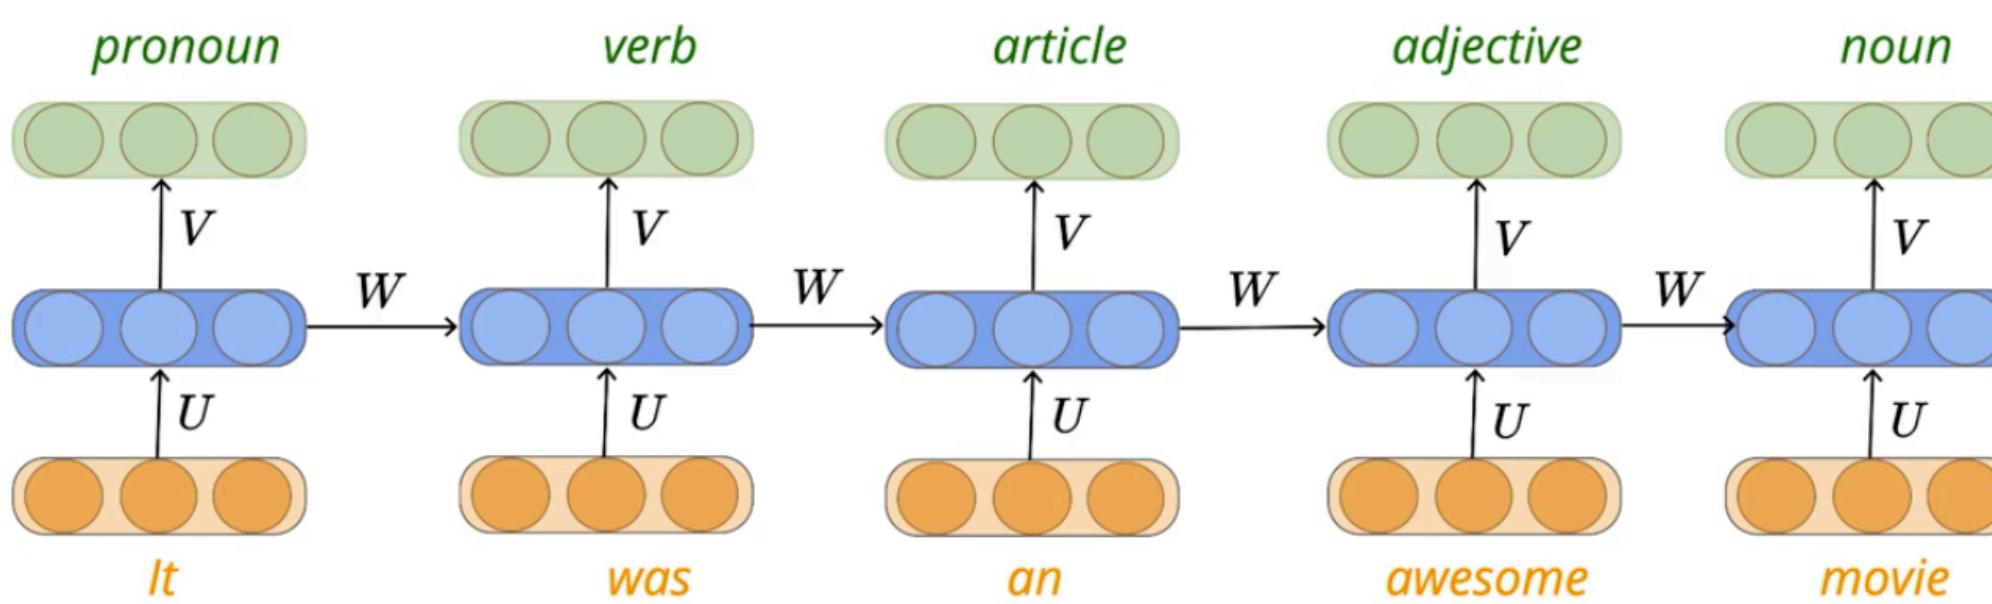
## Fully connected architectures



Convolutional neural network (CNN) architectures  
for image processing (today)



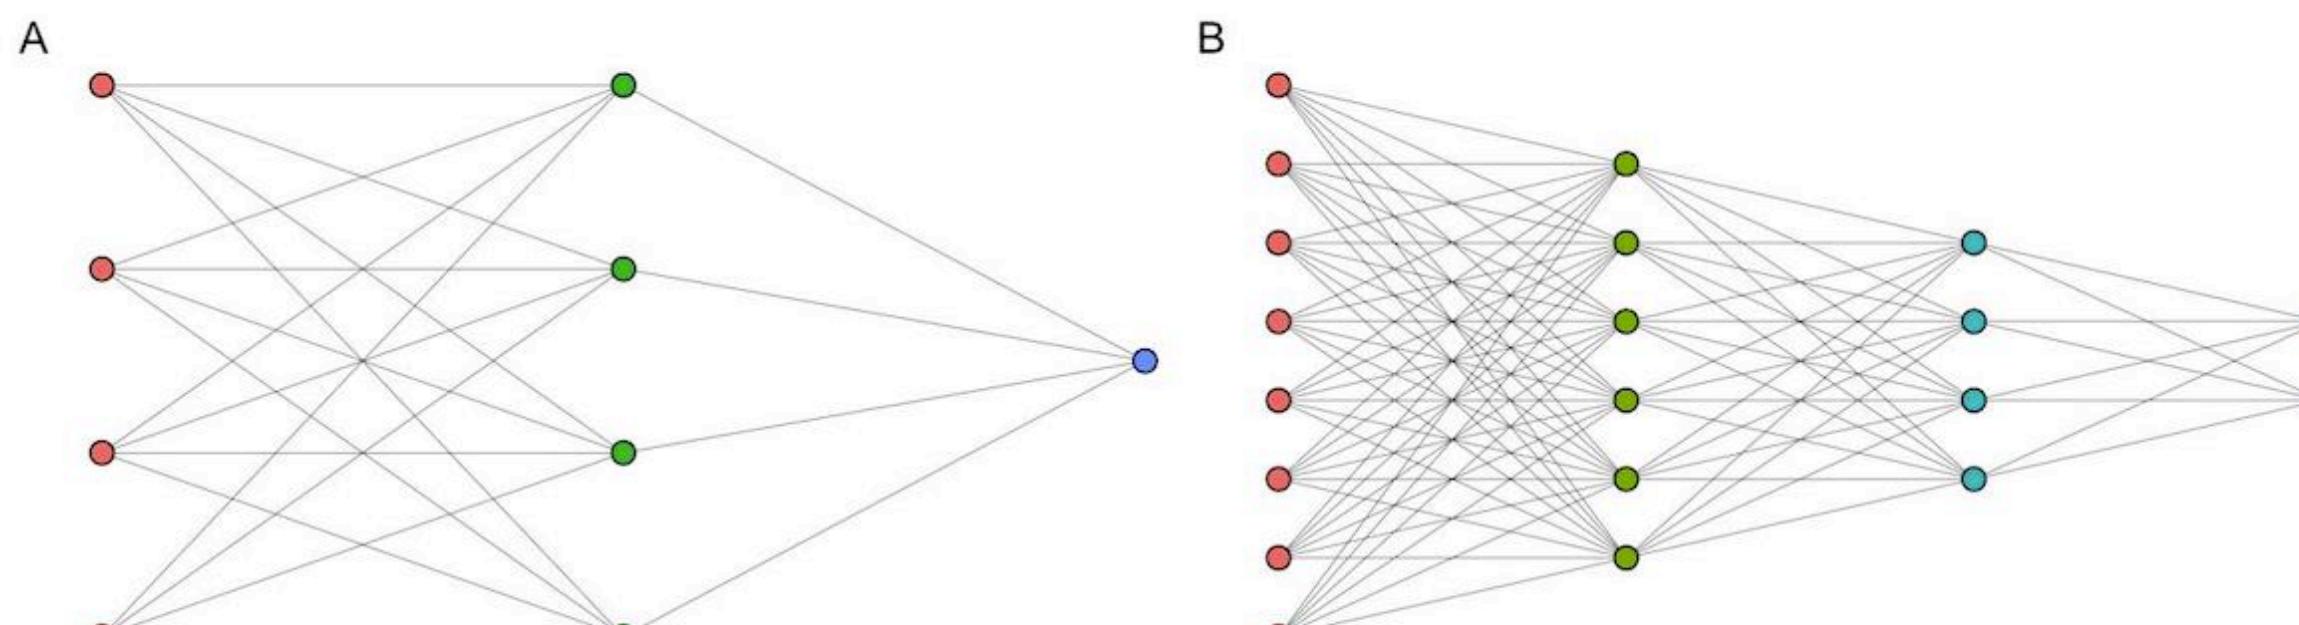
## Recurrent neural network architectures for language processing (Thursday)



<https://towardsdatascience.com/recurrent-neural-networks-rnn-explained-the-eli5-way-3956887e8b75>

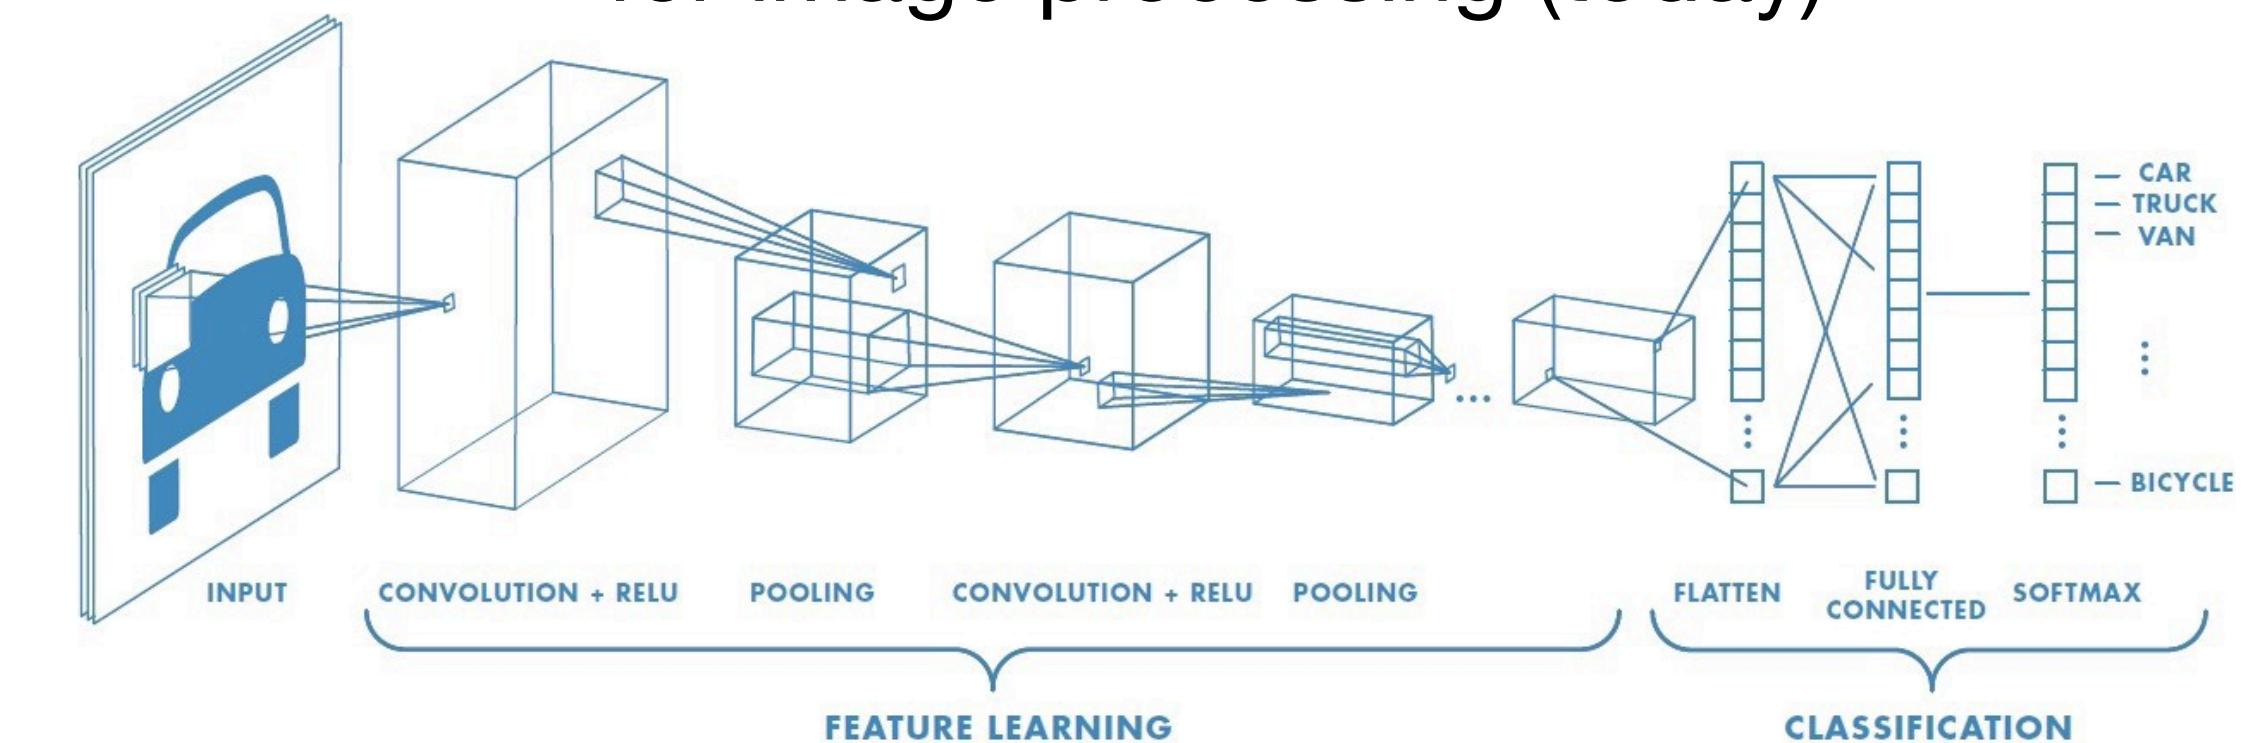
# Network architectures

## Fully connected architectures



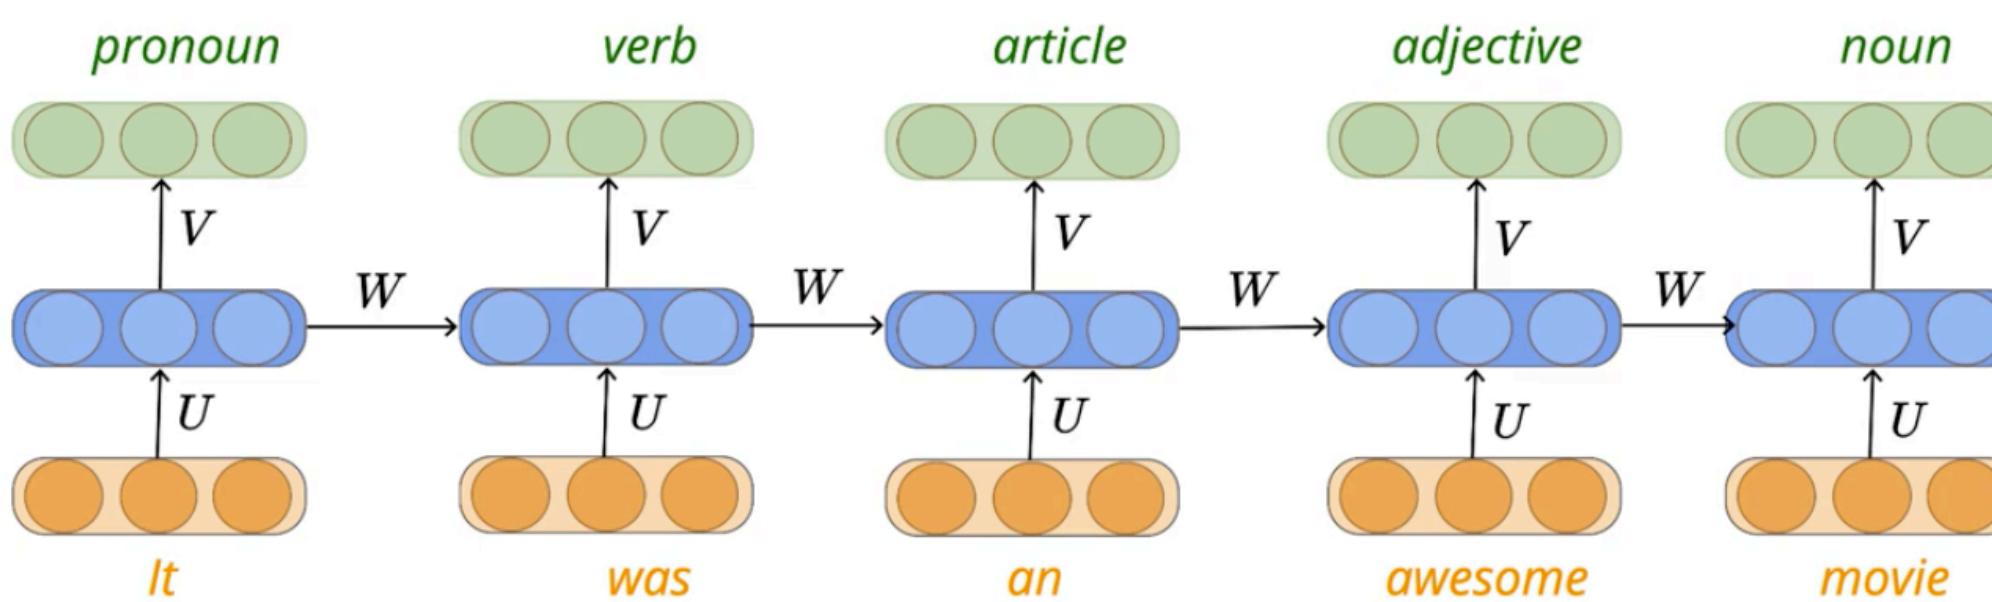
<https://community.rstudio.com/t/visualising-neural-network-architectures/41723>

## Convolutional neural network (CNN) architectures for image processing (today)



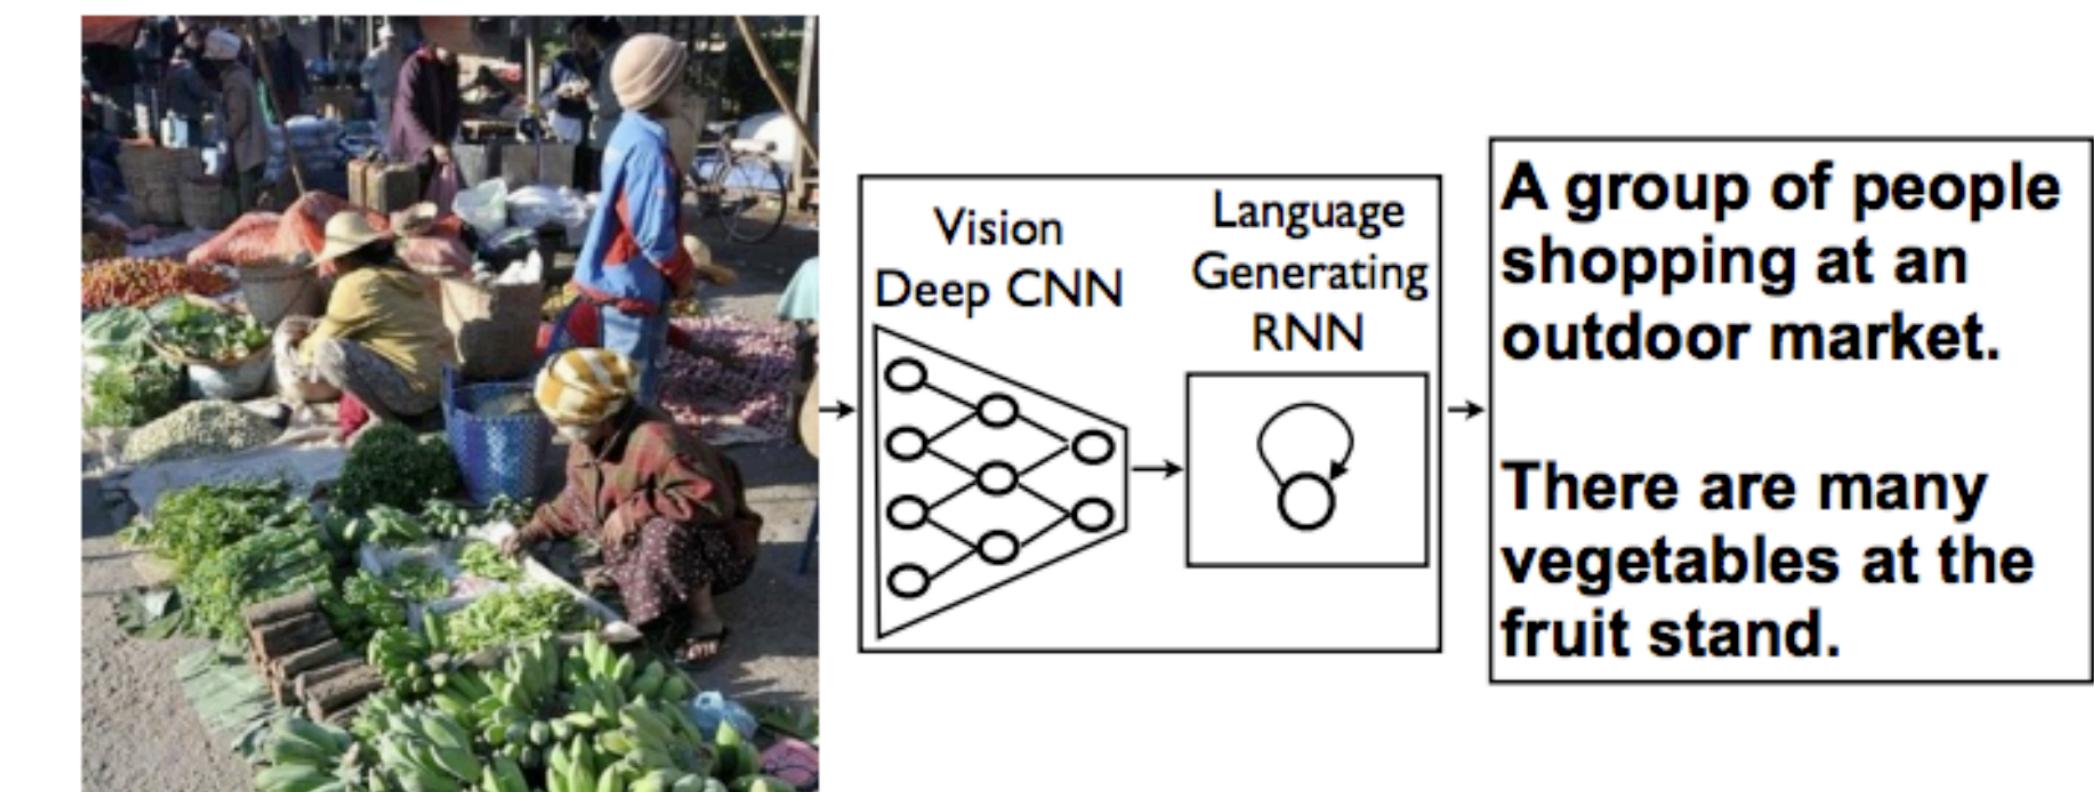
<https://towardsdatascience.com/a-comprehensive-guide-to-convolutional-neural-networks-the-eli5-way-3bd2b1164a53>

## Recurrent neural network architectures for language processing (Thursday)



<https://towardsdatascience.com/recurrent-neural-networks-rnn-explained-the-eli5-way-3956887e8b75>

Architecture components are modular and can be composed, e.g. image captioning



[https://subscription.packtpub.com/book/big\\_data\\_and\\_business\\_intelligence/9781788398060/3/ch03lvl1sec22/what-is-caption-generation](https://subscription.packtpub.com/book/big_data_and_business_intelligence/9781788398060/3/ch03lvl1sec22/what-is-caption-generation)

# Case study: Image classification

# Case study: Image classification

Prototypical computer vision task:

Given an image, classify according to what object it depicts.



<http://ai.stanford.edu/~olga/papers/iccv13-ILSVRCanalysis.pdf>

# Case study: Image classification

Prototypical computer vision task:

Given an image, classify according to what object it depicts.

Challenges:



<http://ai.stanford.edu/~olga/papers/iccv13-ILSVRCanalysis.pdf>

# Case study: Image classification

Prototypical computer vision task:

Given an image, classify according to what object it depicts.

Challenges:

- Viewpoint variation



<http://ai.stanford.edu/~olga/papers/iccv13-ILSVRCanalysis.pdf>

# Case study: Image classification

Prototypical computer vision task:

Given an image, classify according to what object it depicts.

Challenges:

- Viewpoint variation
- Illumination



<http://ai.stanford.edu/~olga/papers/iccv13-ILSVRCanalysis.pdf>

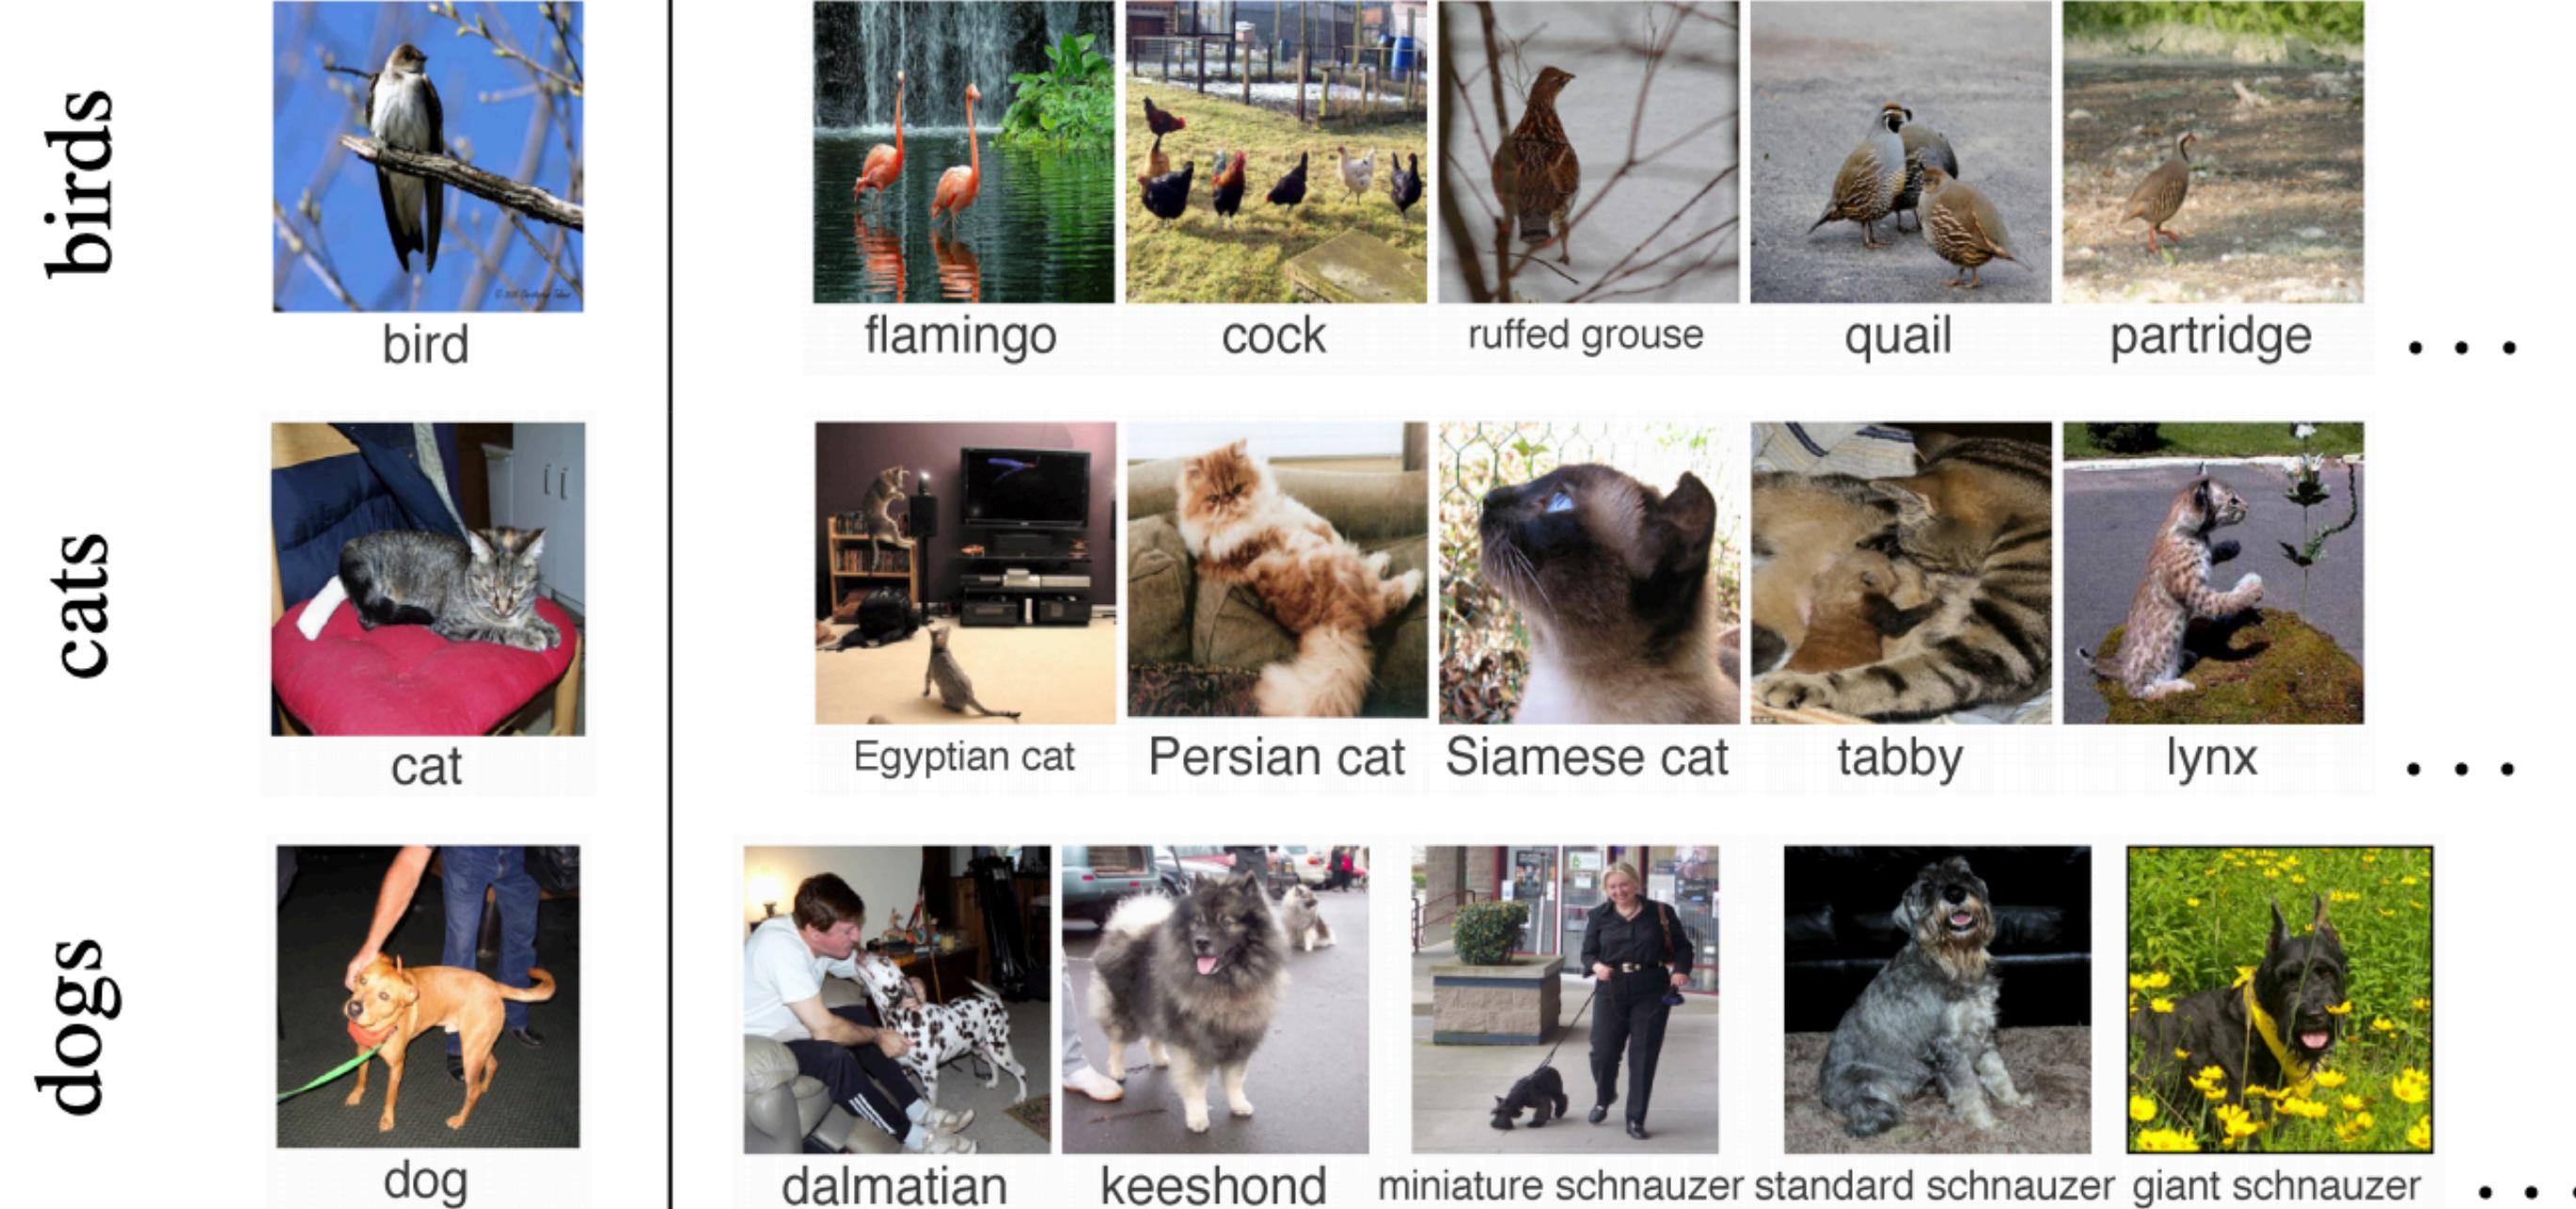
# Case study: Image classification

Prototypical computer vision task:

Given an image, classify according to what object it depicts.

Challenges:

- Viewpoint variation
- Illumination
- Deformation



<http://ai.stanford.edu/~olga/papers/iccv13-ILSVRCanalysis.pdf>

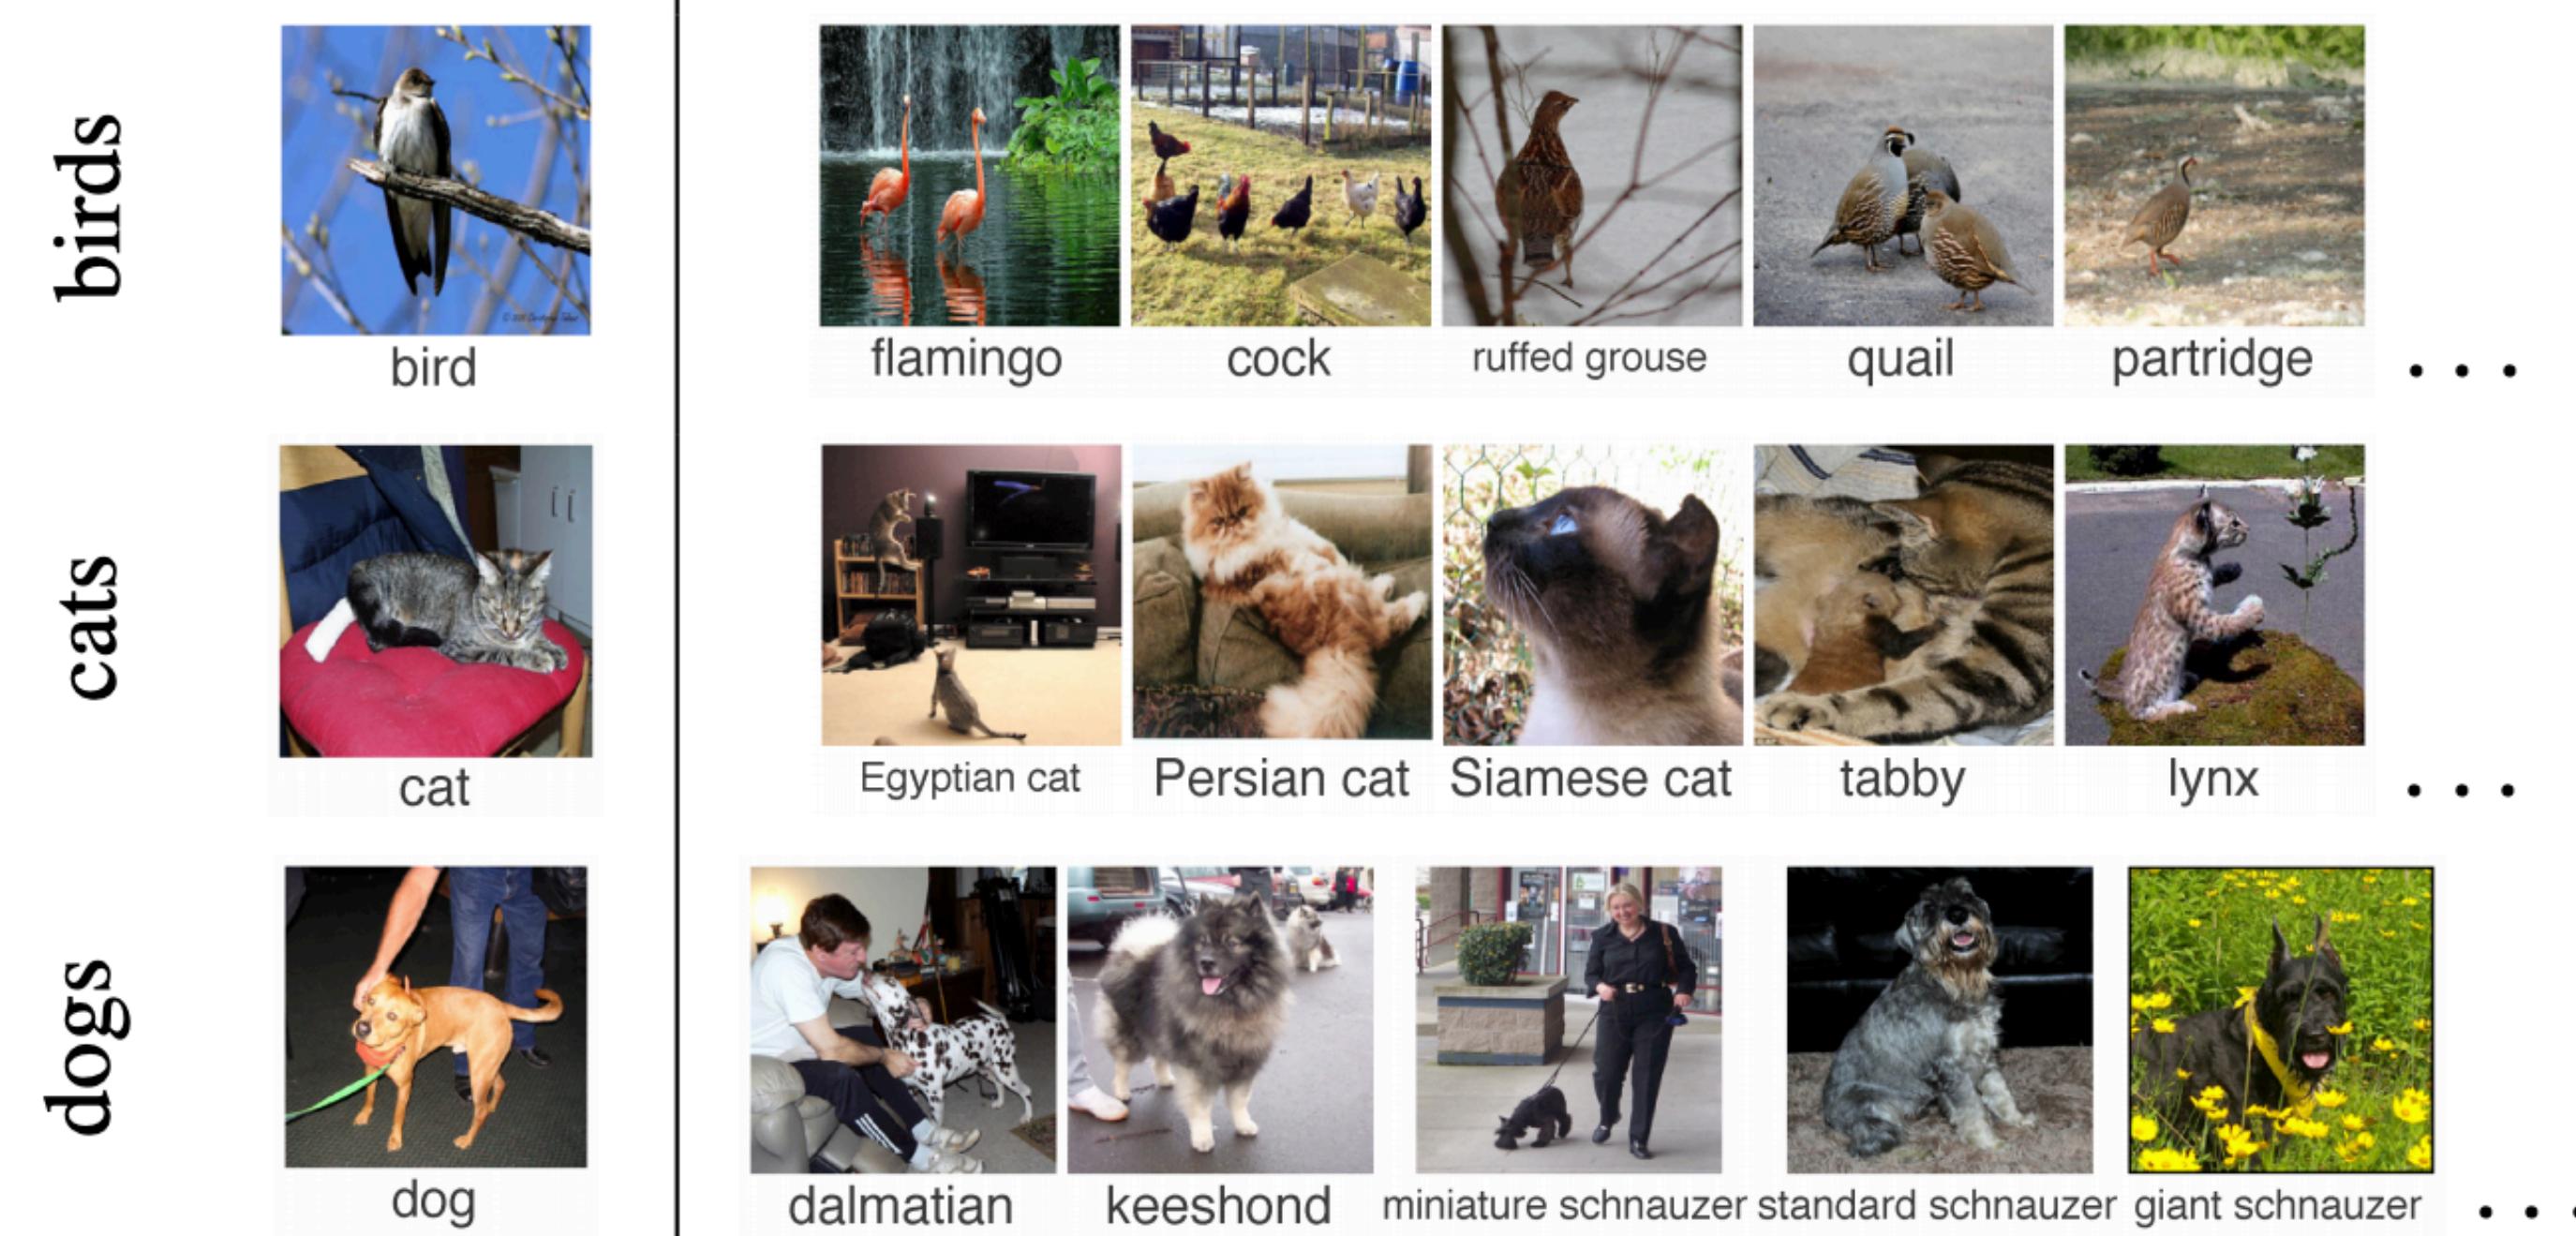
# Case study: Image classification

Prototypical computer vision task:

Given an image, classify according to what object it depicts.

Challenges:

- Viewpoint variation
- Illumination
- Deformation
- Occlusion



<http://ai.stanford.edu/~olga/papers/iccv13-ILSVRCanalysis.pdf>

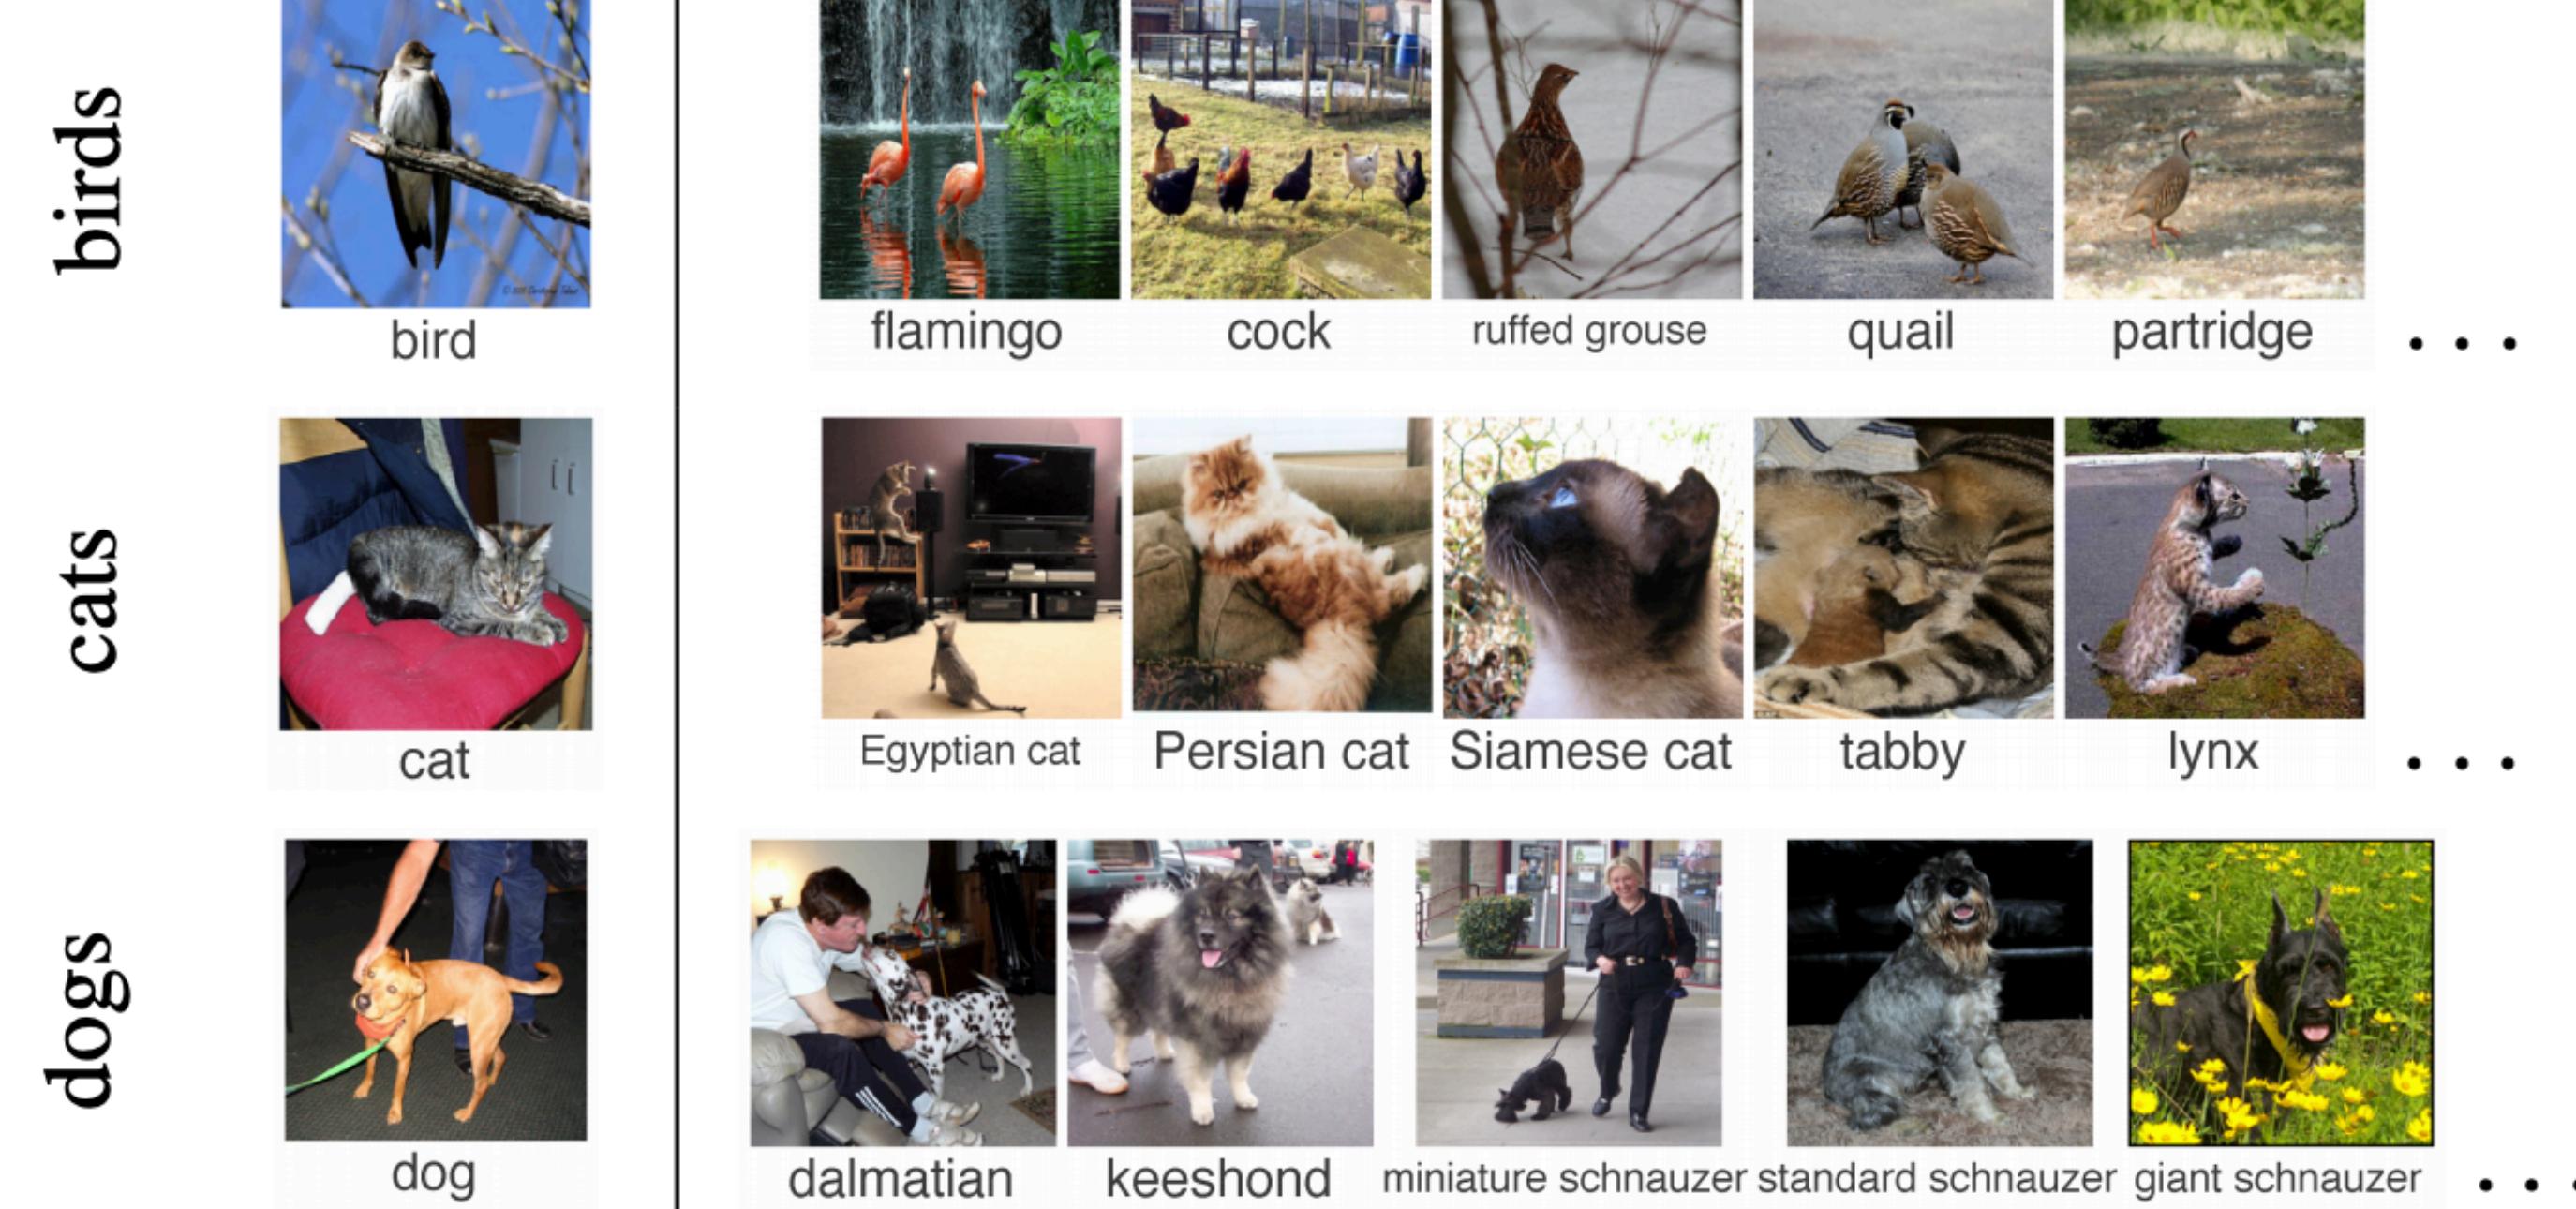
# Case study: Image classification

Prototypical computer vision task:

Given an image, classify according to what object it depicts.

Challenges:

- Viewpoint variation
- Illumination
- Deformation
- Occlusion
- Background clutter



<http://ai.stanford.edu/~olga/papers/iccv13-ILSVRCanalysis.pdf>

# Case study: Image classification

Prototypical computer vision task:

Given an image, classify according to what object it depicts.

Challenges:

- Viewpoint variation
- Illumination
- Deformation
- Occlusion
- Background clutter
- Intraclass variation



<http://ai.stanford.edu/~olga/papers/iccv13-ILSVRCanalysis.pdf>

# ImageNet

## A large dataset for image classification

# ImageNet

## A large dataset for image classification

Assembled in 2009 by downloading lots of images from the web and crowdsourcing their labels.



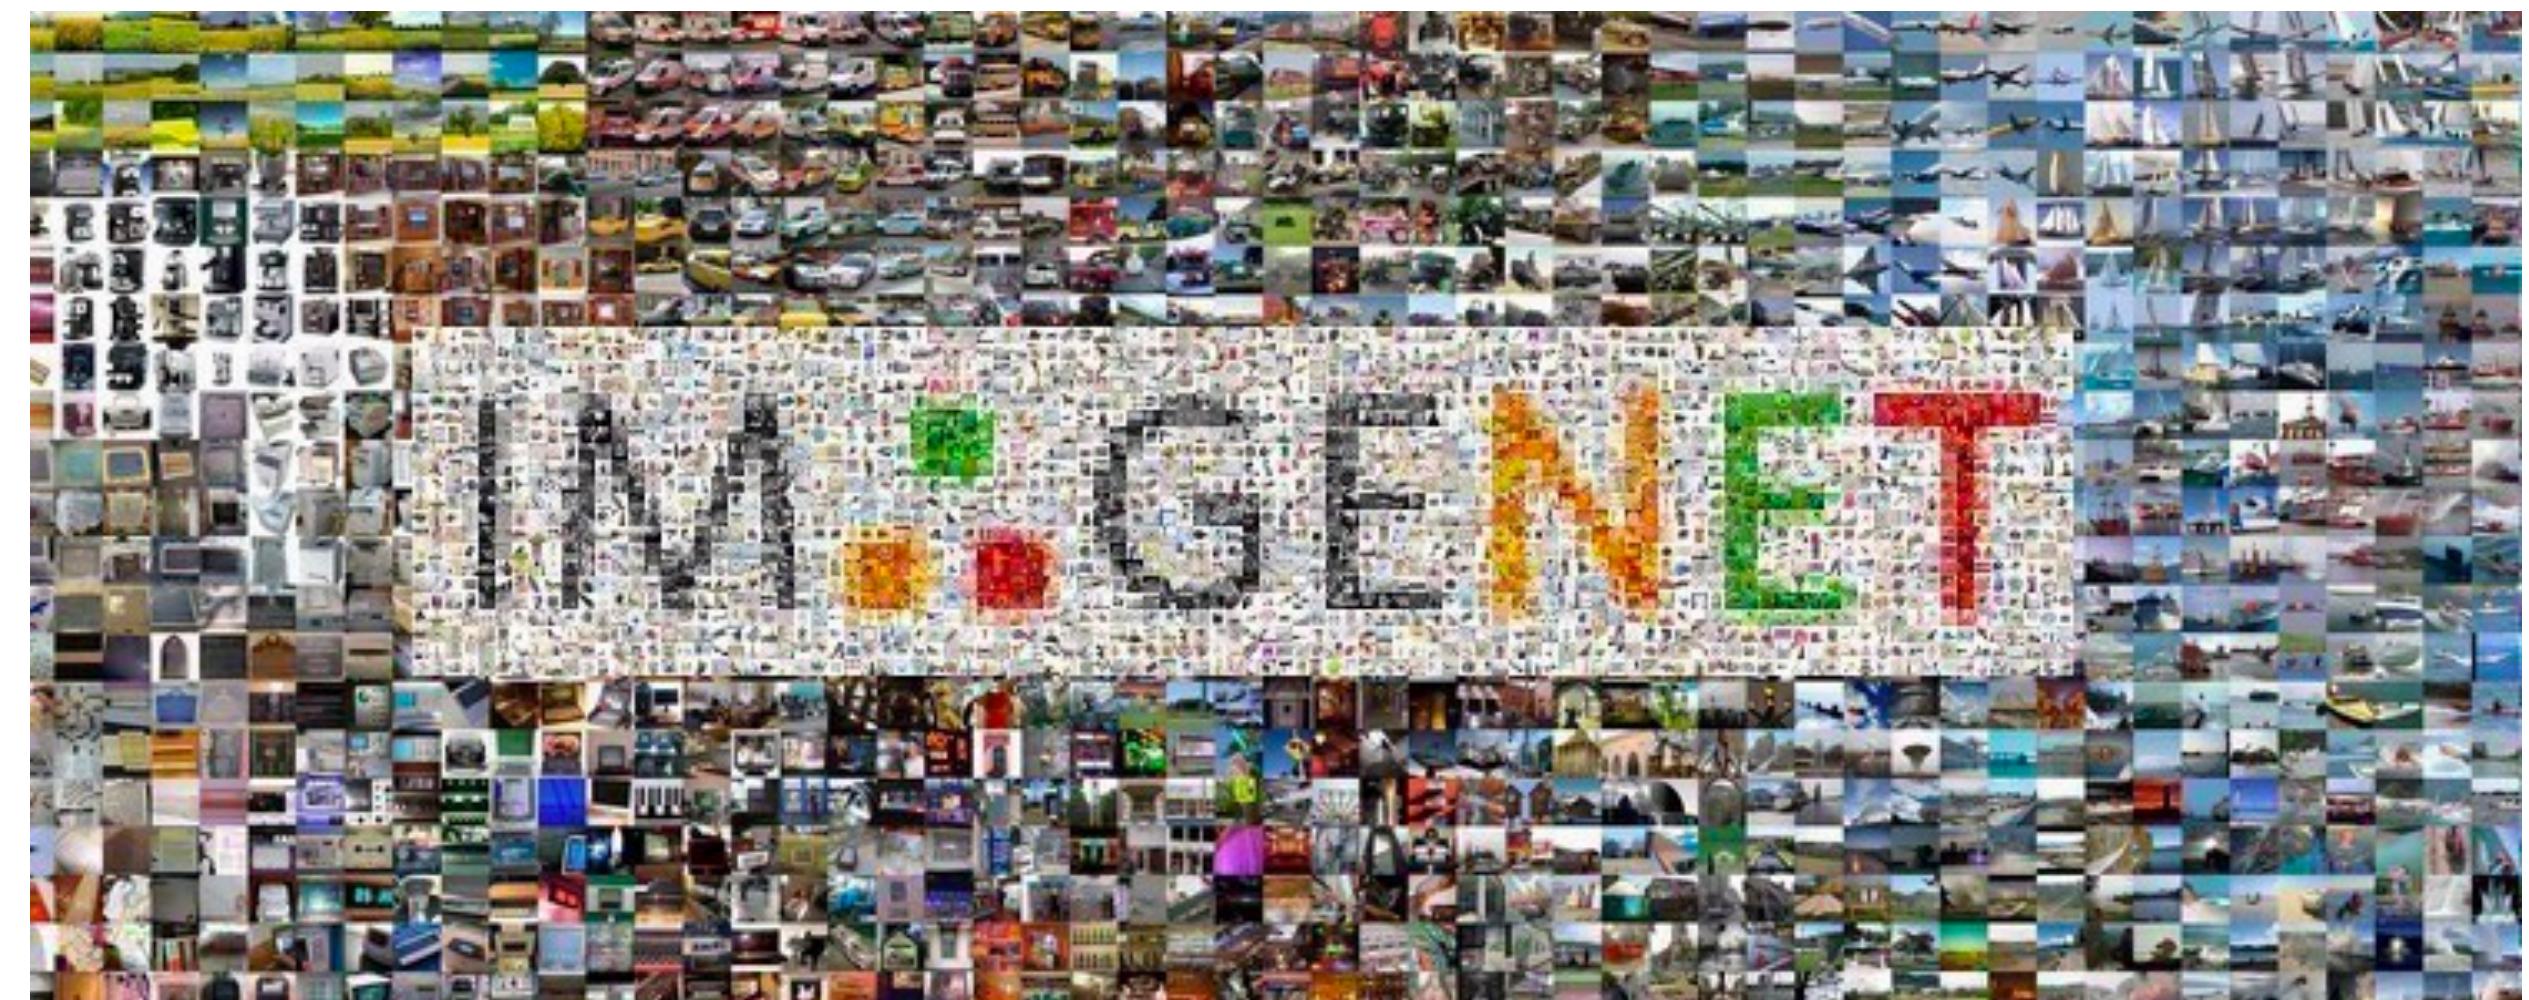
<https://medium.com/syncedreview/sensetime-trains-imagenet-alexnet-in-record-1-5-minutes-e944ab049b2c>

# ImageNet

## A large dataset for image classification

Assembled in 2009 by downloading lots of images from the web and crowdsourcing their labels.

- Training set: 1.2 million images
- Test set: 100,000 images
- 1000 classes



<https://medium.com/syncedreview/sensetime-trains-imagenet-alexnet-in-record-1-5-minutes-e944ab049b2c>

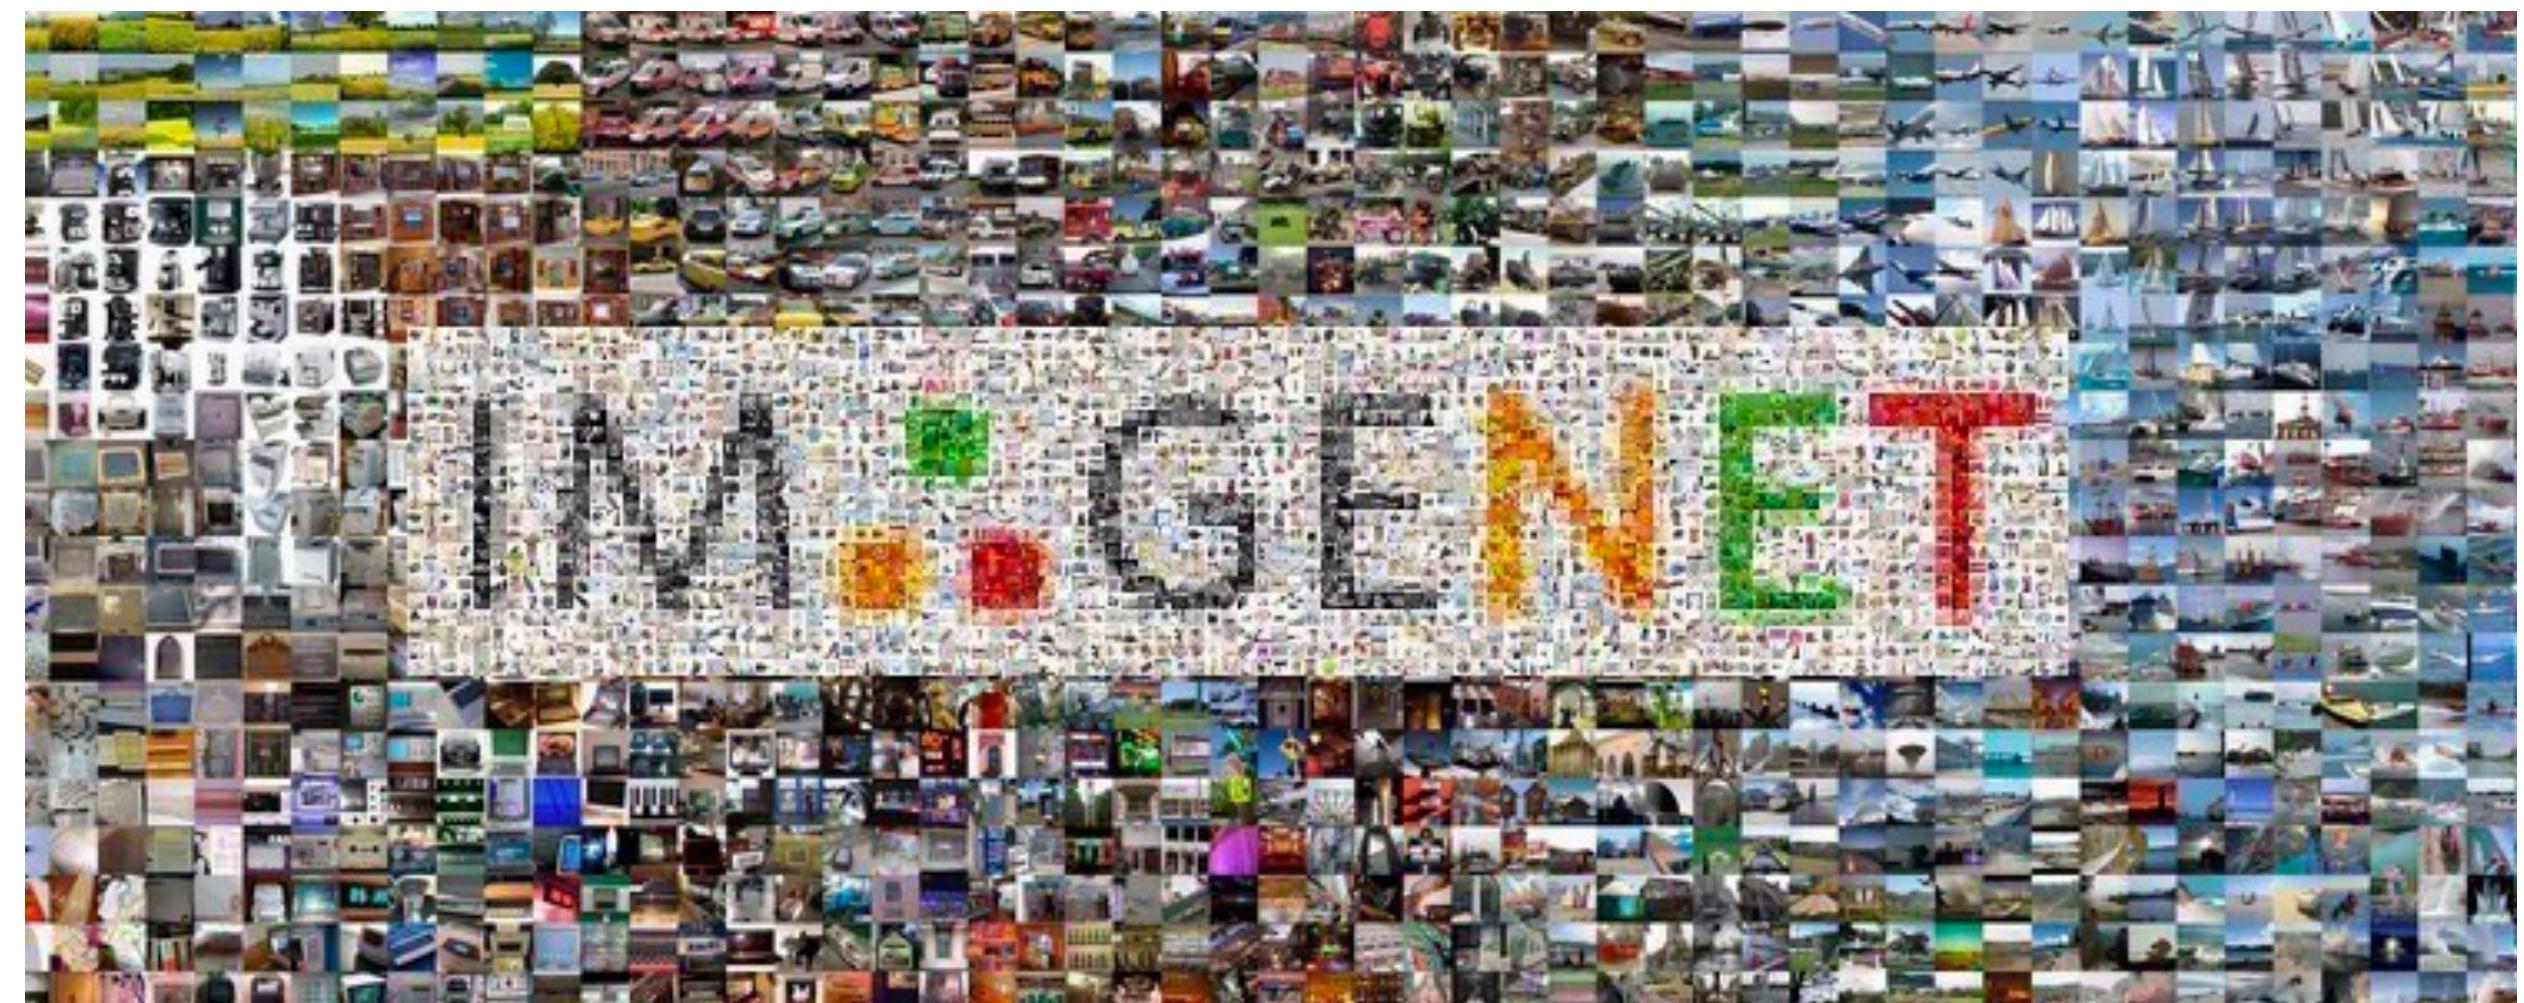
# ImageNet

## A large dataset for image classification

Assembled in 2009 by downloading lots of images from the web and crowdsourcing their labels.

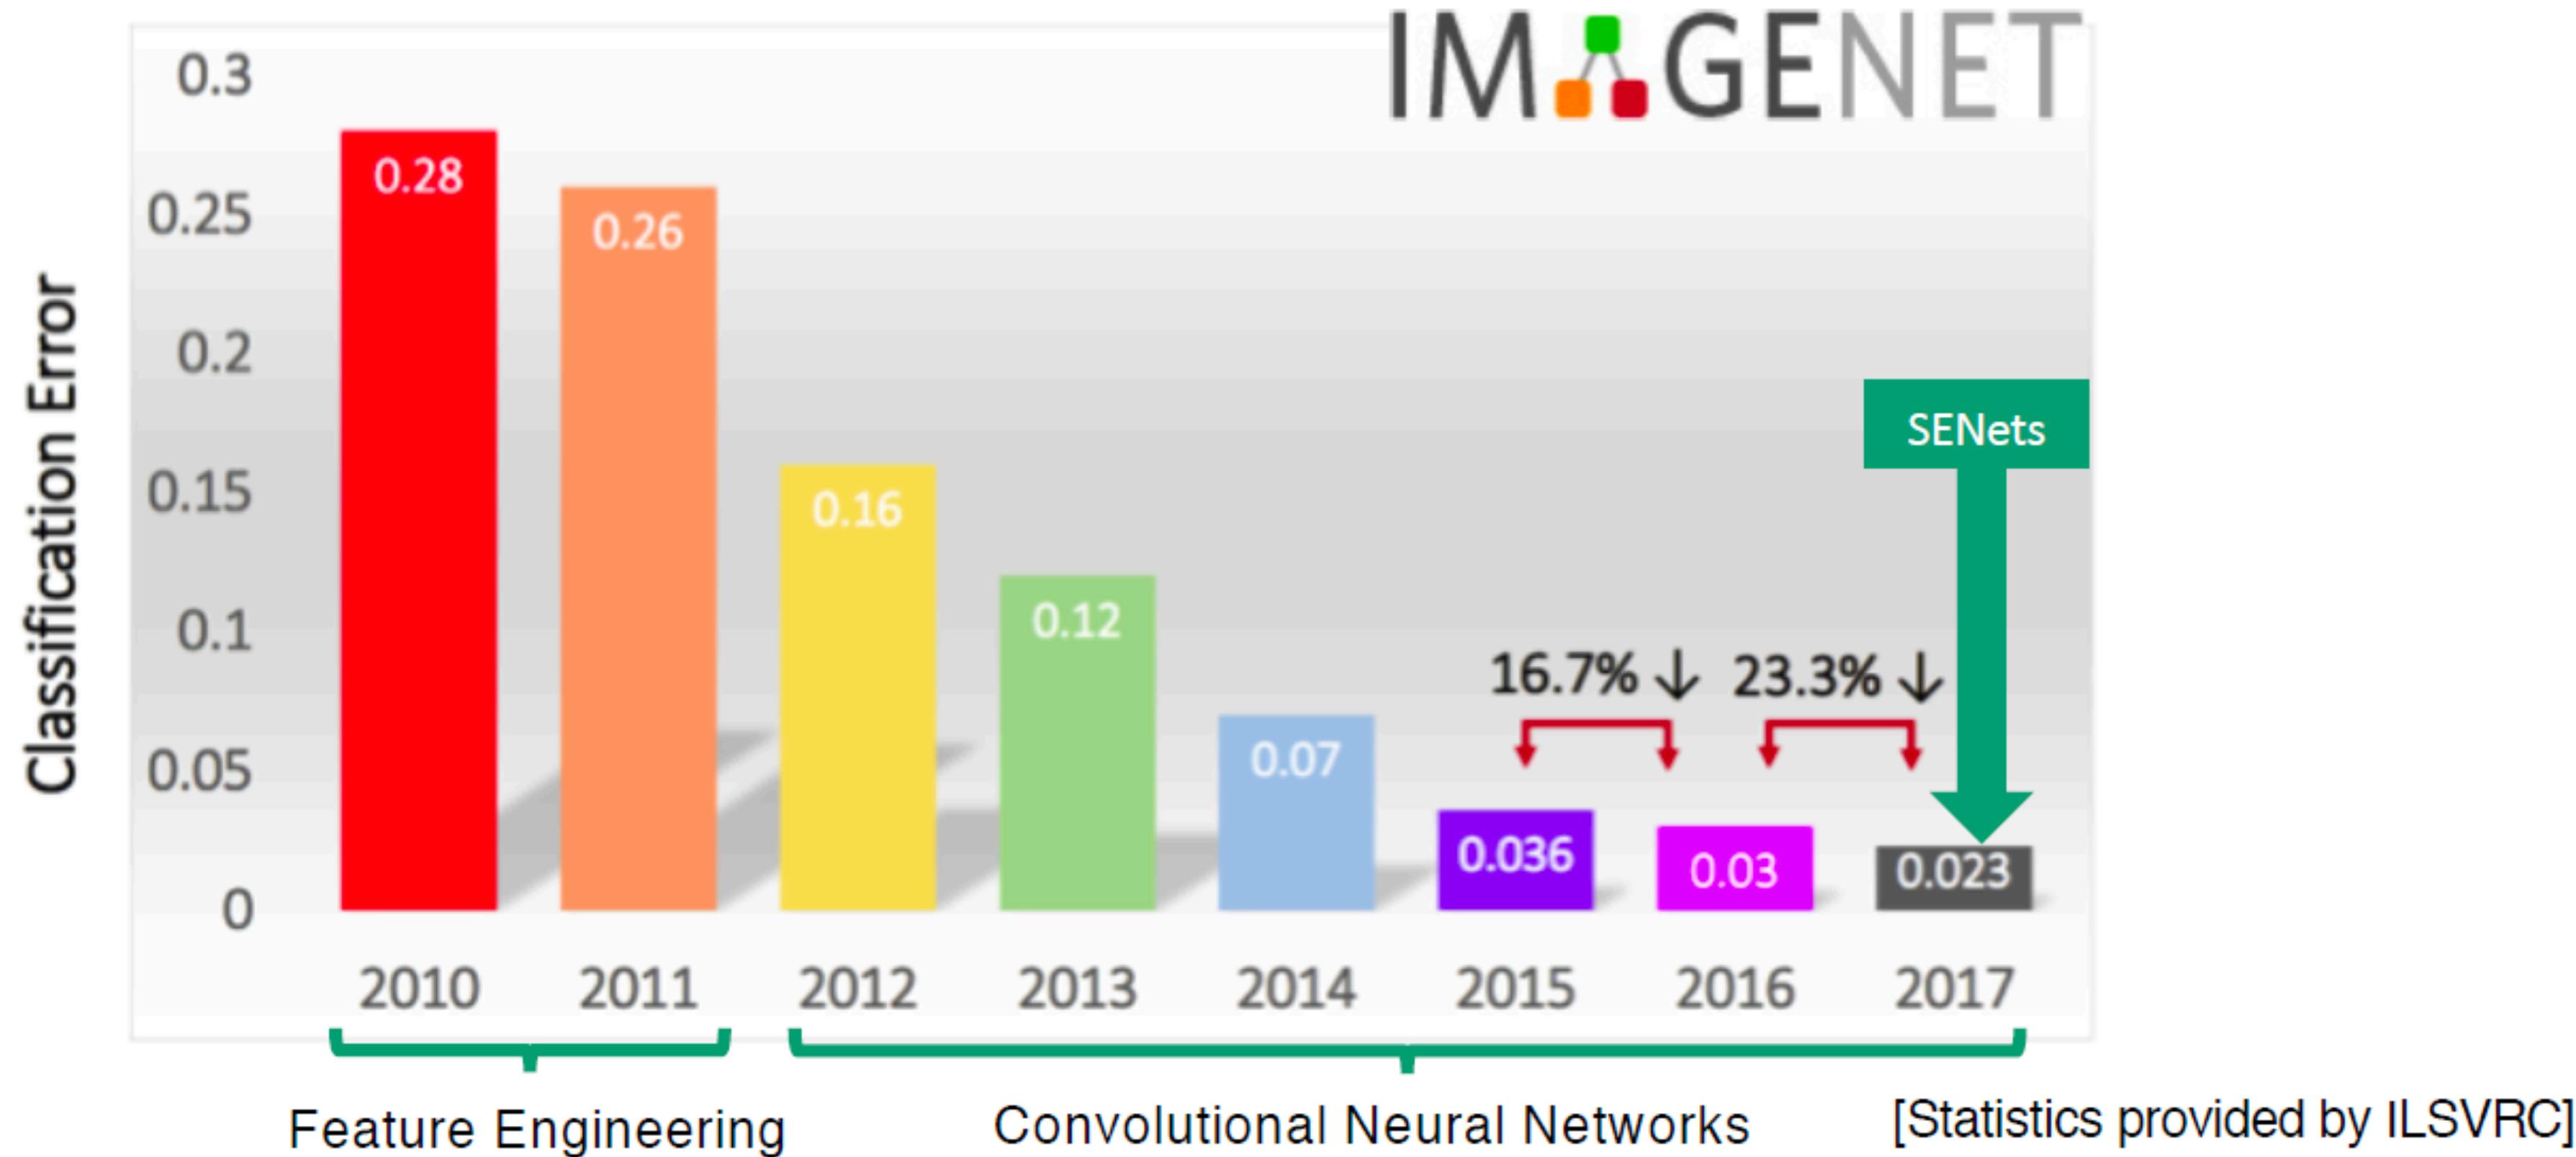
- Training set: 1.2 million images
- Test set: 100,000 images
- 1000 classes

ImageNet Large Scale Visual Recognition Challenge (ILSVRC) held annually between 2010-2017.

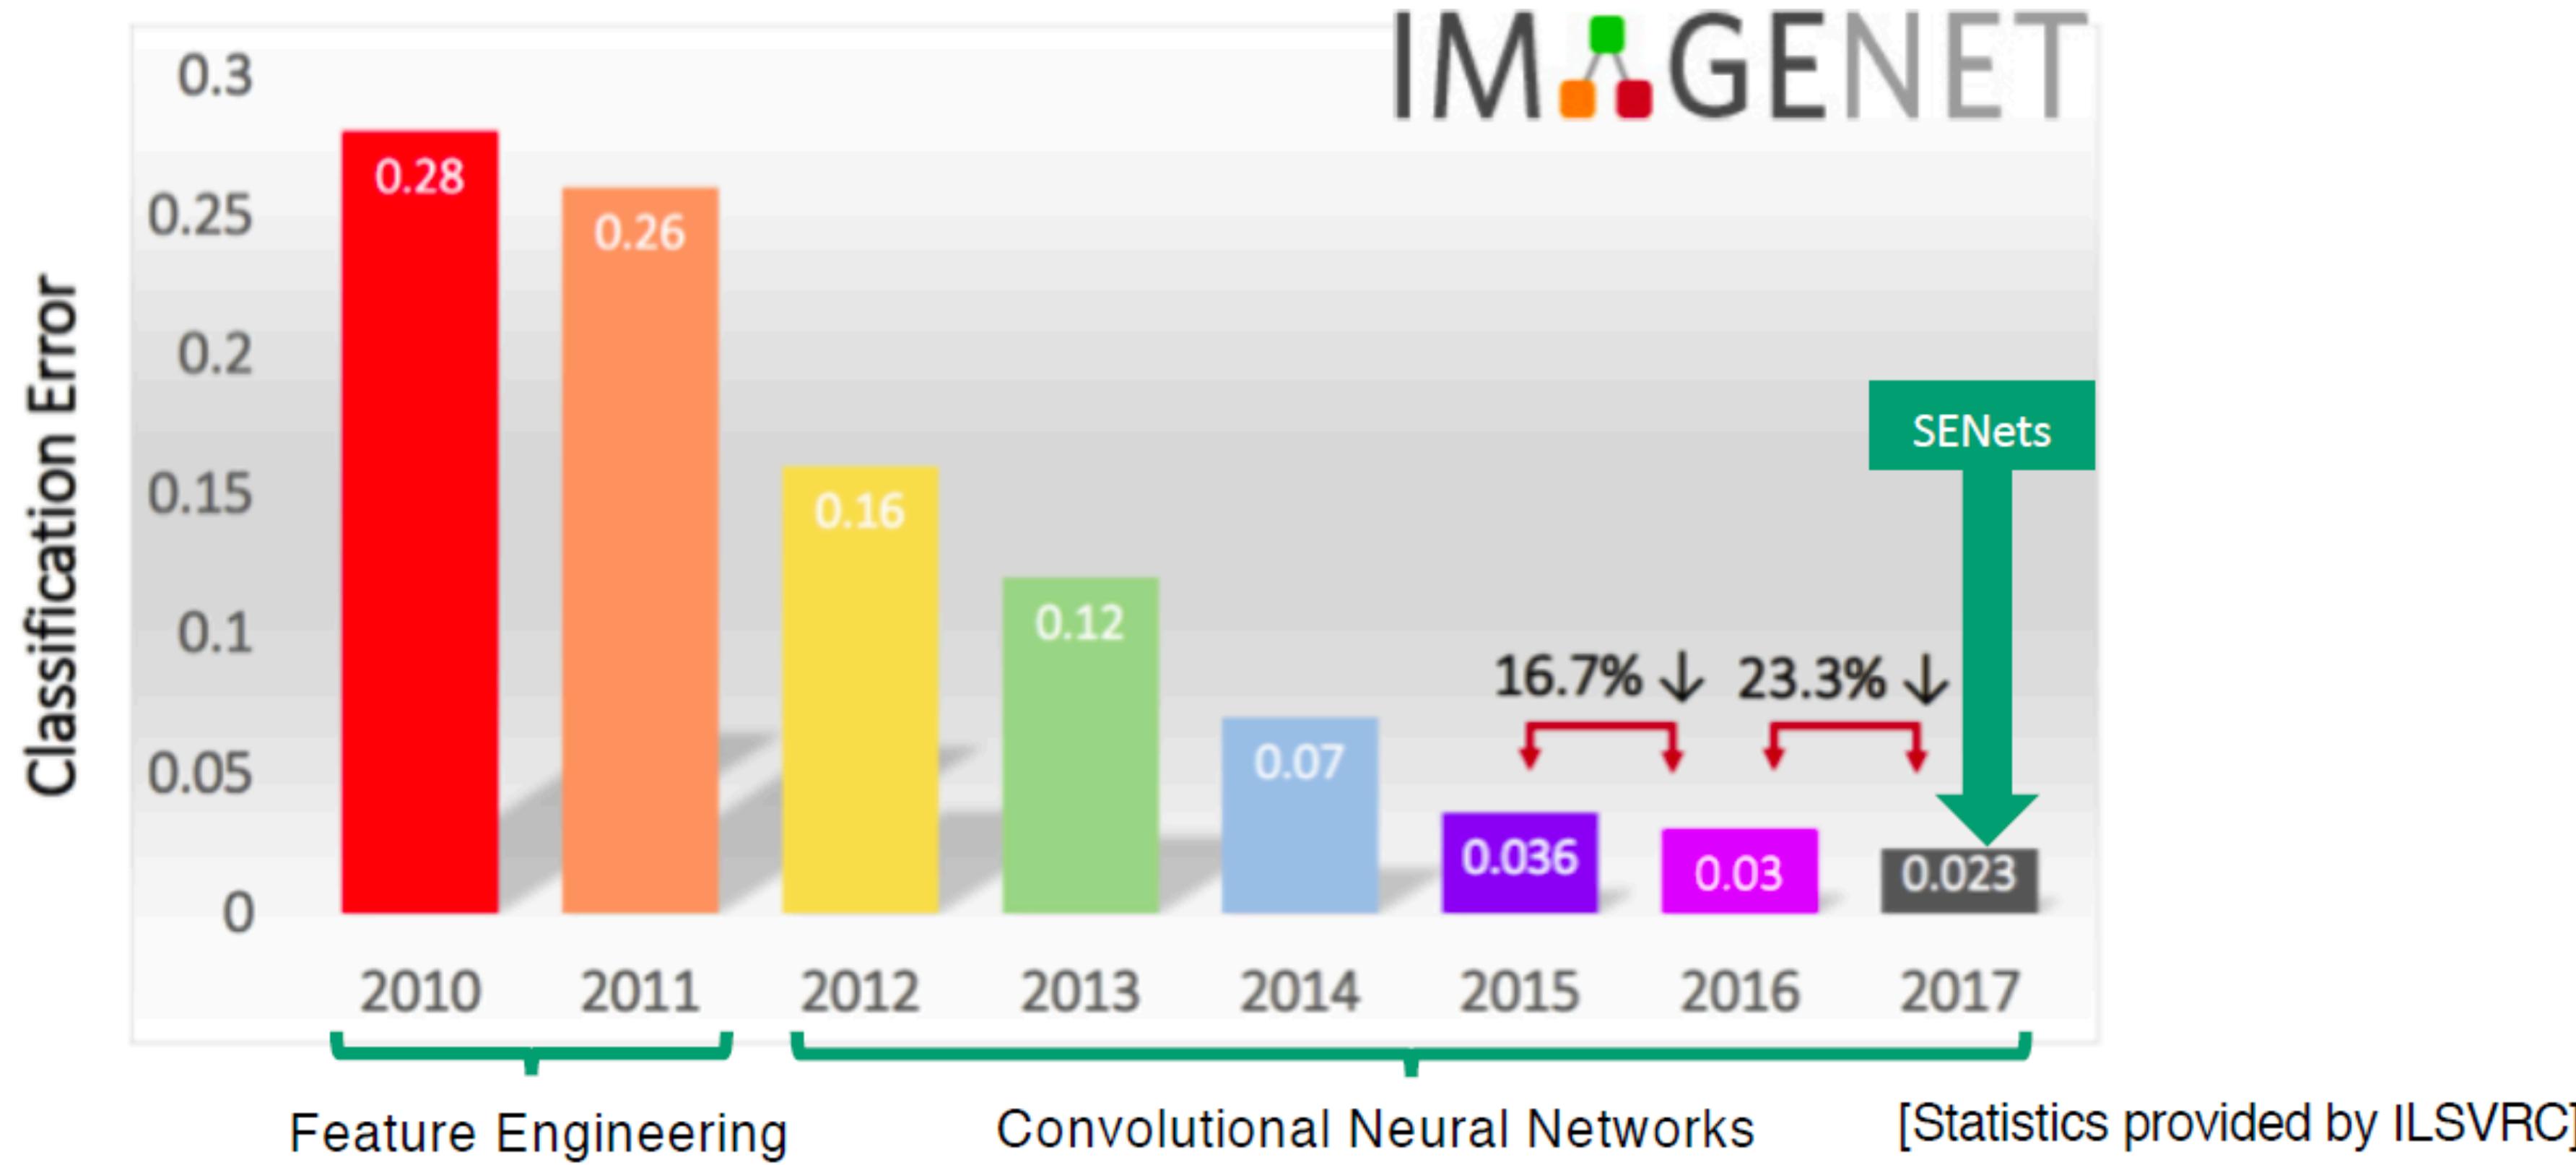


<https://medium.com/syncedreview/sensetime-trains-imagenet-alexnet-in-record-1-5-minutes-e944ab049b2c>

# ILSVRC results over the years



# ILSVRC results over the years

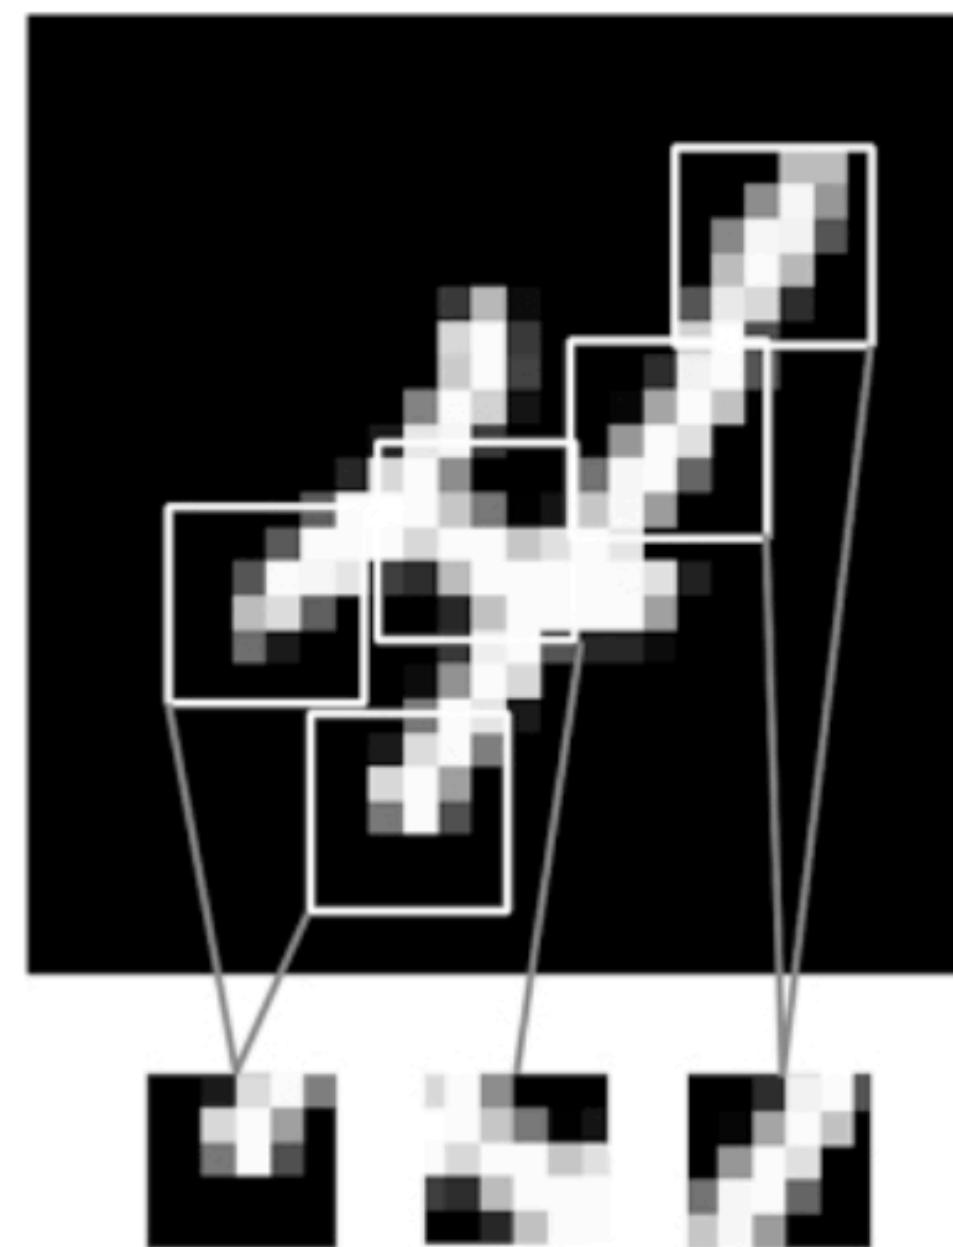


Convolutional neural networks (CNNs) have dominated since 2012.

# CNNs are built on image-specific properties

# CNNs are built on image-specific properties

Figure 5.1. Images can be broken into local patterns such as edges, textures, and so on.



# CNNs are built on image-specific properties

Figure 5.1. Images can be broken into local patterns such as edges, textures, and so on.

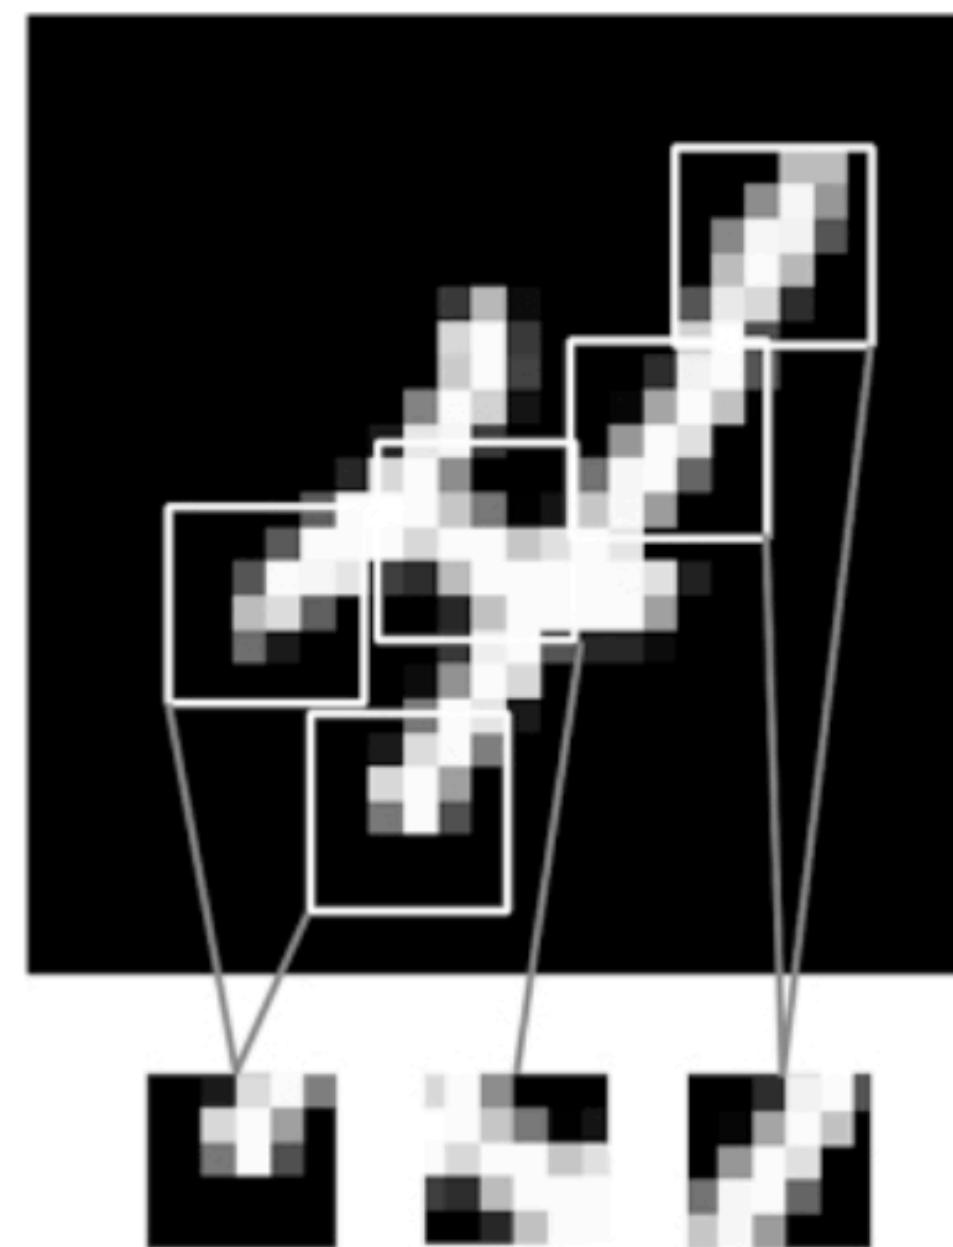
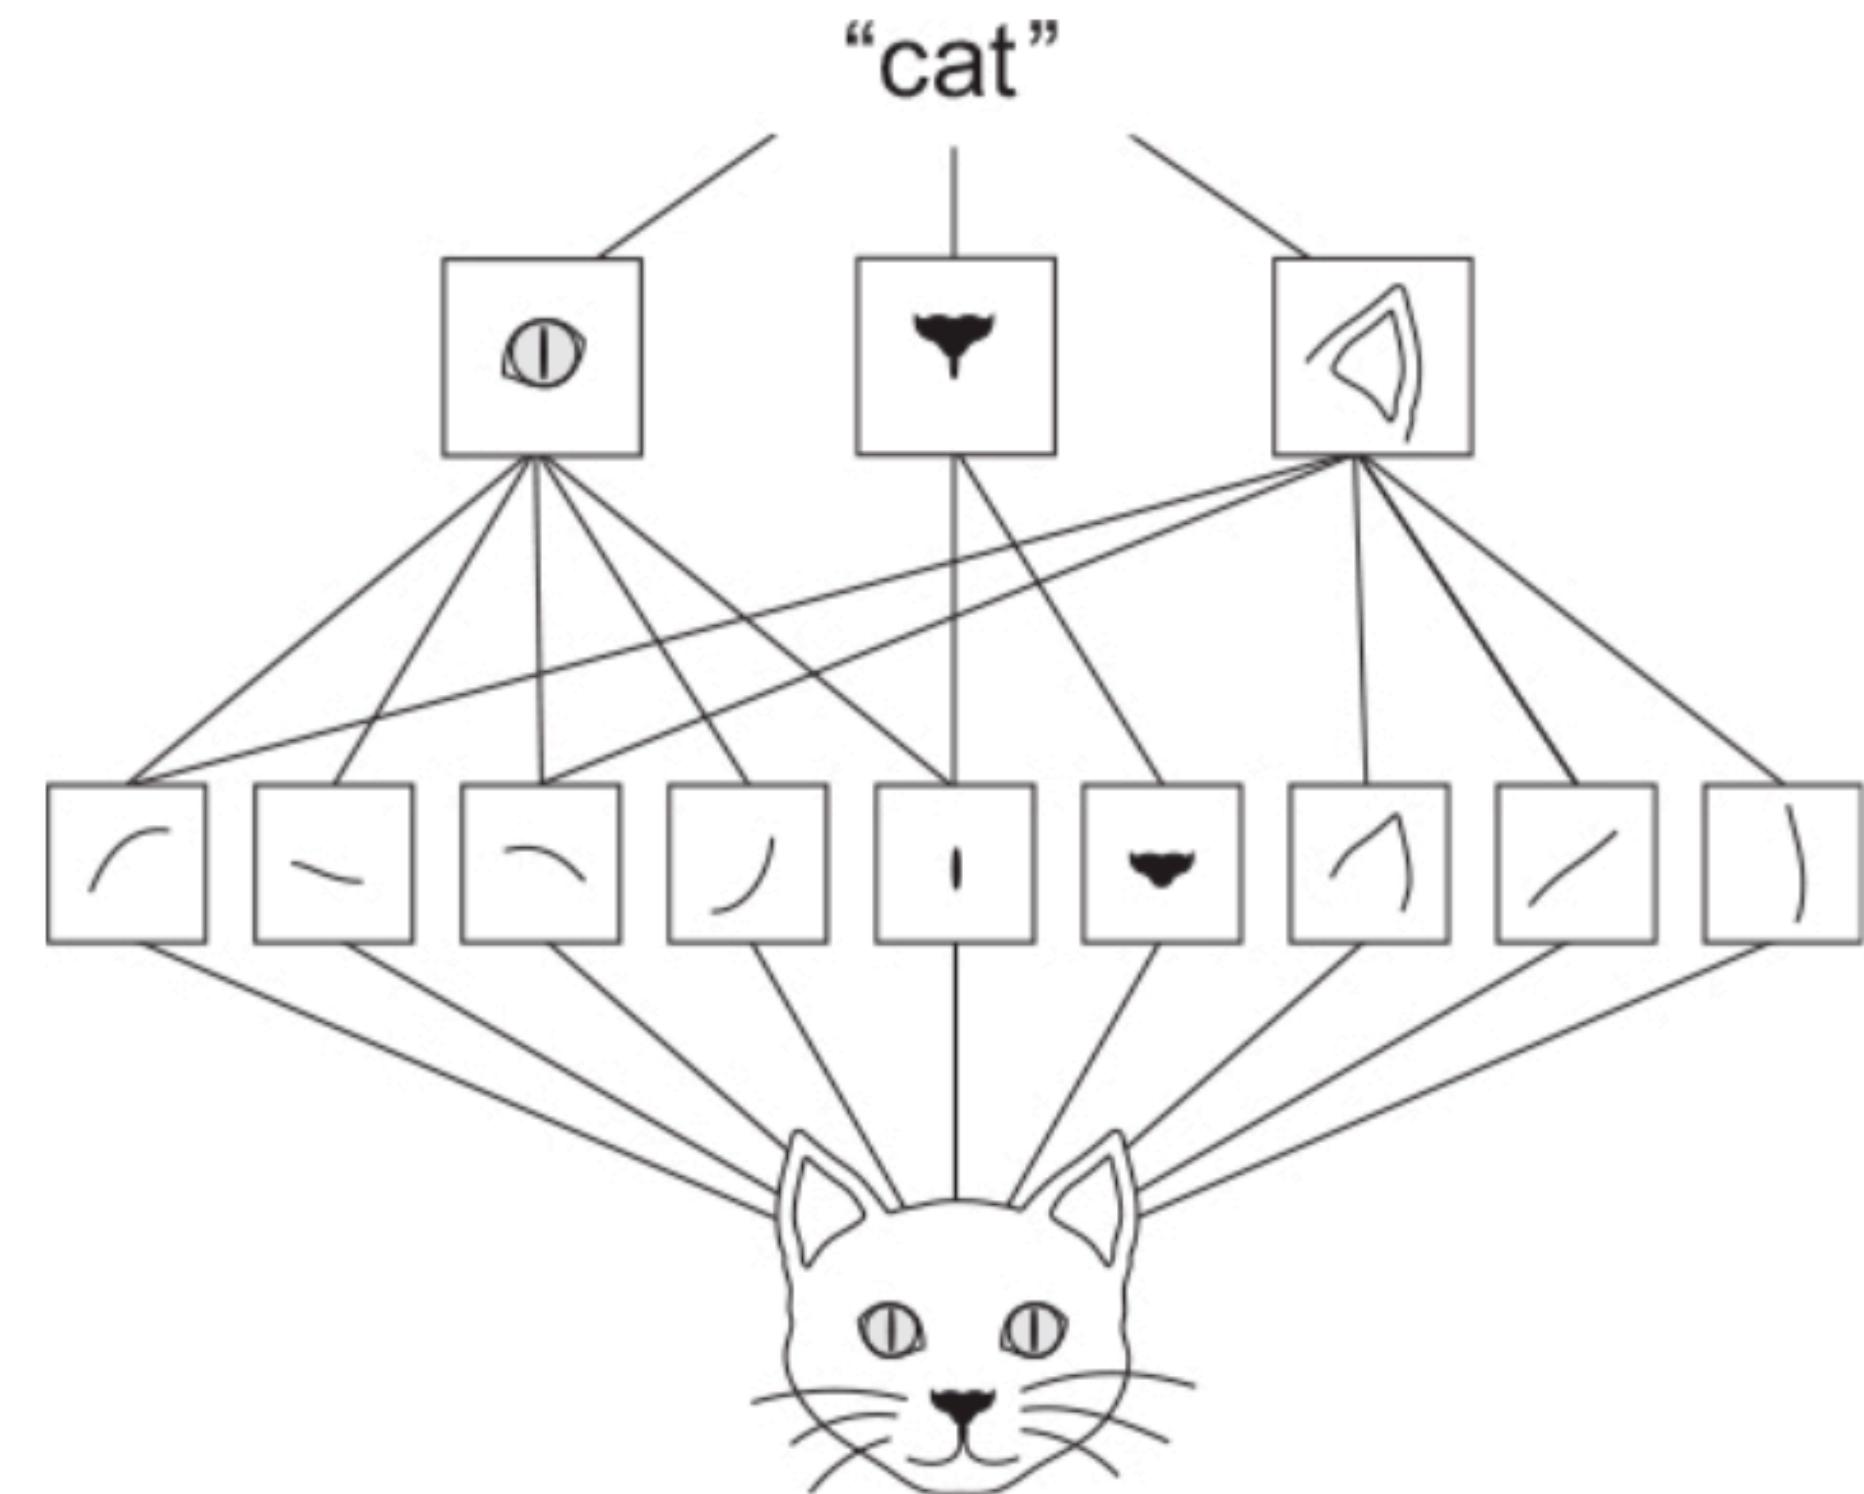


Figure 5.2. The visual world forms a spatial hierarchy of visual modules: hyperlocal edges combine into local objects such as eyes or ears, which combine into high-level concepts such as “cat.”



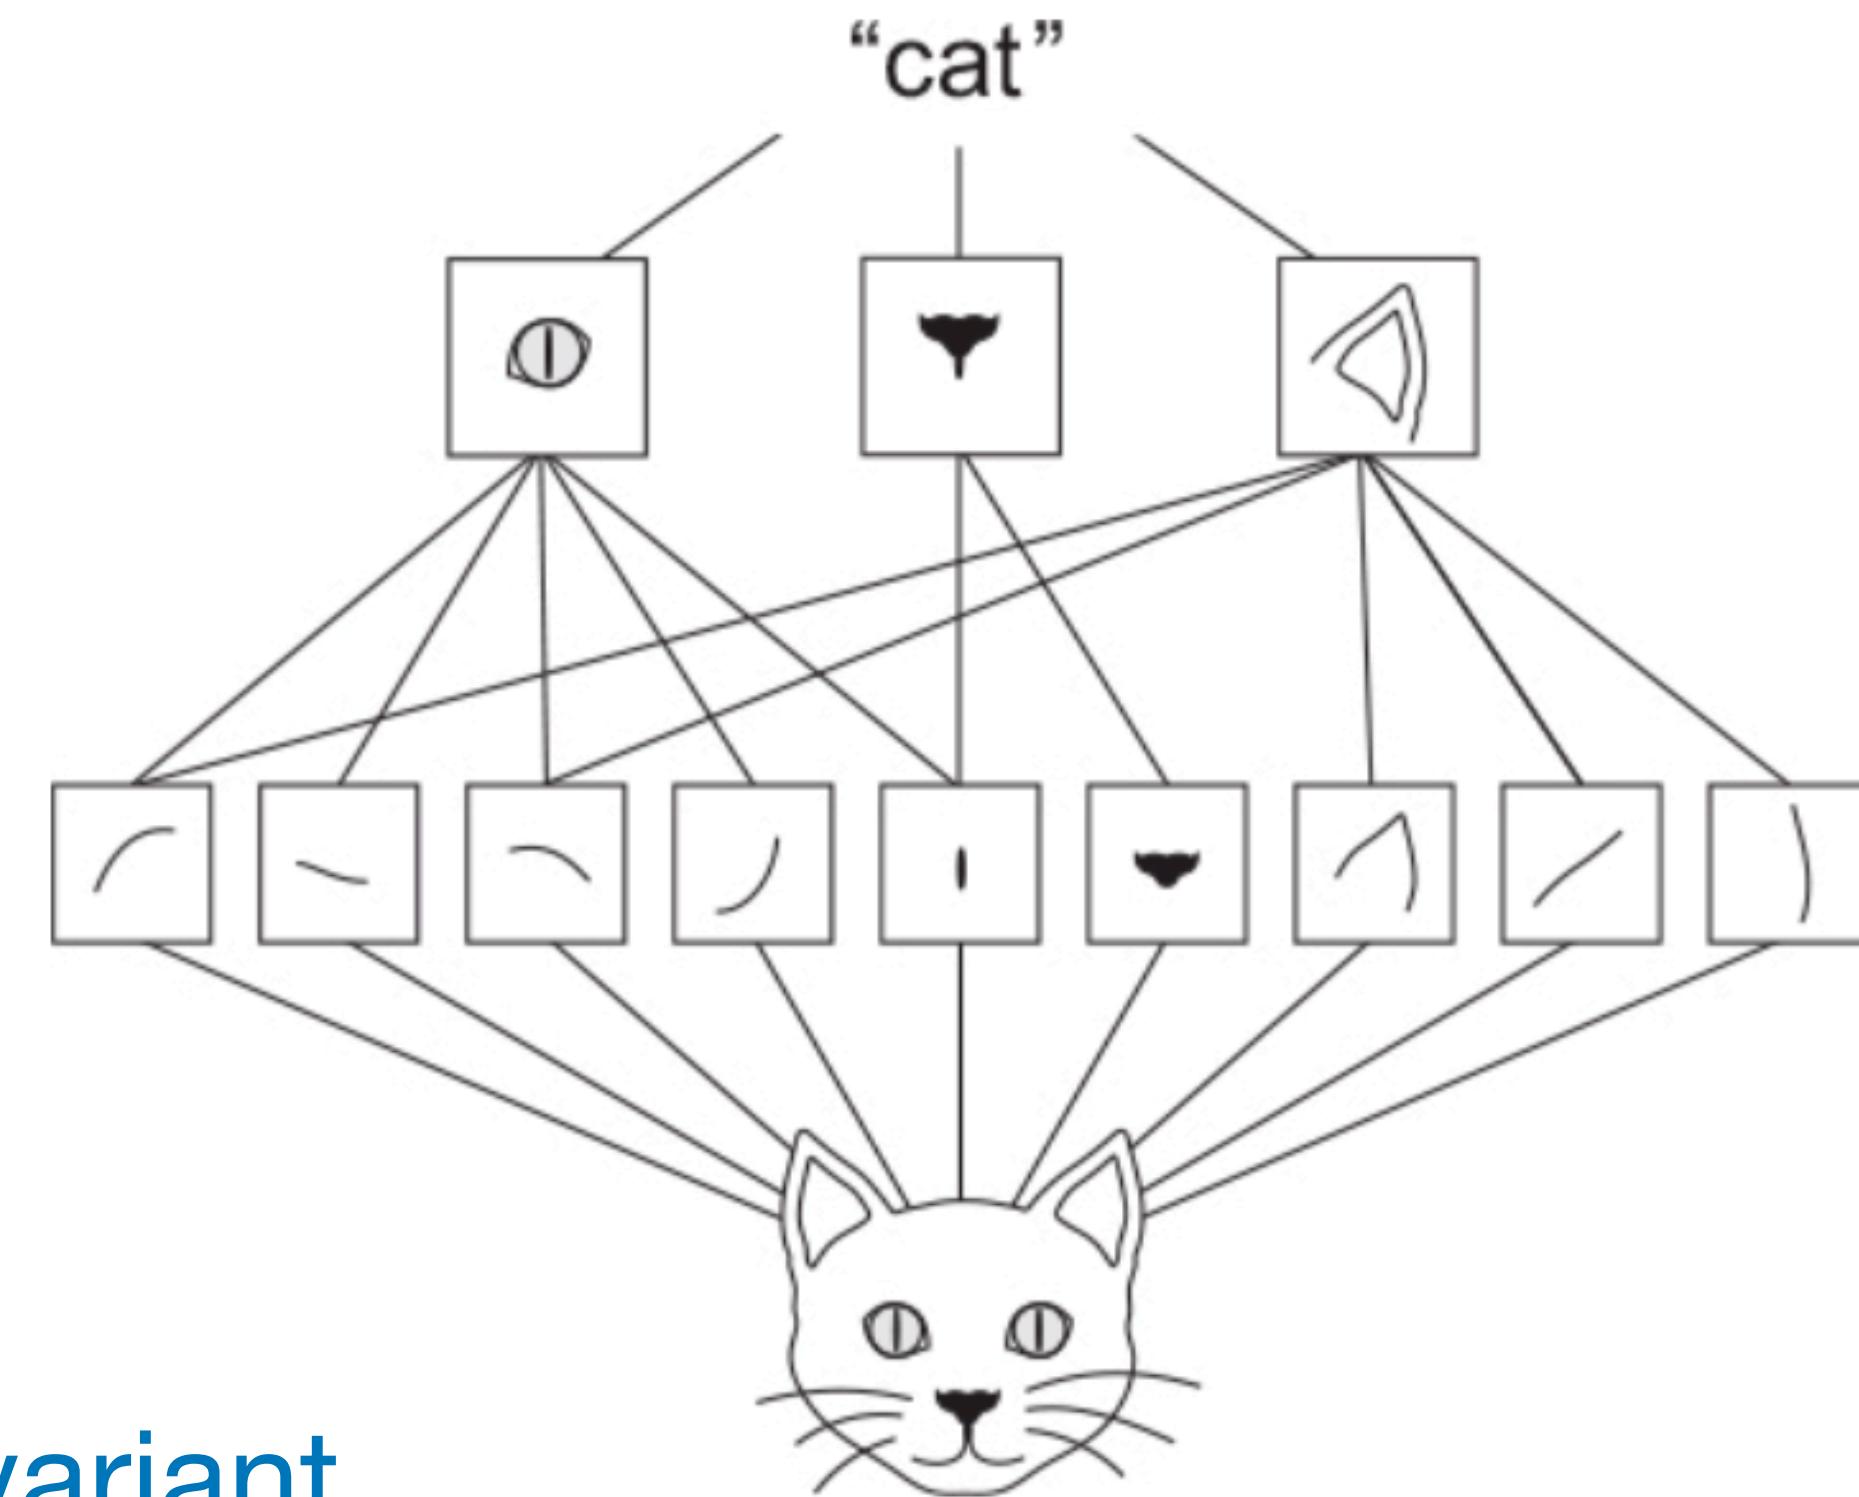
# CNNs are built on image-specific properties

Figure 5.1. Images can be broken into local patterns such as edges, textures, and so on.



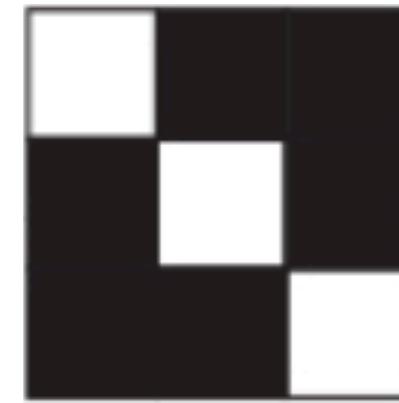
- Patterns are
- Local
  - Translation-invariant
  - Hierarchical

Figure 5.2. The visual world forms a spatial hierarchy of visual modules: hyperlocal edges combine into local objects such as eyes or ears, which combine into high-level concepts such as “cat.”



# Convolution: Searching for patterns

Filter (3x3)



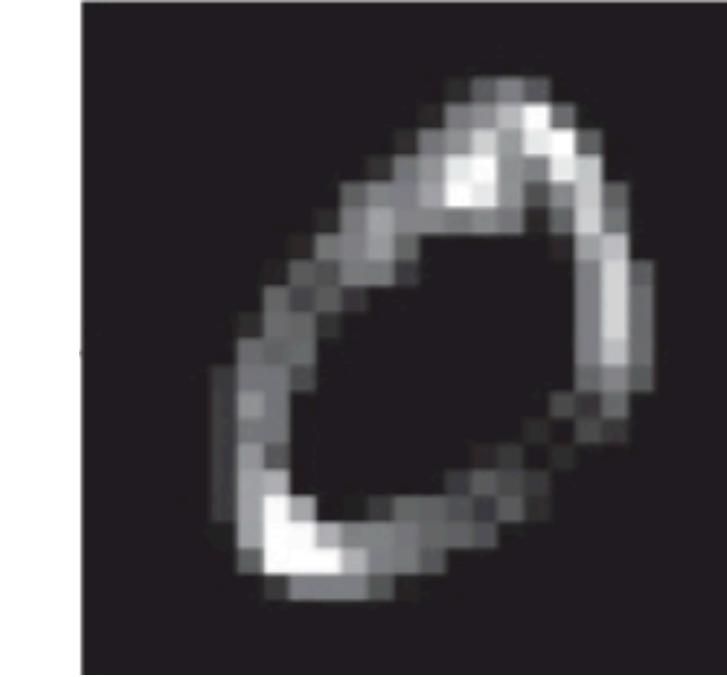
(pattern)

0	1	2
2	2	0
0	1	2

Input image



Activation map



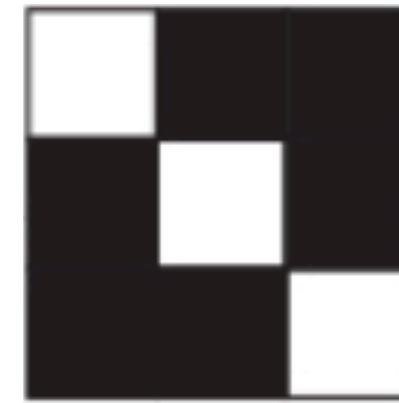
(presence of  
pattern)

$3_0$	$3_1$	$2_2$	1	0
$0_2$	$0_2$	$1_0$	3	1
$3_0$	$1_1$	$2_2$	2	3
2	0	0	2	2
2	0	0	0	1

12.0	12.0	17.0
10.0	17.0	19.0
9.0	6.0	14.0

# Convolution: Searching for patterns

Filter (3x3)



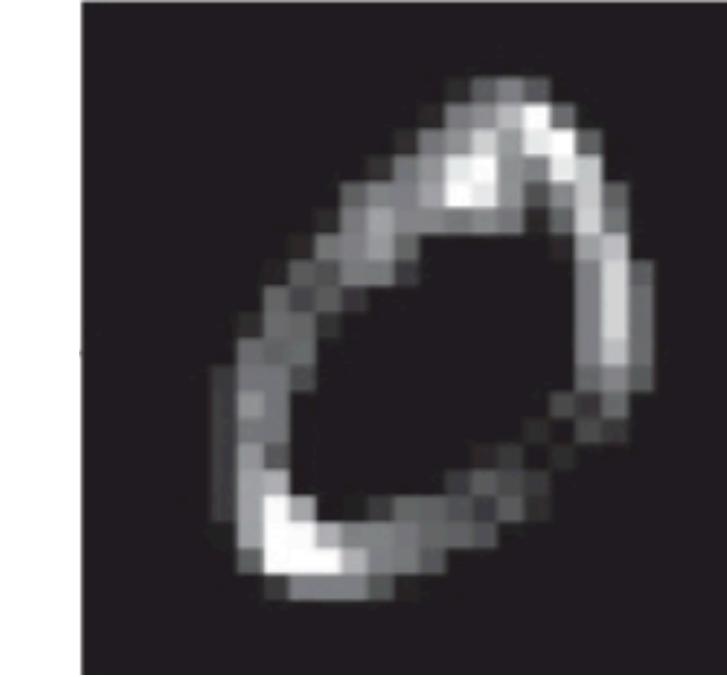
(pattern)

0	1	2
2	2	0
0	1	2

Input image



Activation map



(presence of  
pattern)

$3_0$	$3_1$	$2_2$	1	0
$0_2$	$0_2$	$1_0$	3	1
$3_0$	$1_1$	$2_2$	2	3
2	0	0	2	2
2	0	0	0	1

12.0	12.0	17.0
10.0	17.0	19.0
9.0	6.0	14.0

# Convolution: Searching for patterns

Filter (3x3)



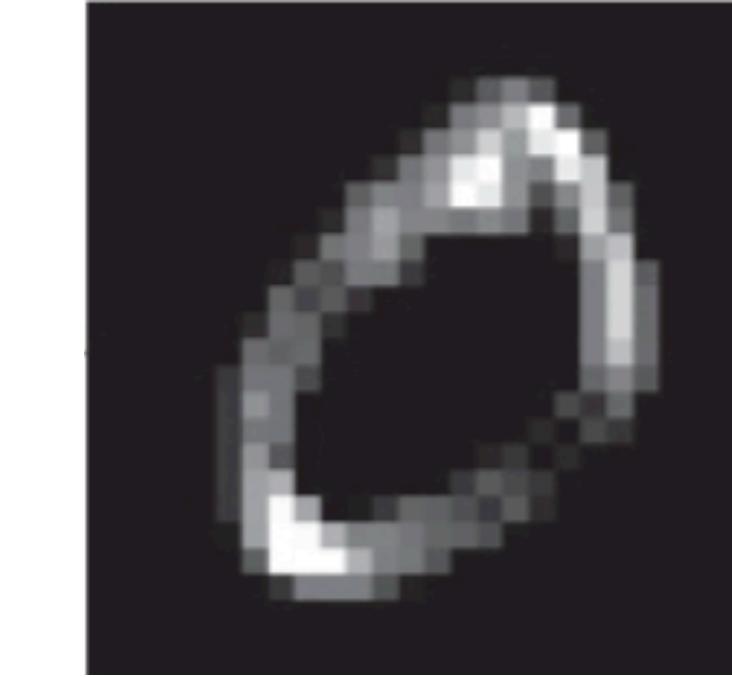
(pattern)

0	1	2
2	2	0
0	1	2

Input image



Activation map



(presence of  
pattern)

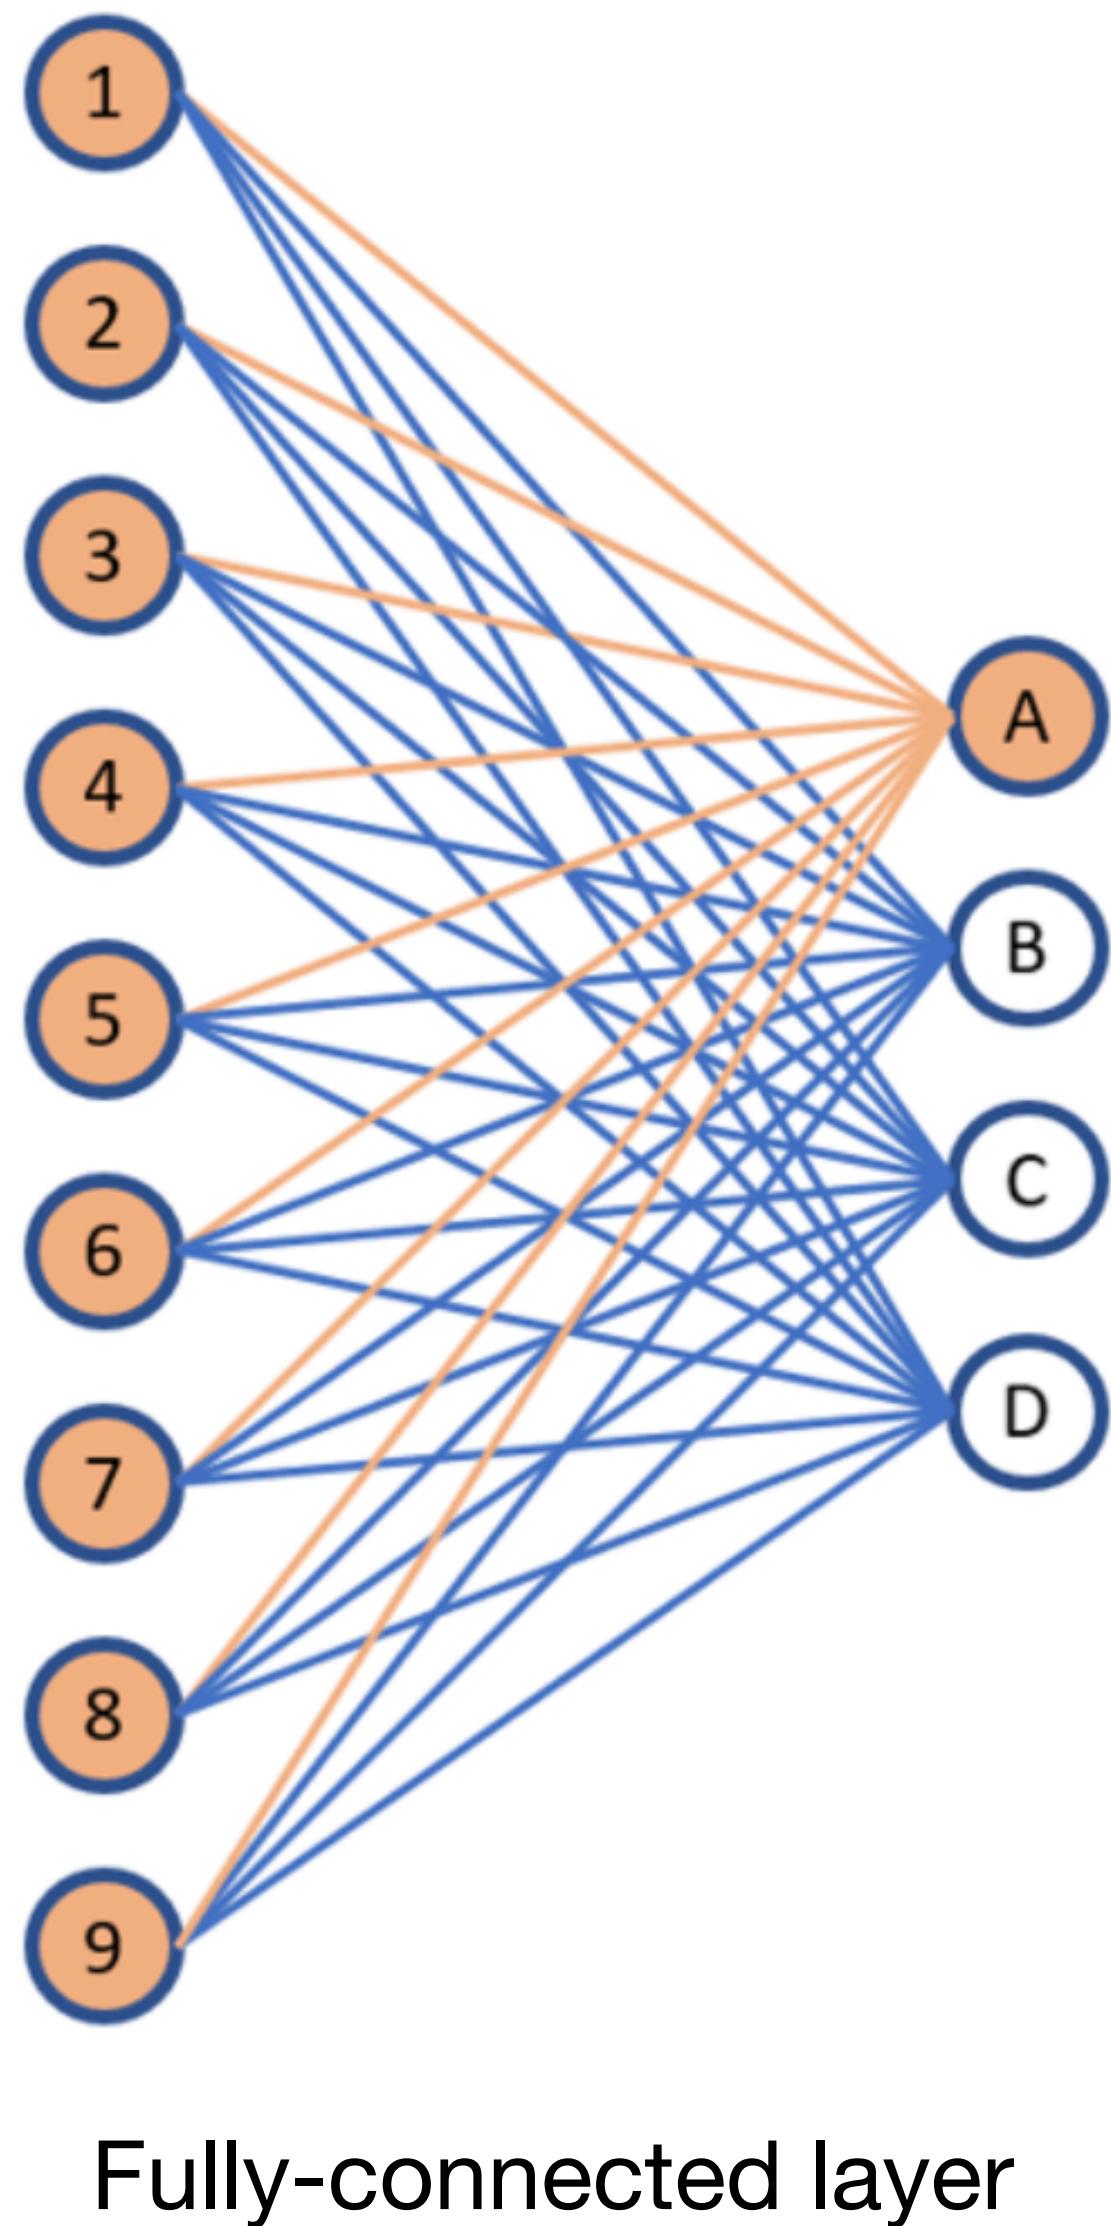
$3_0$	$3_1$	$2_2$	1	0
$0_2$	$0_2$	$1_0$	3	1
$3_0$	$1_1$	$2_2$	2	3
2	0	0	2	2
2	0	0	0	1

12.0	12.0	17.0
10.0	17.0	19.0
9.0	6.0	14.0

We want to use many filters, each sensitive to a different kind of pattern.

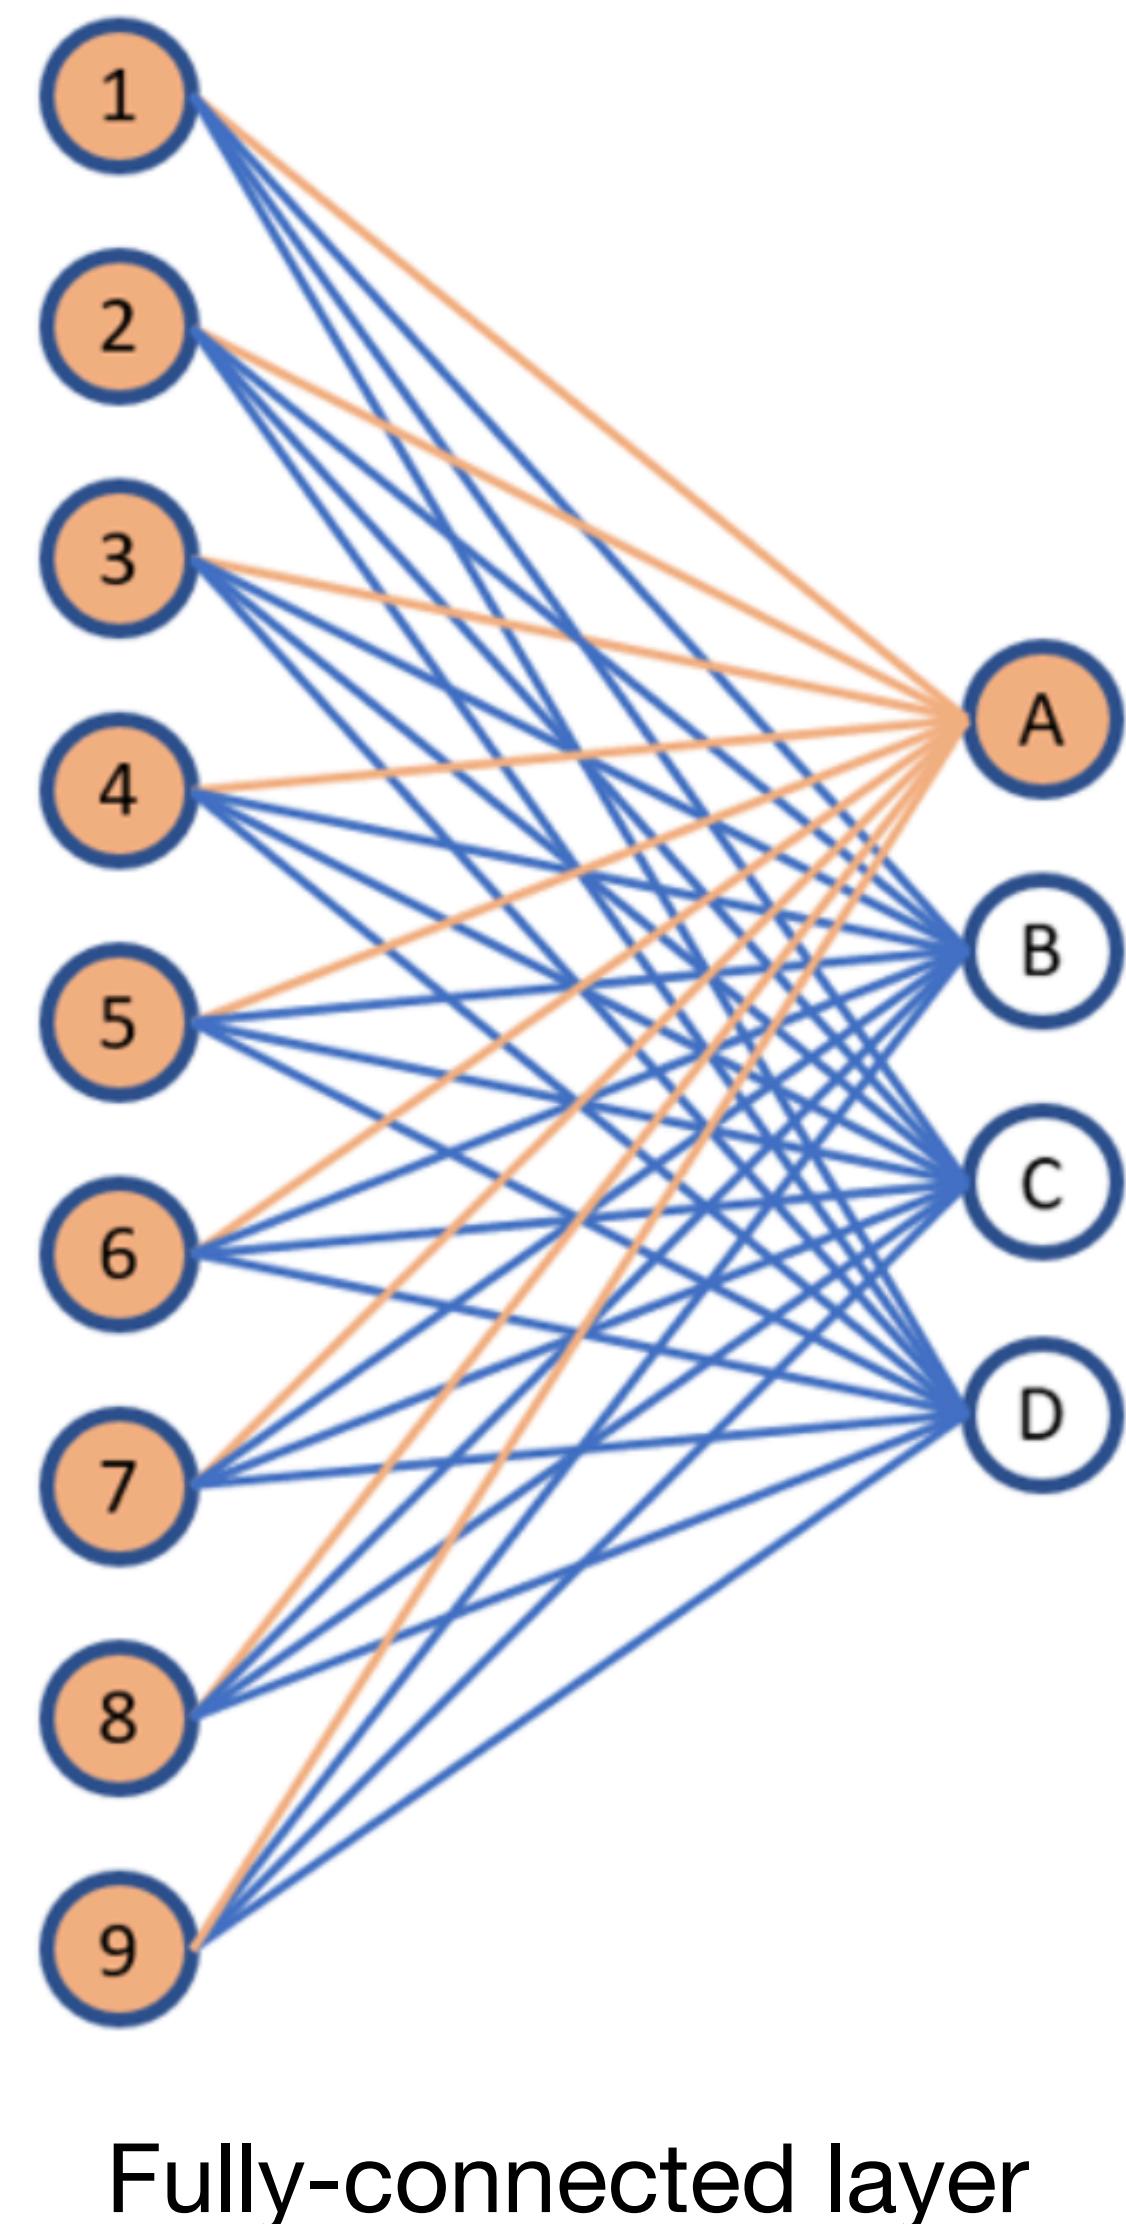
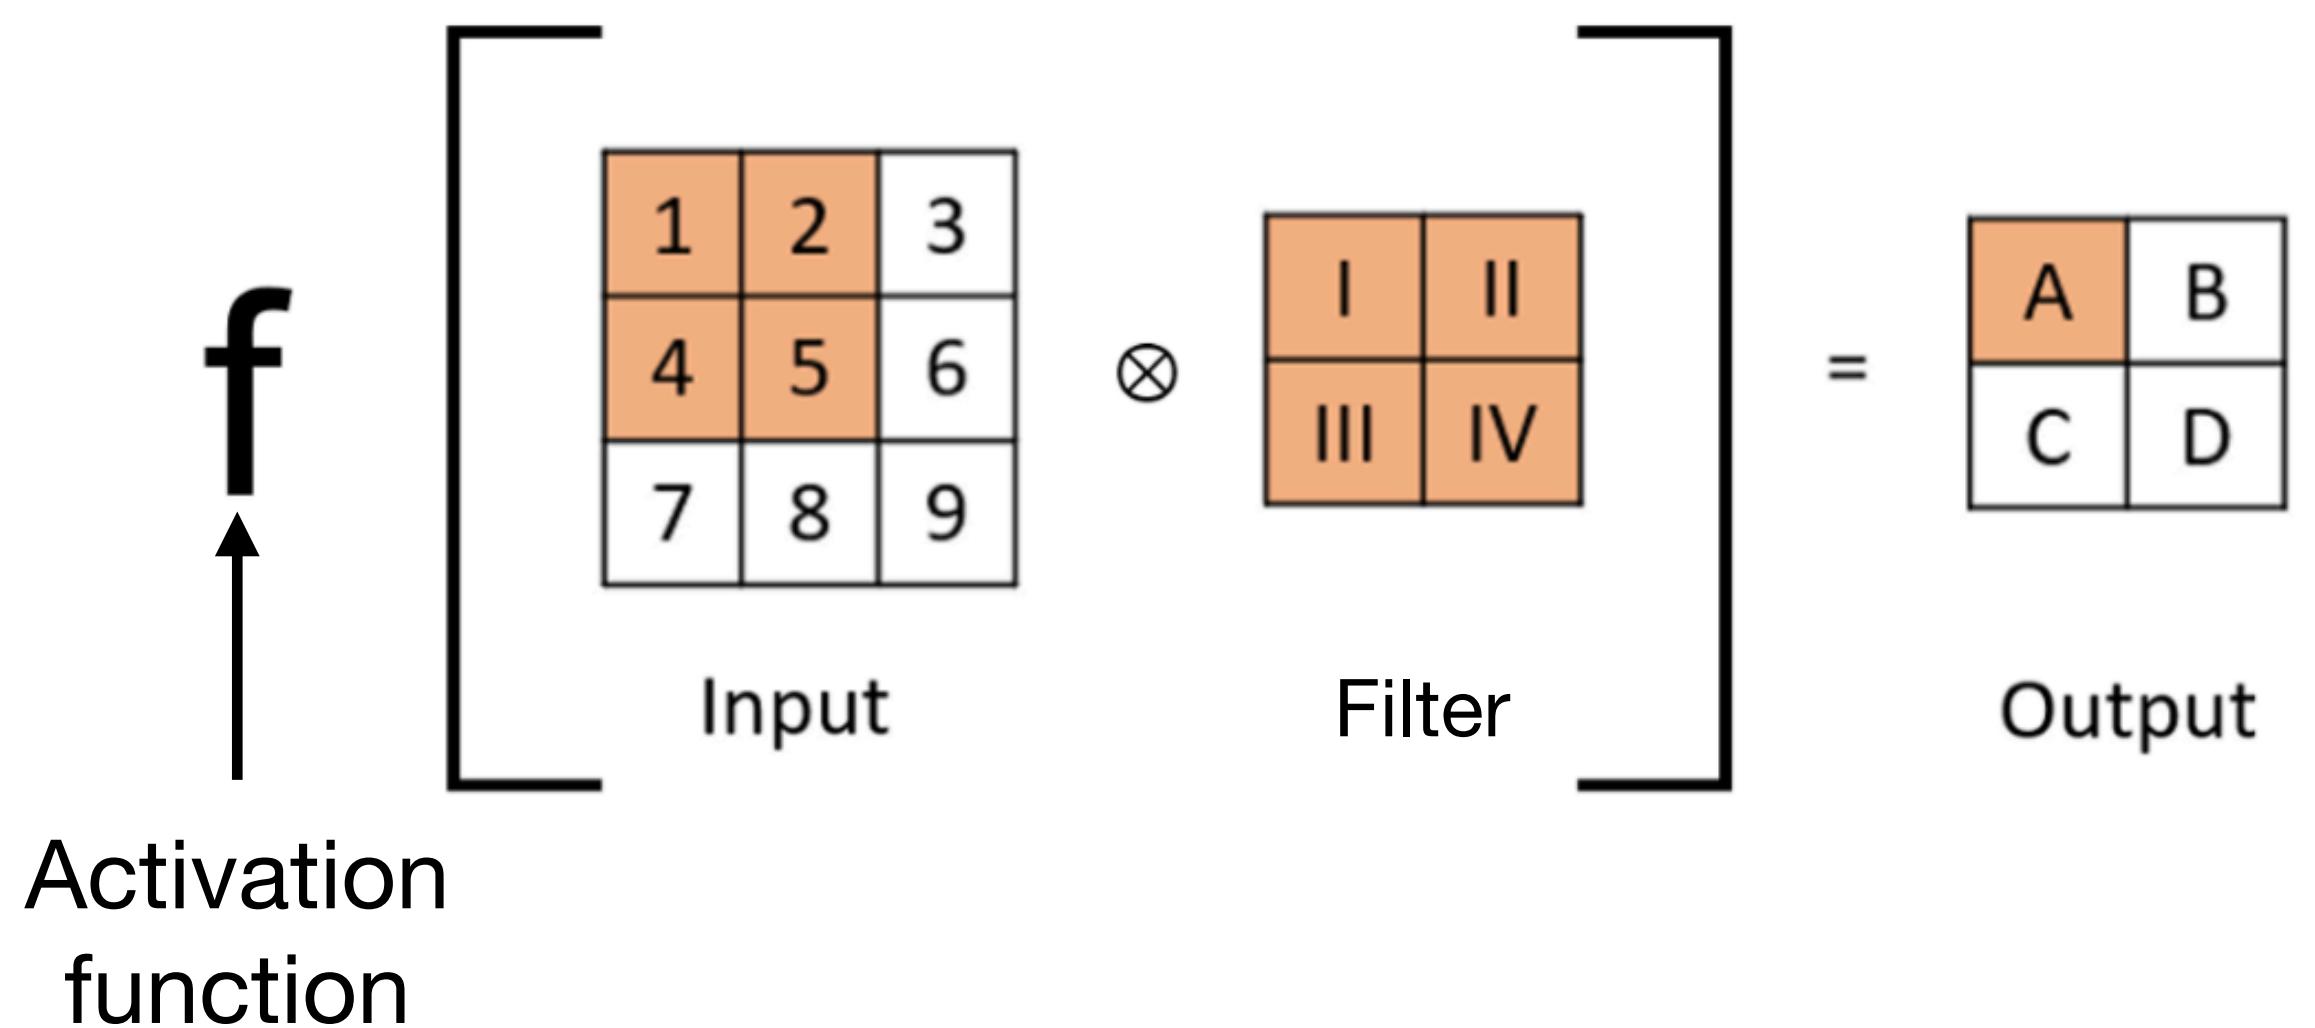
# Convolutional layer versus fully-connected layer

A convolutional layer can be visualized similarly to a fully-connected layer.



# Convolutional layer versus fully-connected layer

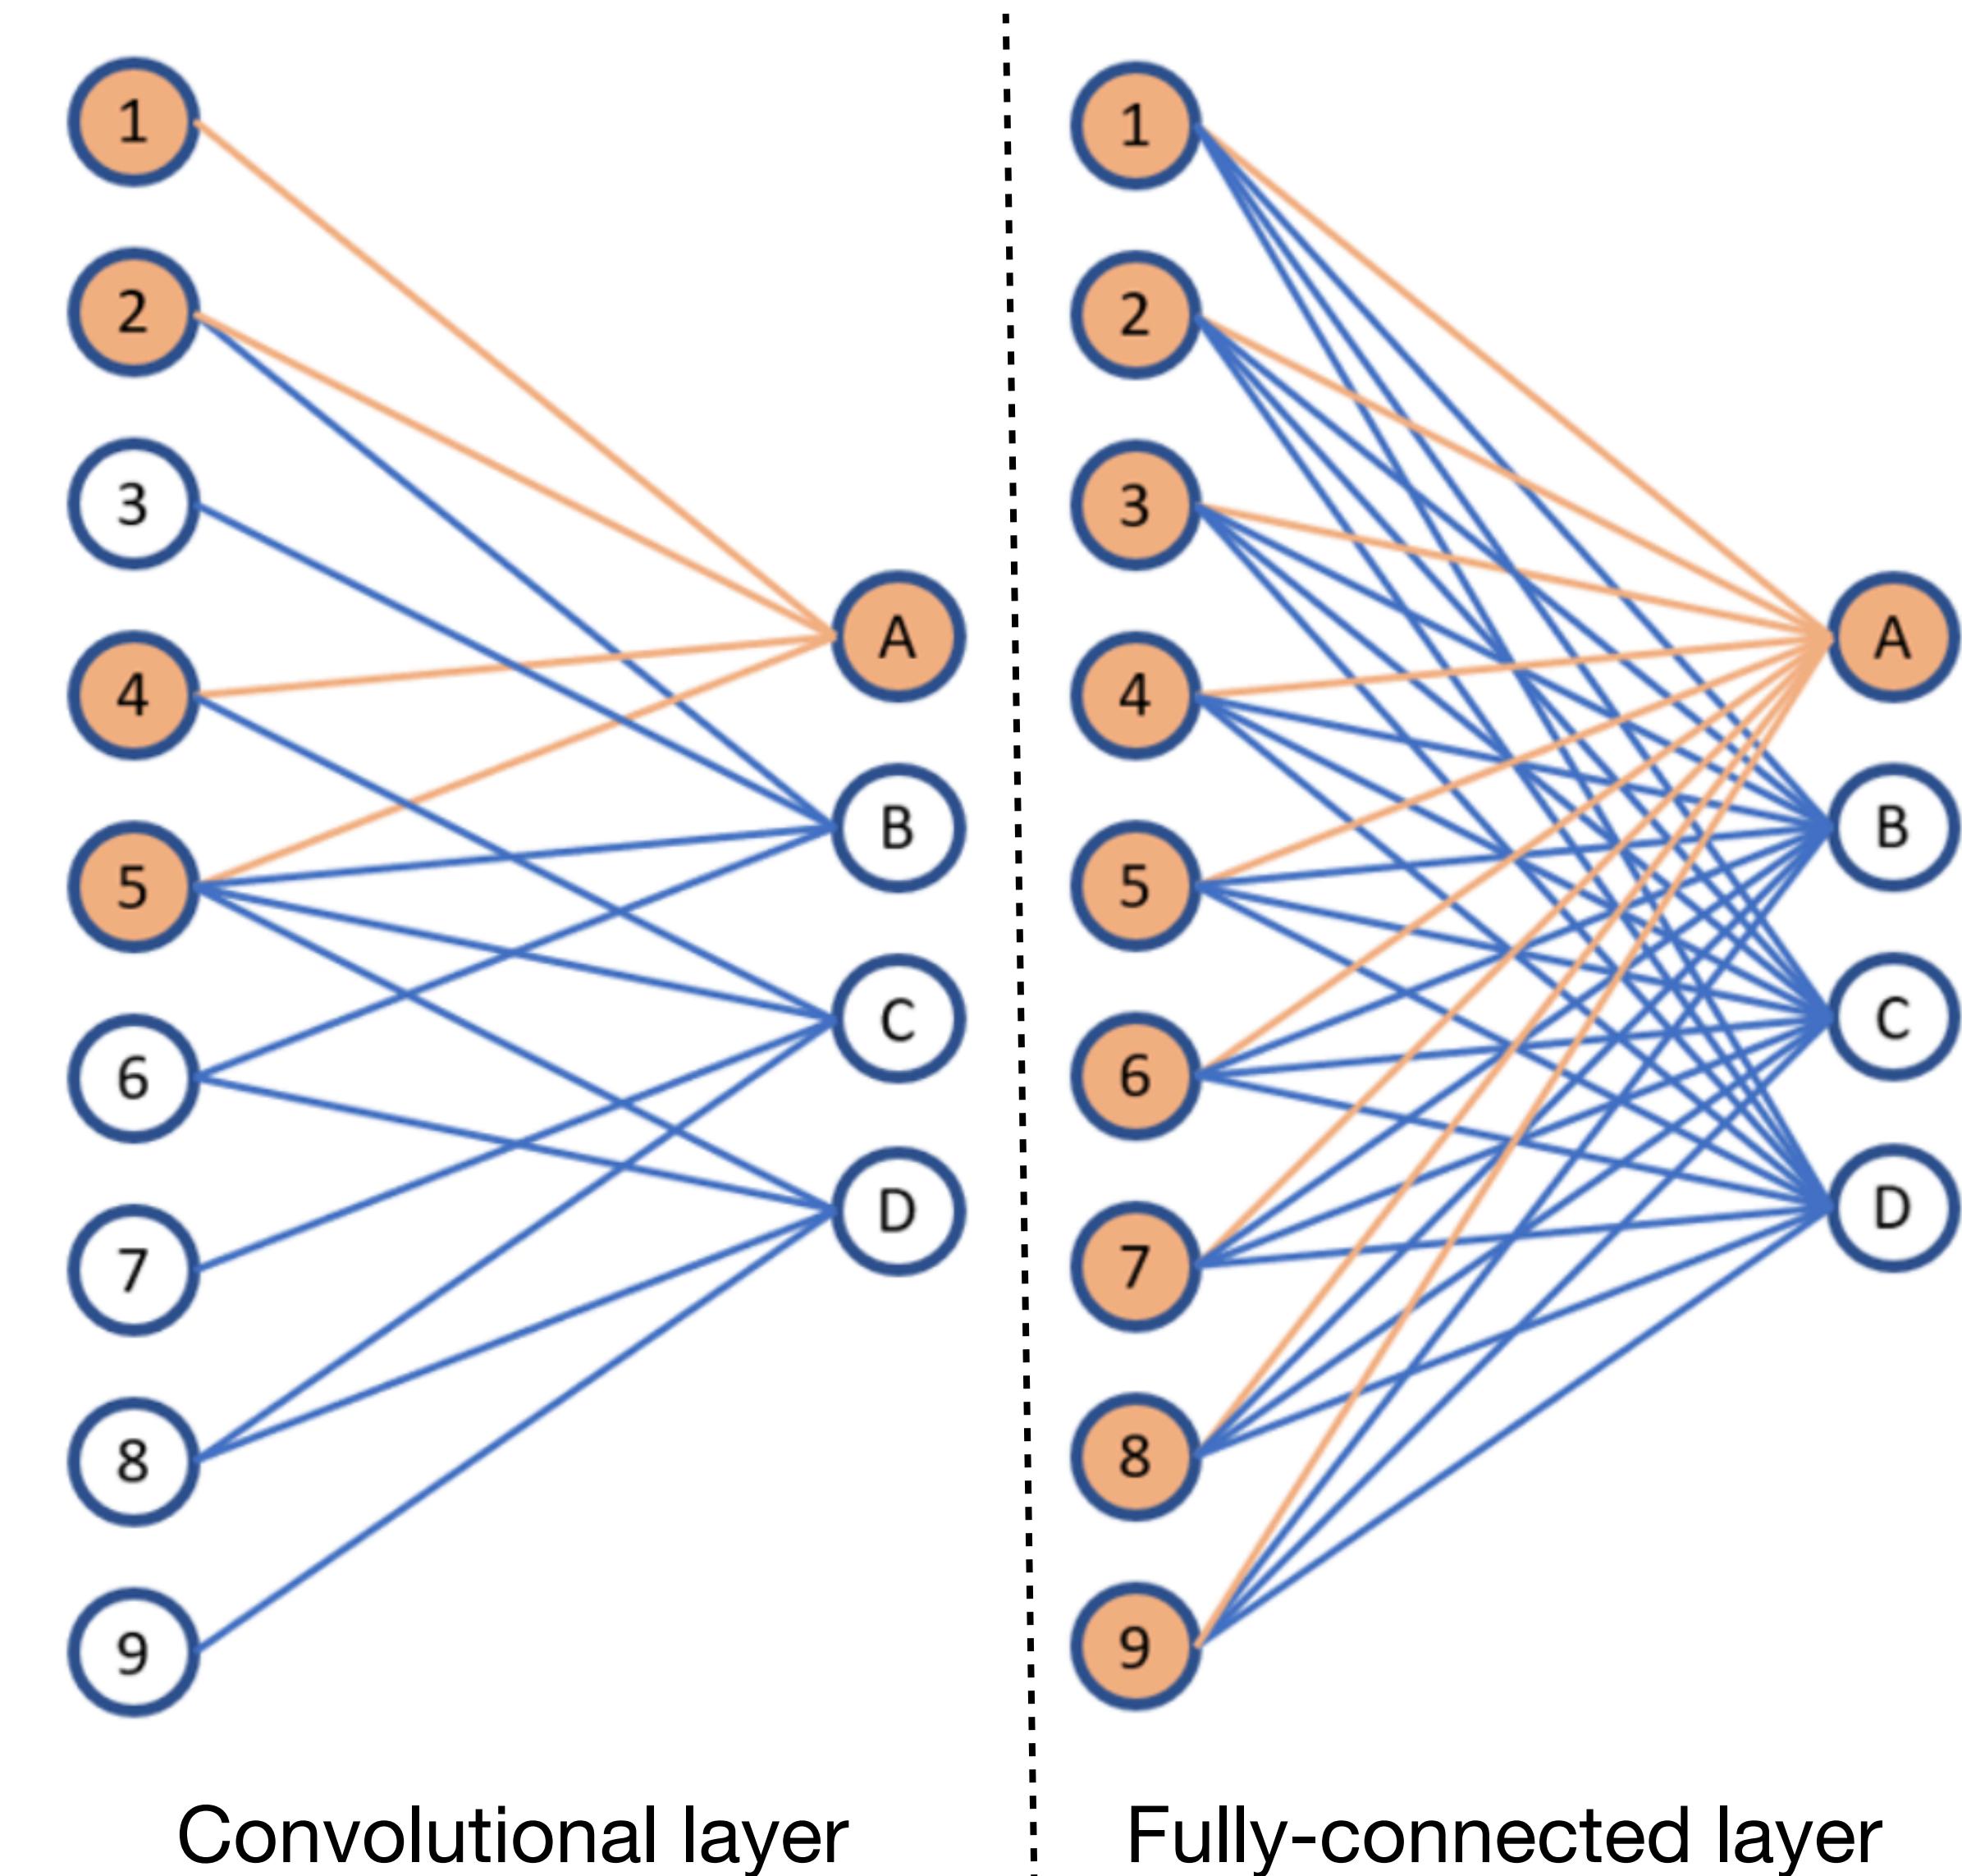
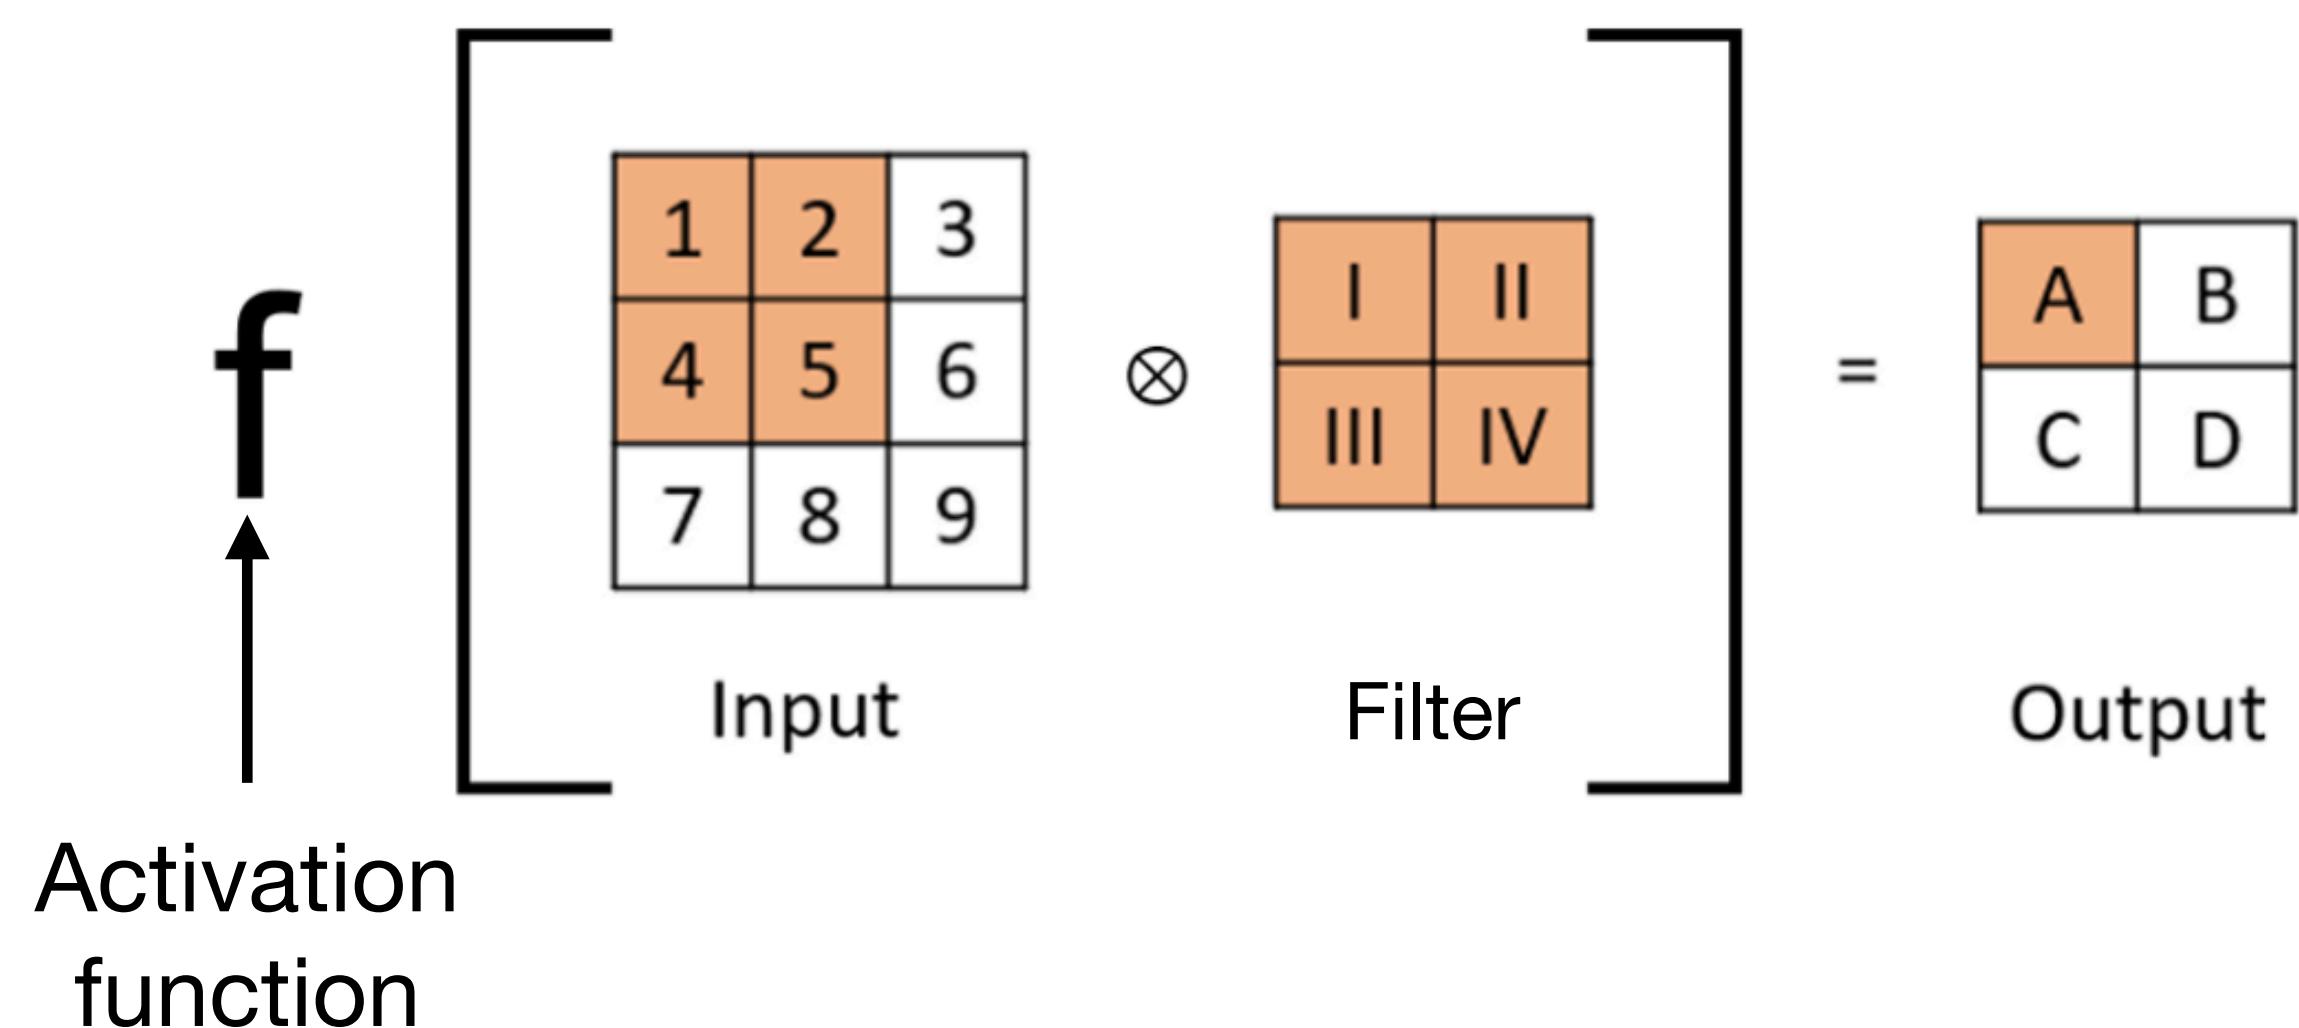
A convolutional layer can be visualized similarly to a fully-connected layer.



Fully-connected layer

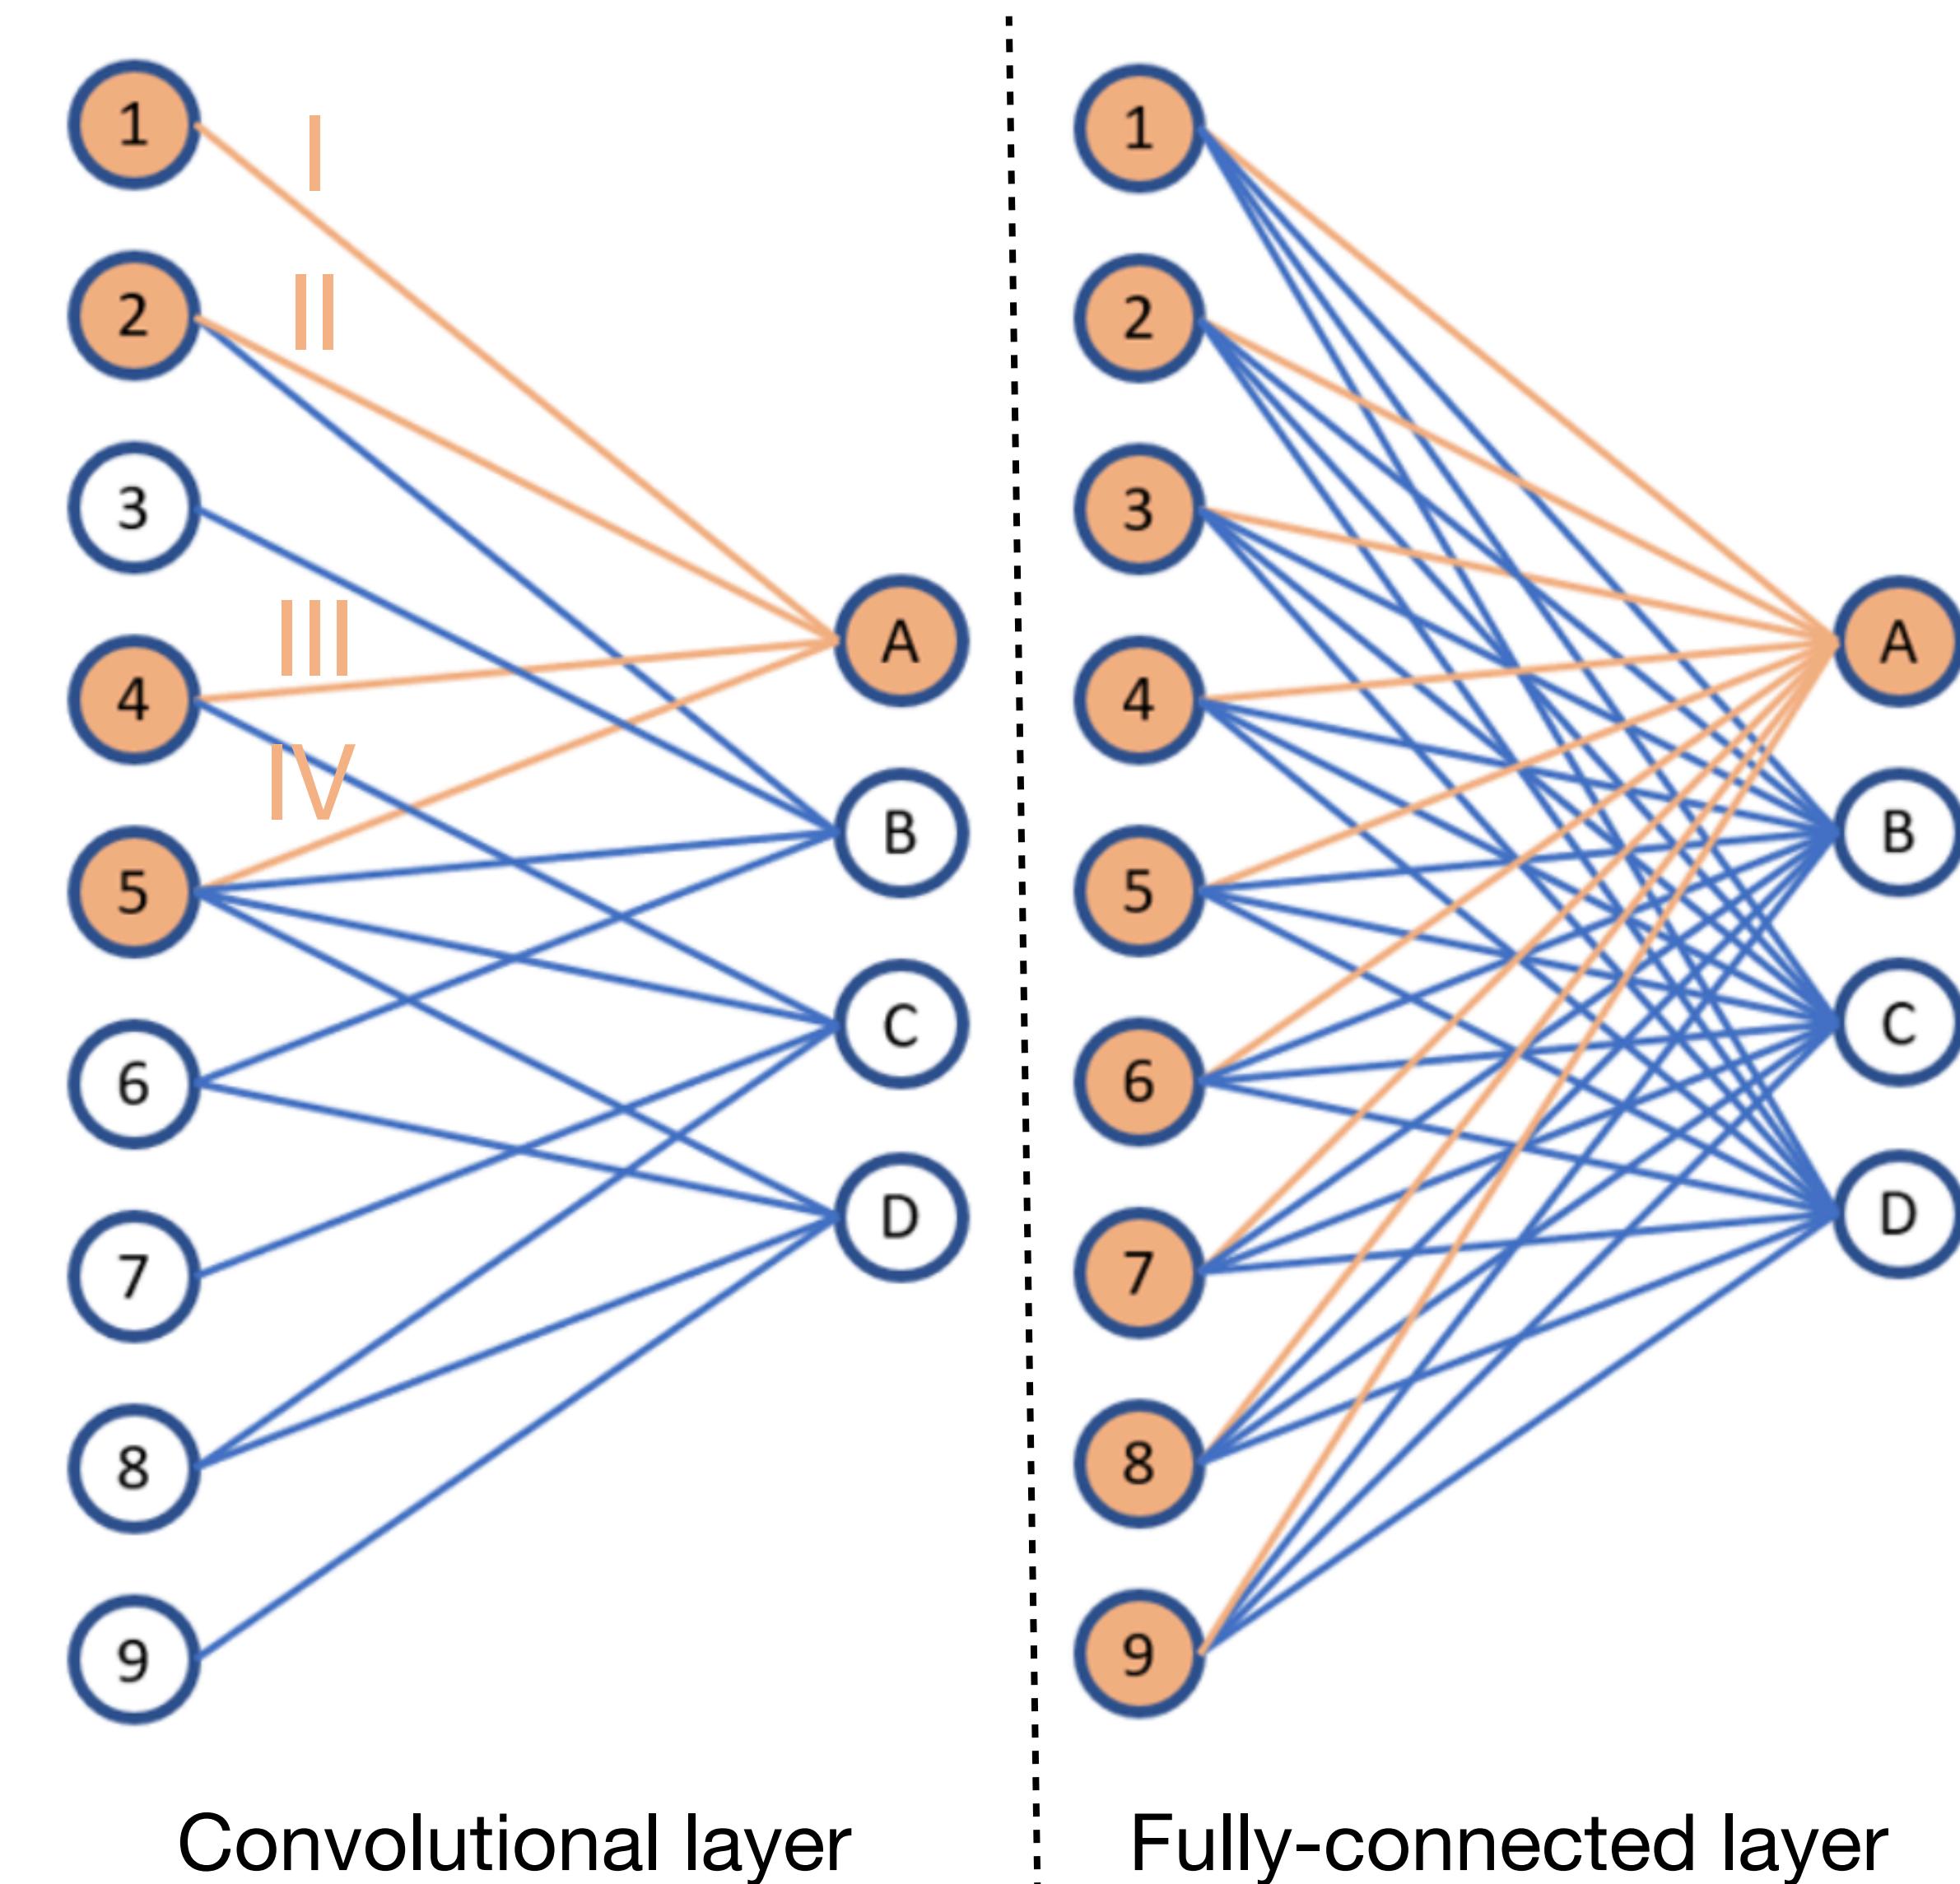
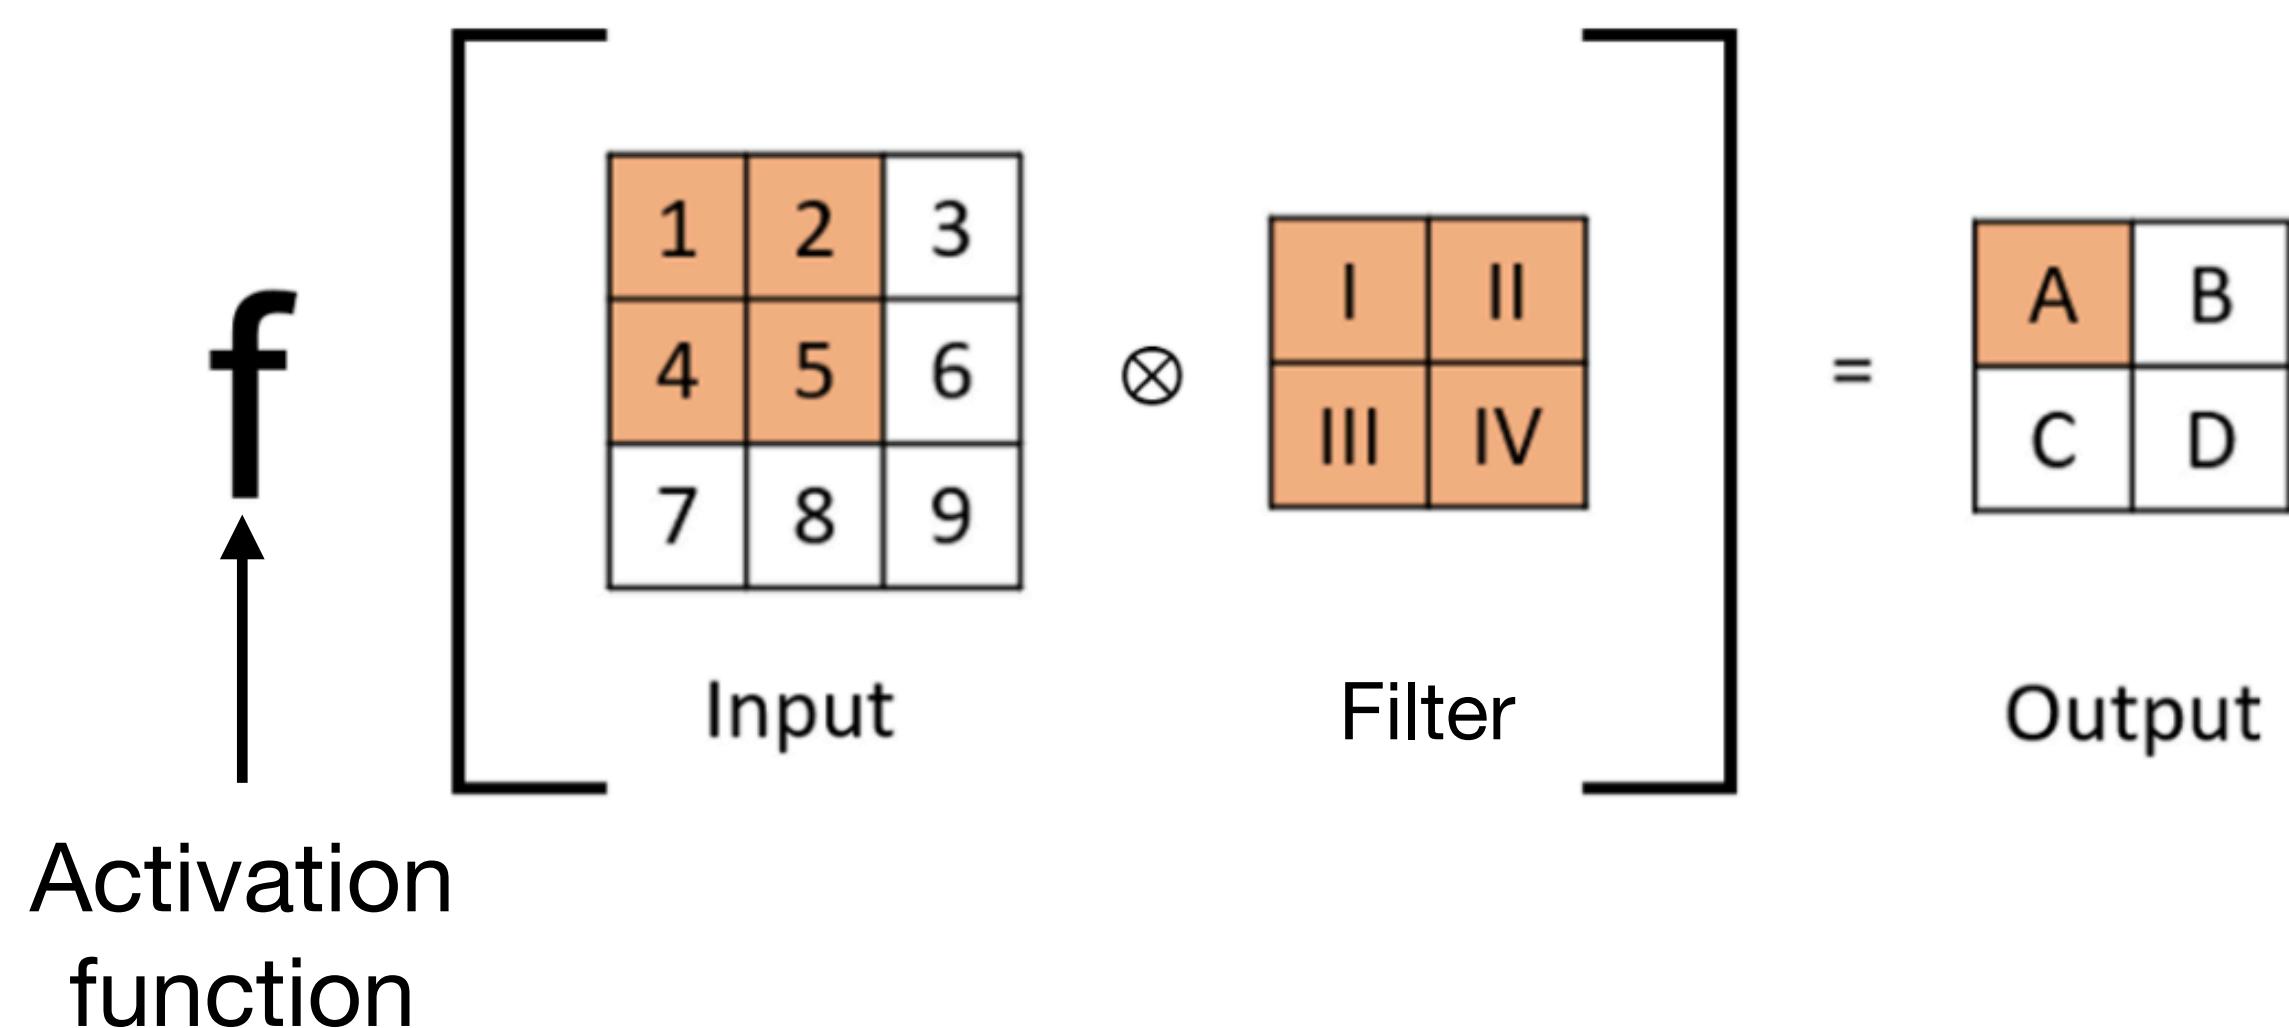
# Convolutional layer versus fully-connected layer

A convolutional layer can be visualized similarly to a fully-connected layer.



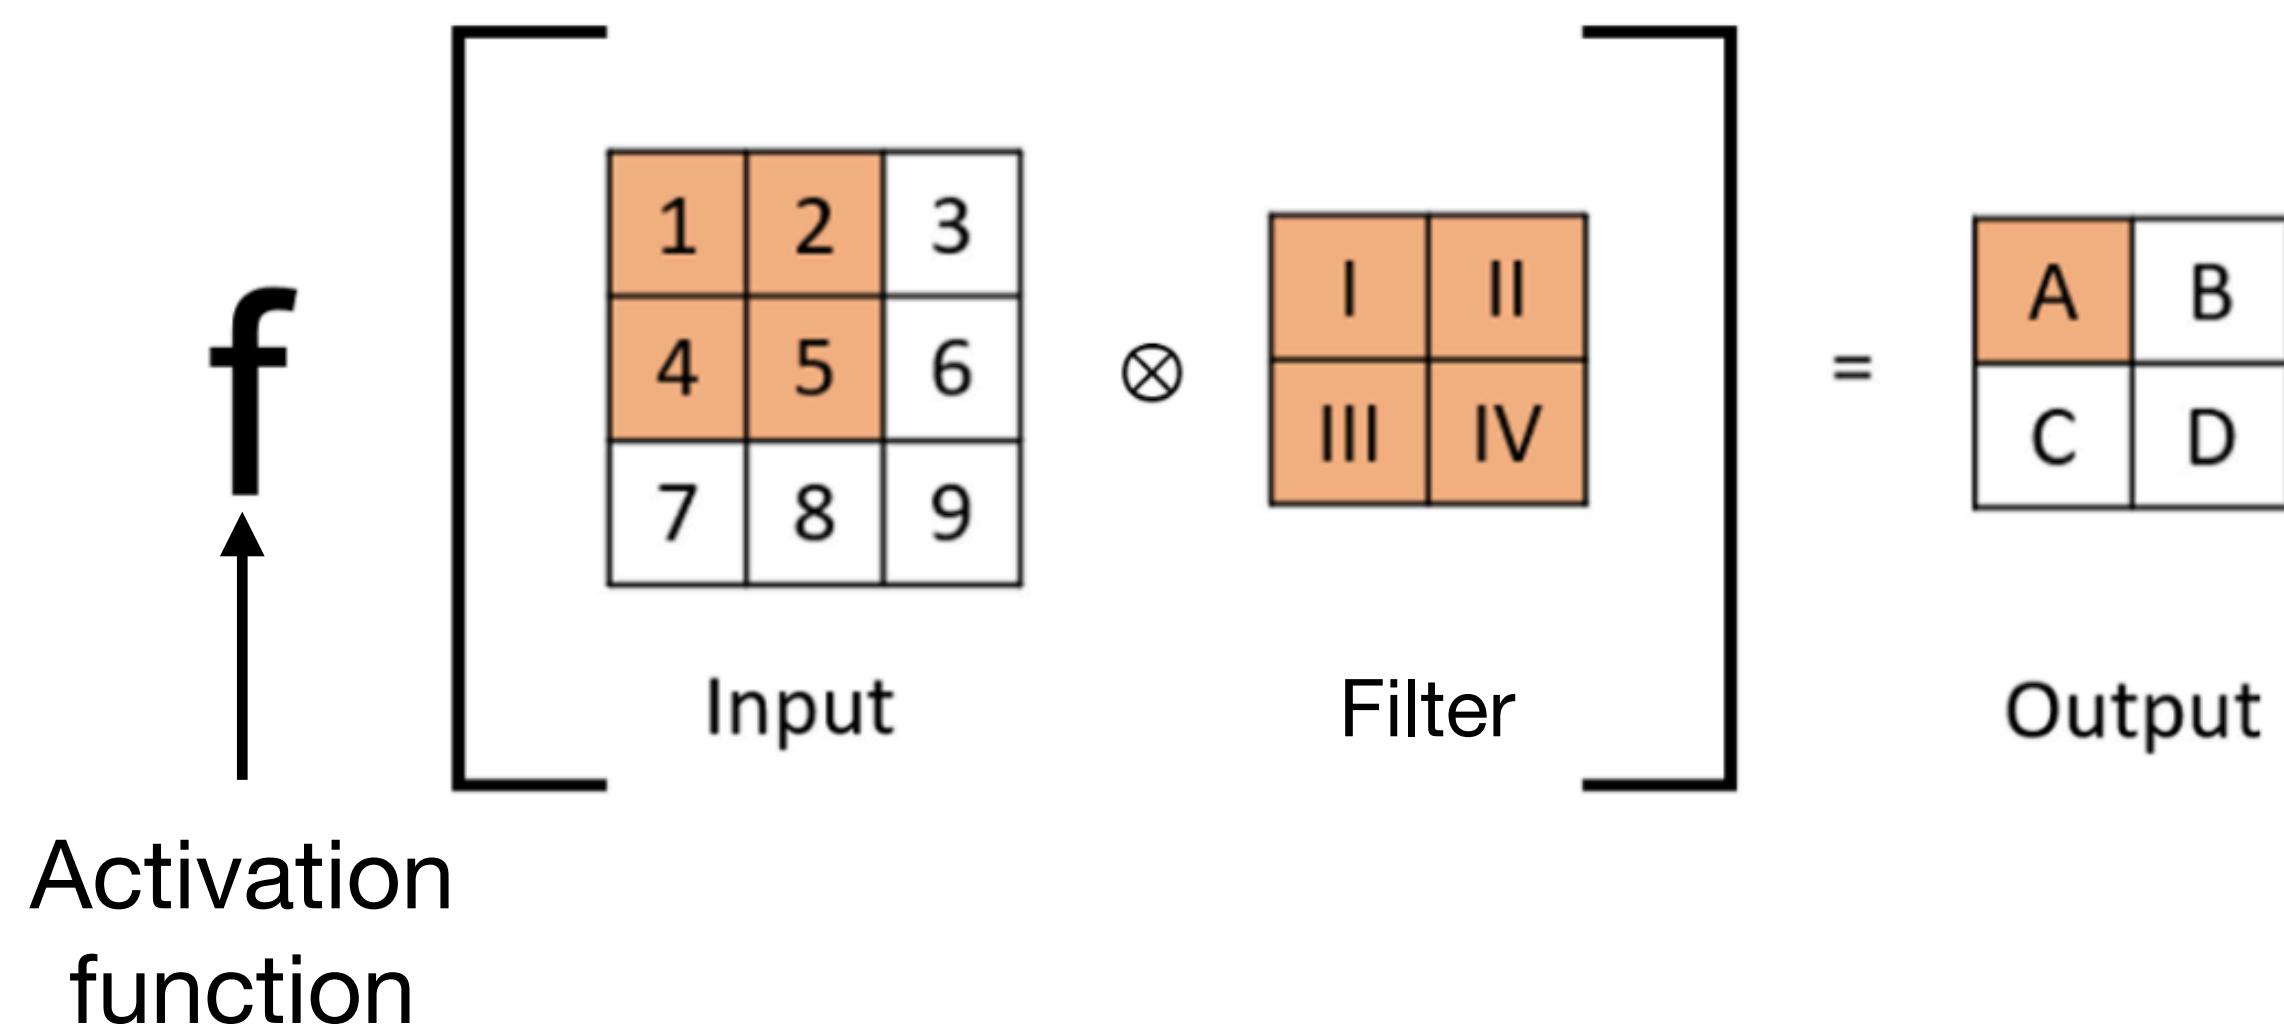
# Convolutional layer versus fully-connected layer

A convolutional layer can be visualized similarly to a fully-connected layer.

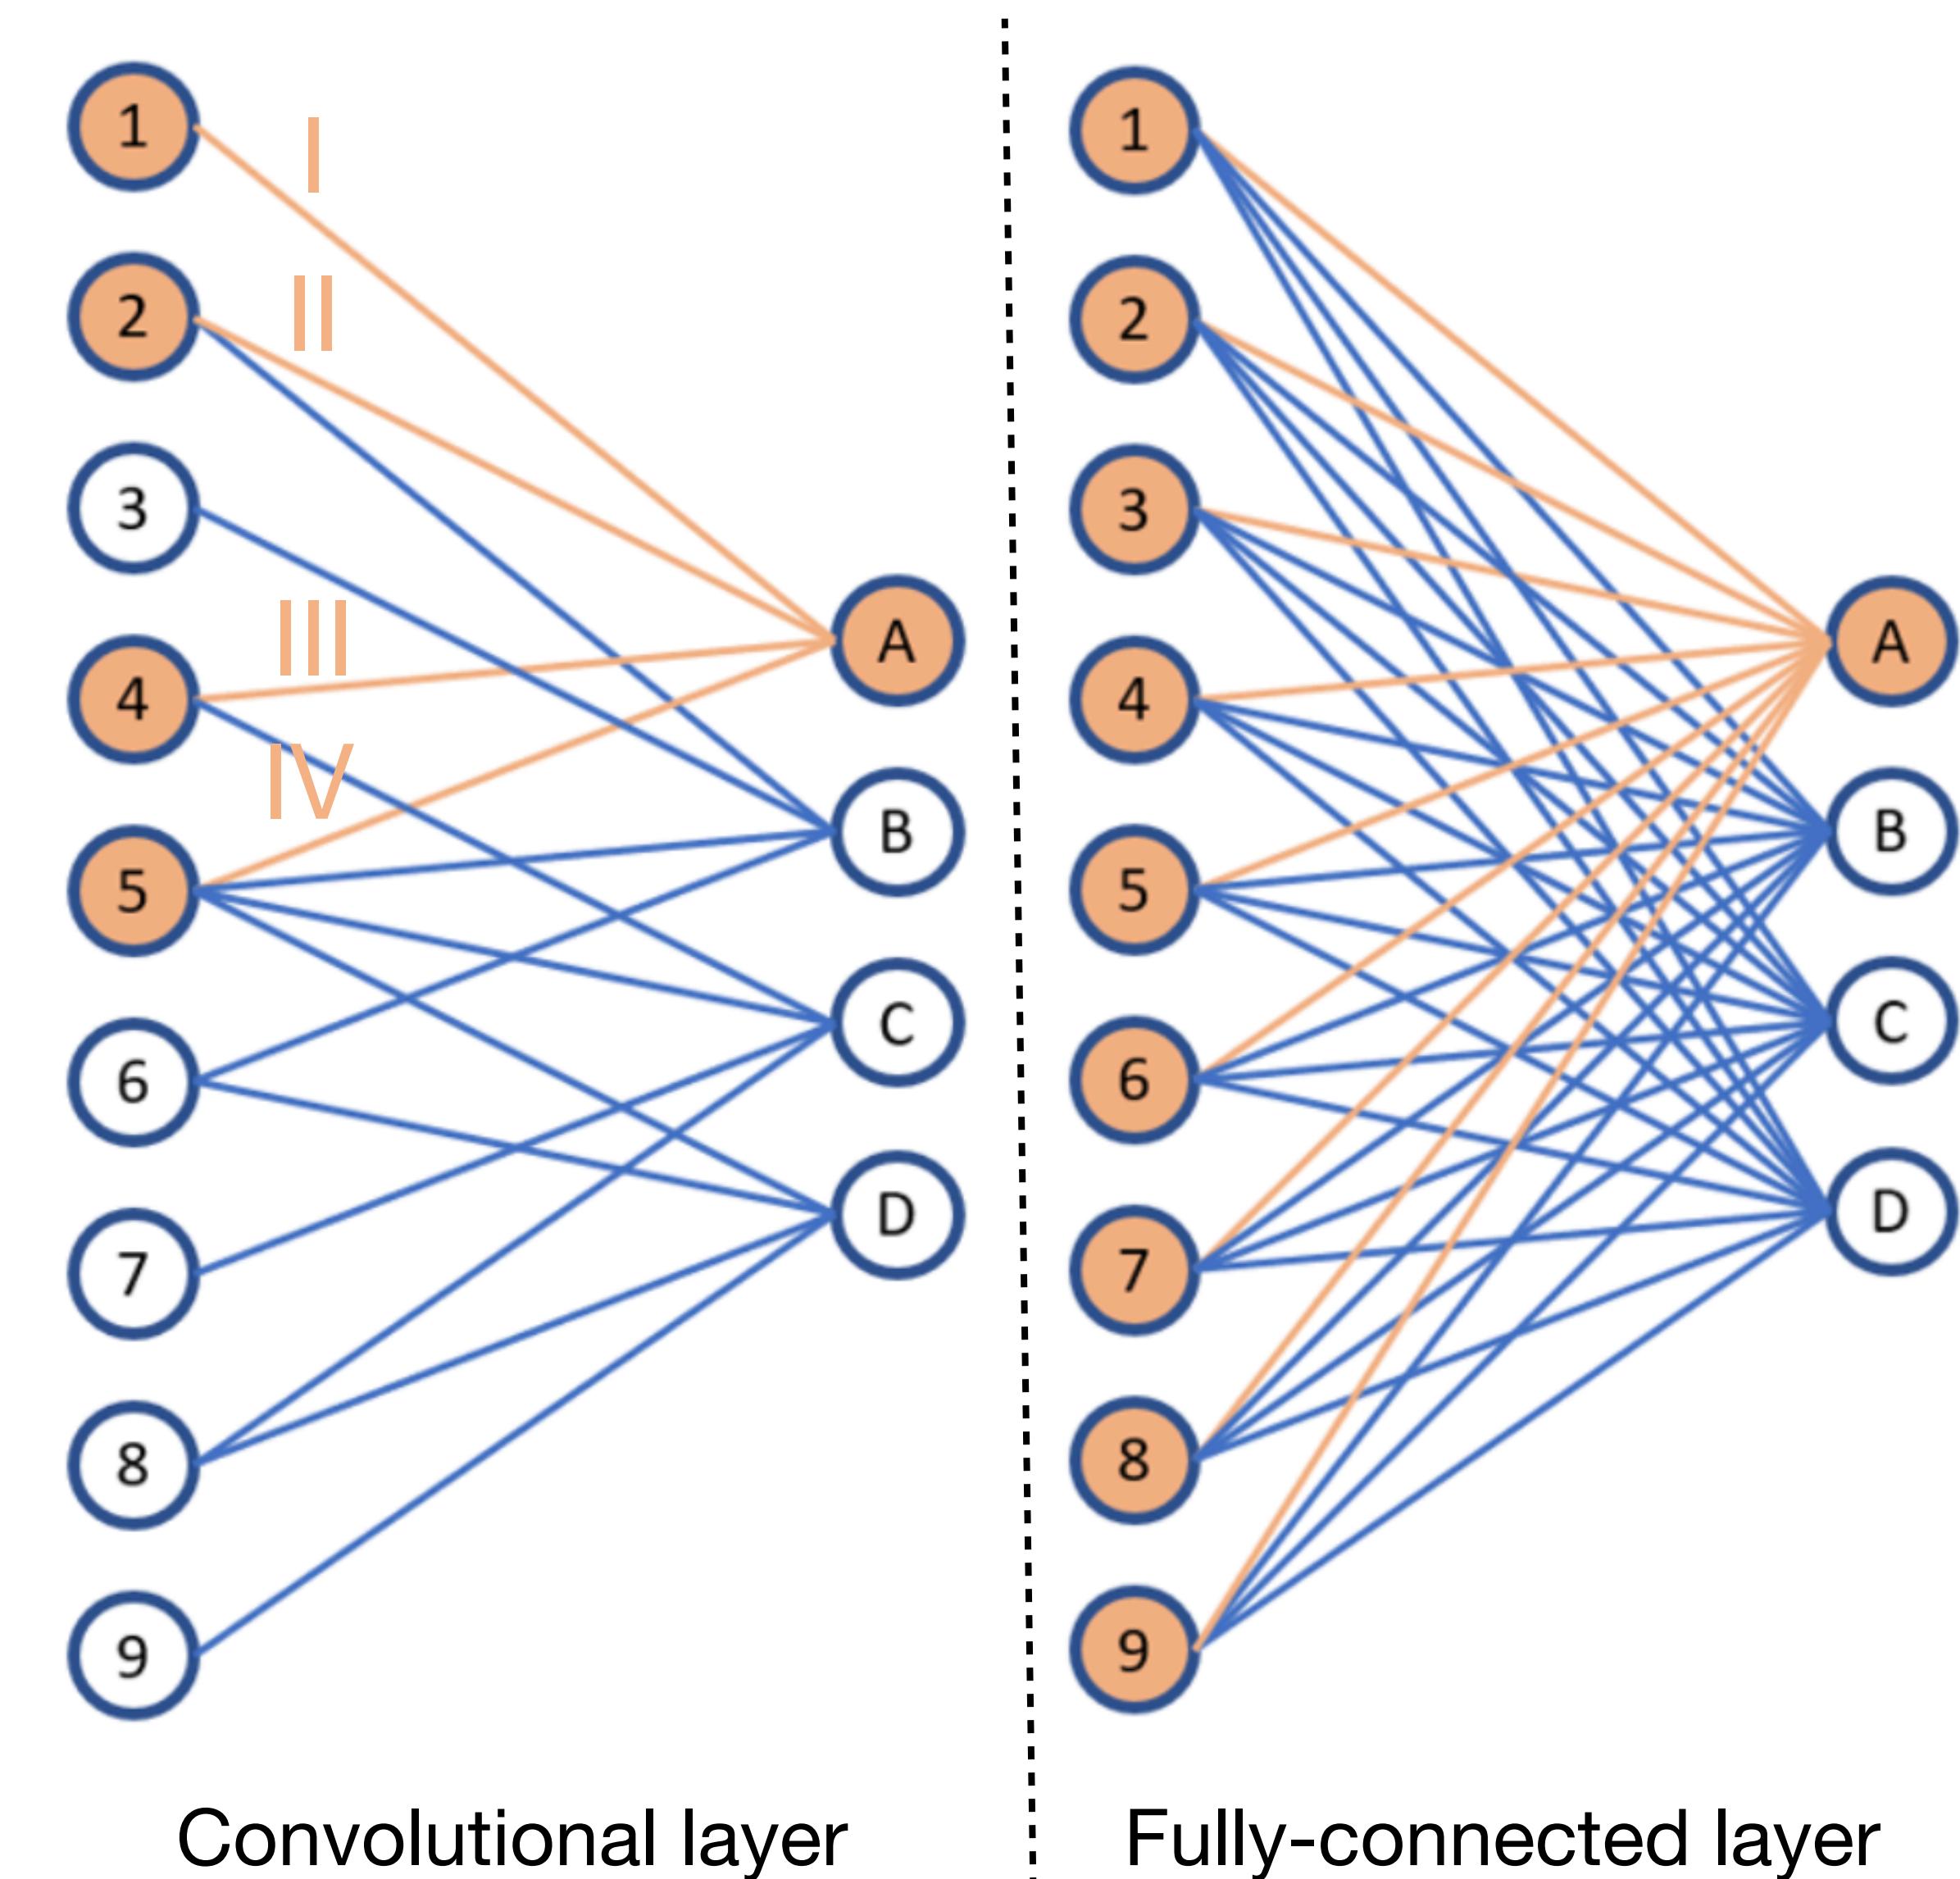


# Convolutional layer versus fully-connected layer

A convolutional layer can be visualized similarly to a fully-connected layer.

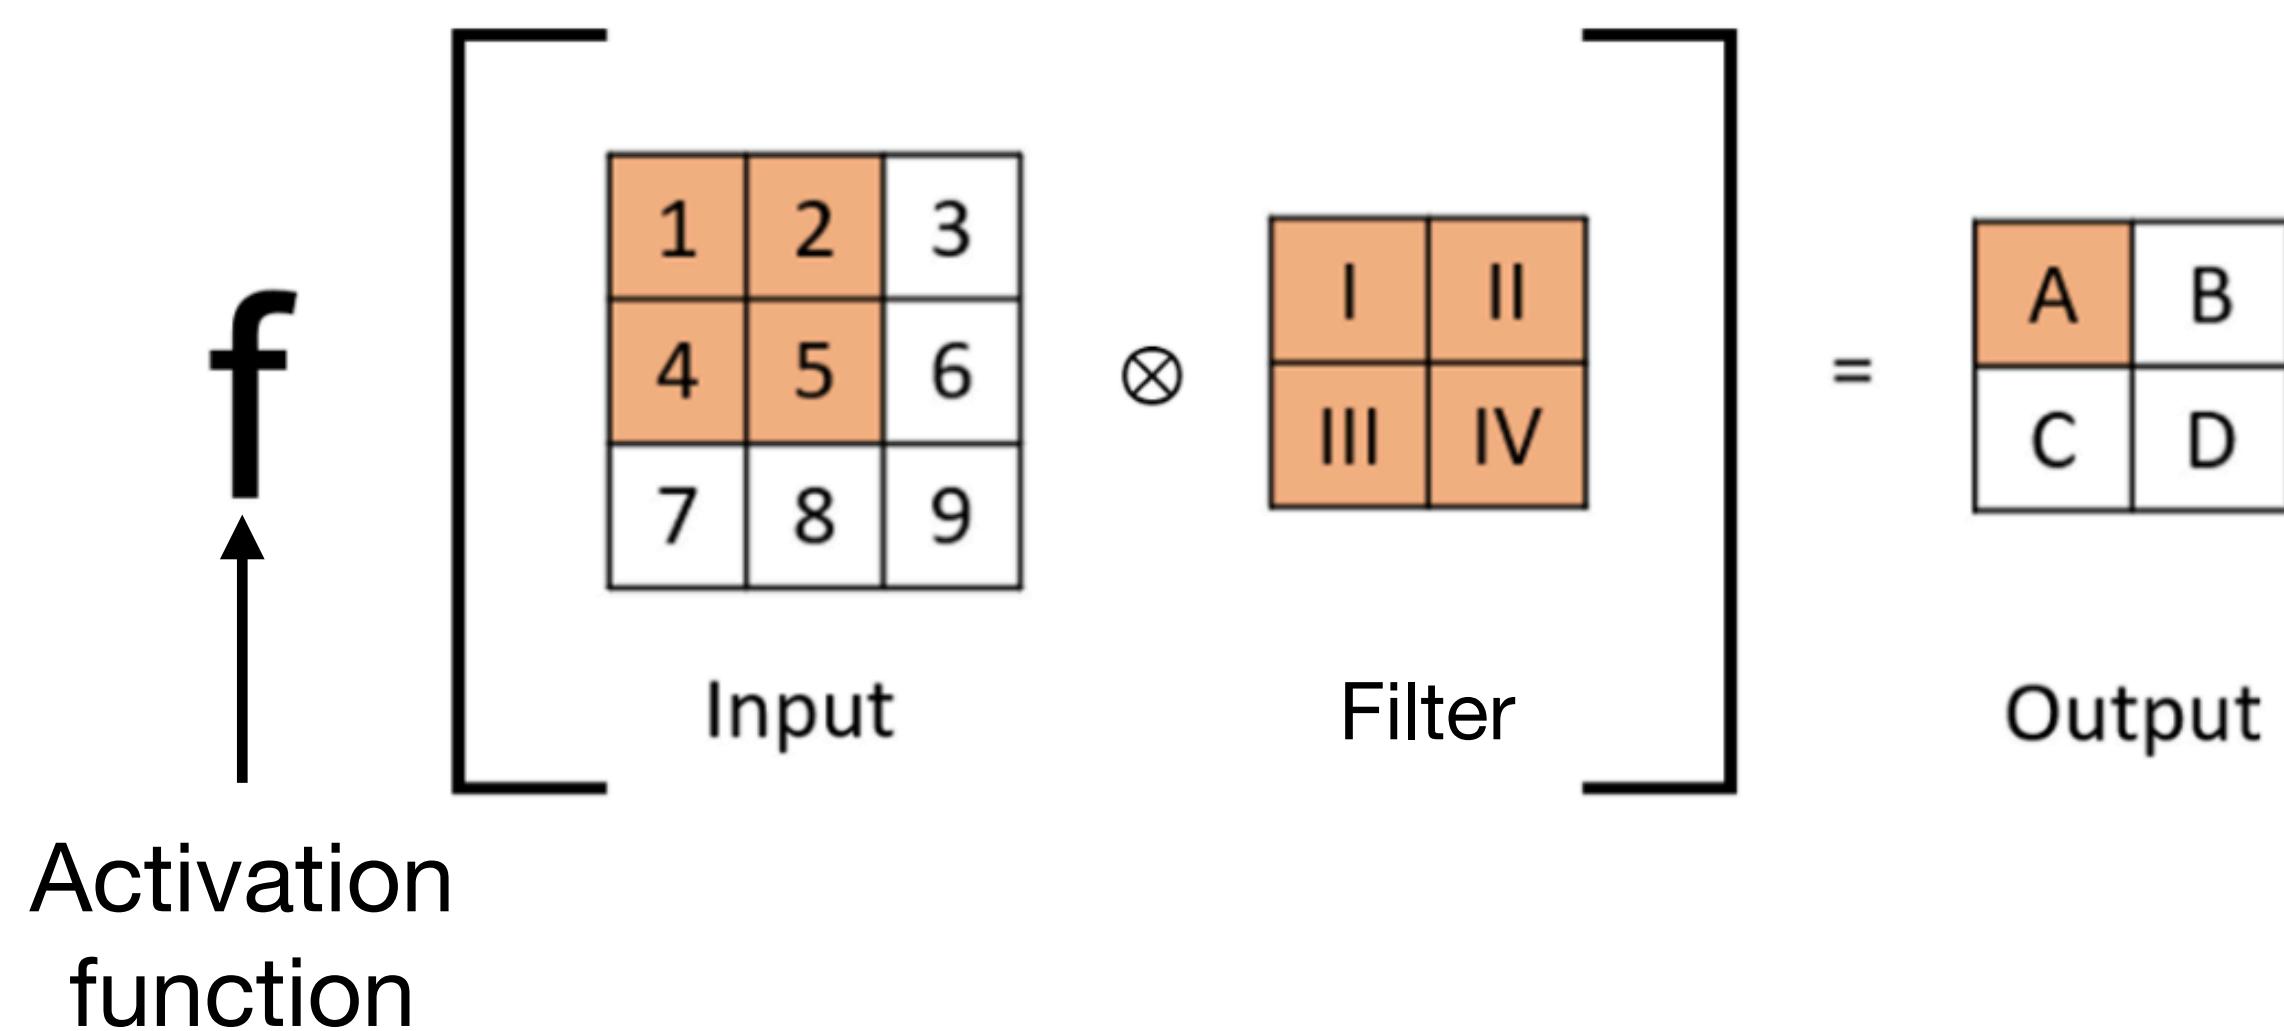


In a convolutional layer:



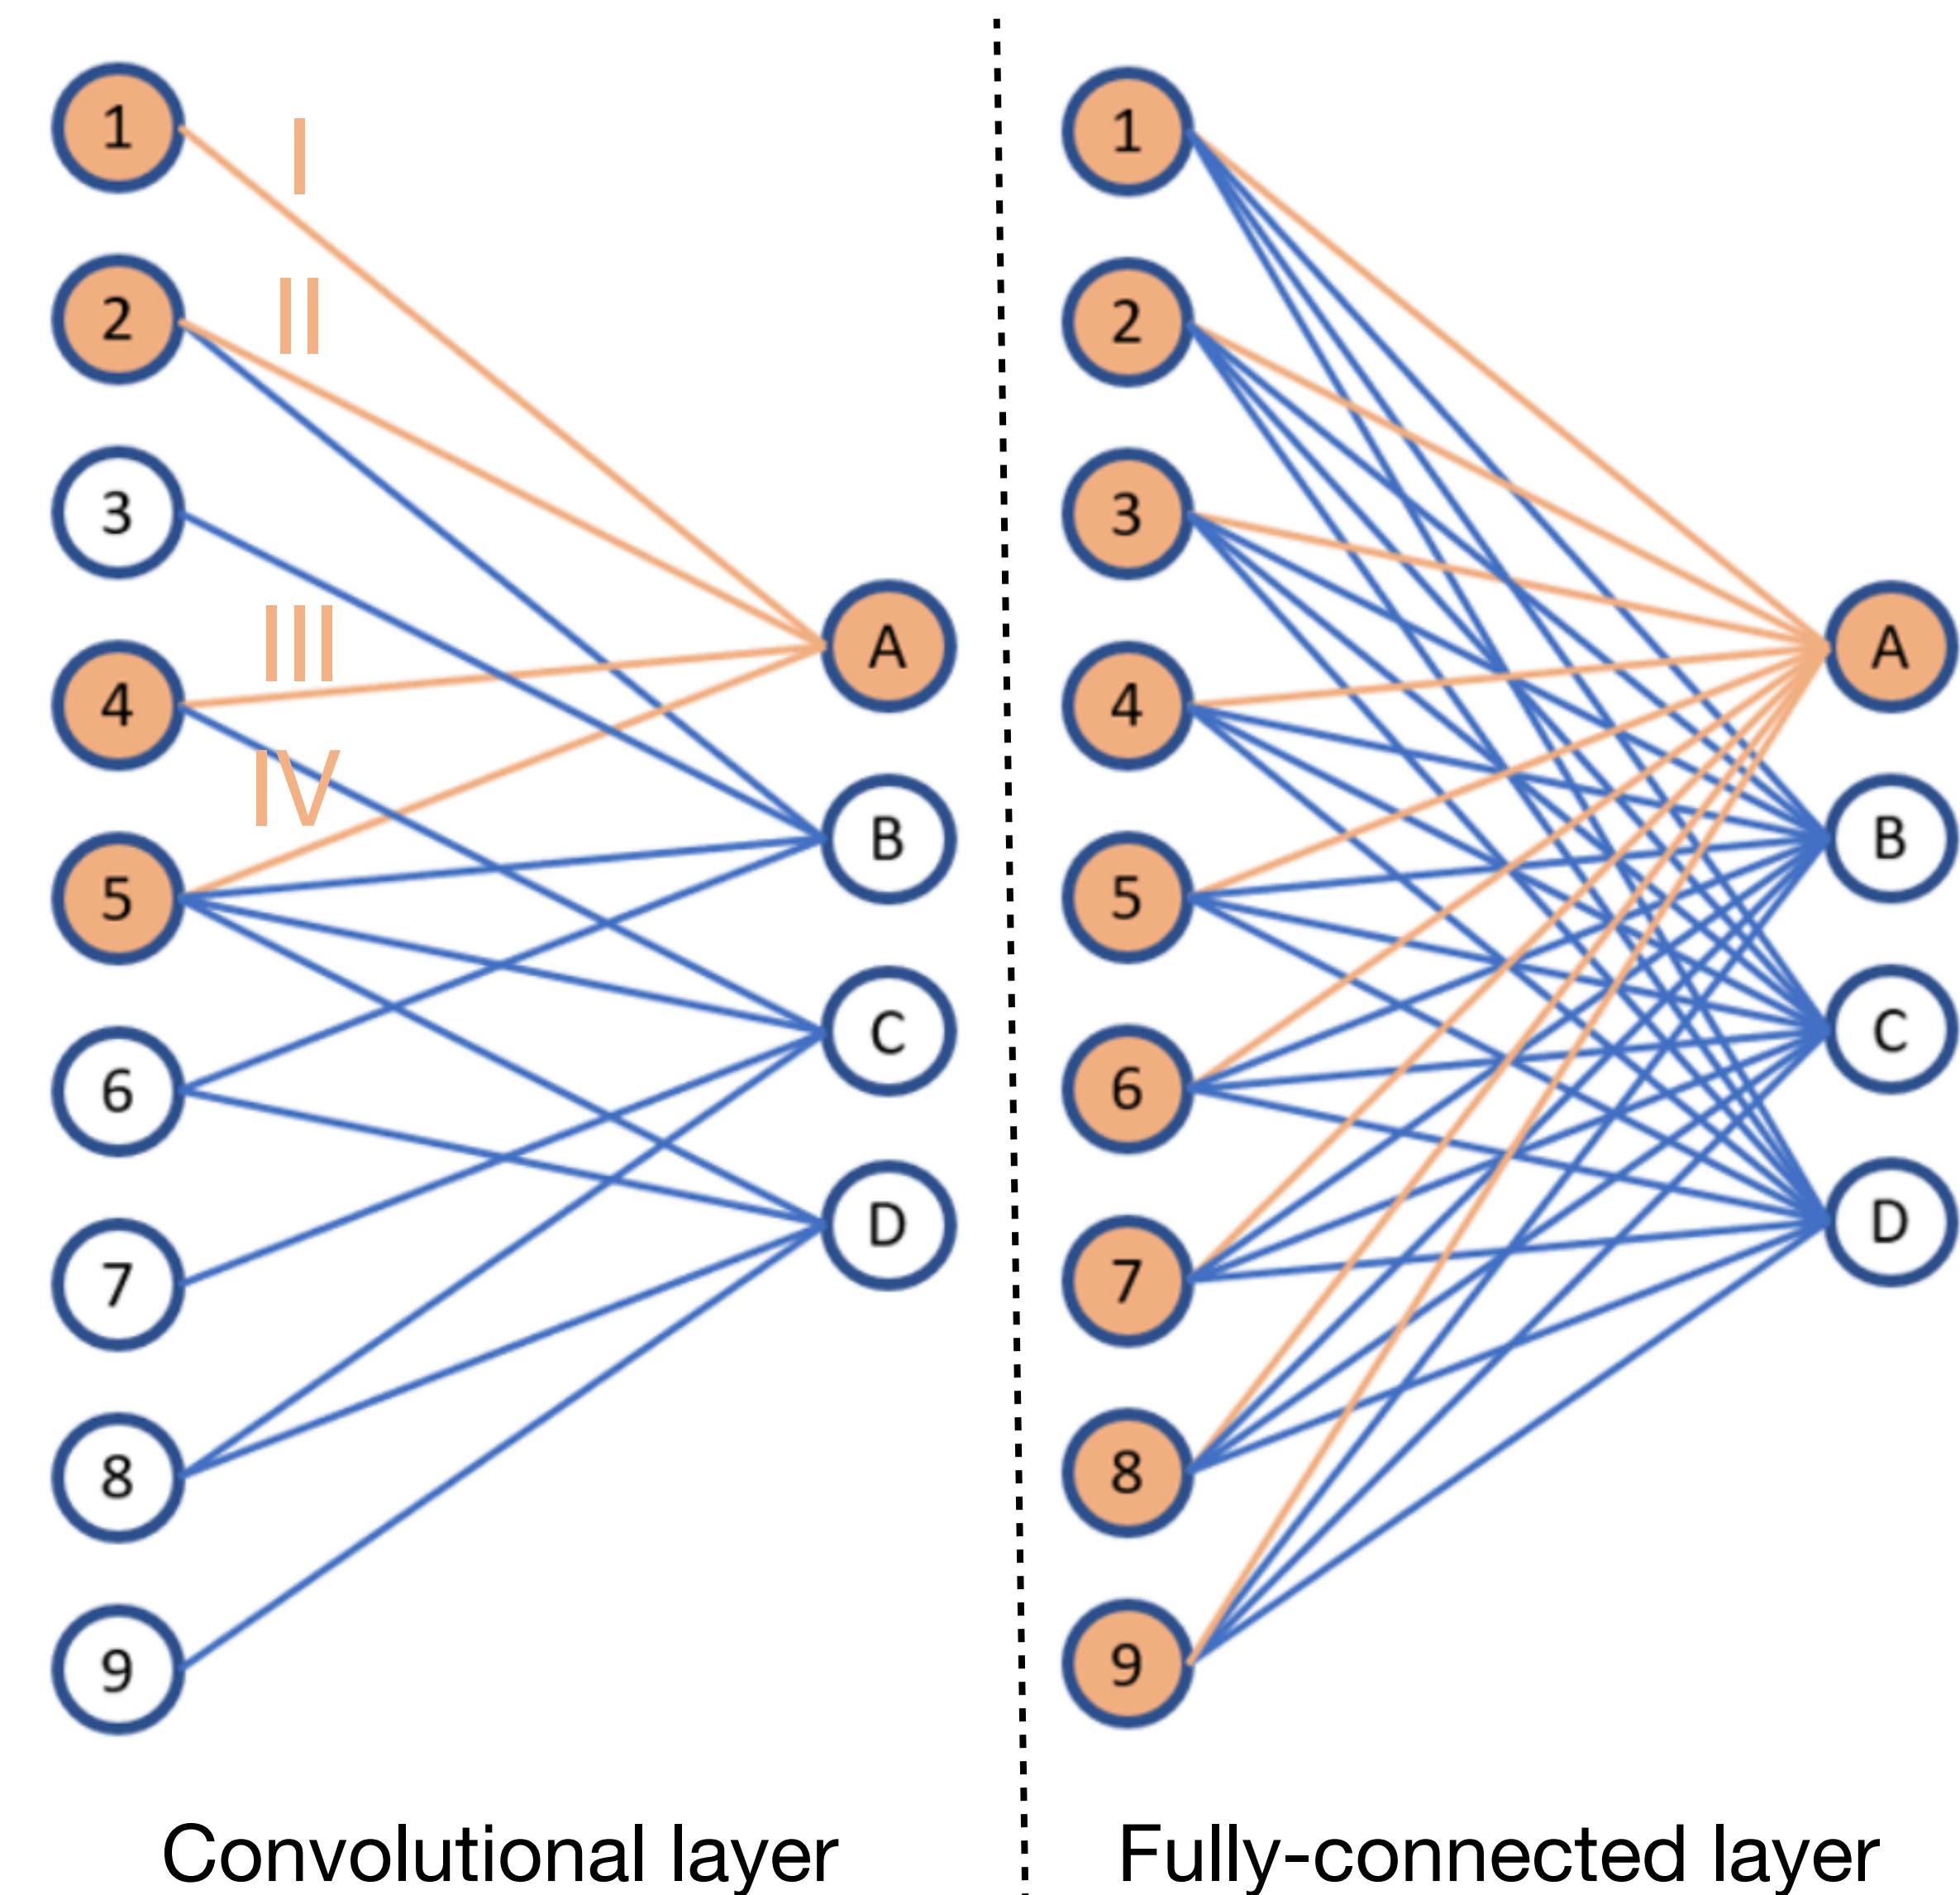
# Convolutional layer versus fully-connected layer

A convolutional layer can be visualized similarly to a fully-connected layer.



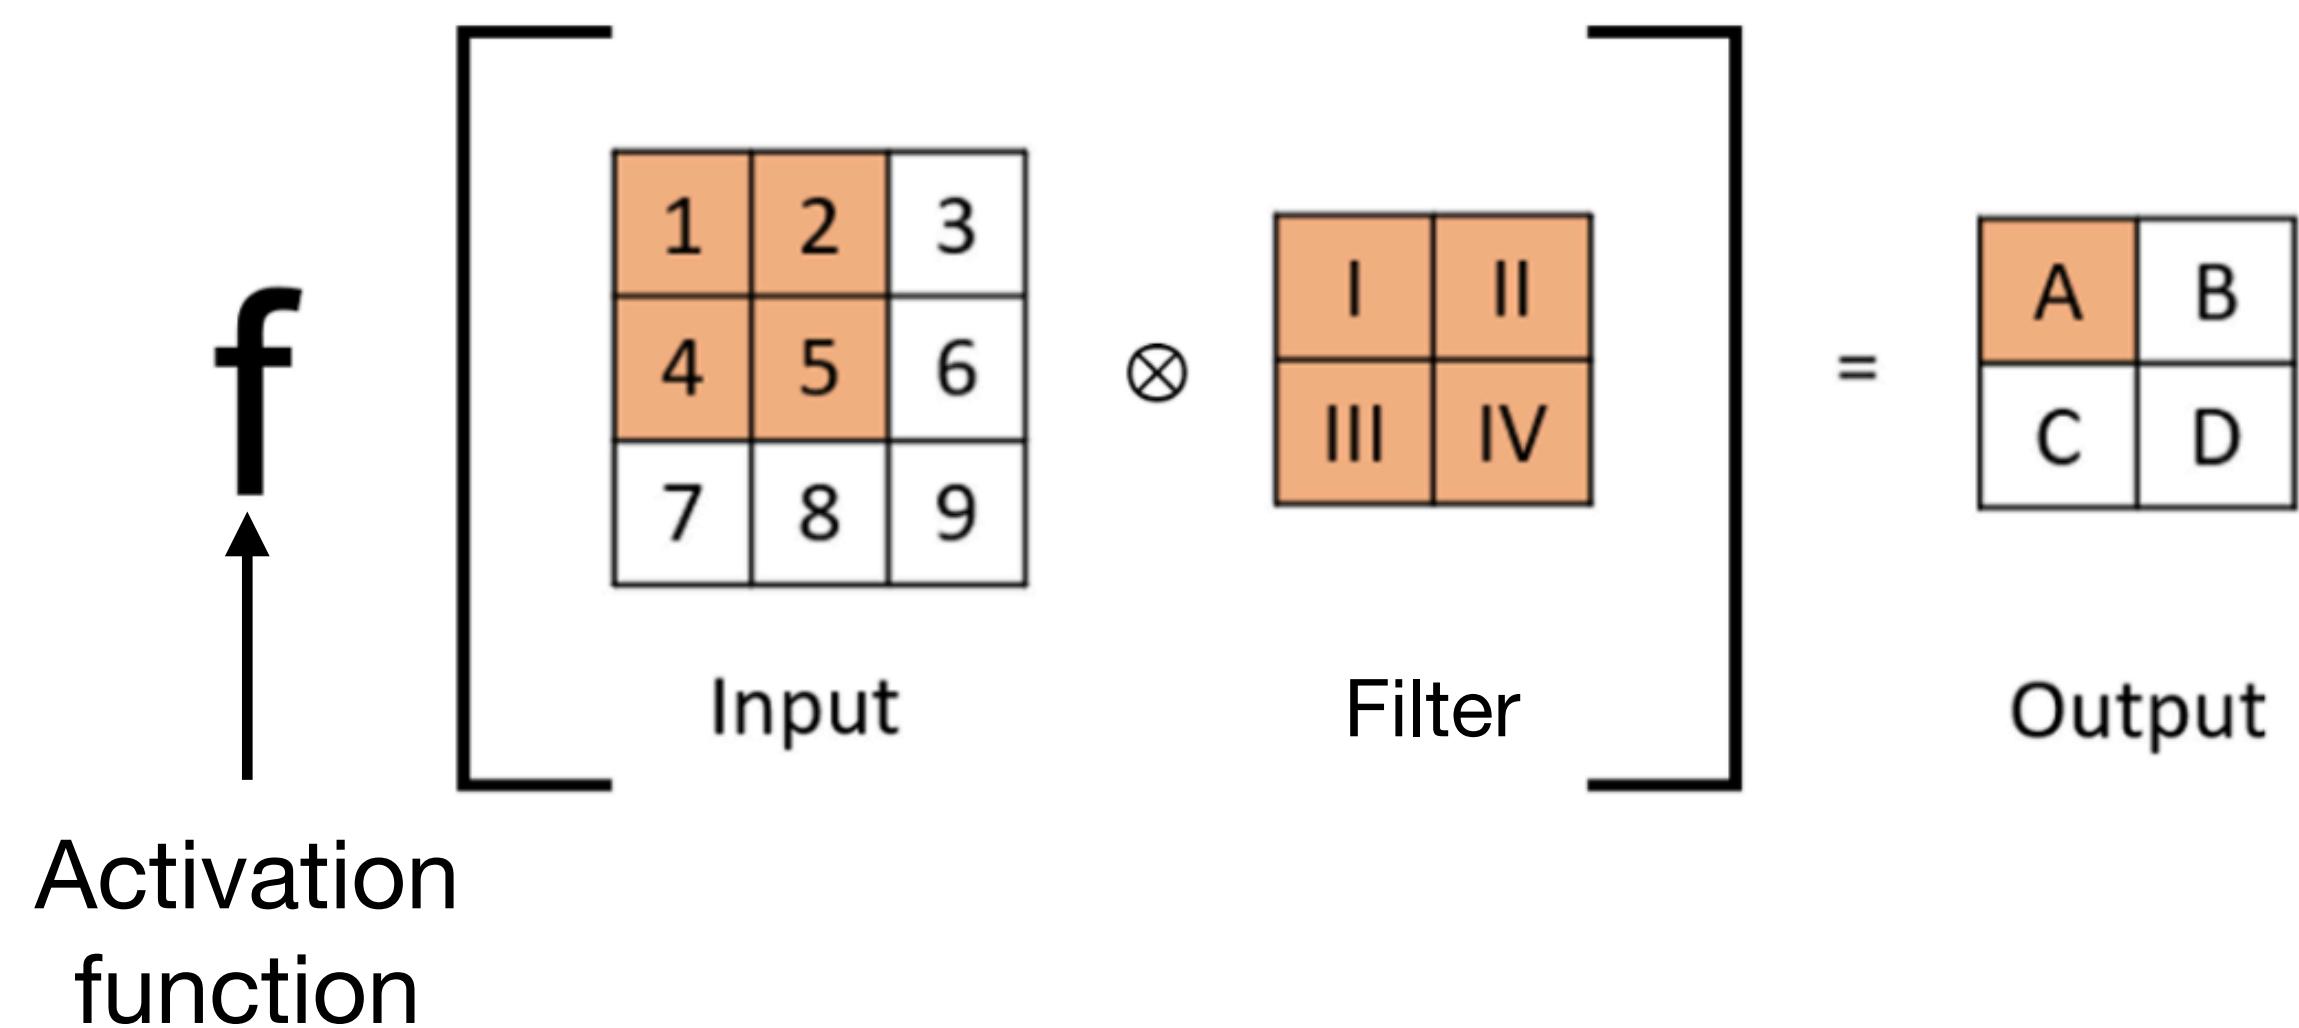
In a convolutional layer:

- Not all node pairs are connected with edges

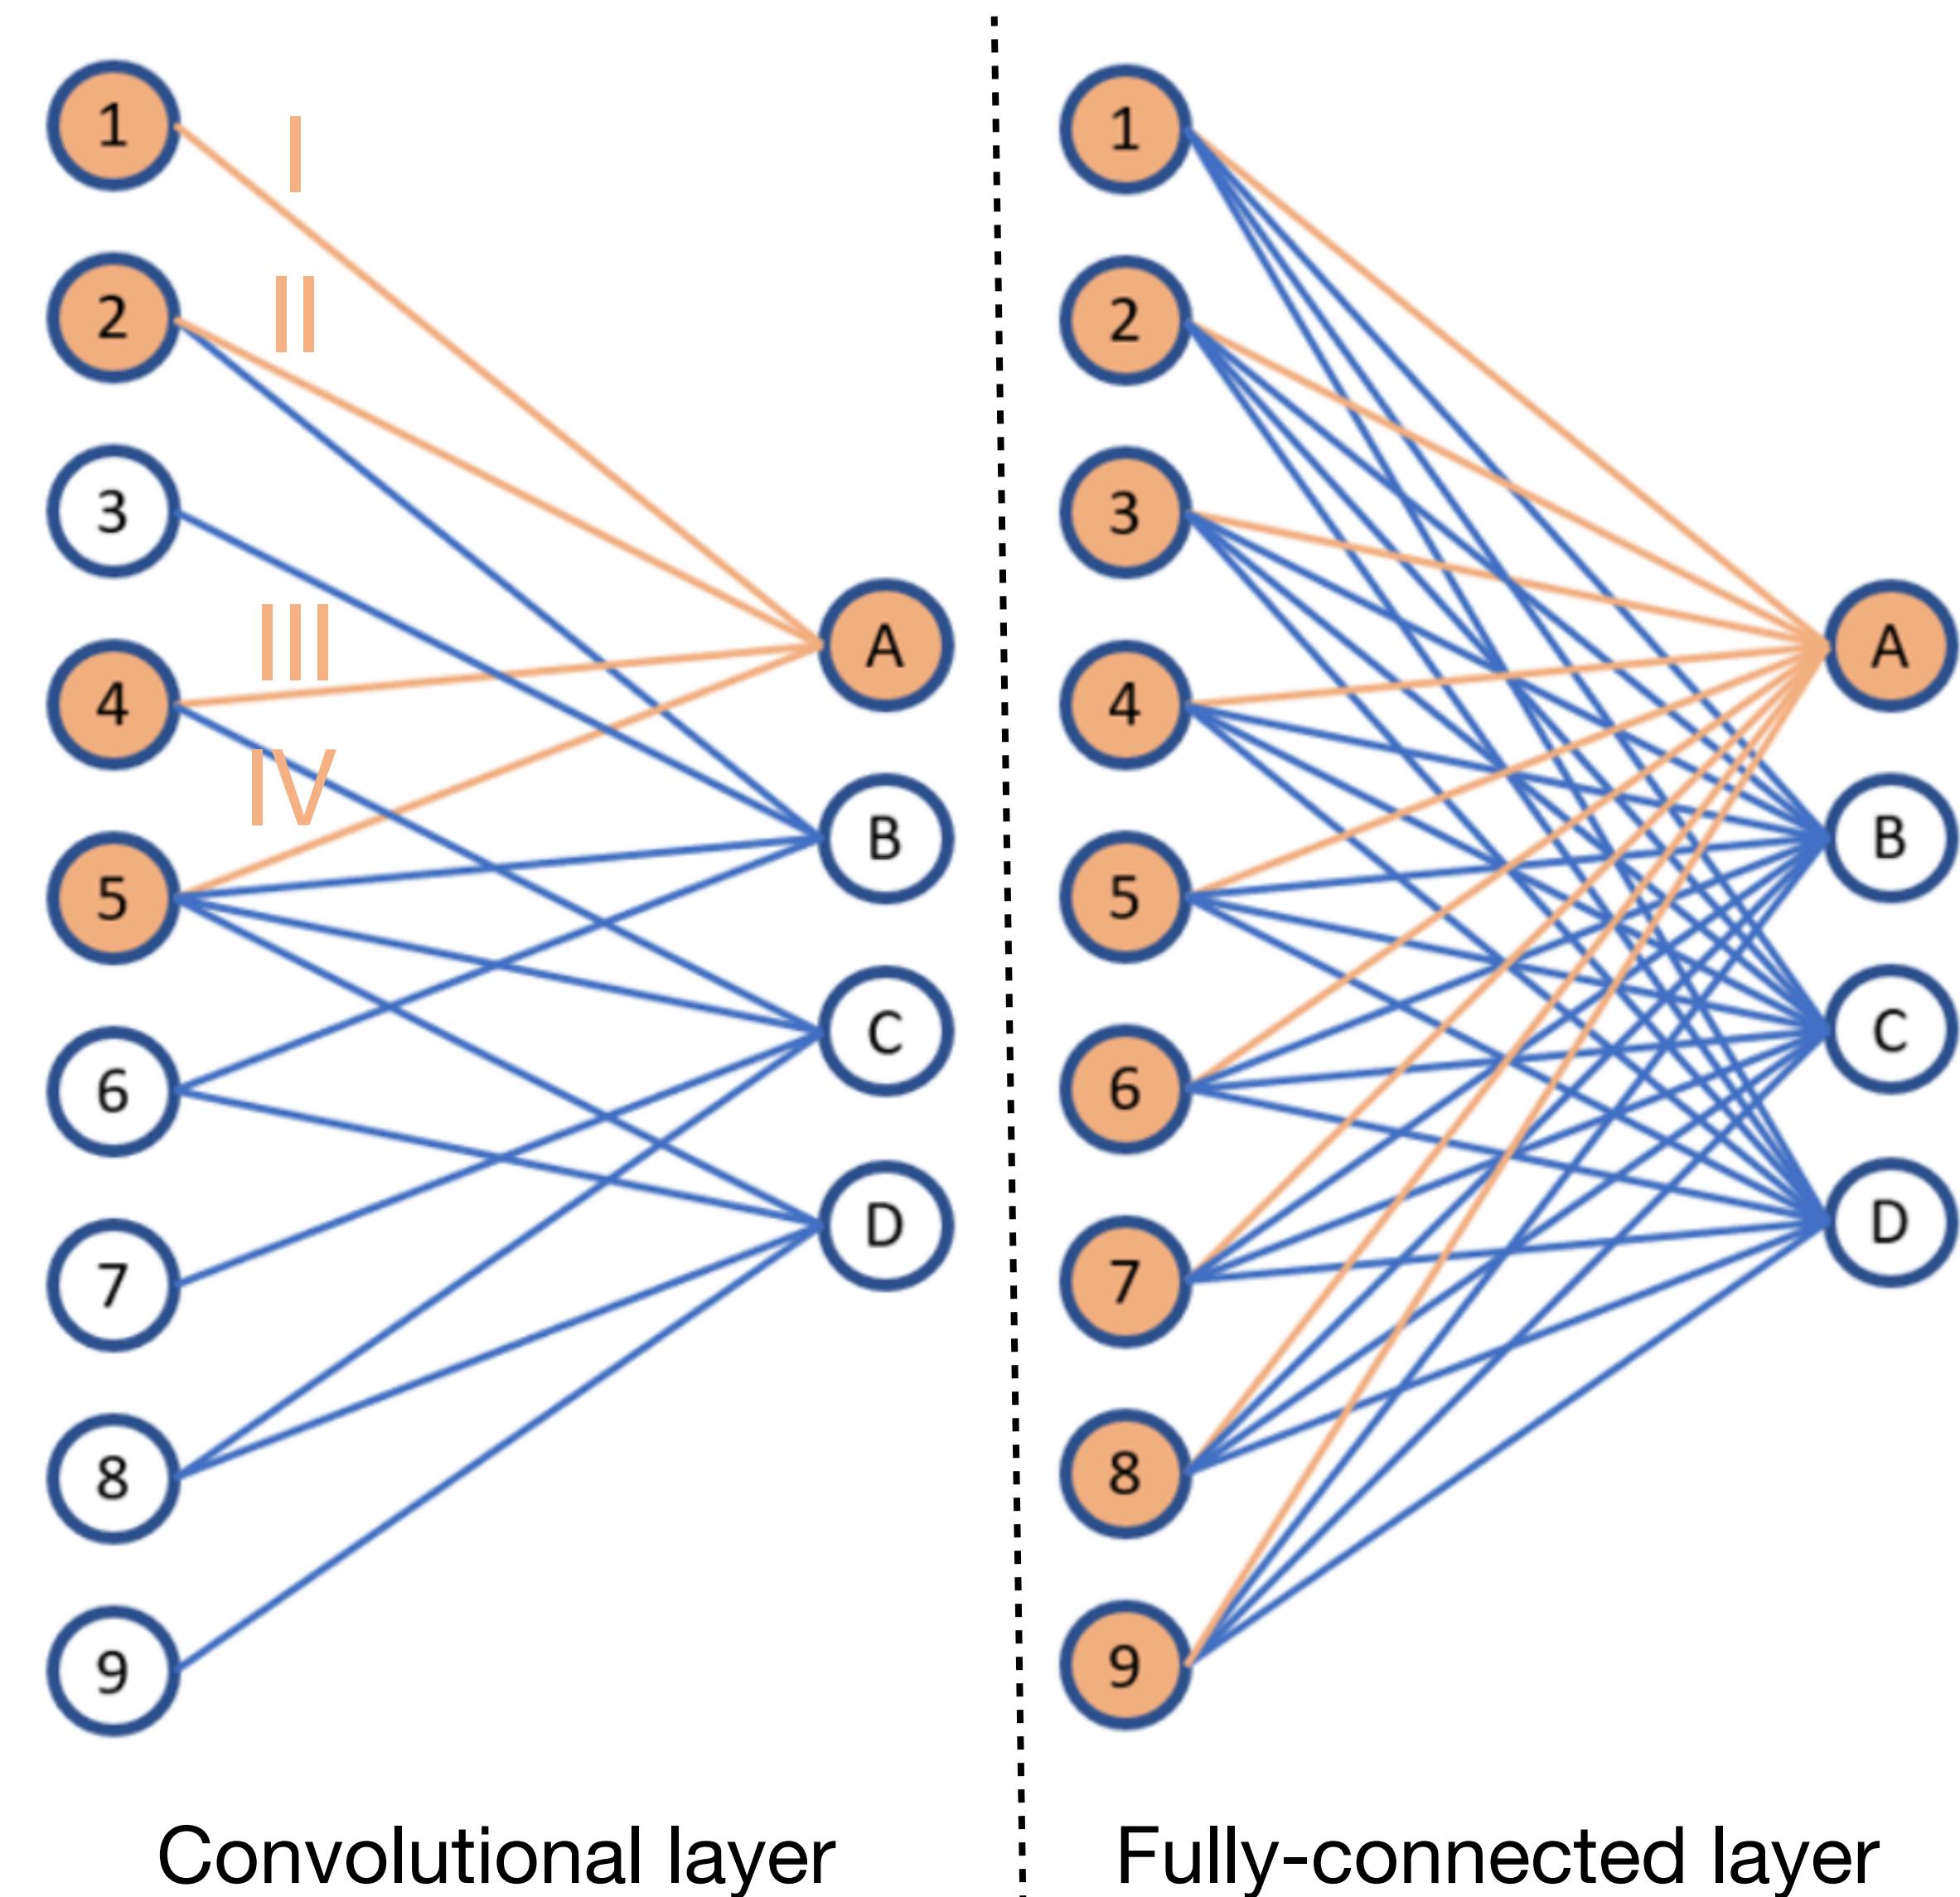


# Convolutional layer versus fully-connected layer

A convolutional layer can be visualized similarly to a fully-connected layer.

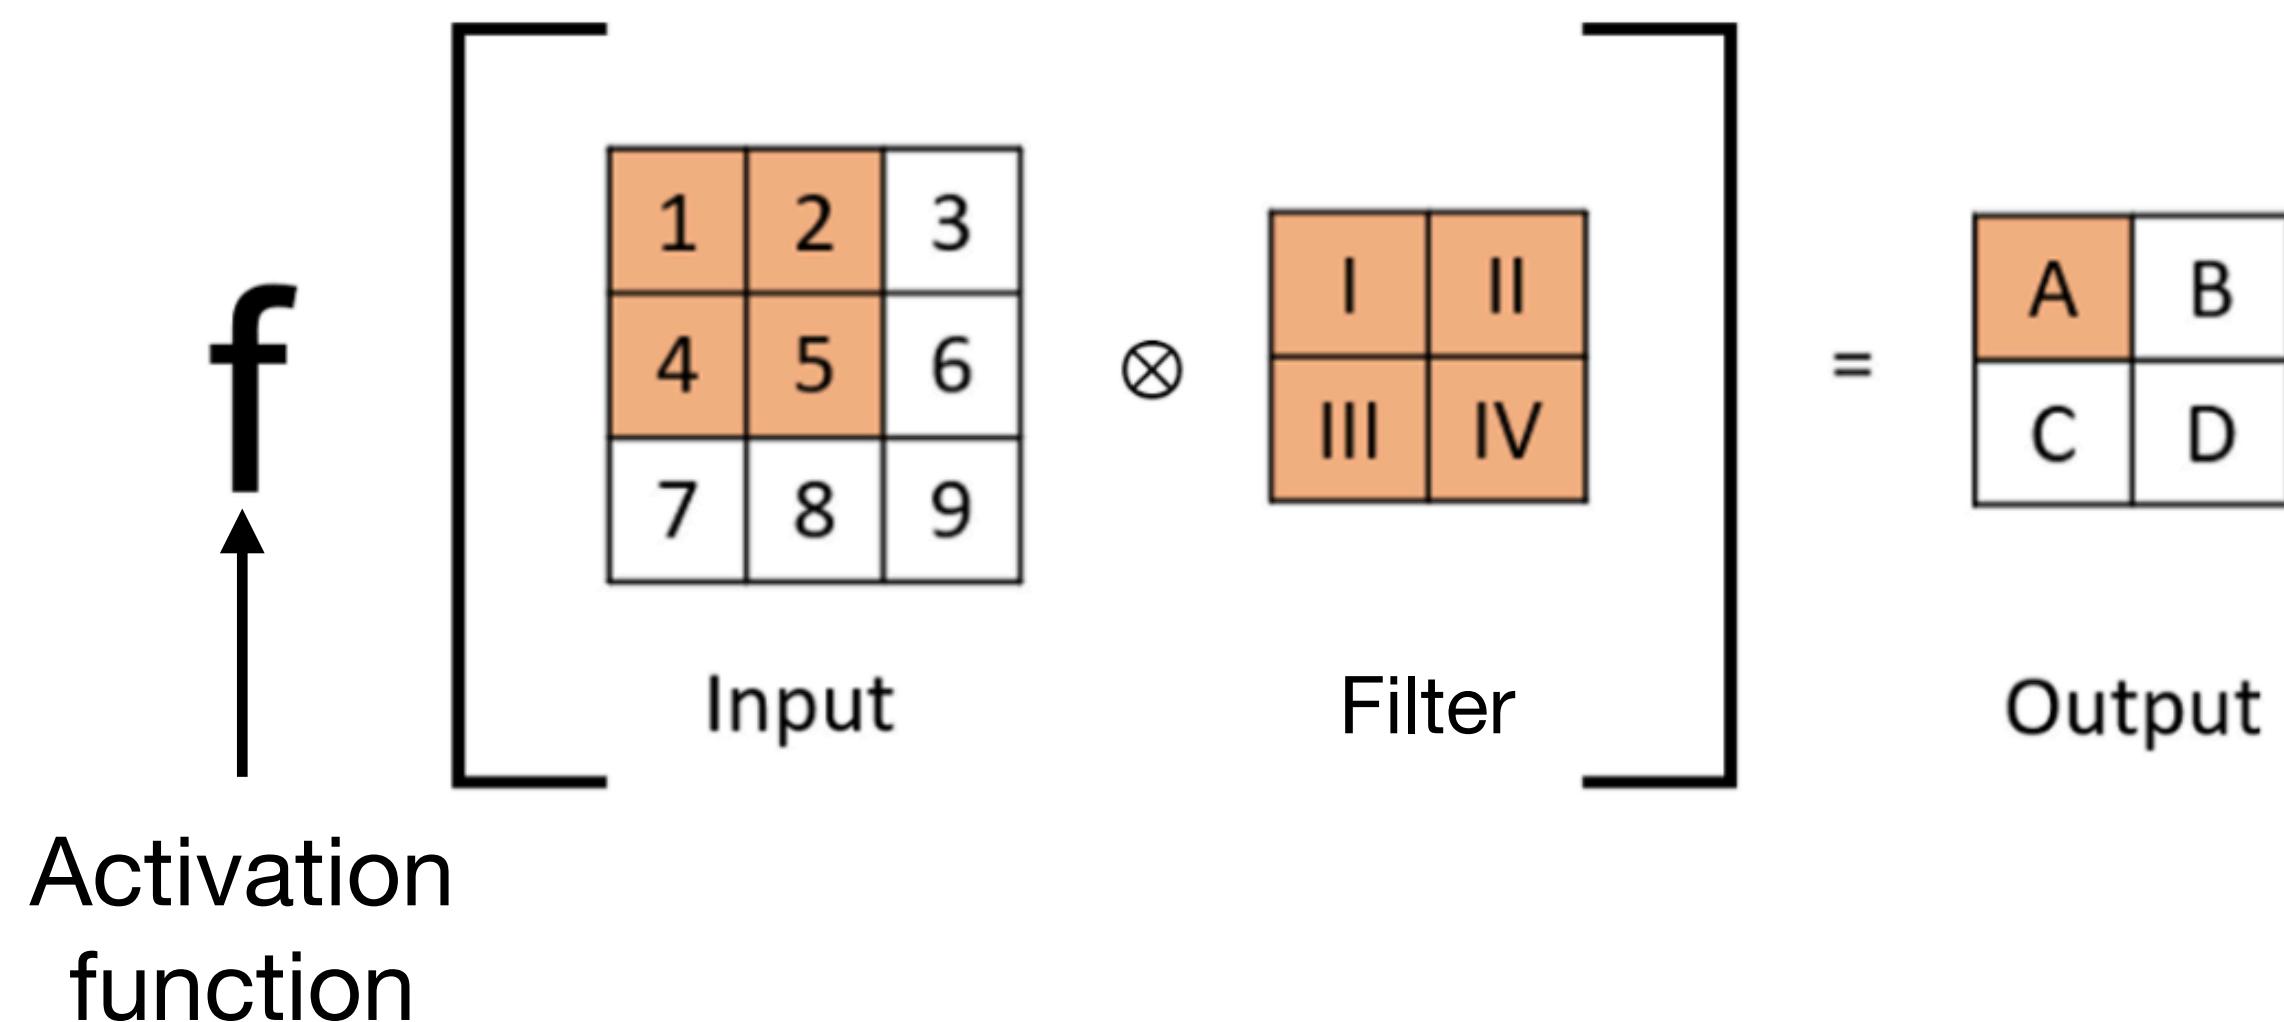


- In a convolutional layer:
- Not all node pairs are connected with edges
  - Weights (from filter) reused across edges



# Convolutional layer versus fully-connected layer

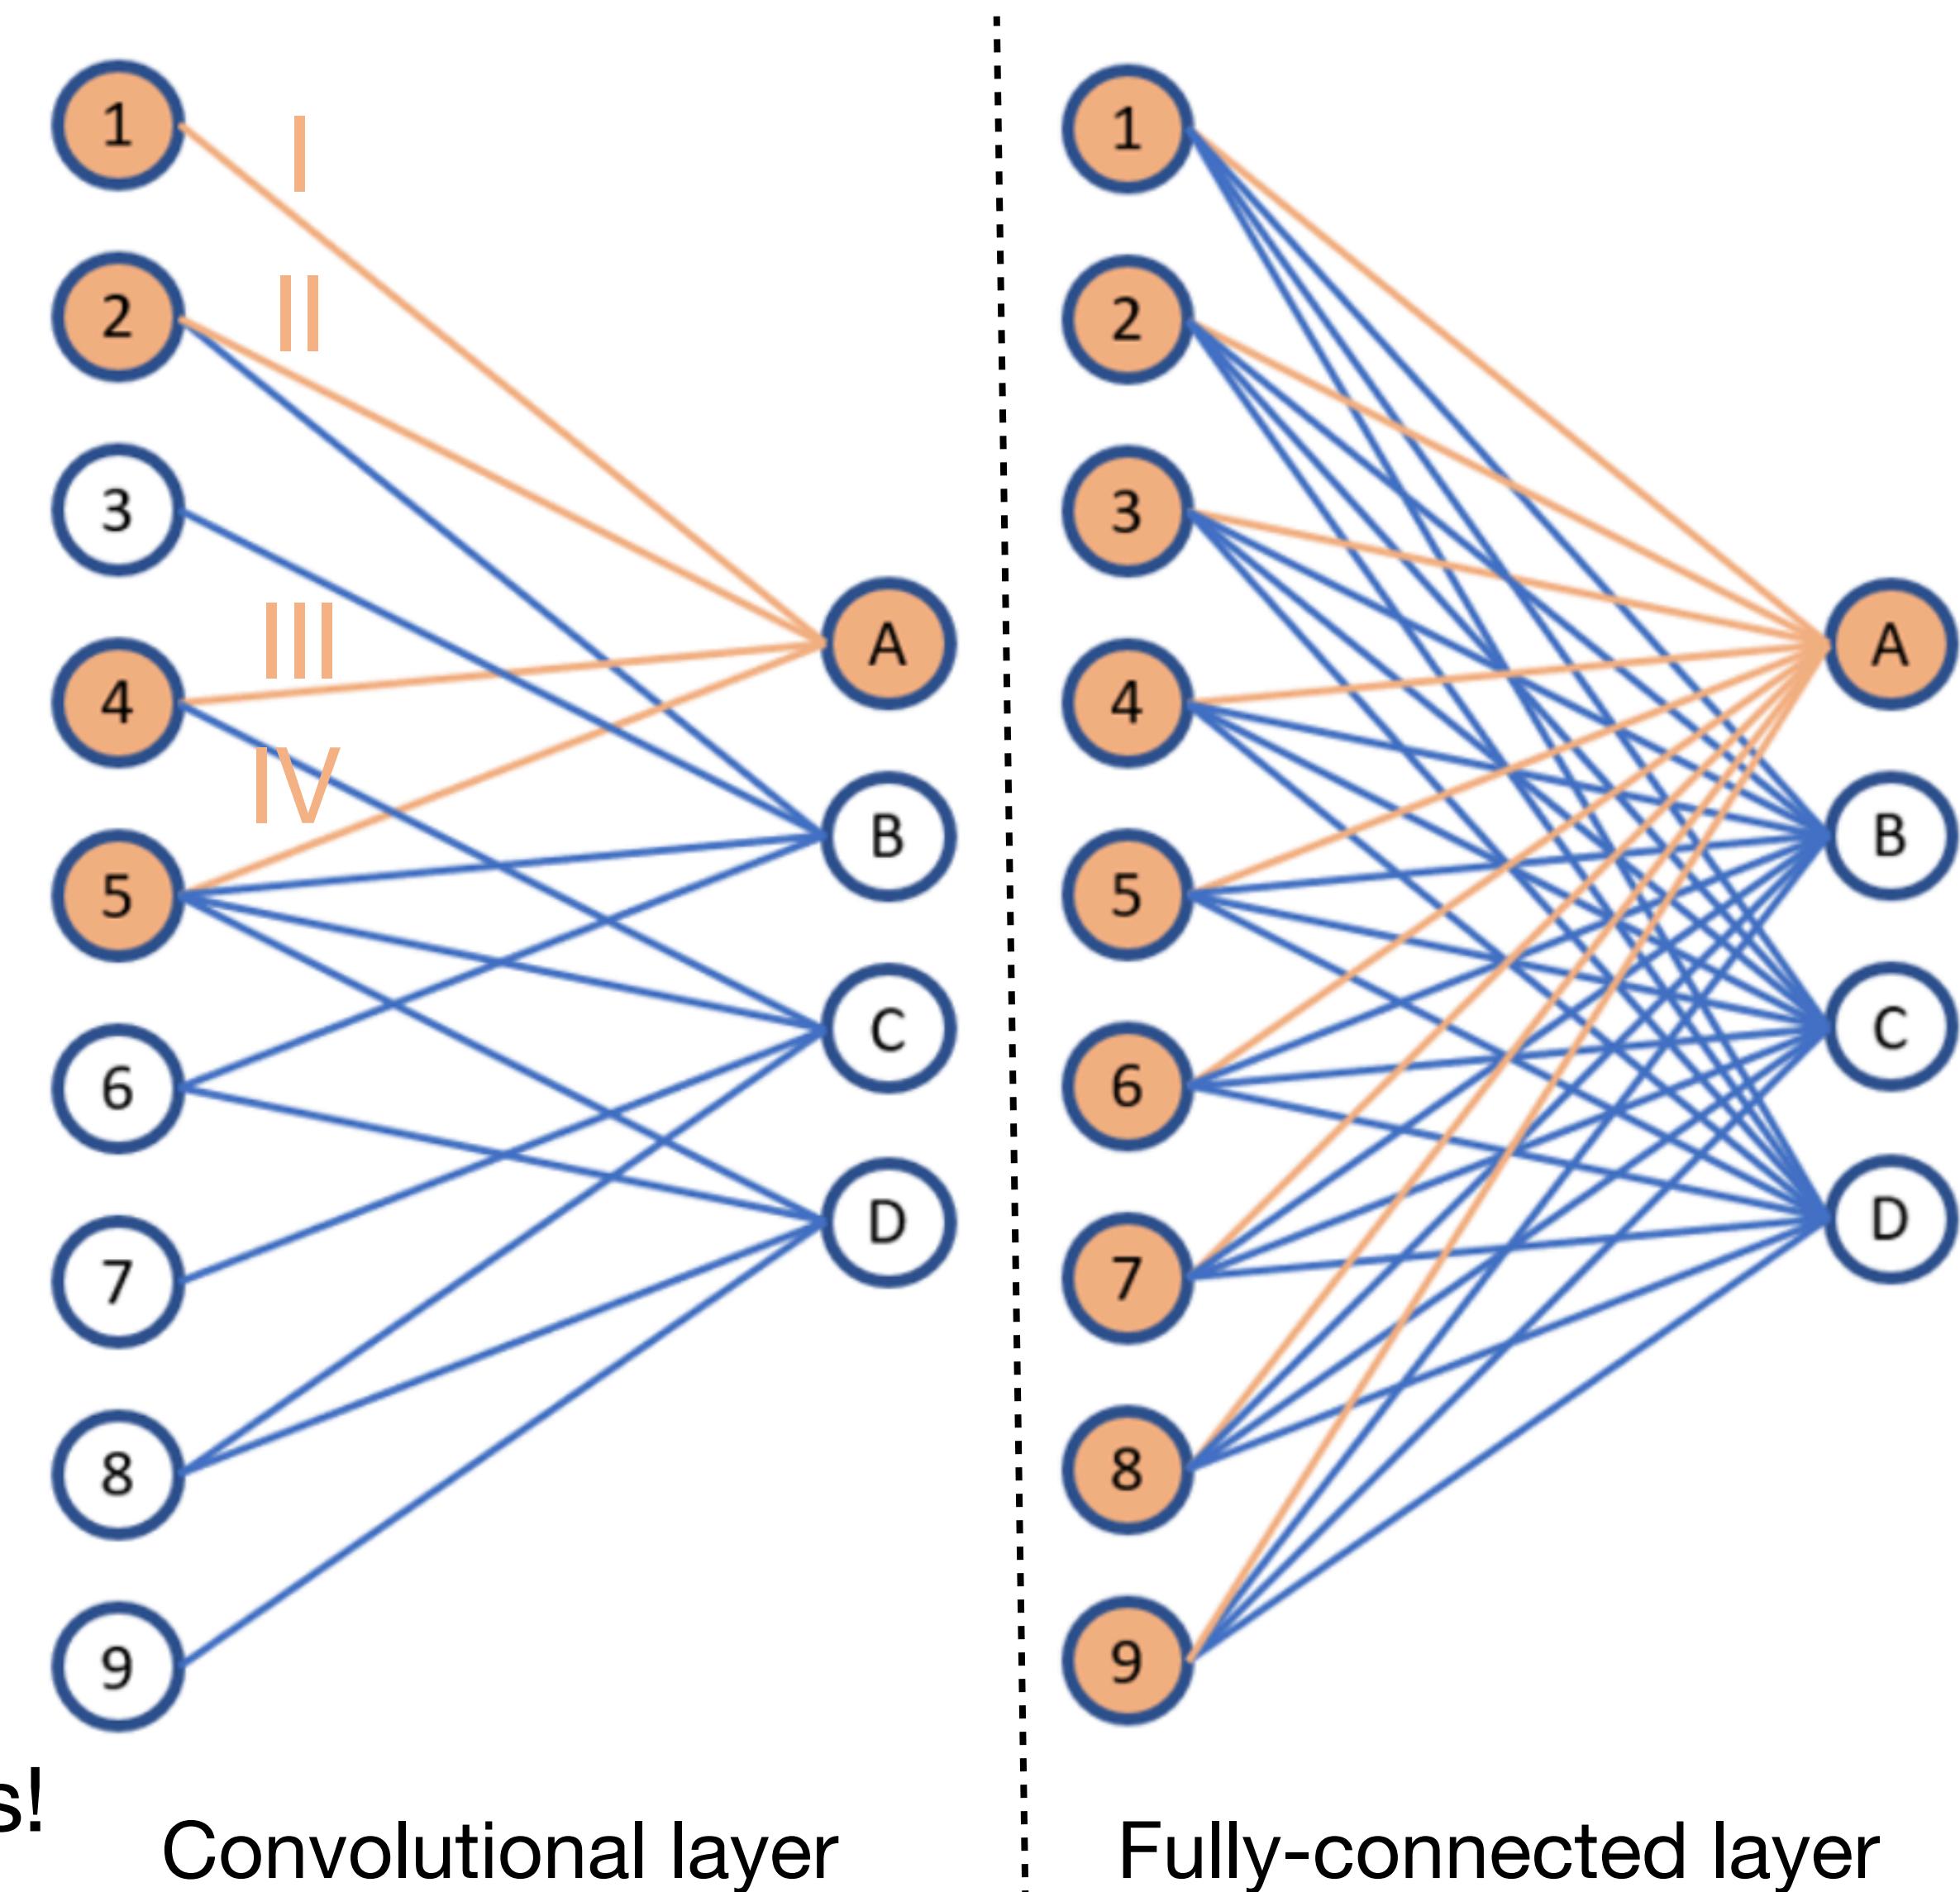
A convolutional layer can be visualized similarly to a fully-connected layer.



In a convolutional layer:

- Not all node pairs are connected with edges
- Weights (from filter) reused across edges

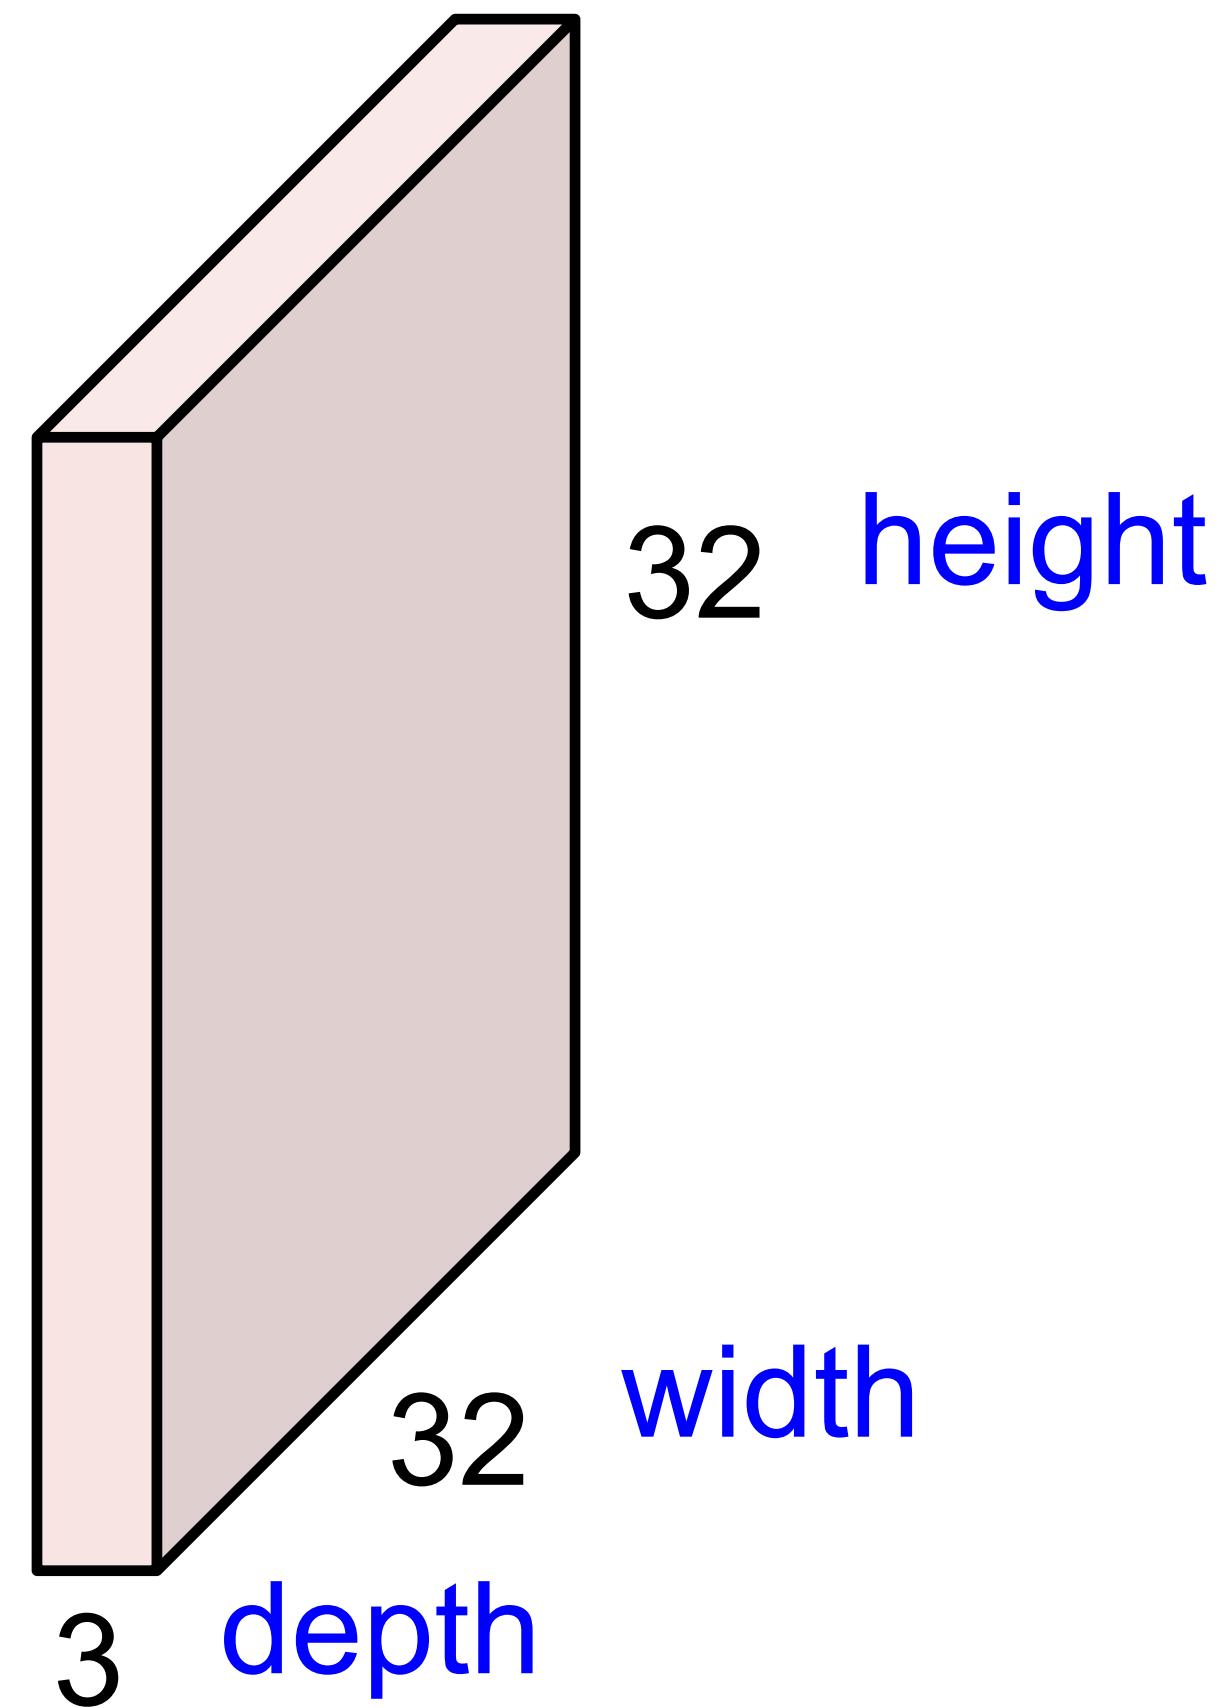
Consequence: Conv layers have fewer parameters!



# Convolution Layer

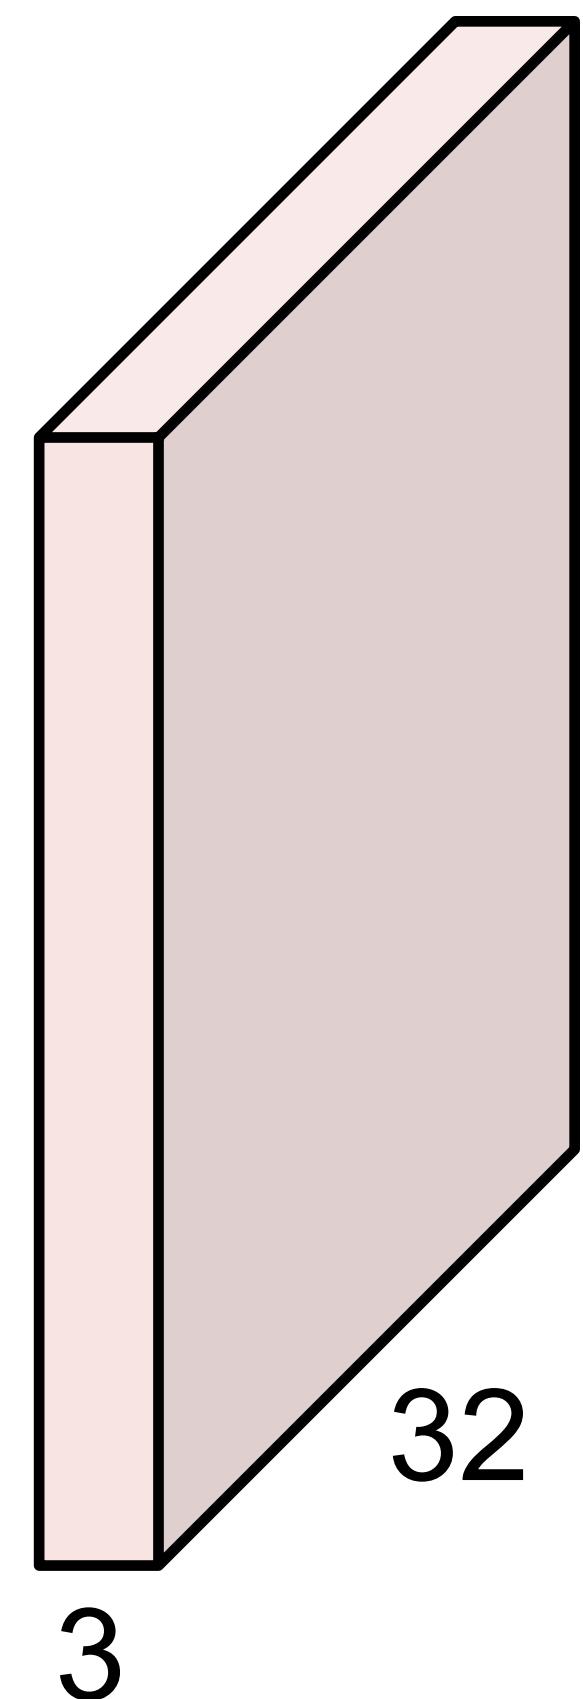
Note: This slide and several following ones are borrowed from Stanford's [CS231n](#).

32x32x3 image (images typically have red, green, and blue **channels**.)

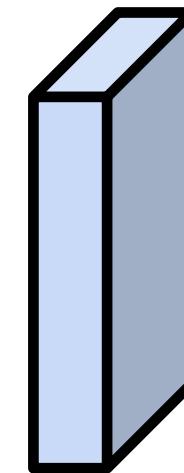


# Convolution Layer

32x32x3 image

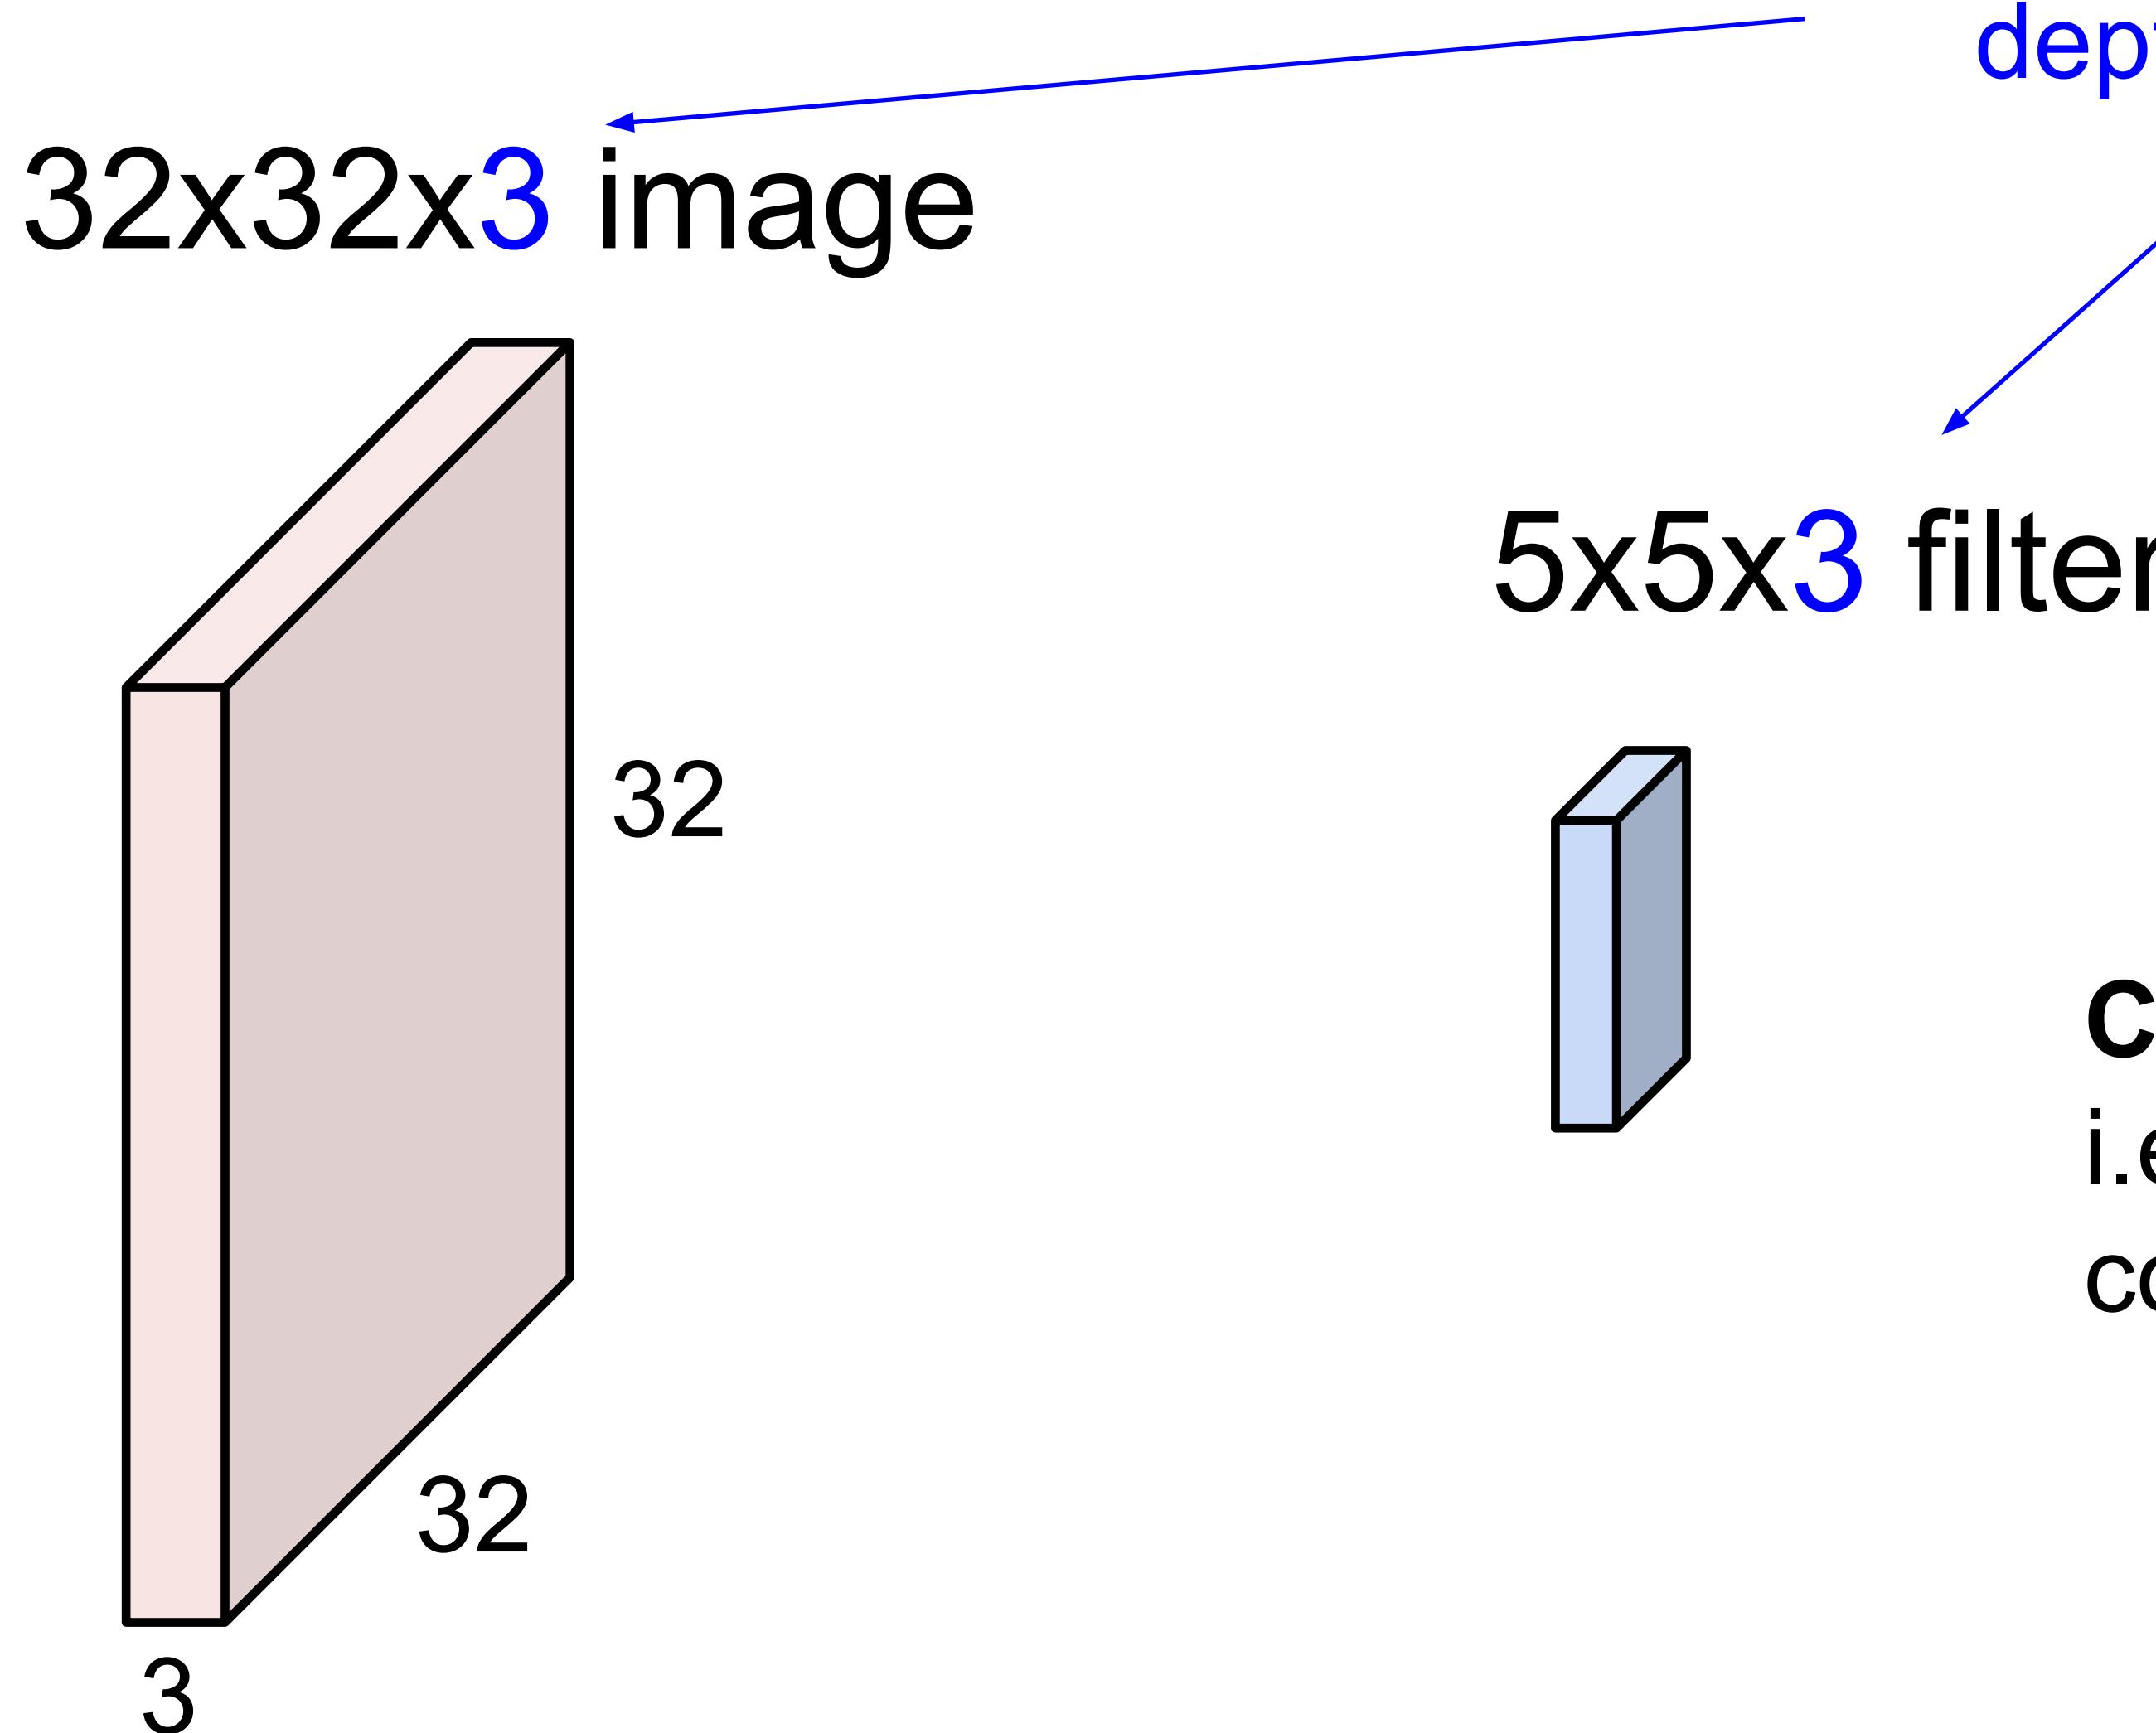


5x5x3 filter



**Convolve** the filter with the image  
i.e. “slide over the image spatially,  
computing dot products”

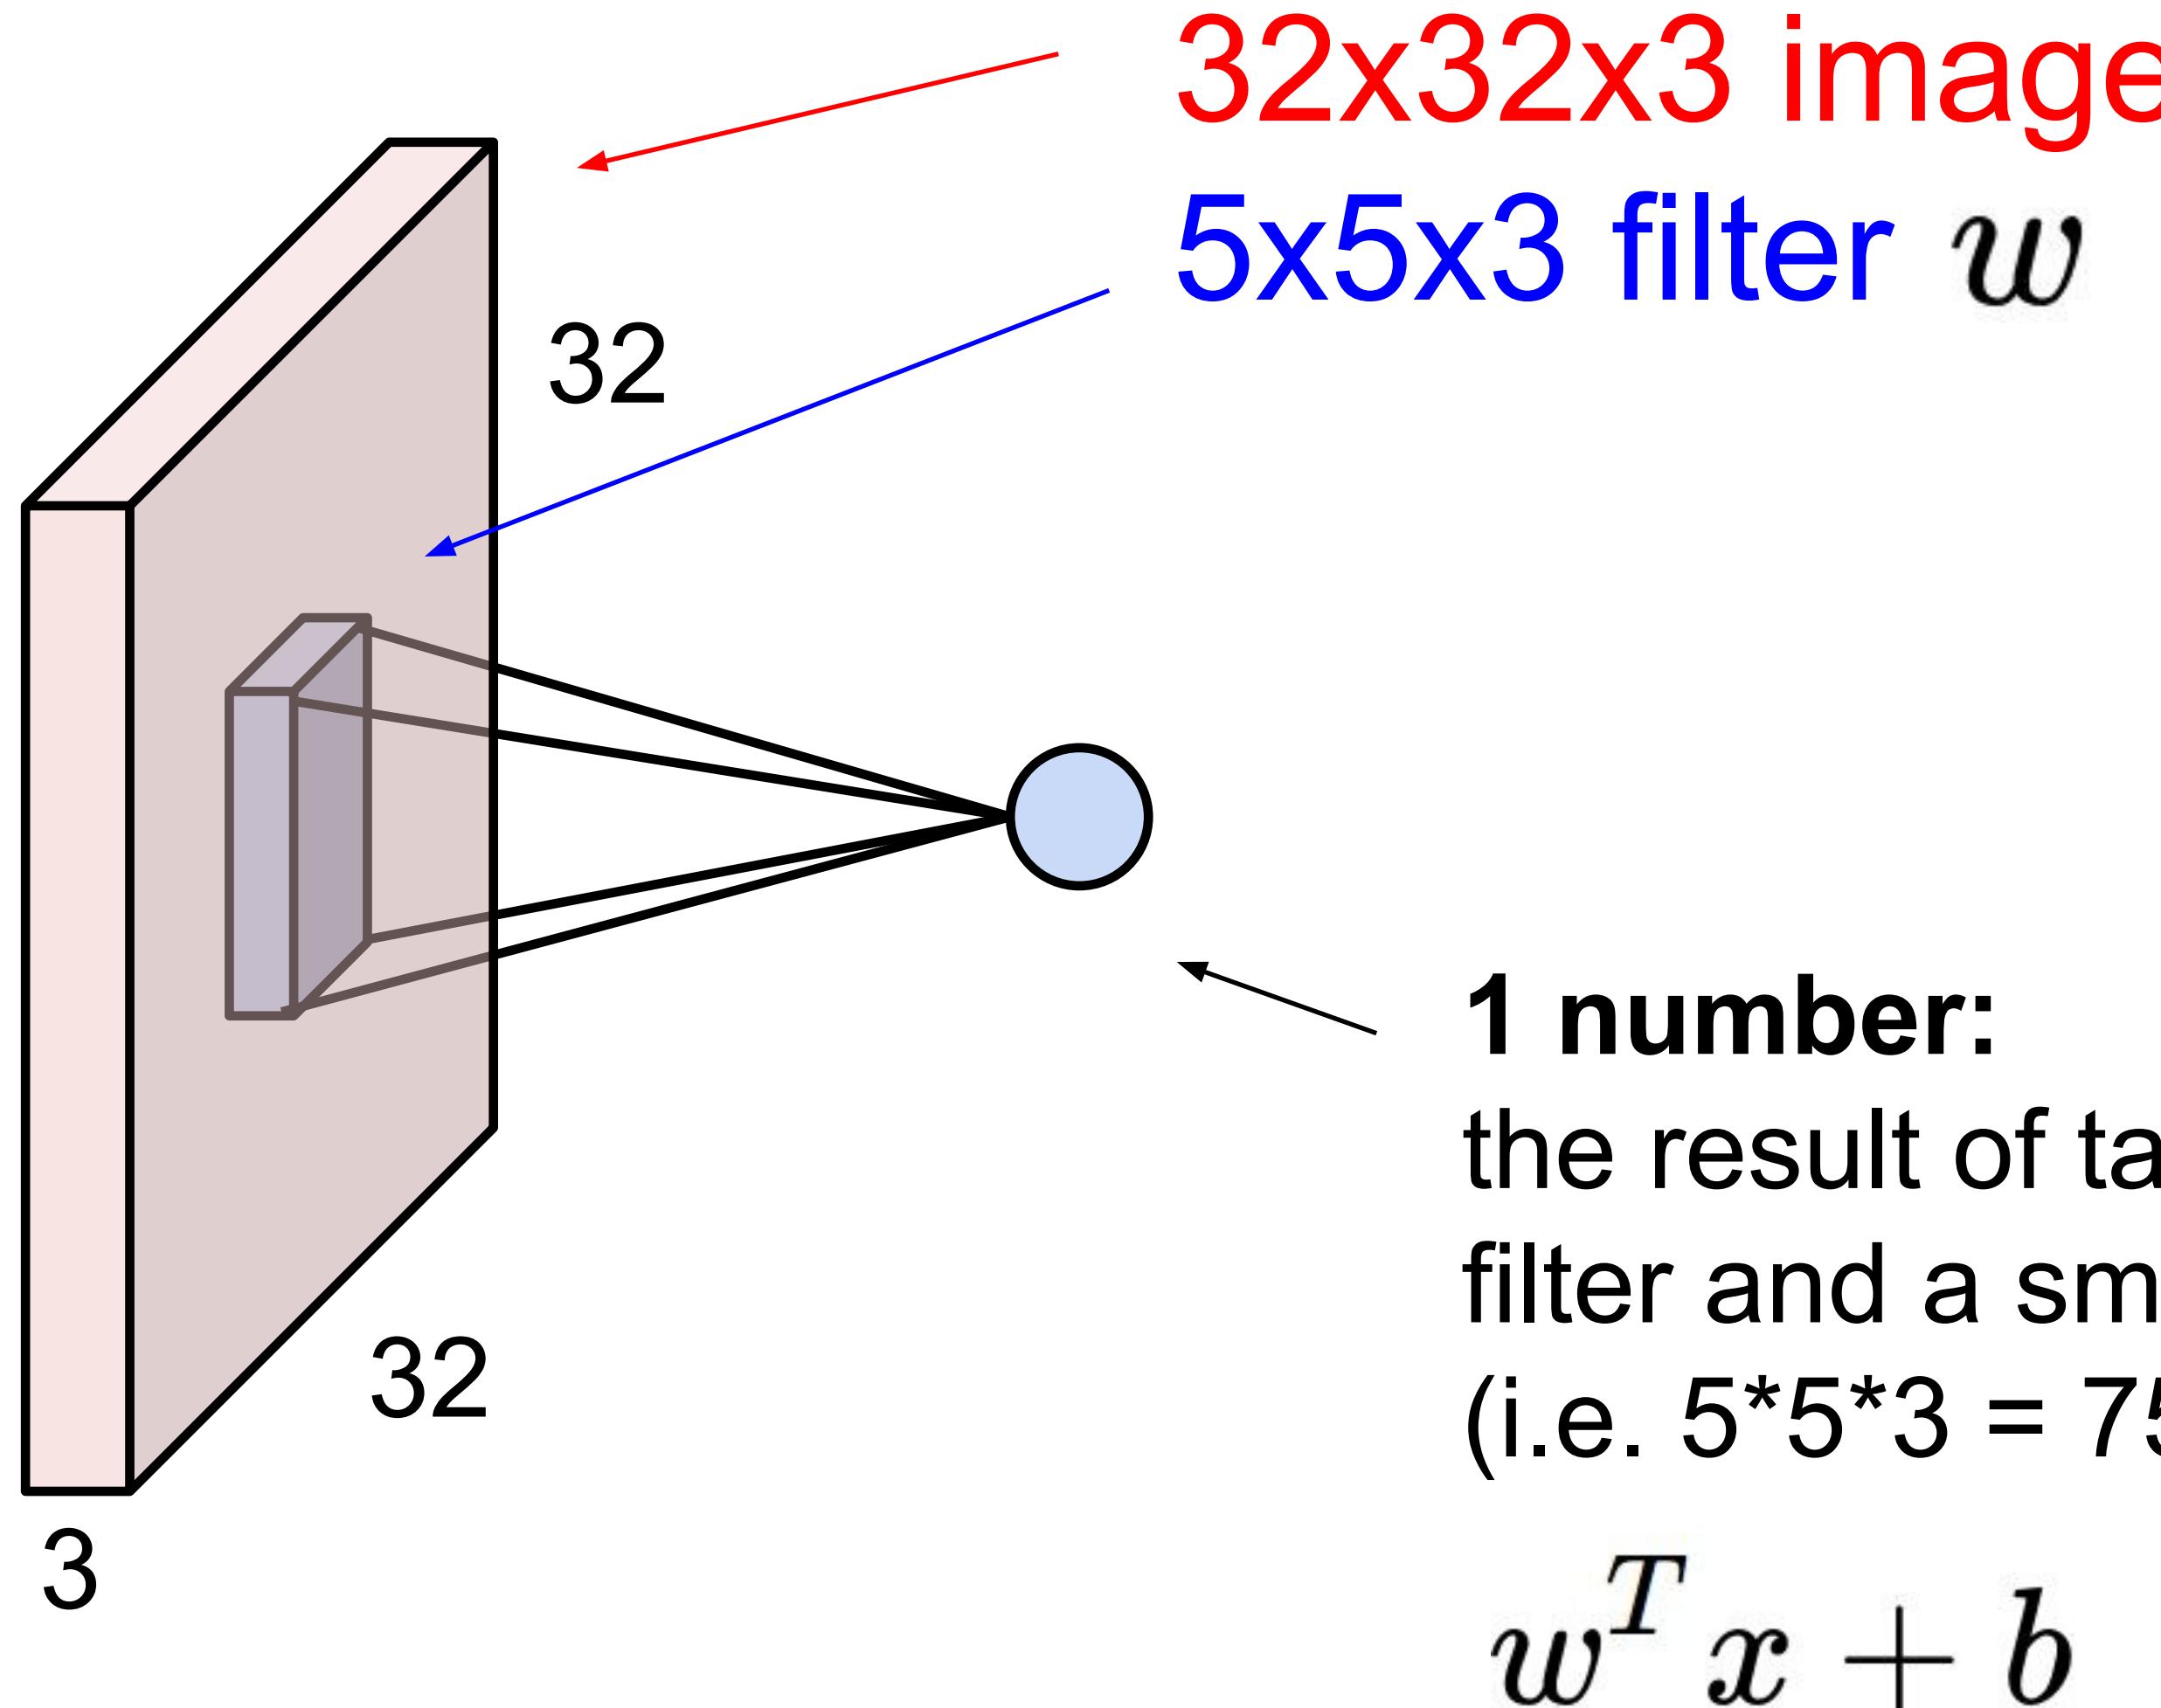
# Convolution Layer



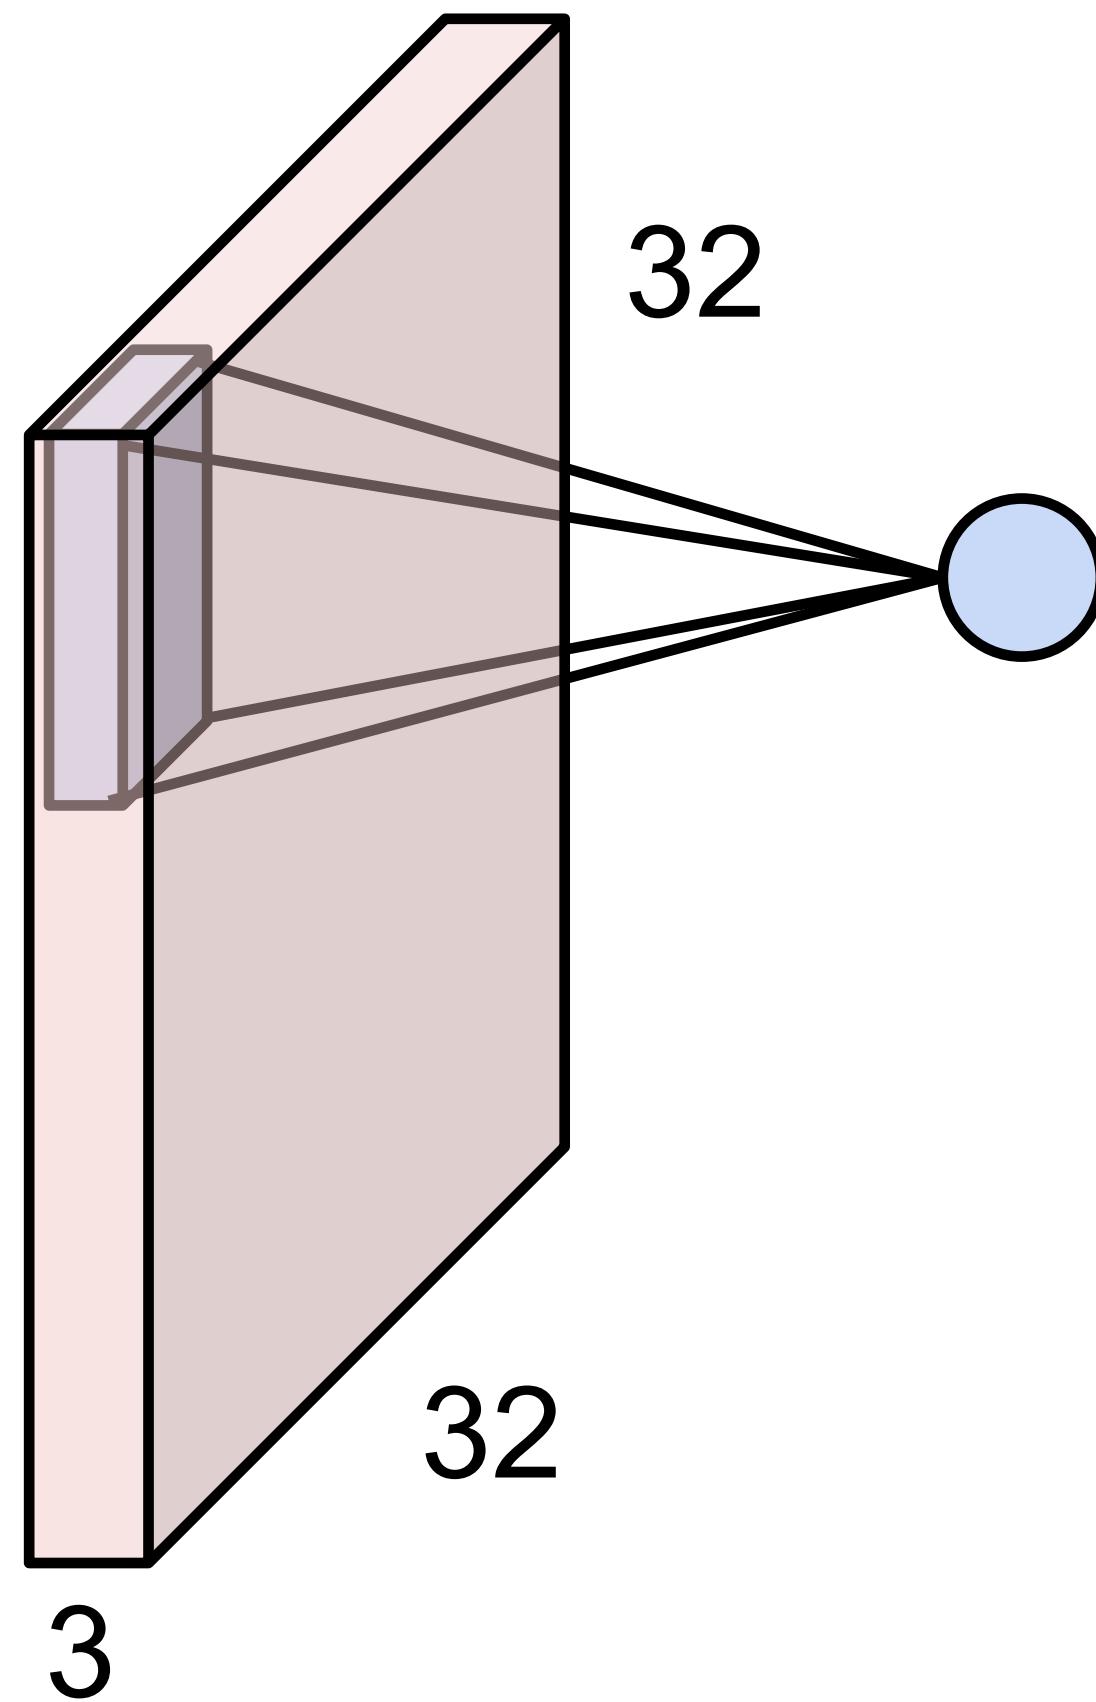
Filters always extend the full depth of the input volume

**Convolve** the filter with the image  
i.e. “slide over the image spatially,  
computing dot products”

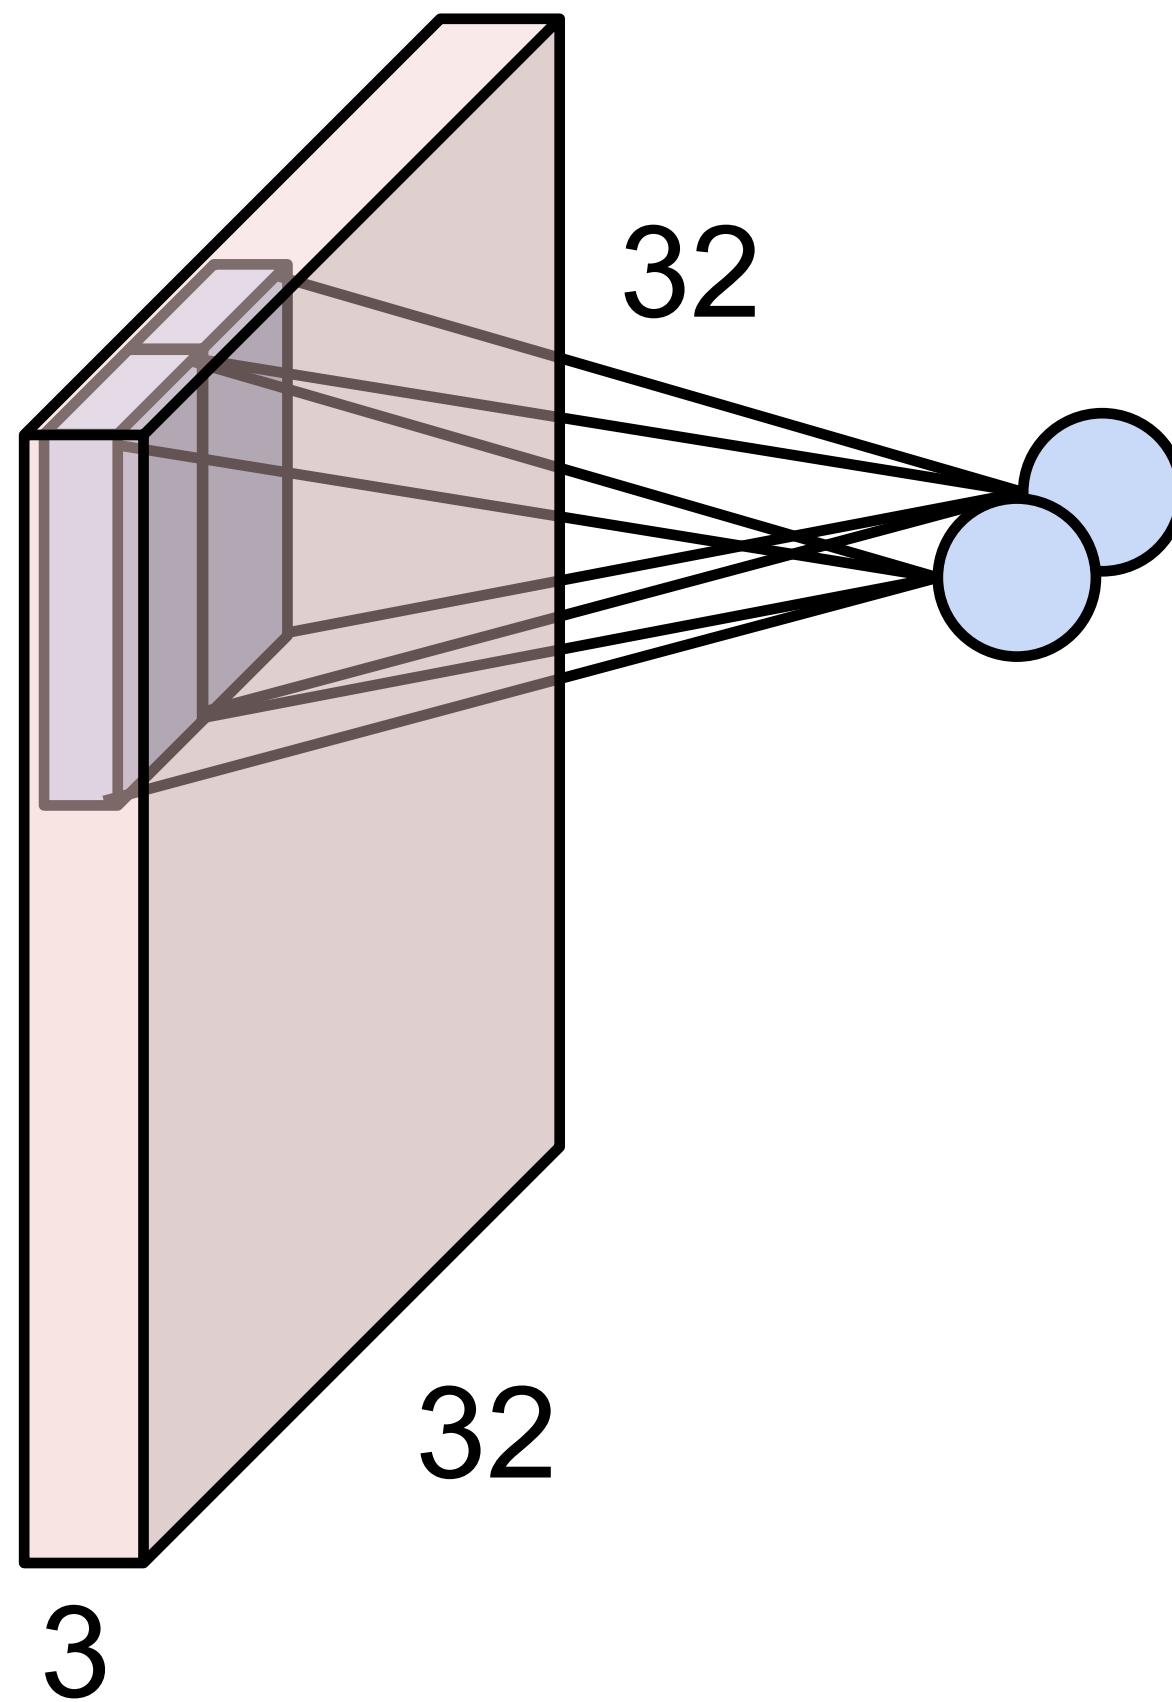
# Convolution Layer



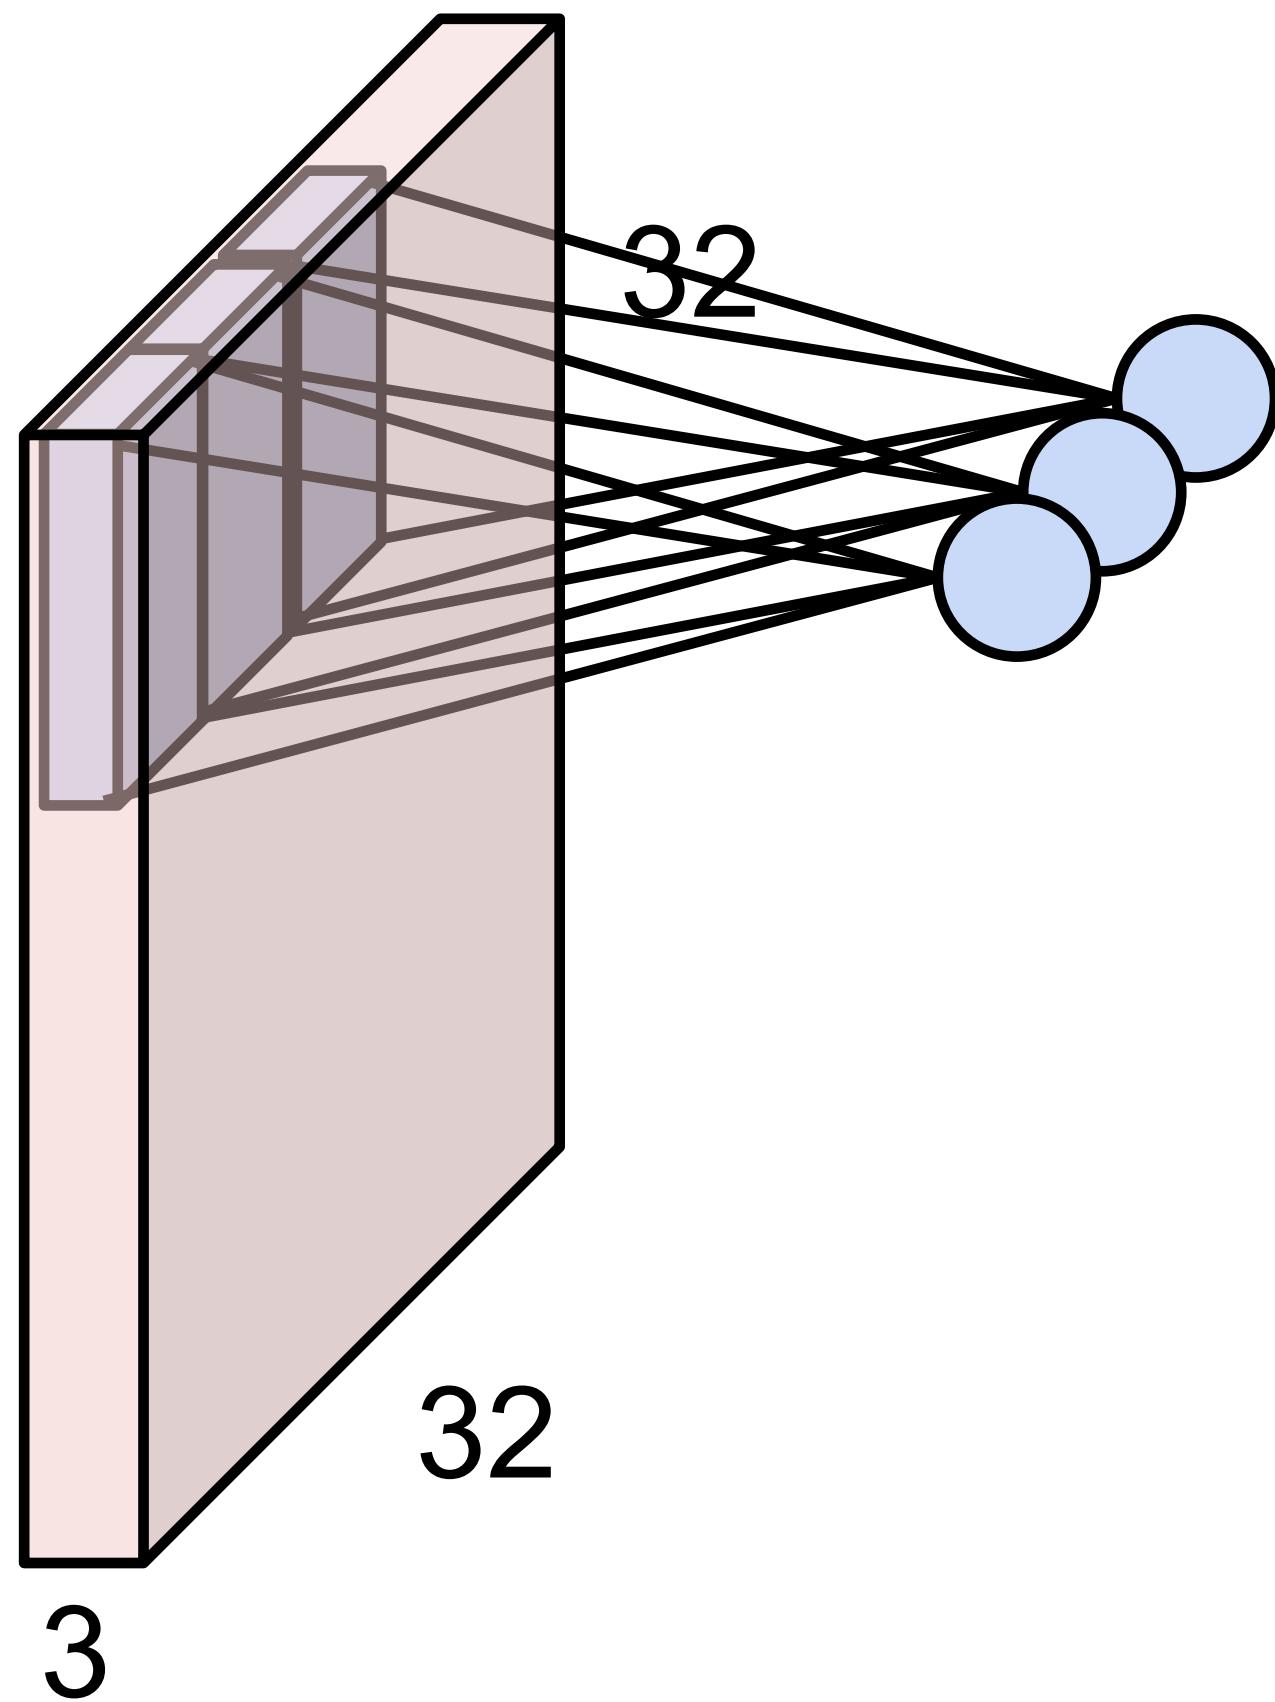
# Convolution Layer



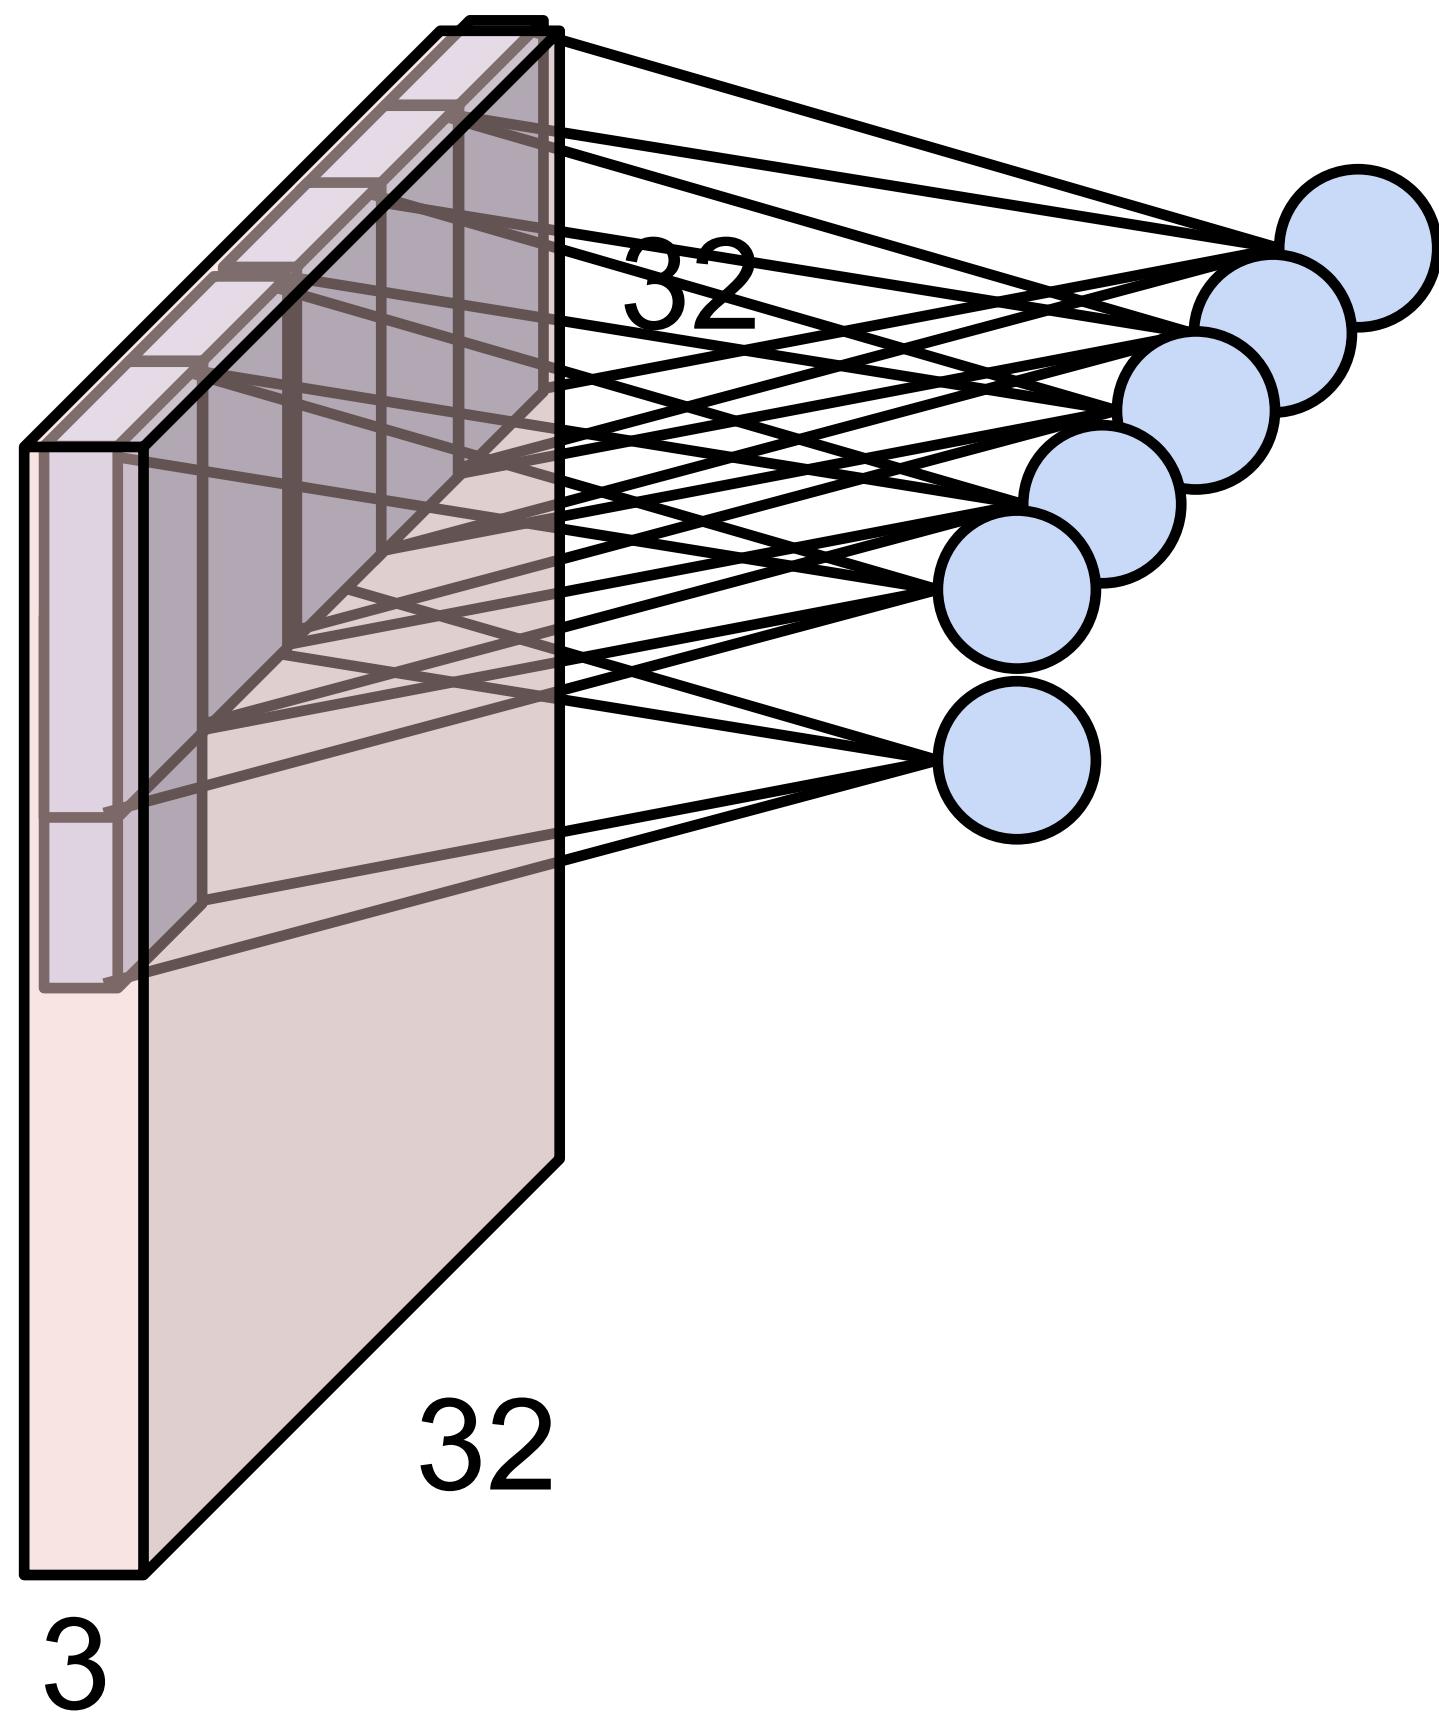
# Convolution Layer



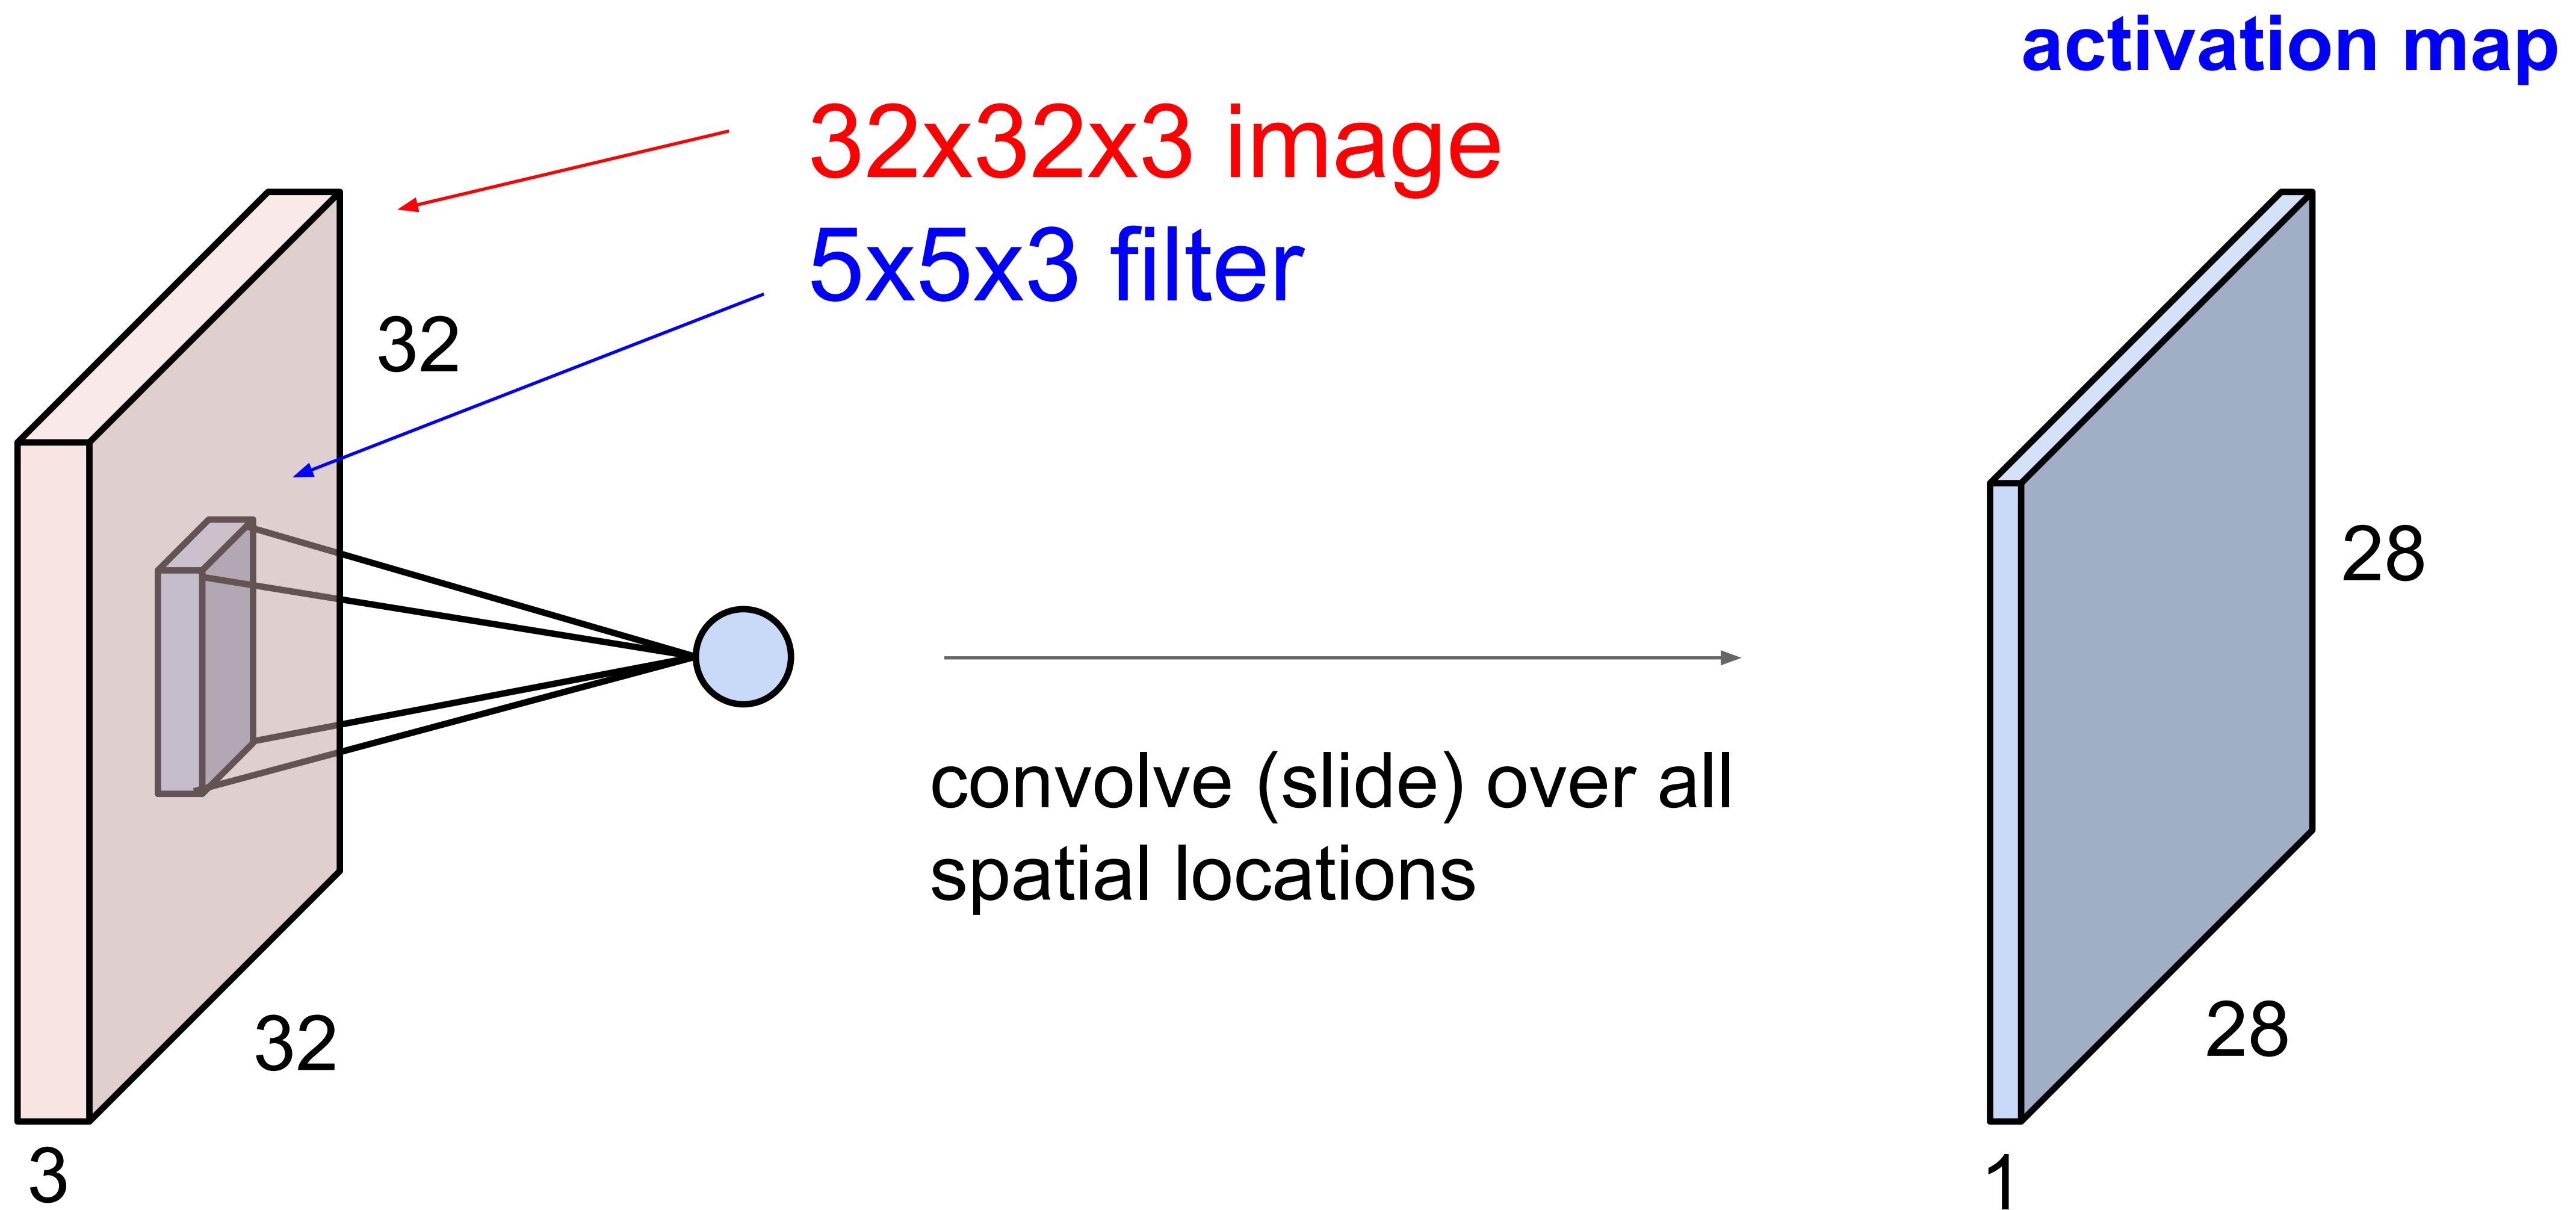
# Convolution Layer



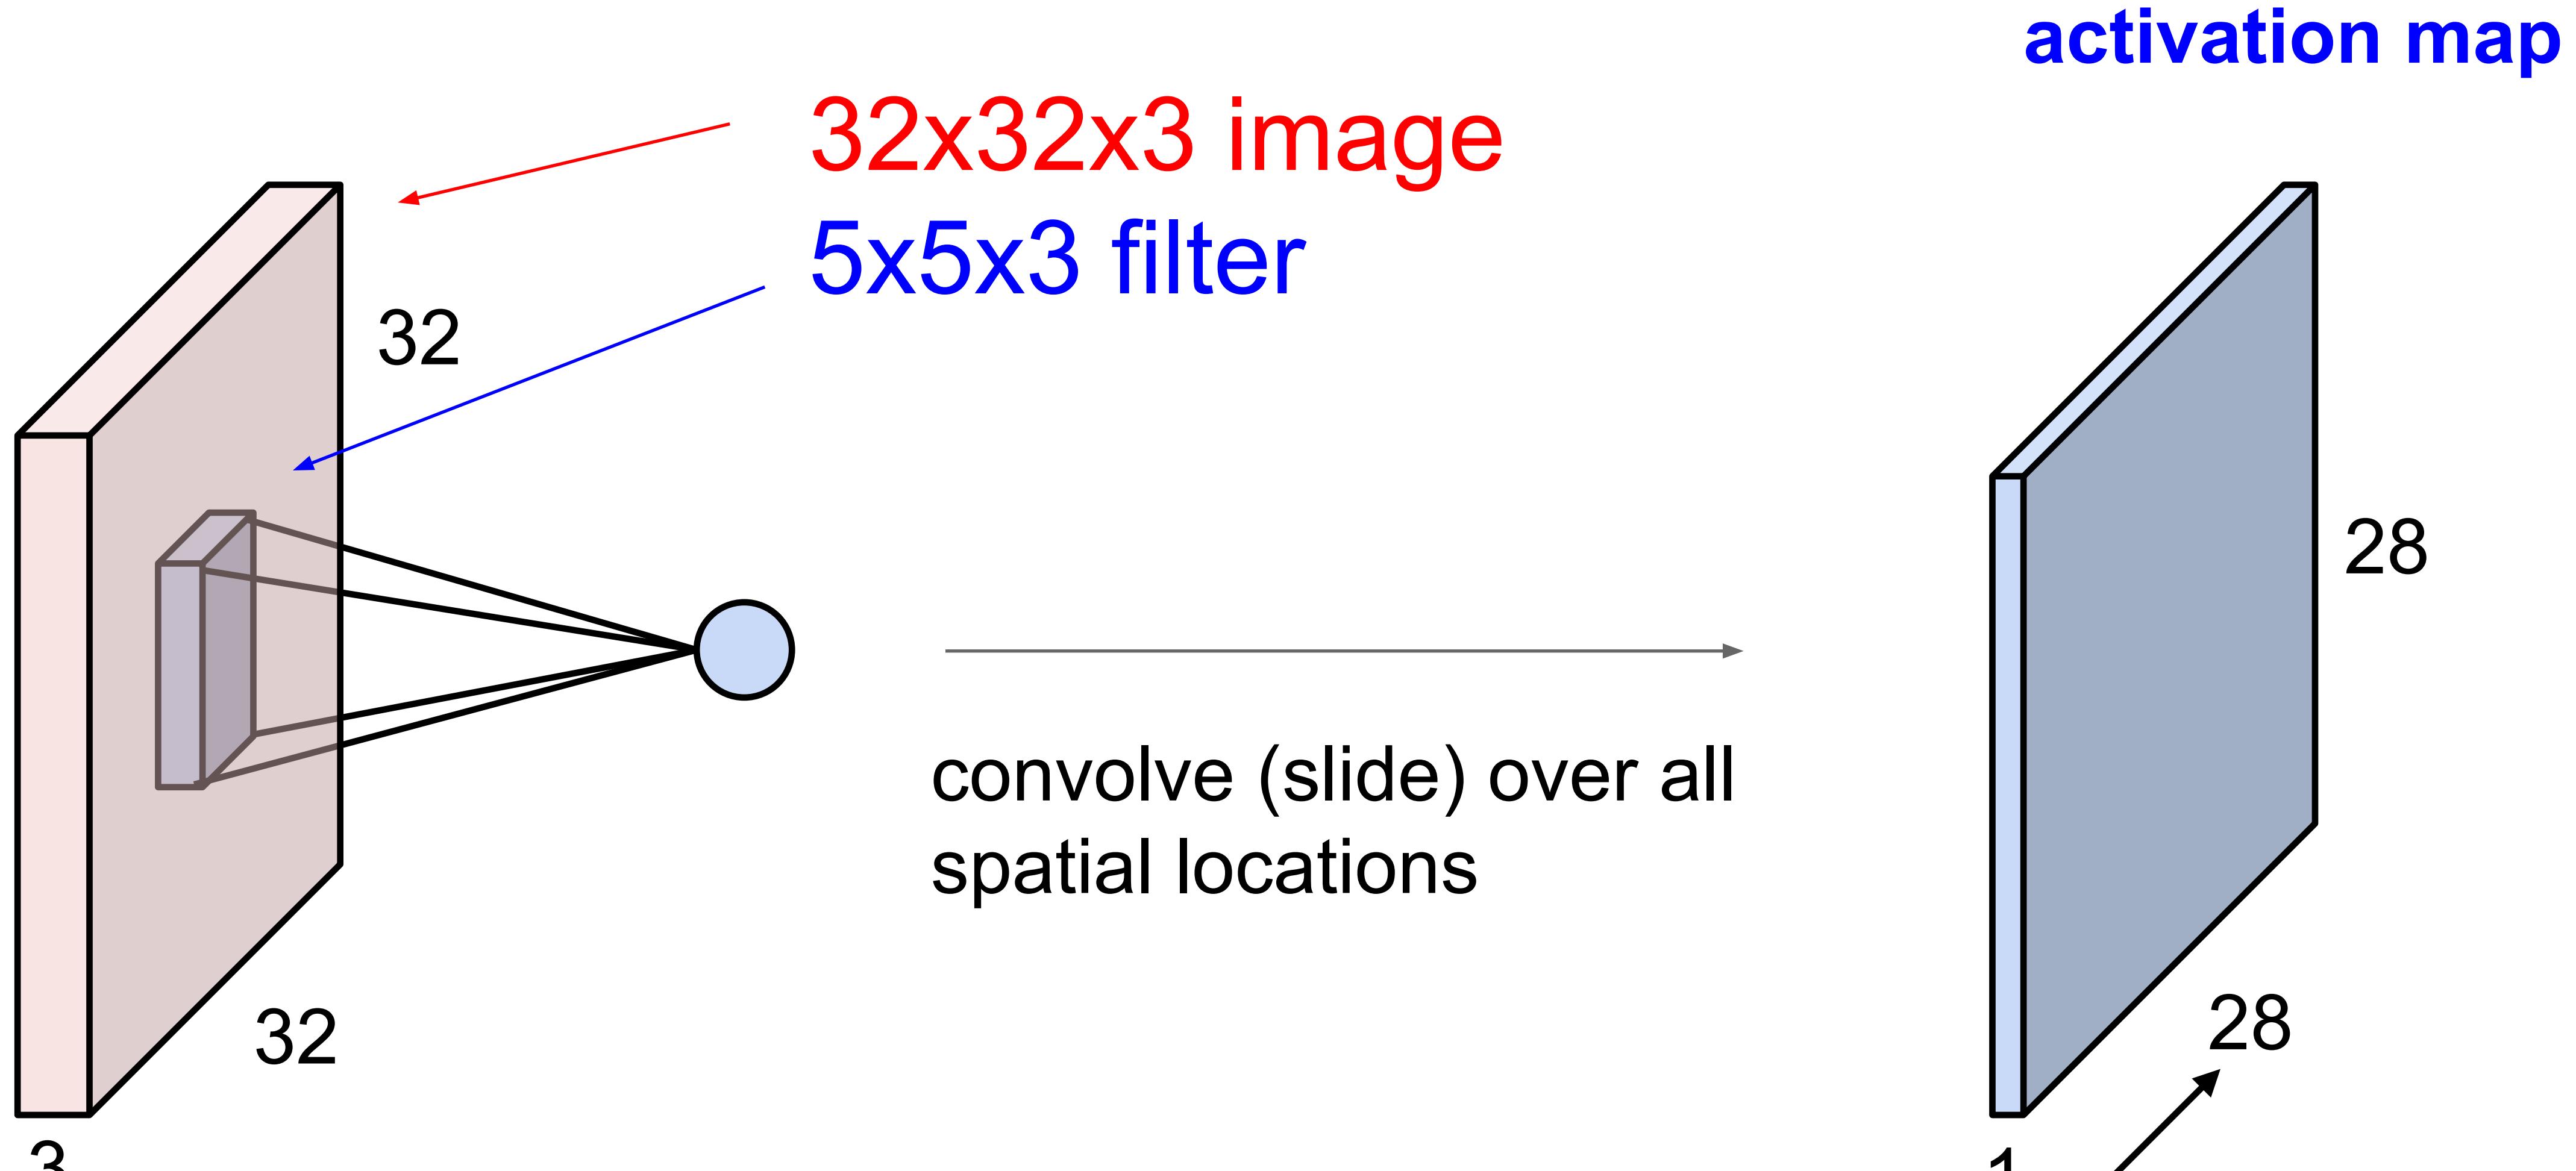
# Convolution Layer



# Convolution Layer



# Convolution Layer



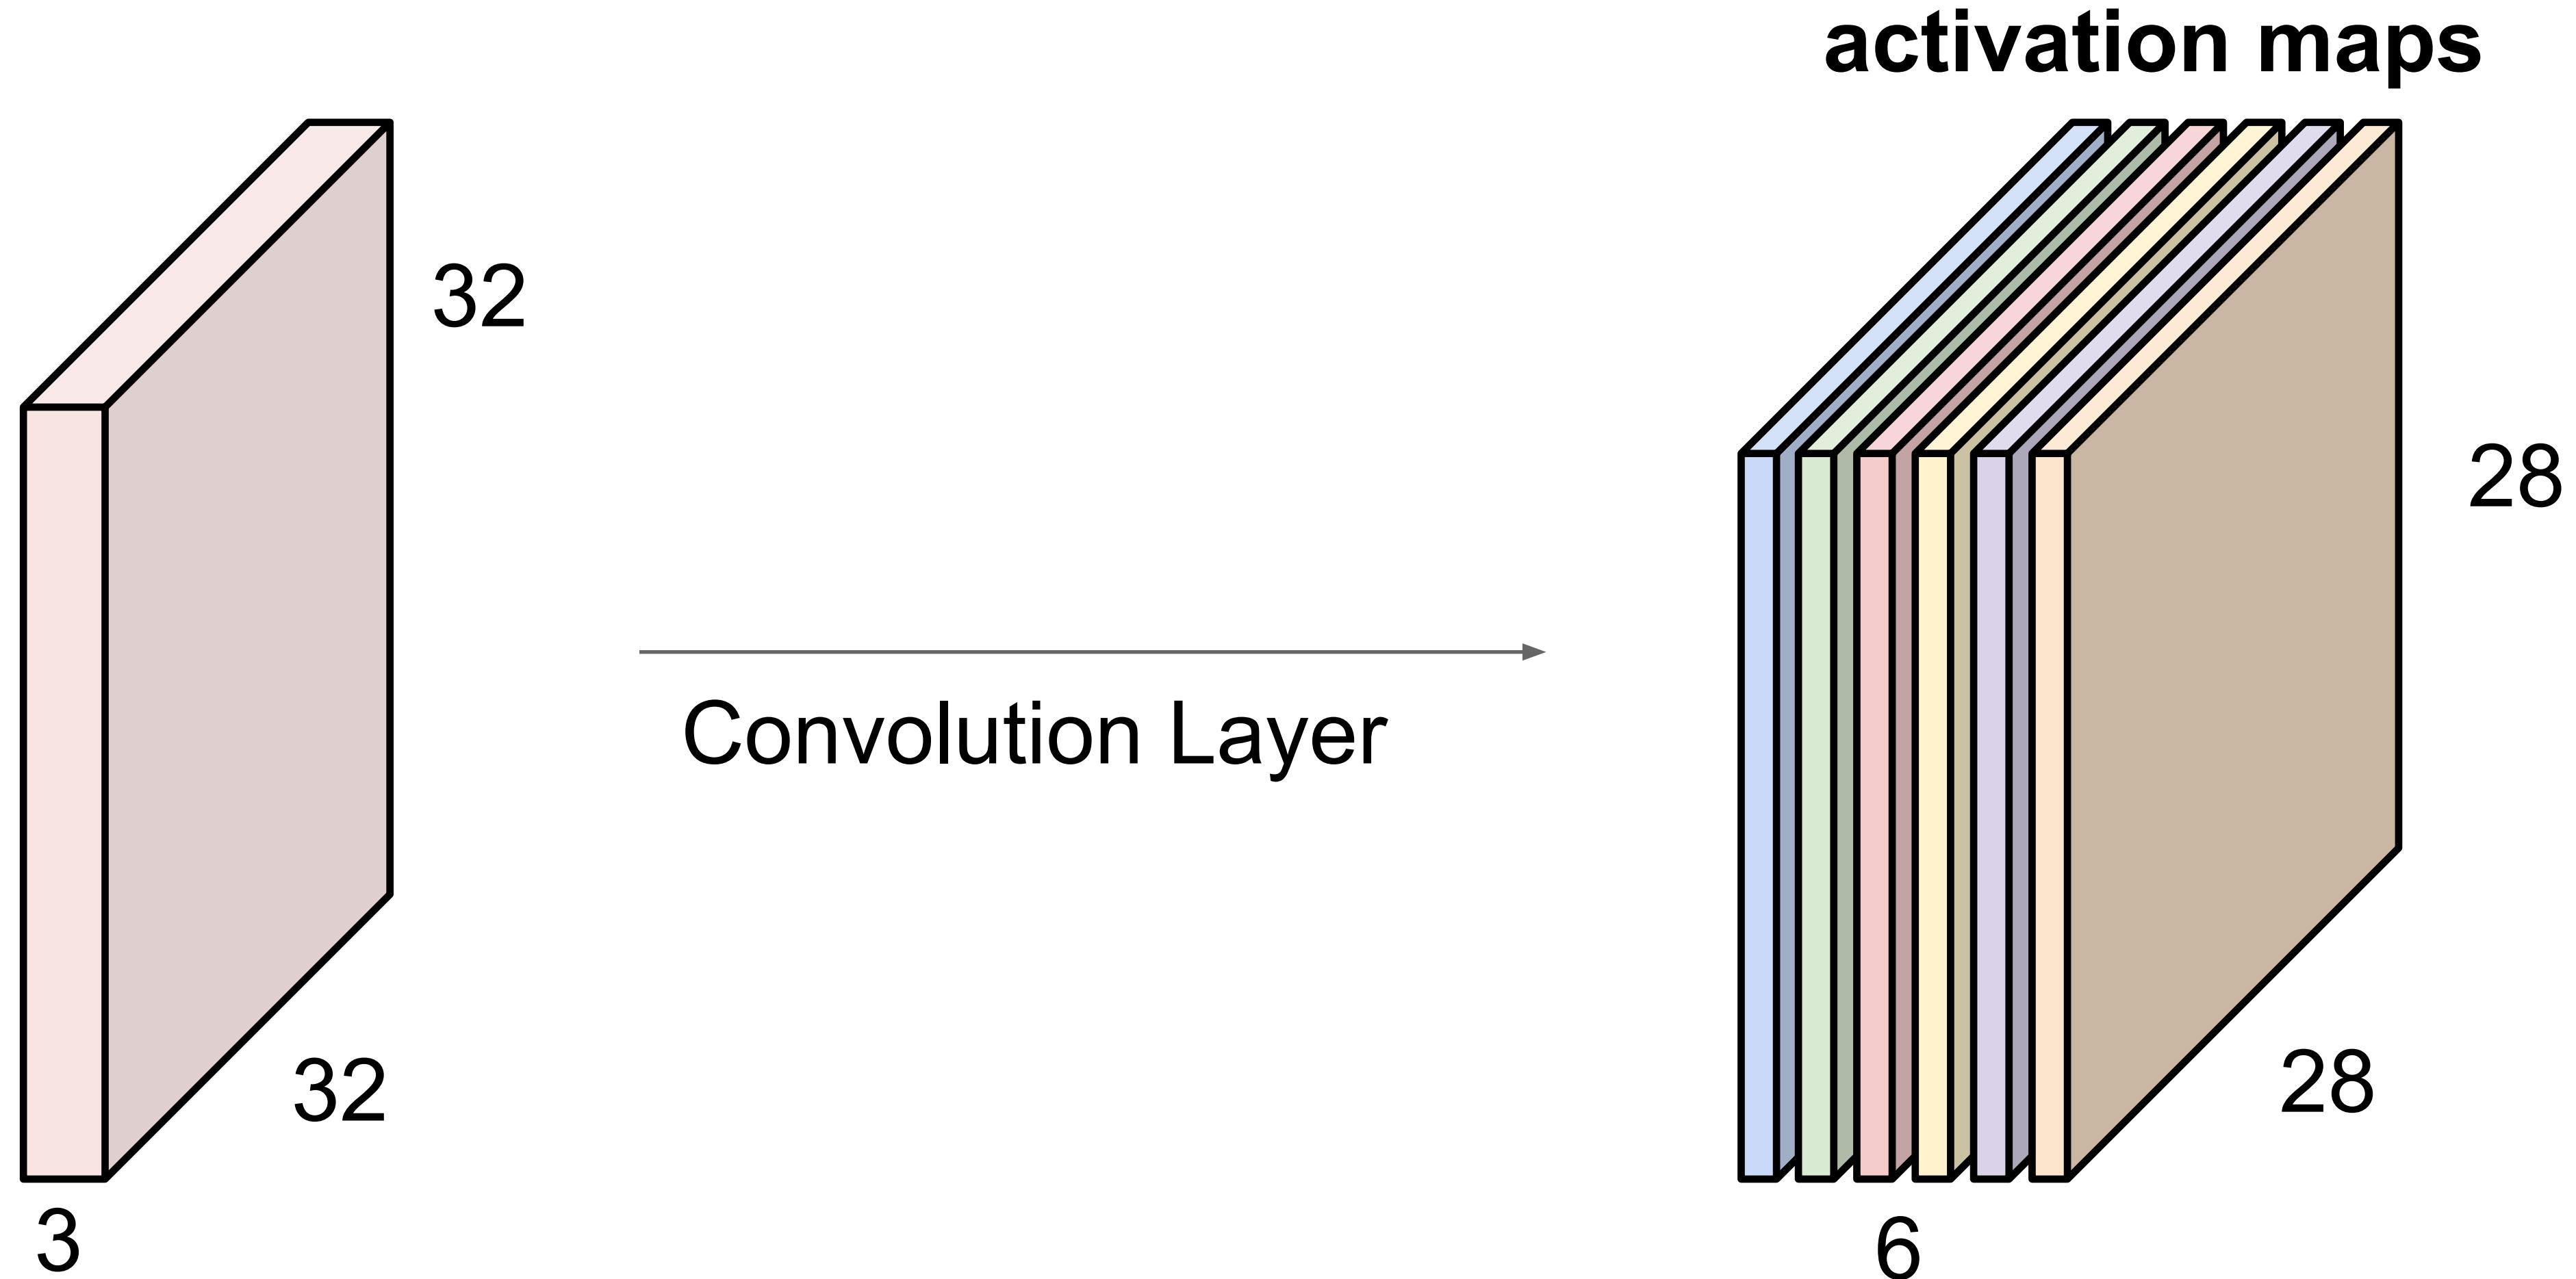
Activation map dimension =  
Input image dimension - Filter dimension + 1

# Convolution Layer

consider a second, green filter

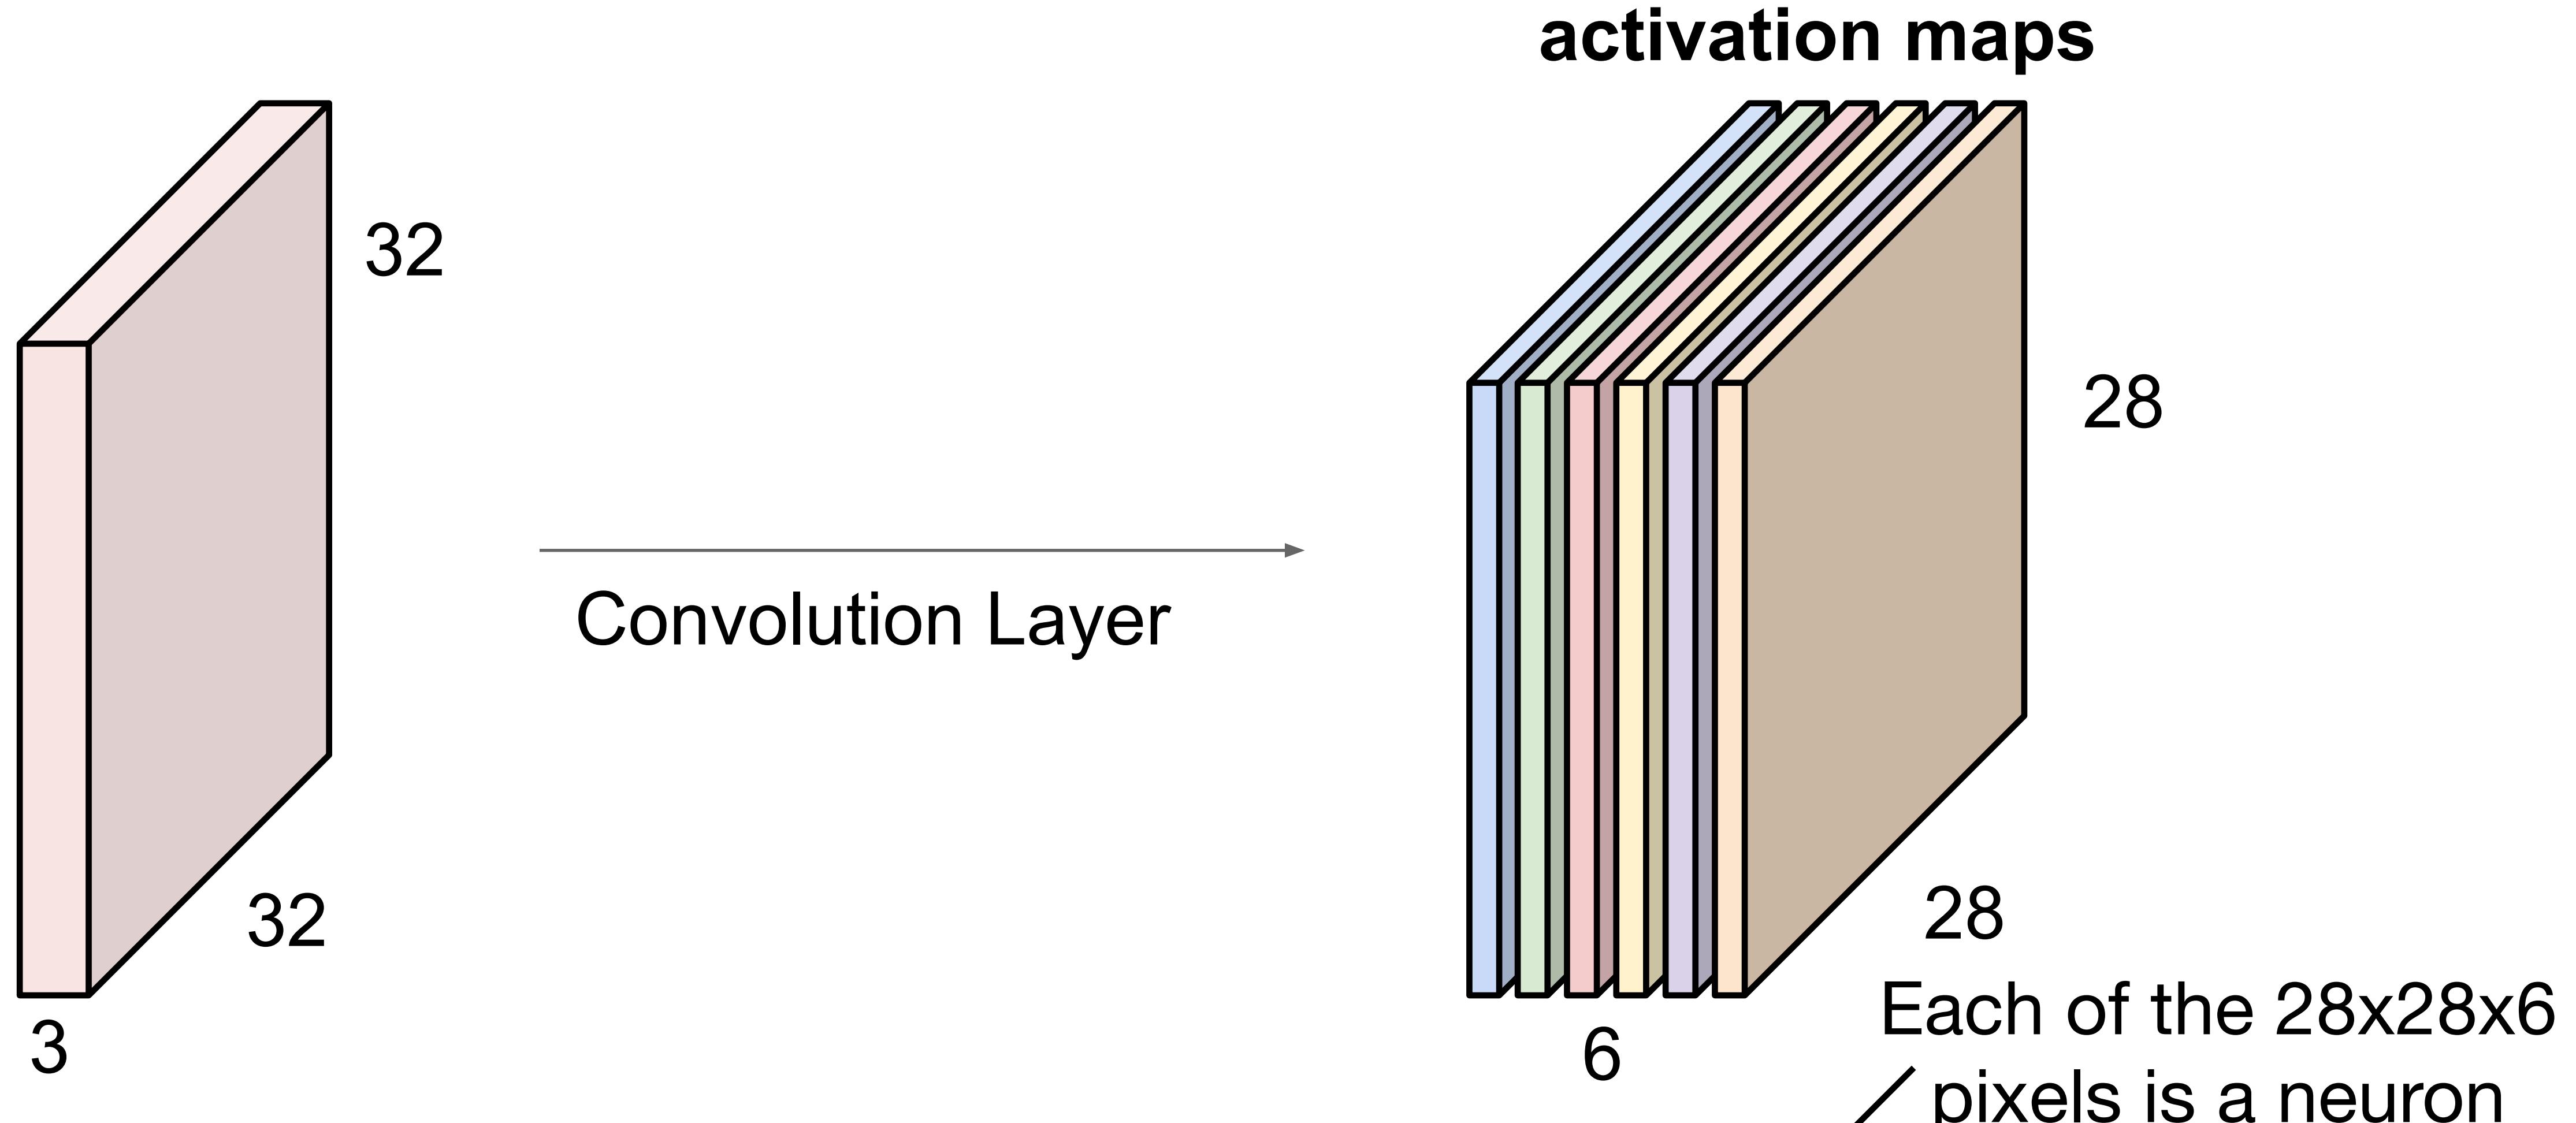


For example, if we had 6 5x5 filters, we'll get 6 separate activation maps:



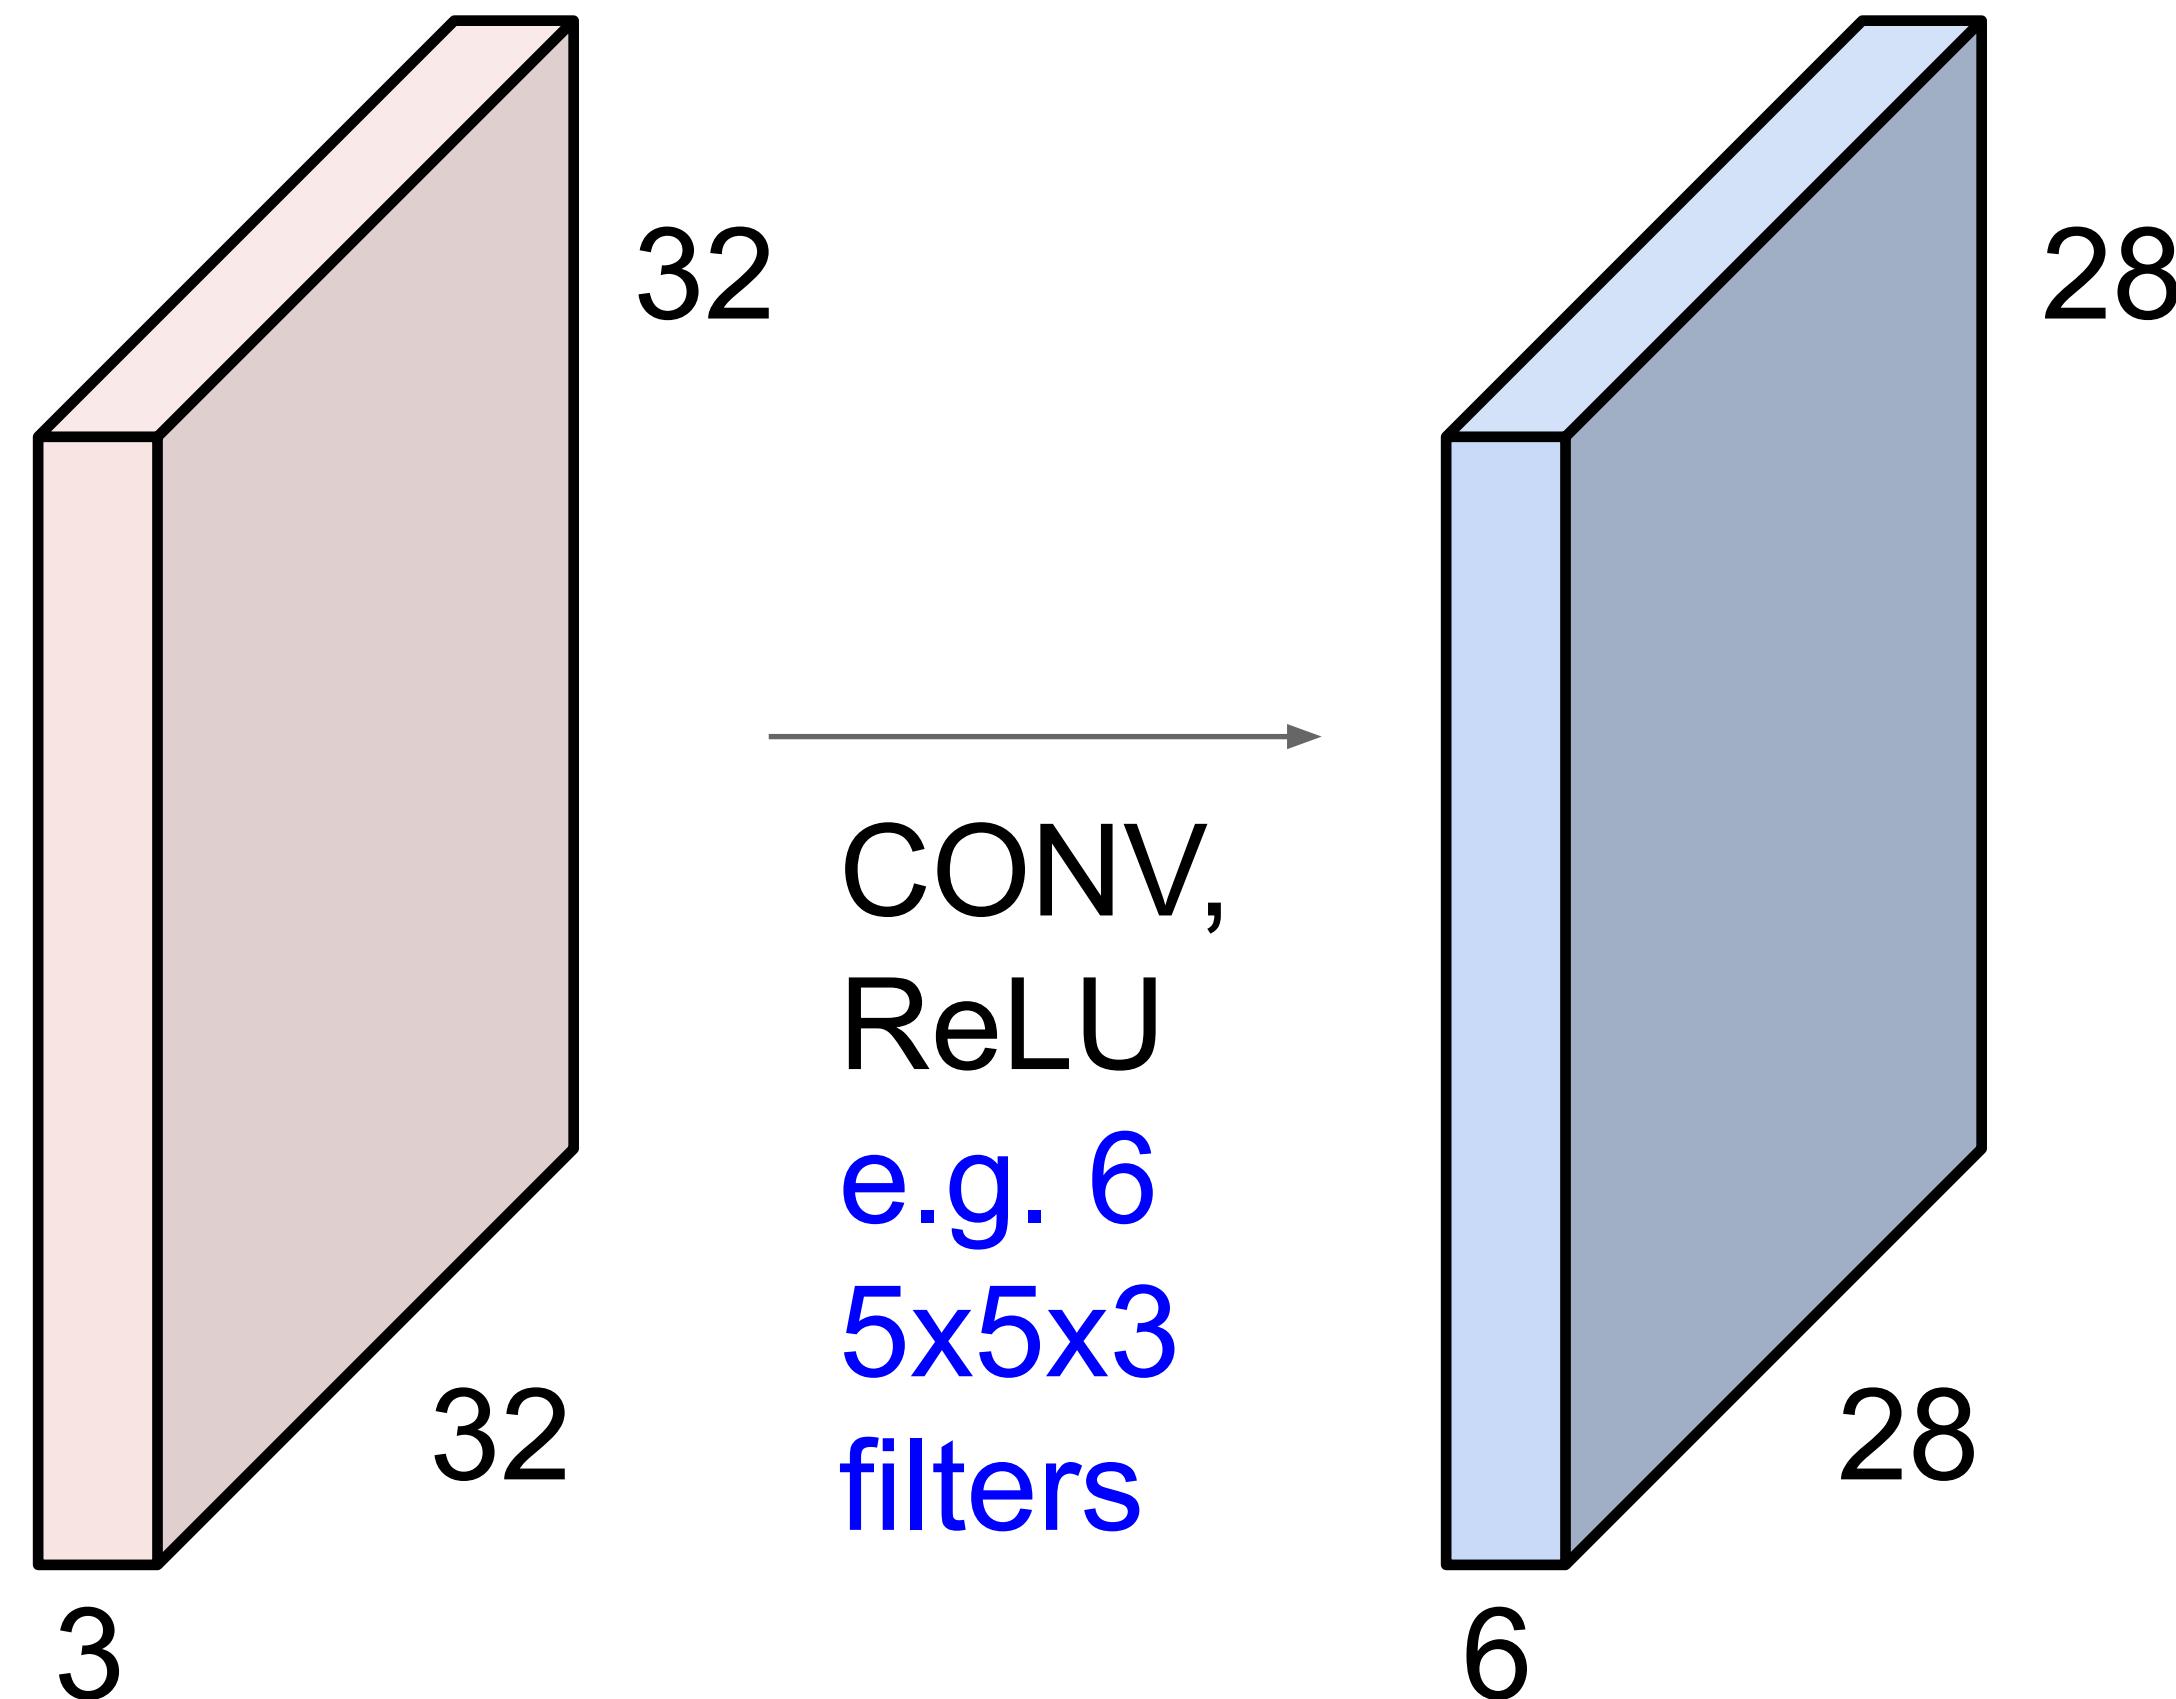
We stack these up to get a “new image” of size  $28 \times 28 \times 6$ !

For example, if we had 6 5x5 filters, we'll get 6 separate activation maps:

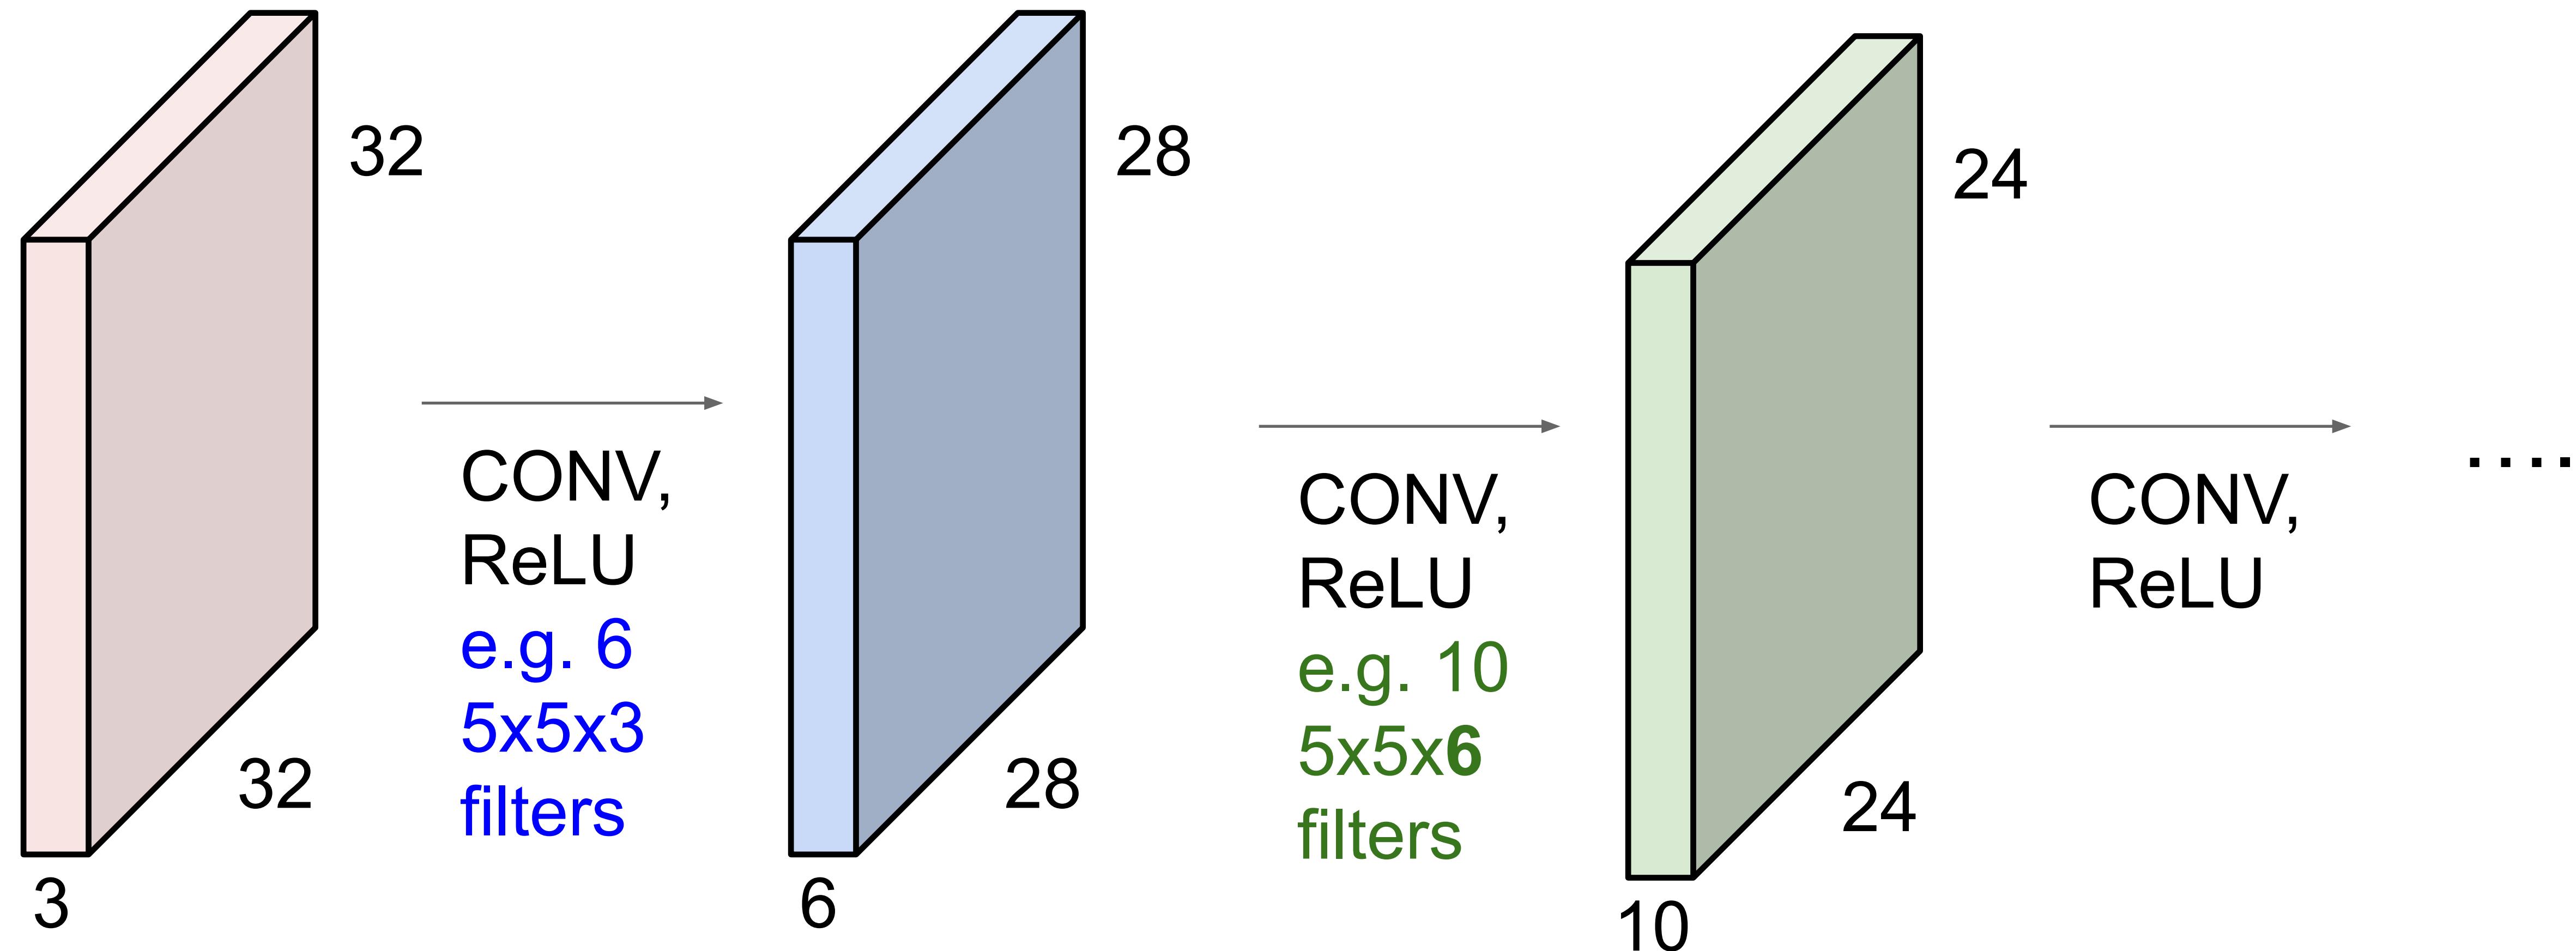


We stack these up to get a “new image” of size  $28 \times 28 \times 6$ !

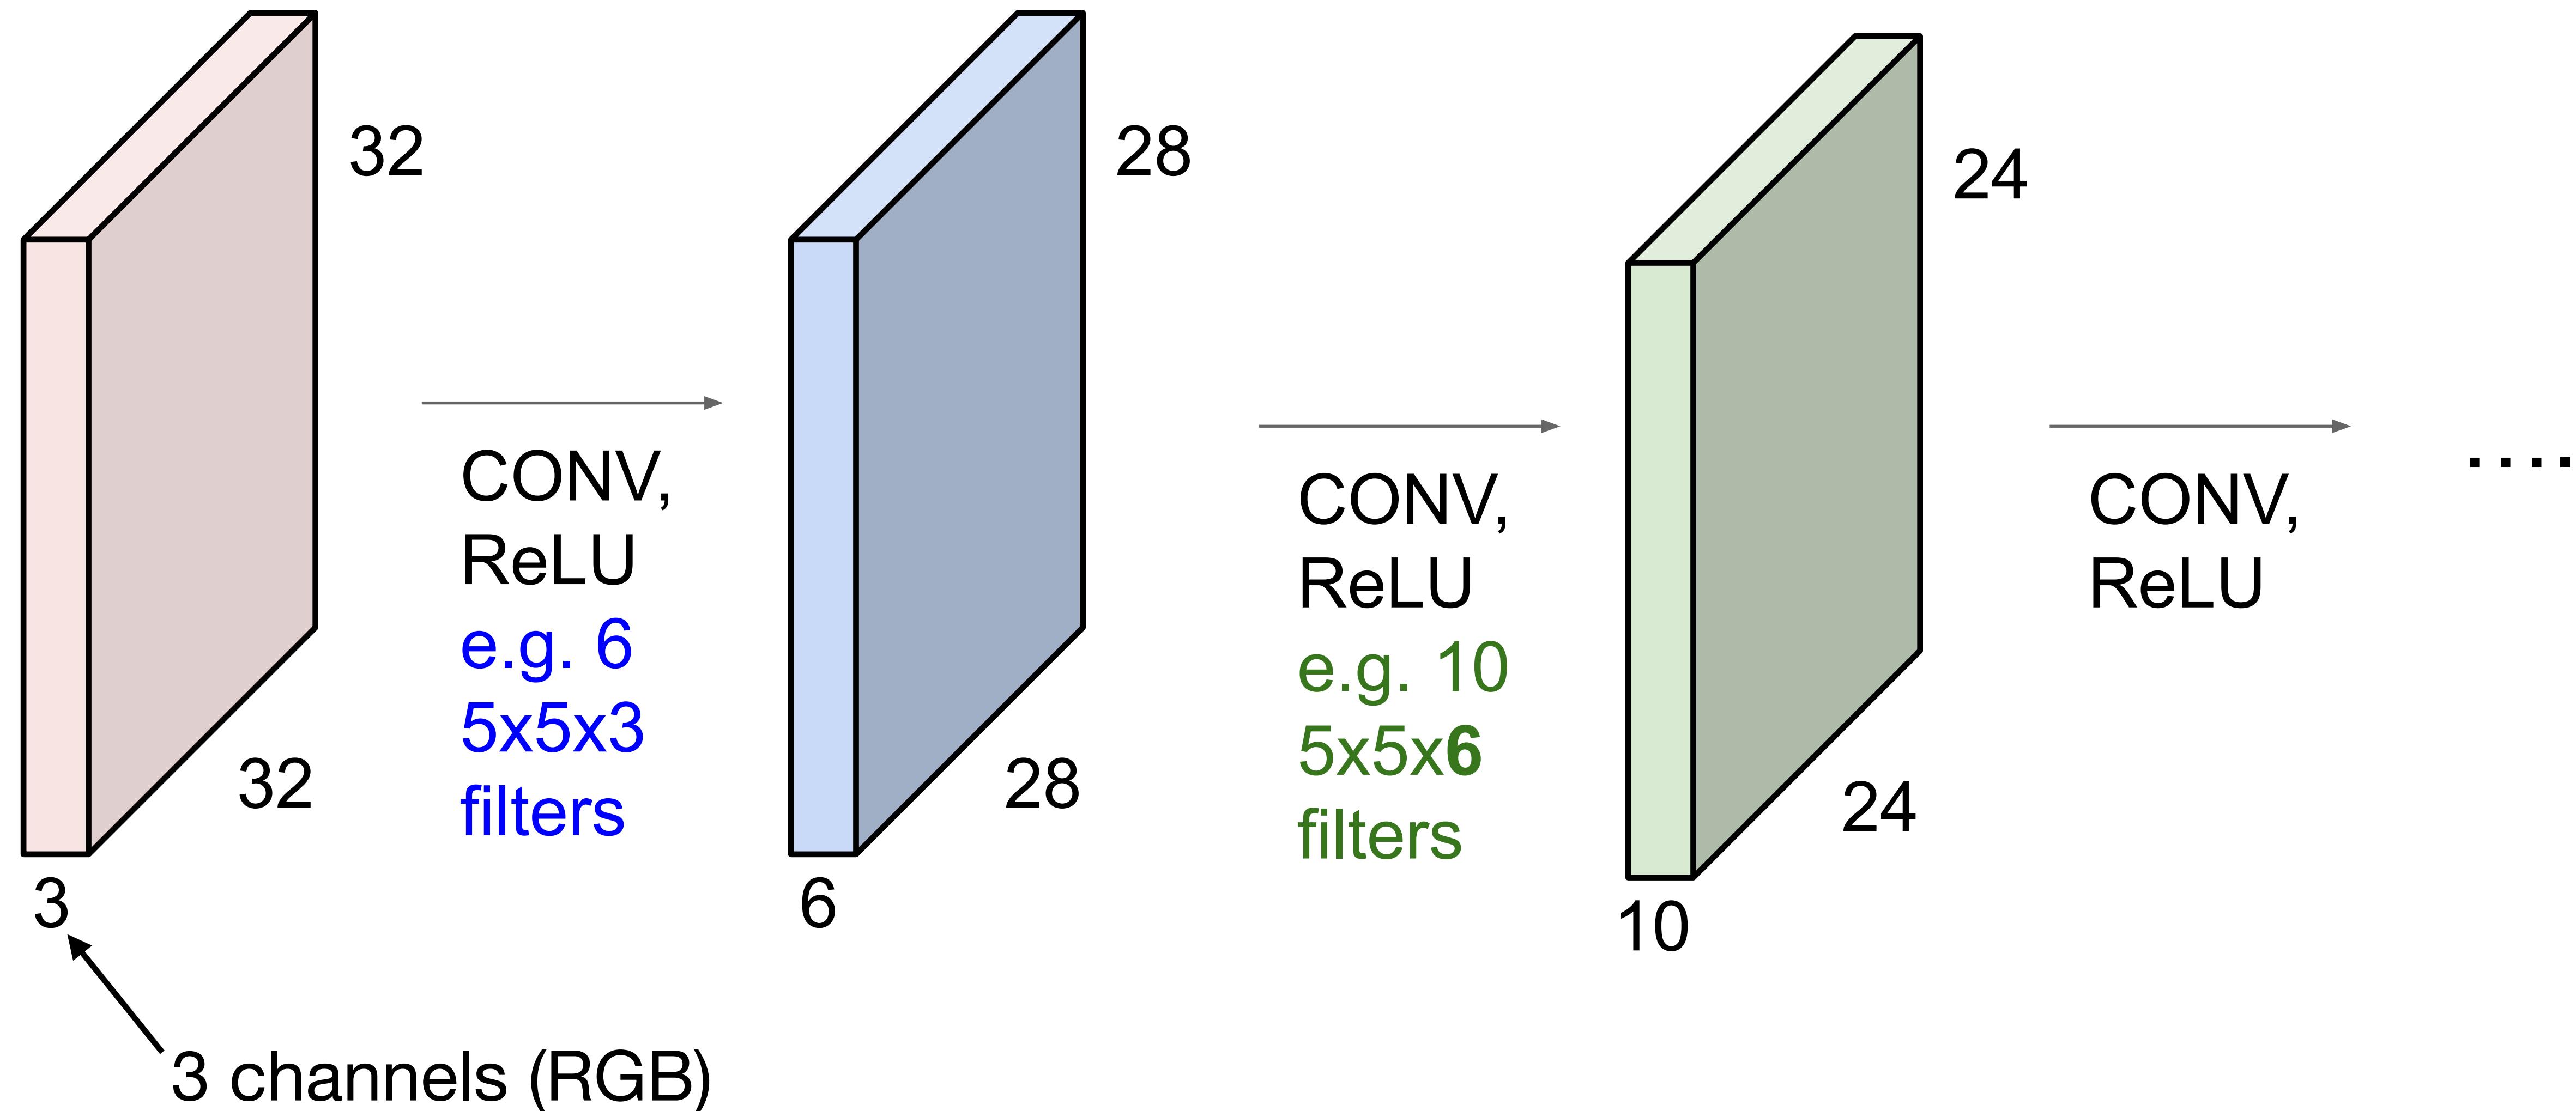
**Preview:** ConvNet is a sequence of Convolution Layers, interspersed with activation functions



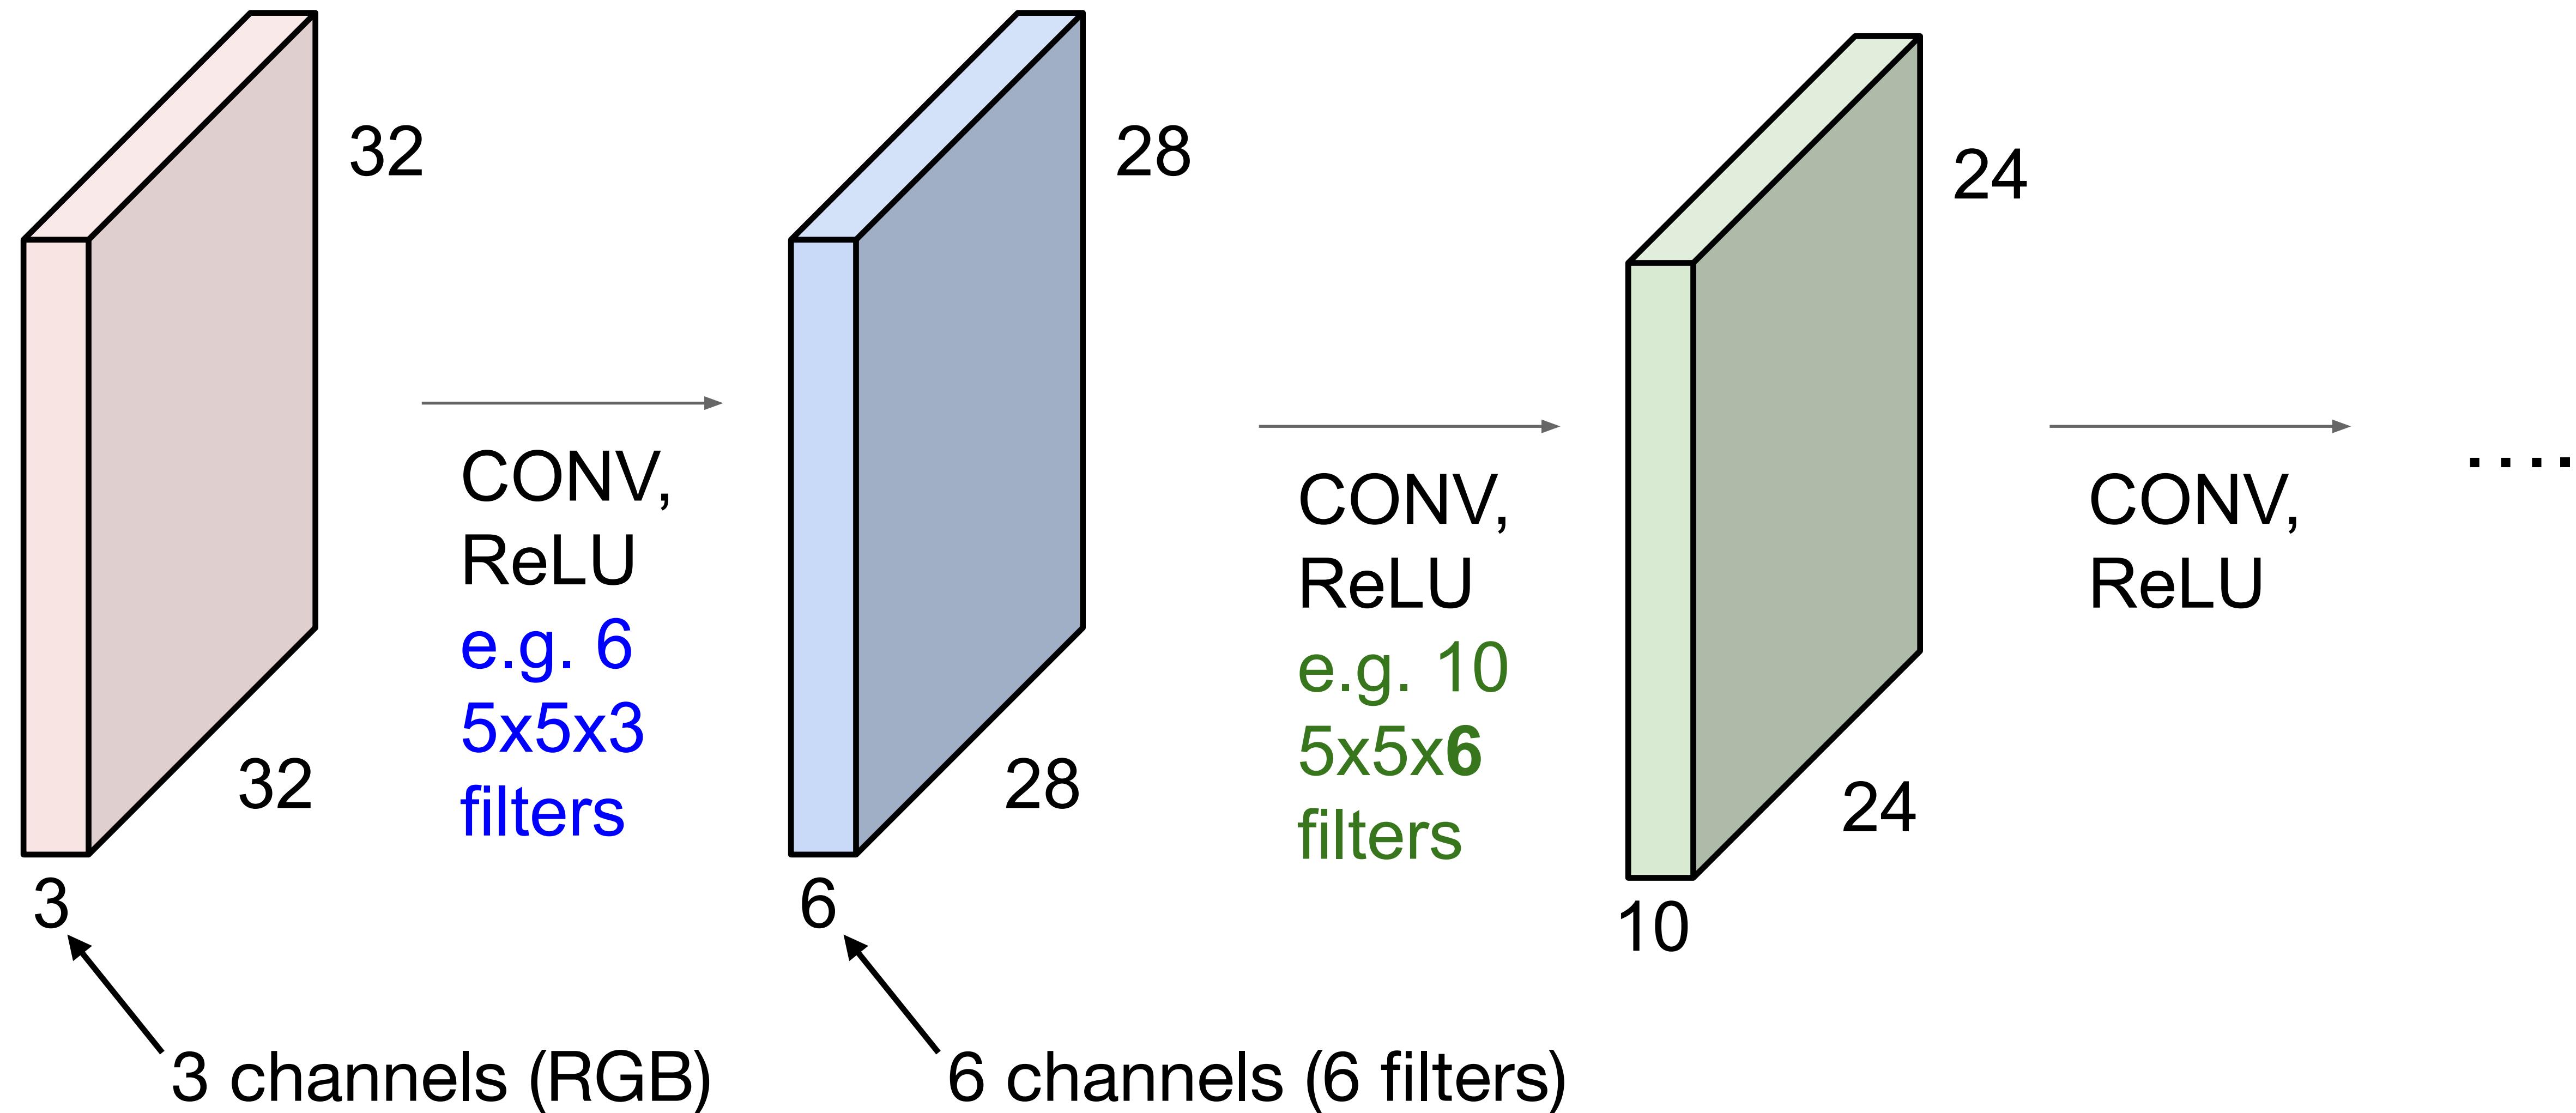
**Preview:** ConvNet is a sequence of Convolution Layers, interspersed with activation functions



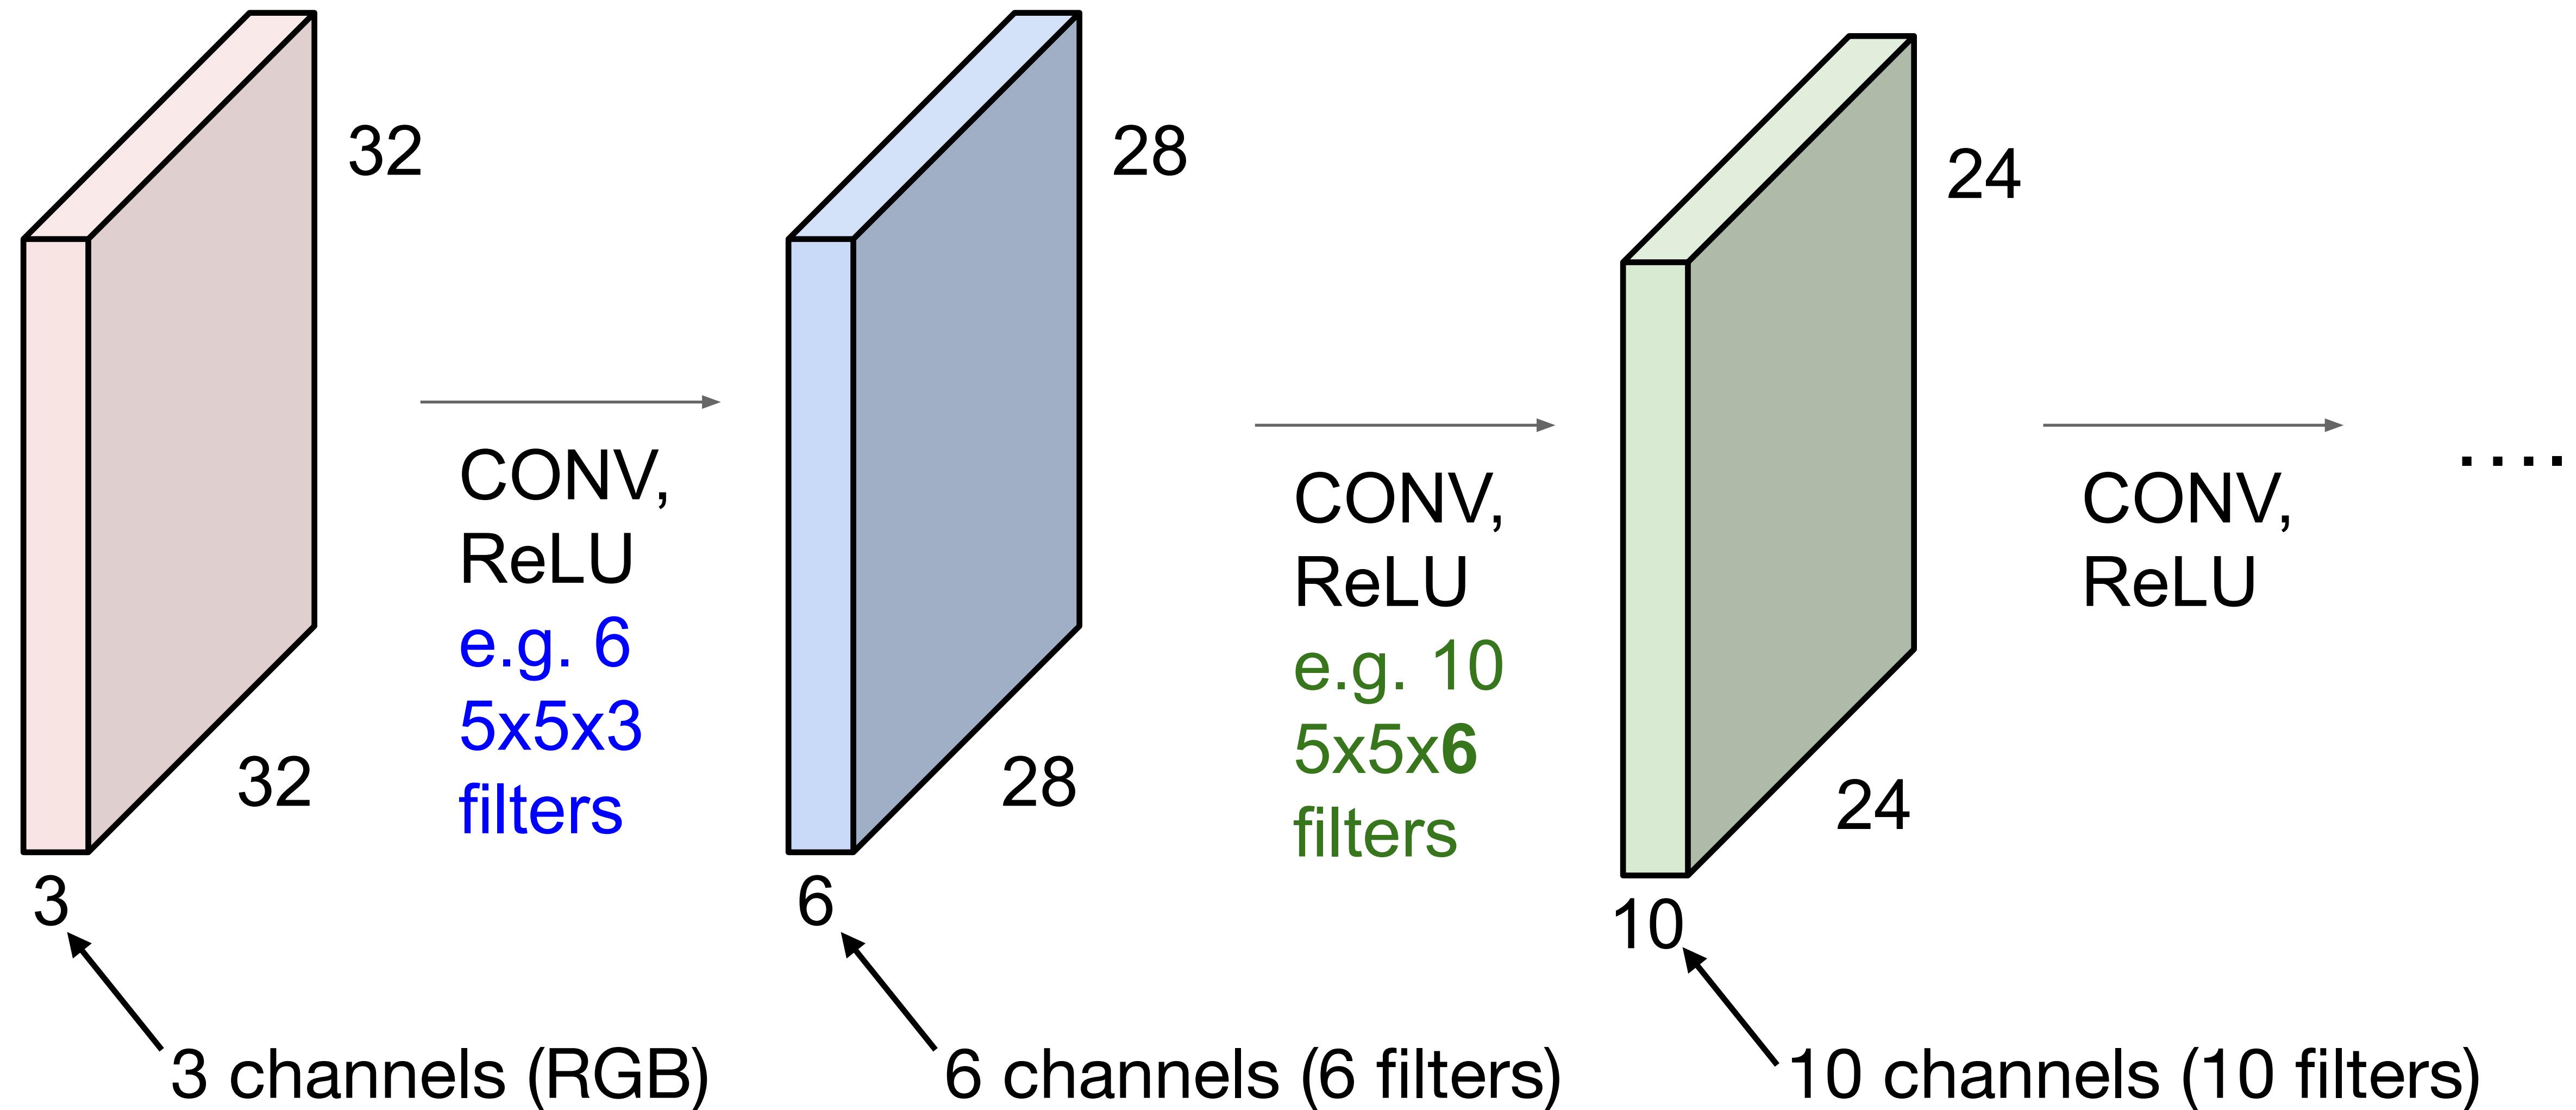
**Preview:** ConvNet is a sequence of Convolution Layers, interspersed with activation functions



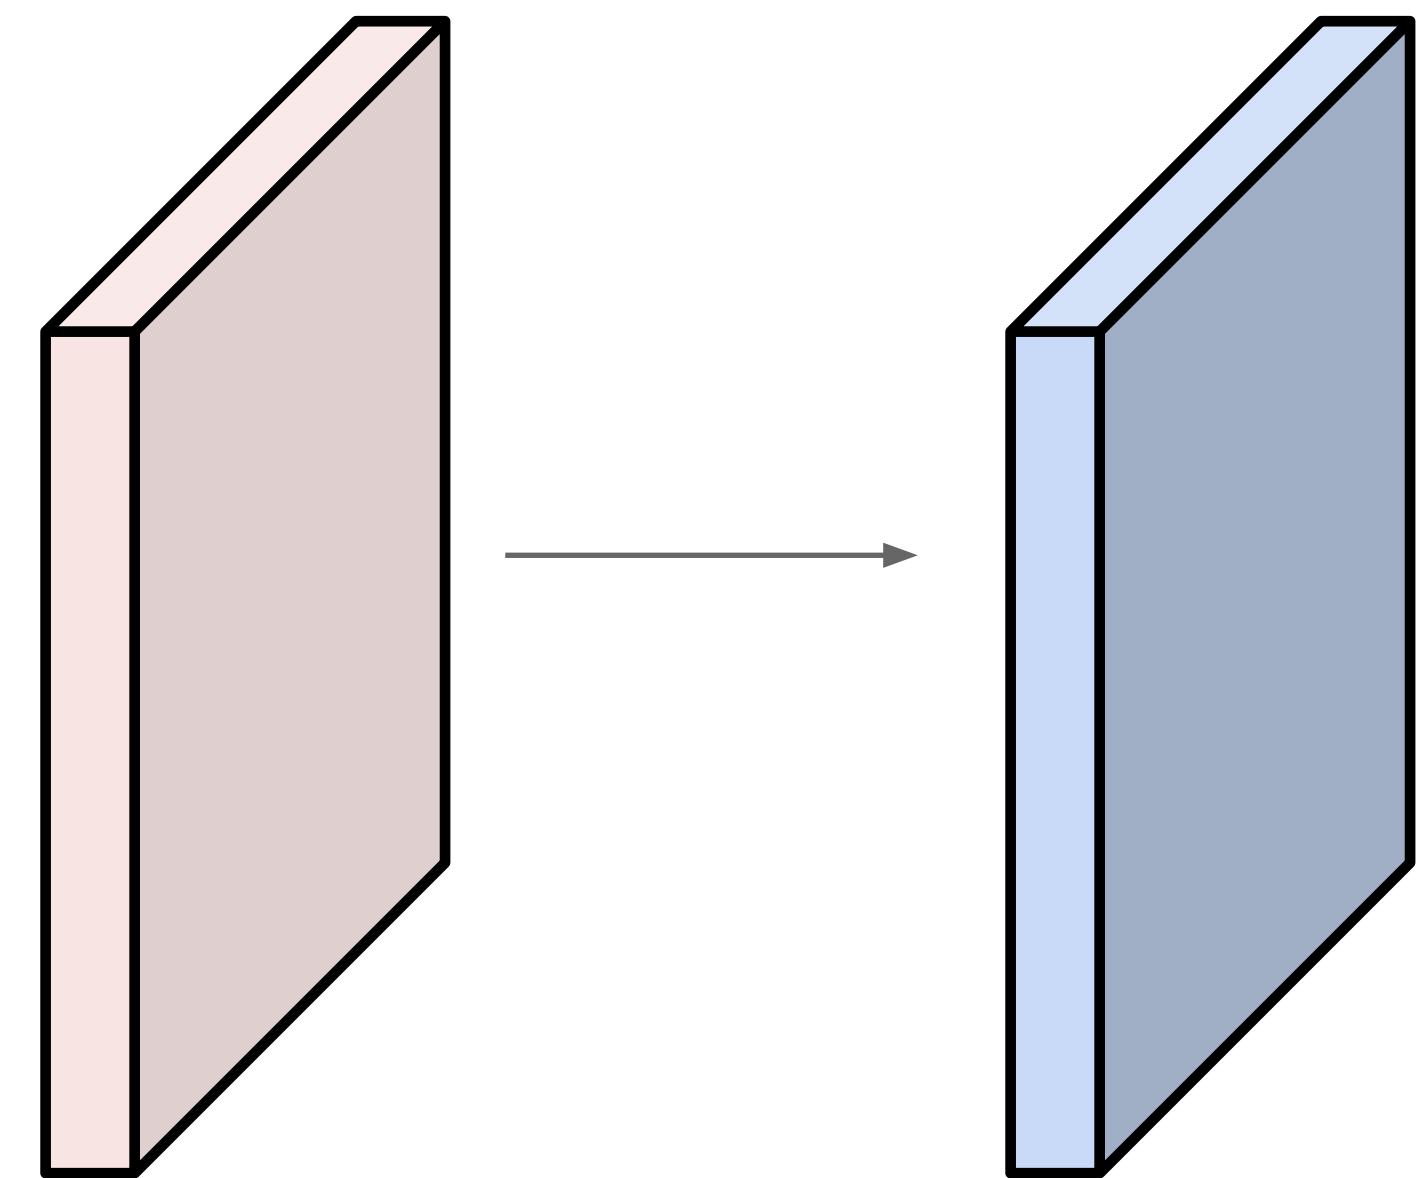
**Preview:** ConvNet is a sequence of Convolution Layers, interspersed with activation functions



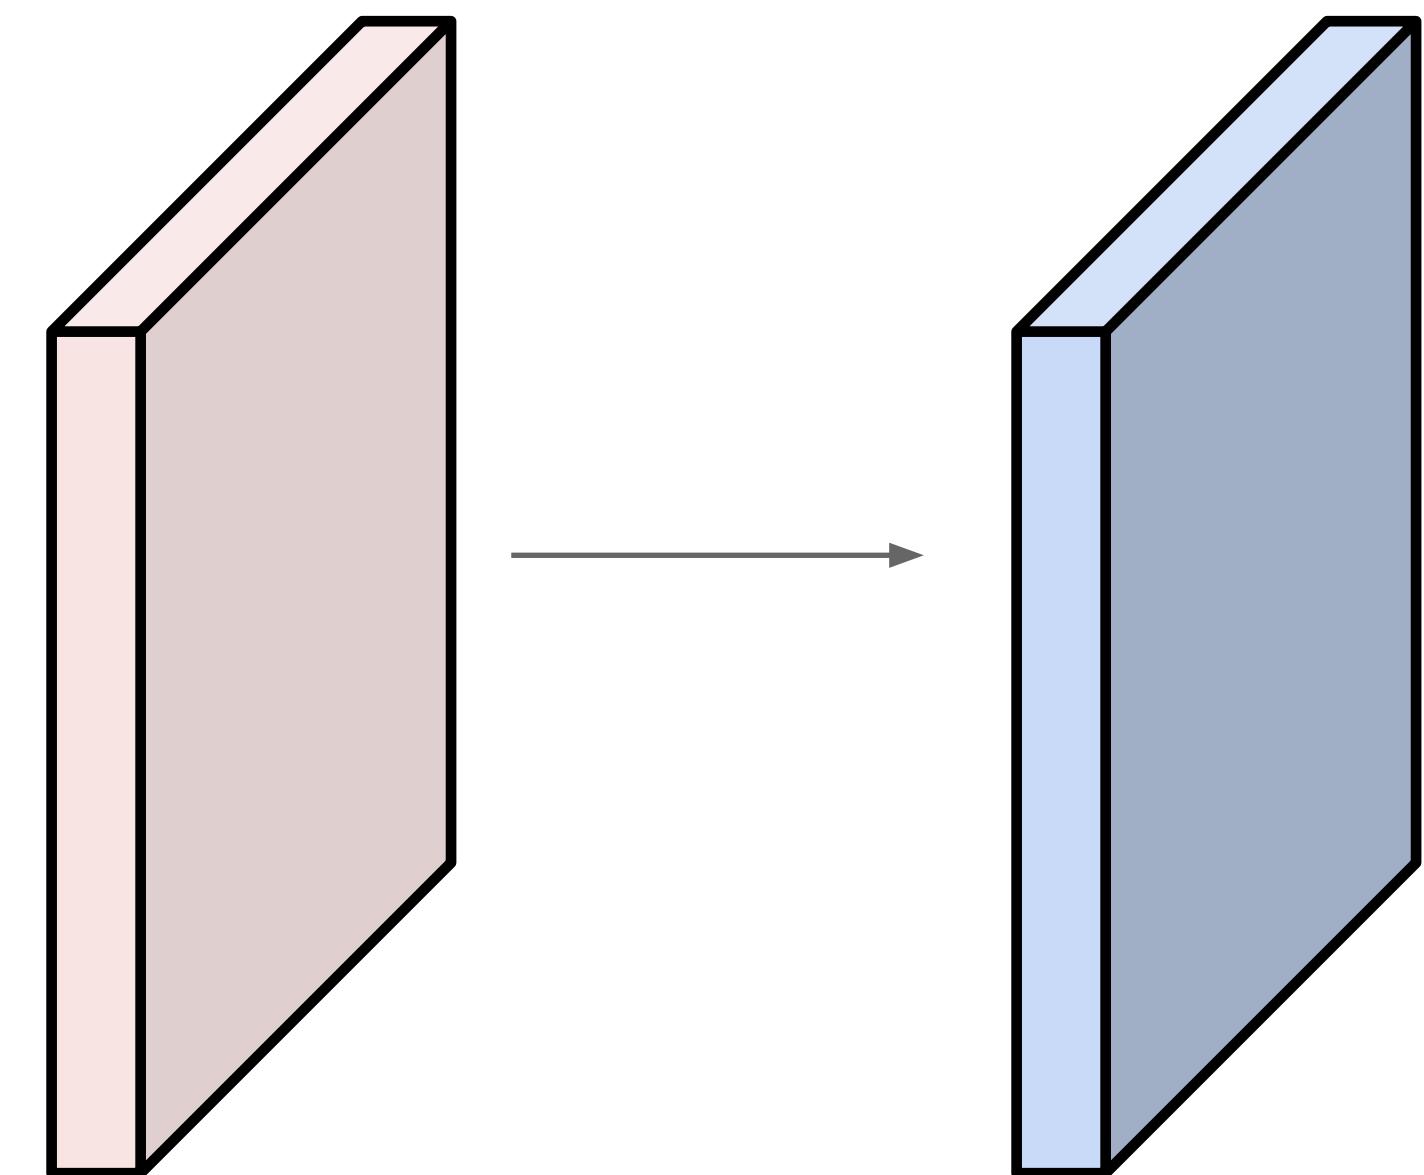
**Preview:** ConvNet is a sequence of Convolution Layers, interspersed with activation functions



**Input volume: 32x32x3  
10 5x5 filters**



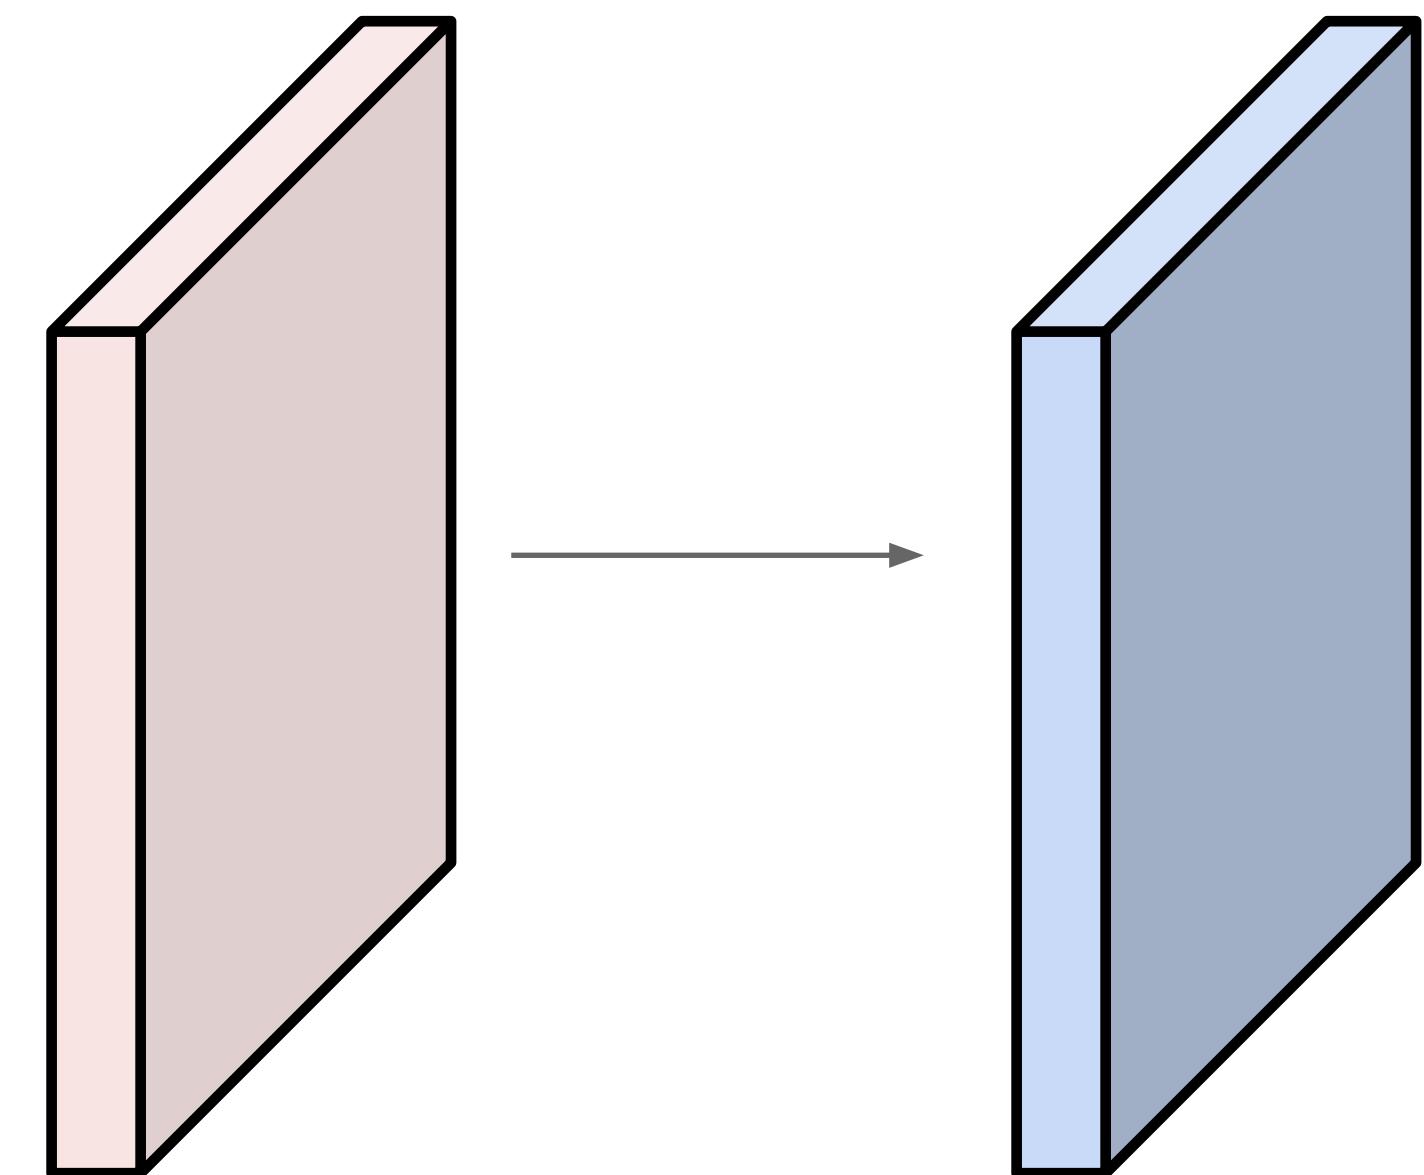
**Number of parameters in this layer?**



Input volume: **32x32x3**  
**10 5x5 filters**

Number of parameters in this layer?

each filter has  $5*5*3 + 1 = 76$  params (+1 for bias)  
 $\Rightarrow 76*10 = 760$



Input volume: **32x32x3**  
**10 5x5 filters**

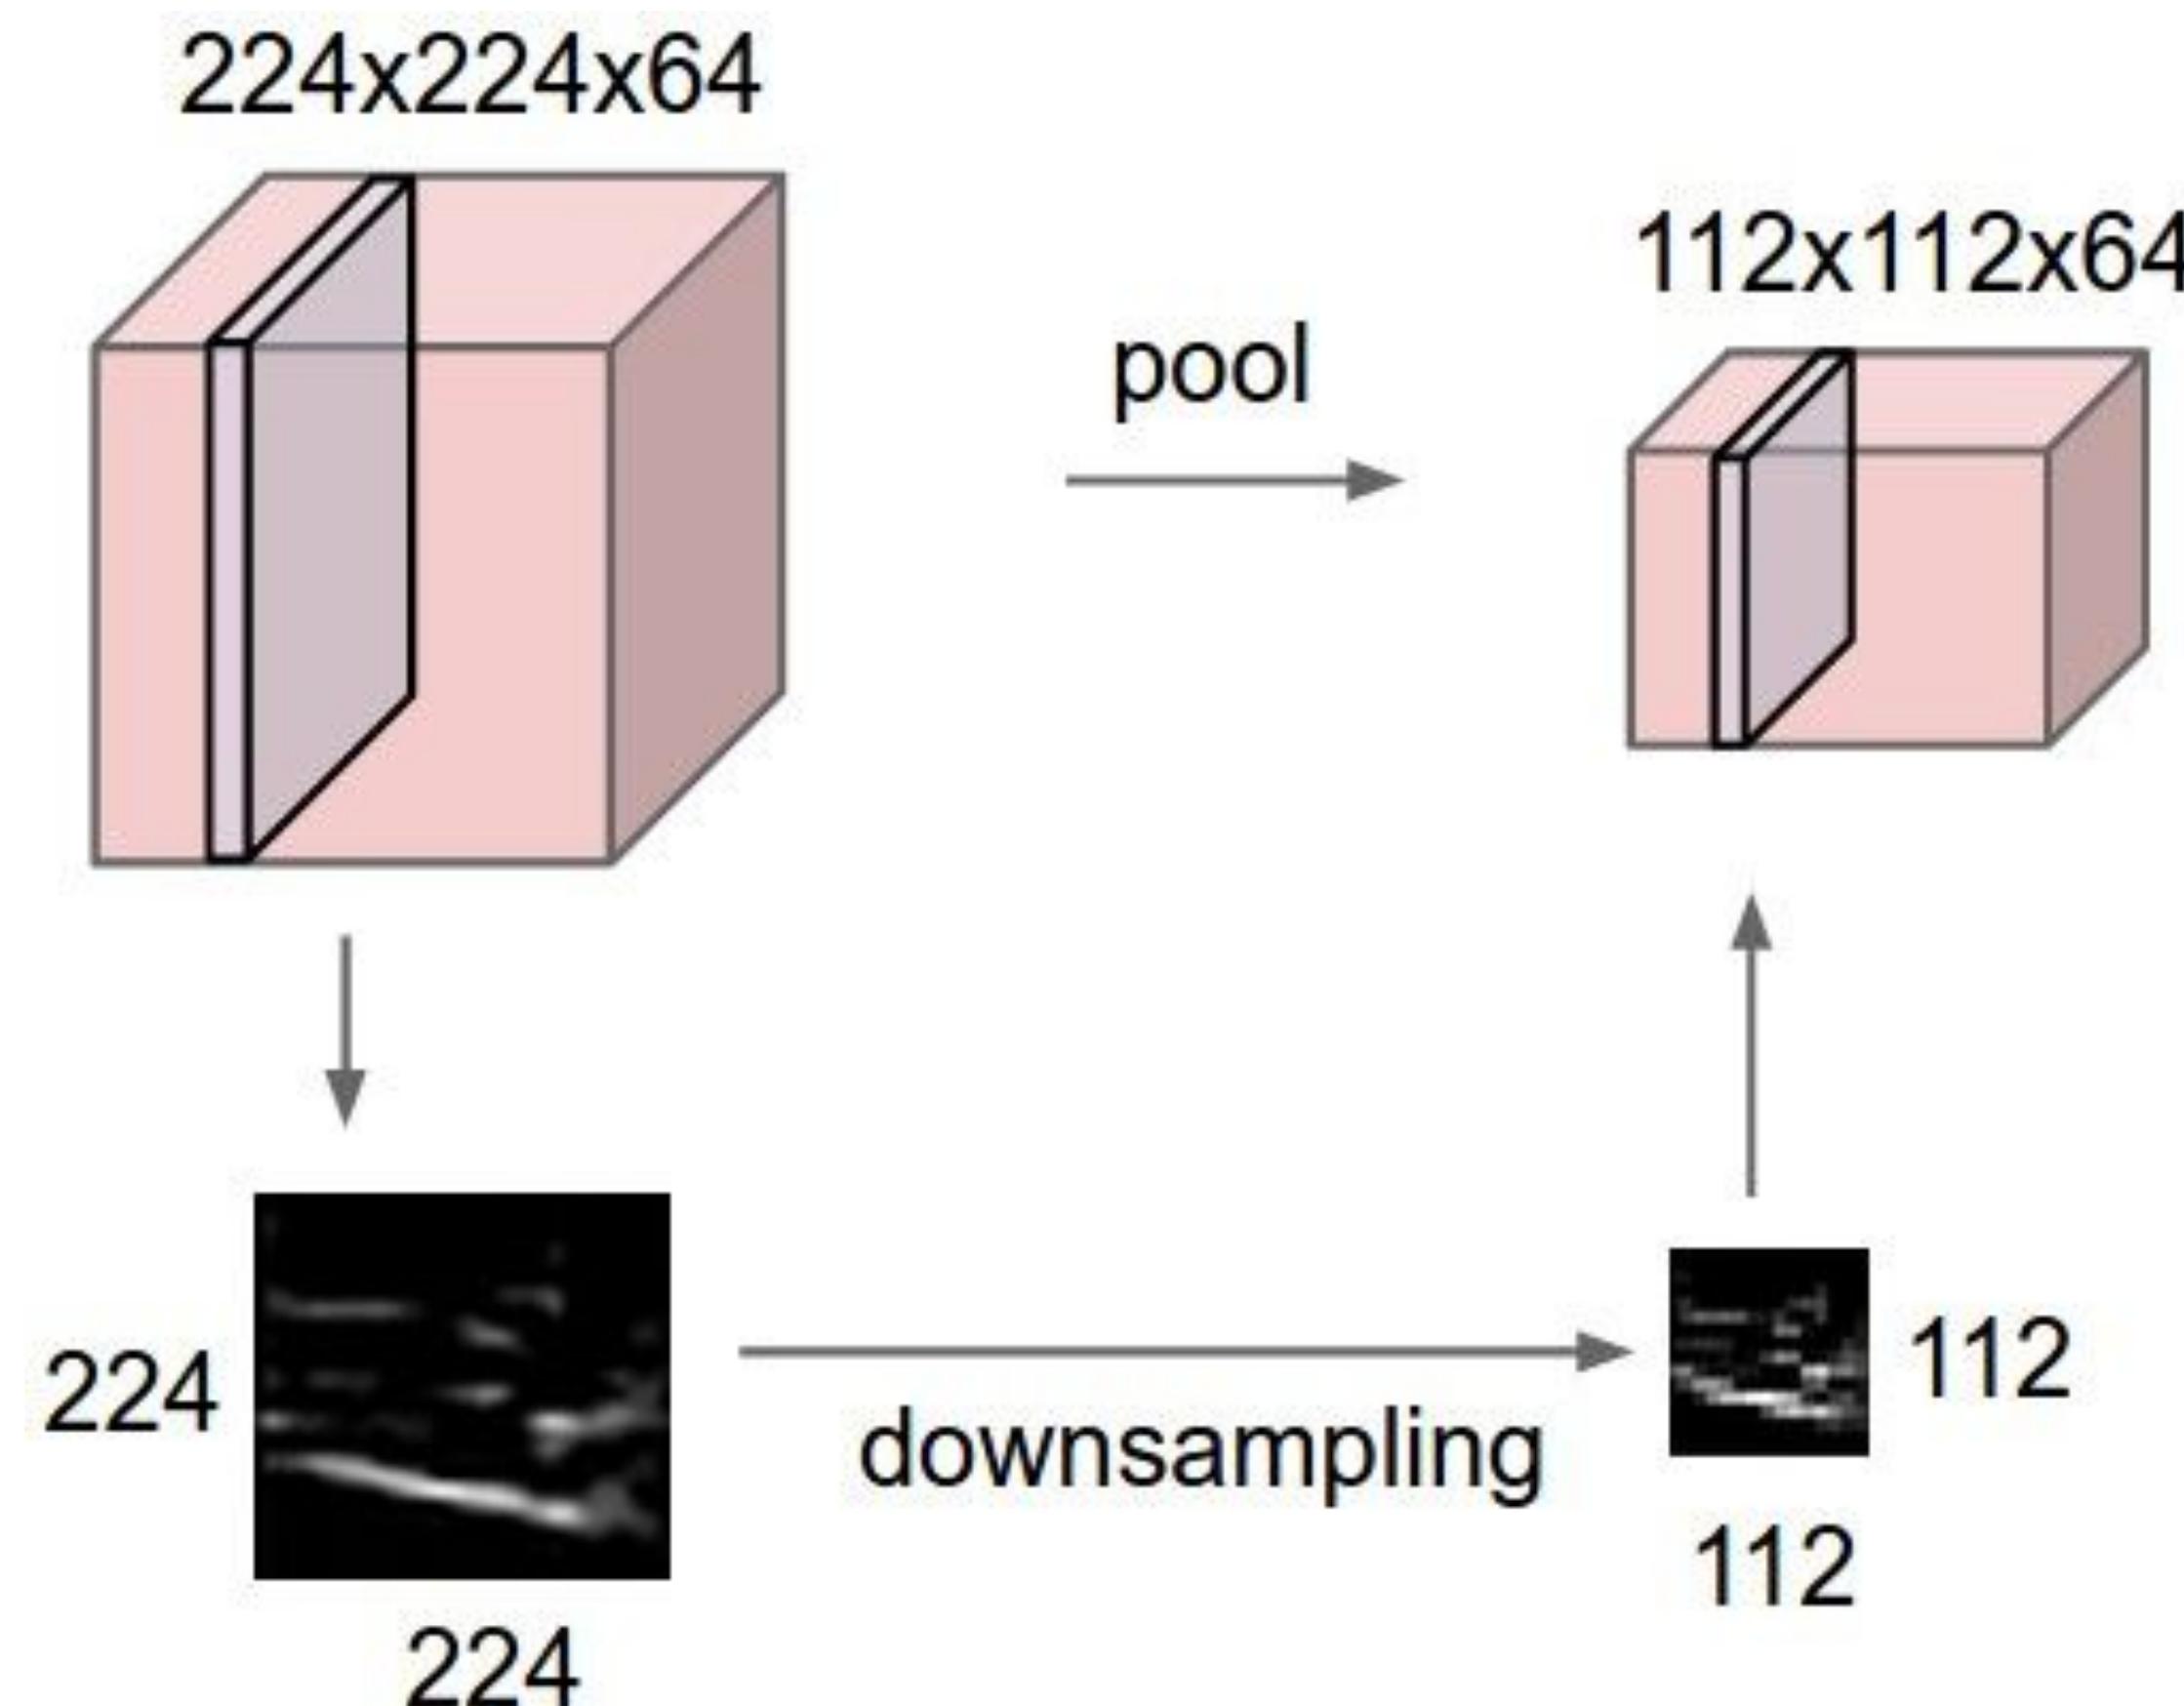
Number of parameters in this layer?

each filter has  $5*5*3 + 1 = 76$  params (+1 for bias)  
 $\Rightarrow 76*10 = 760$

In general, parameters in conv layer =  
**(filter width x filter height x input channels + 1) x number of filters.**

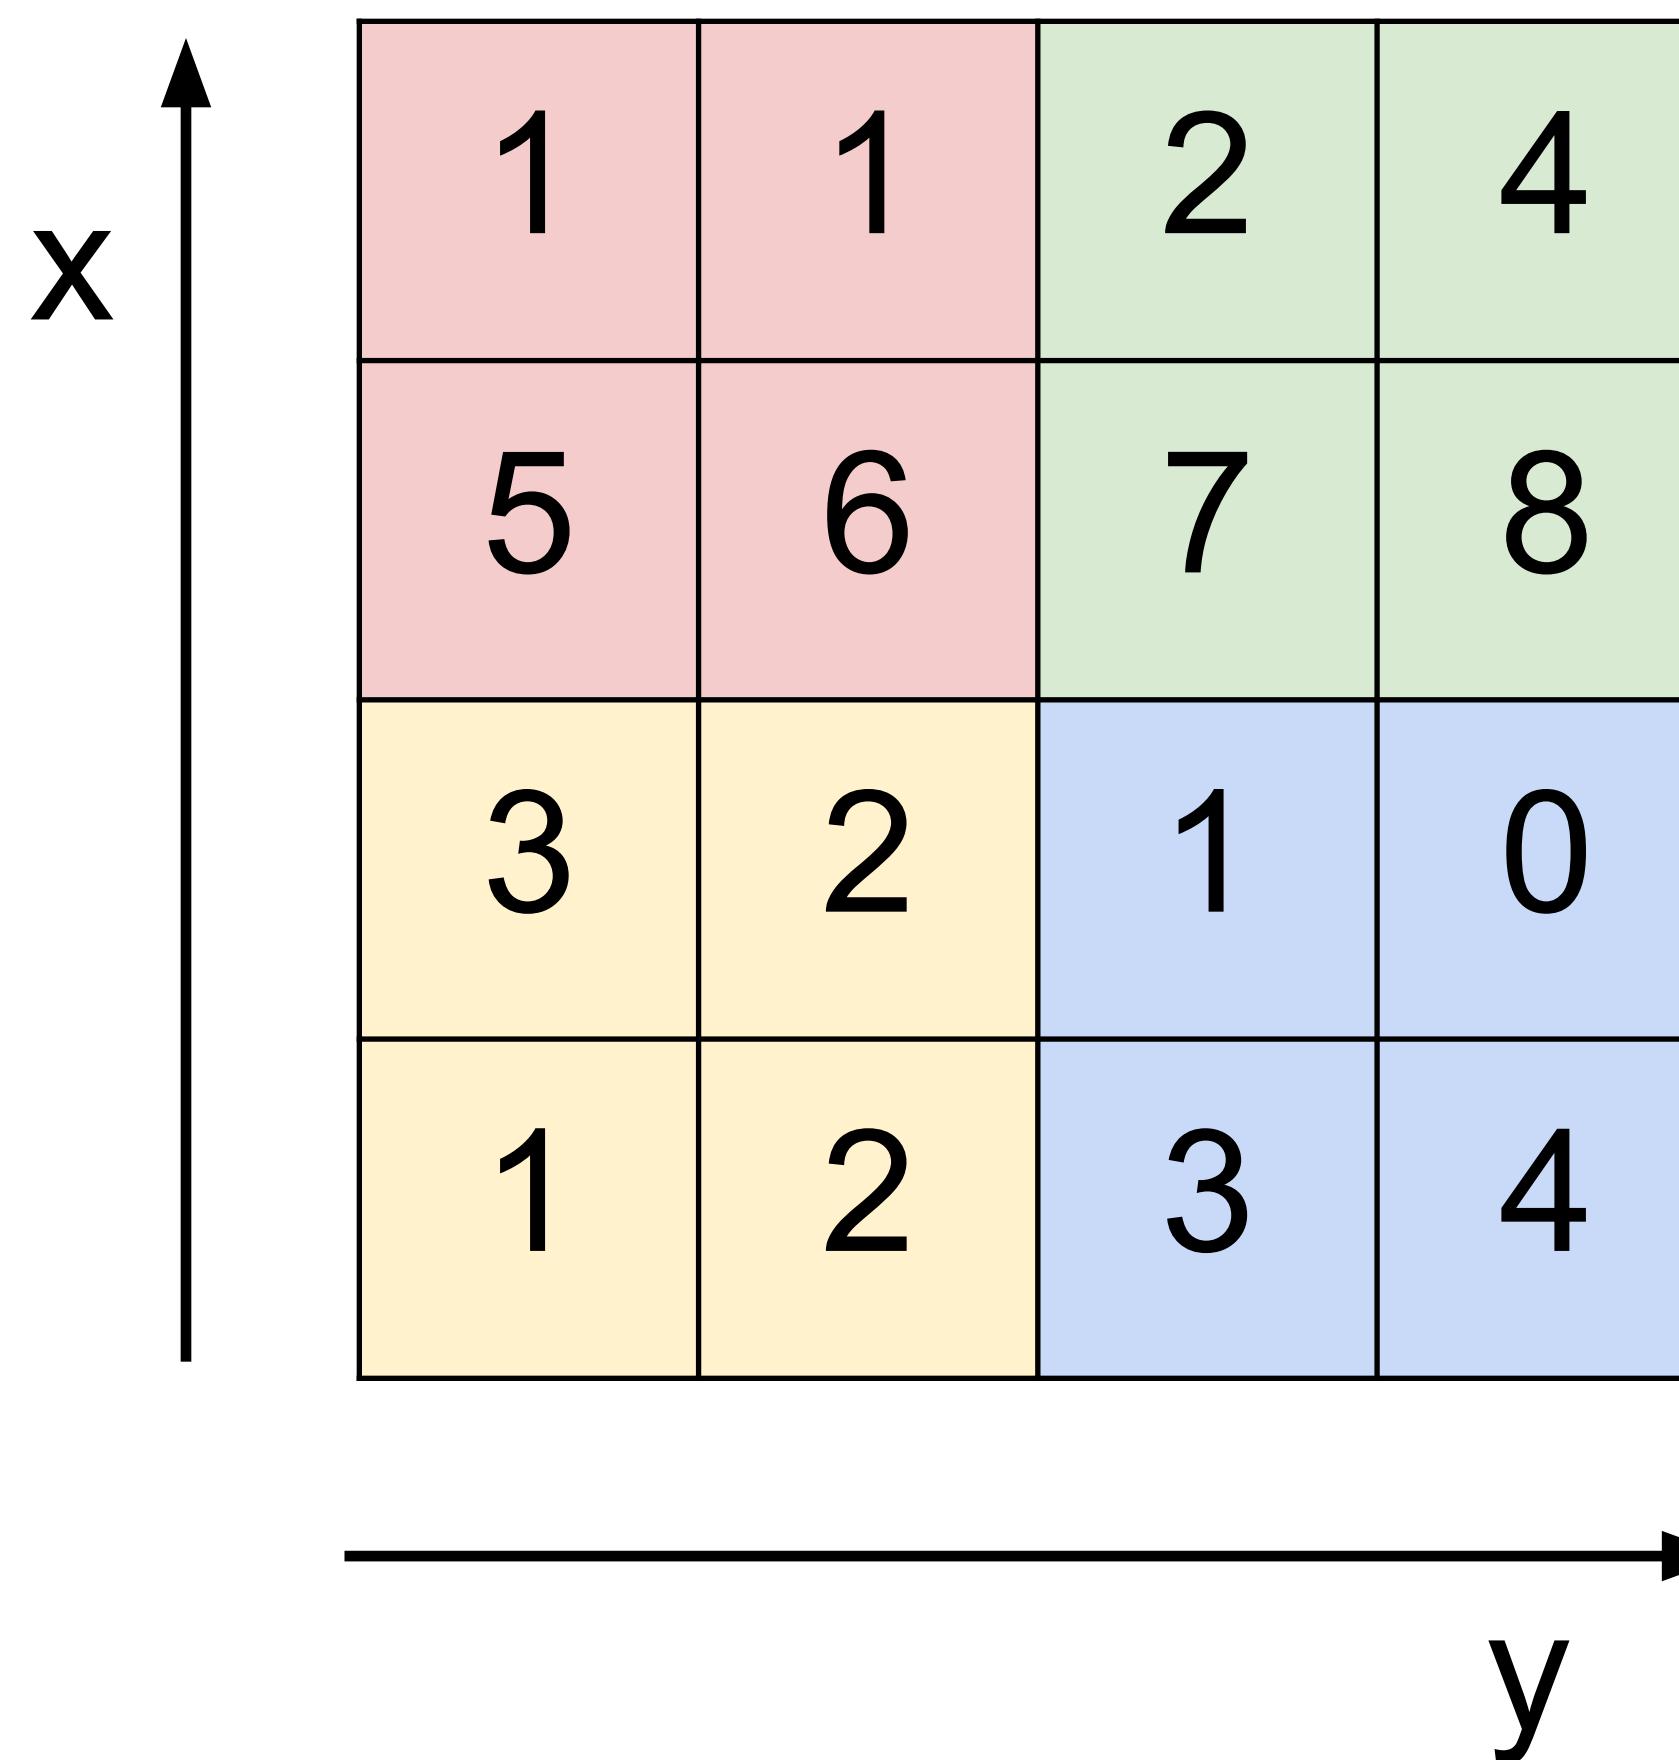
# Pooling layer

- makes the representations smaller and more manageable
- operates over each activation map independently:

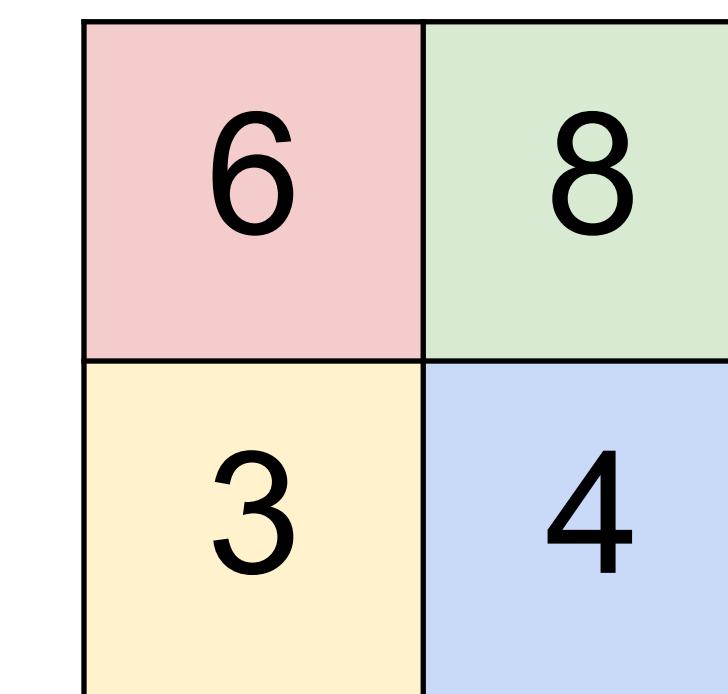


# MAX POOLING

Single depth slice

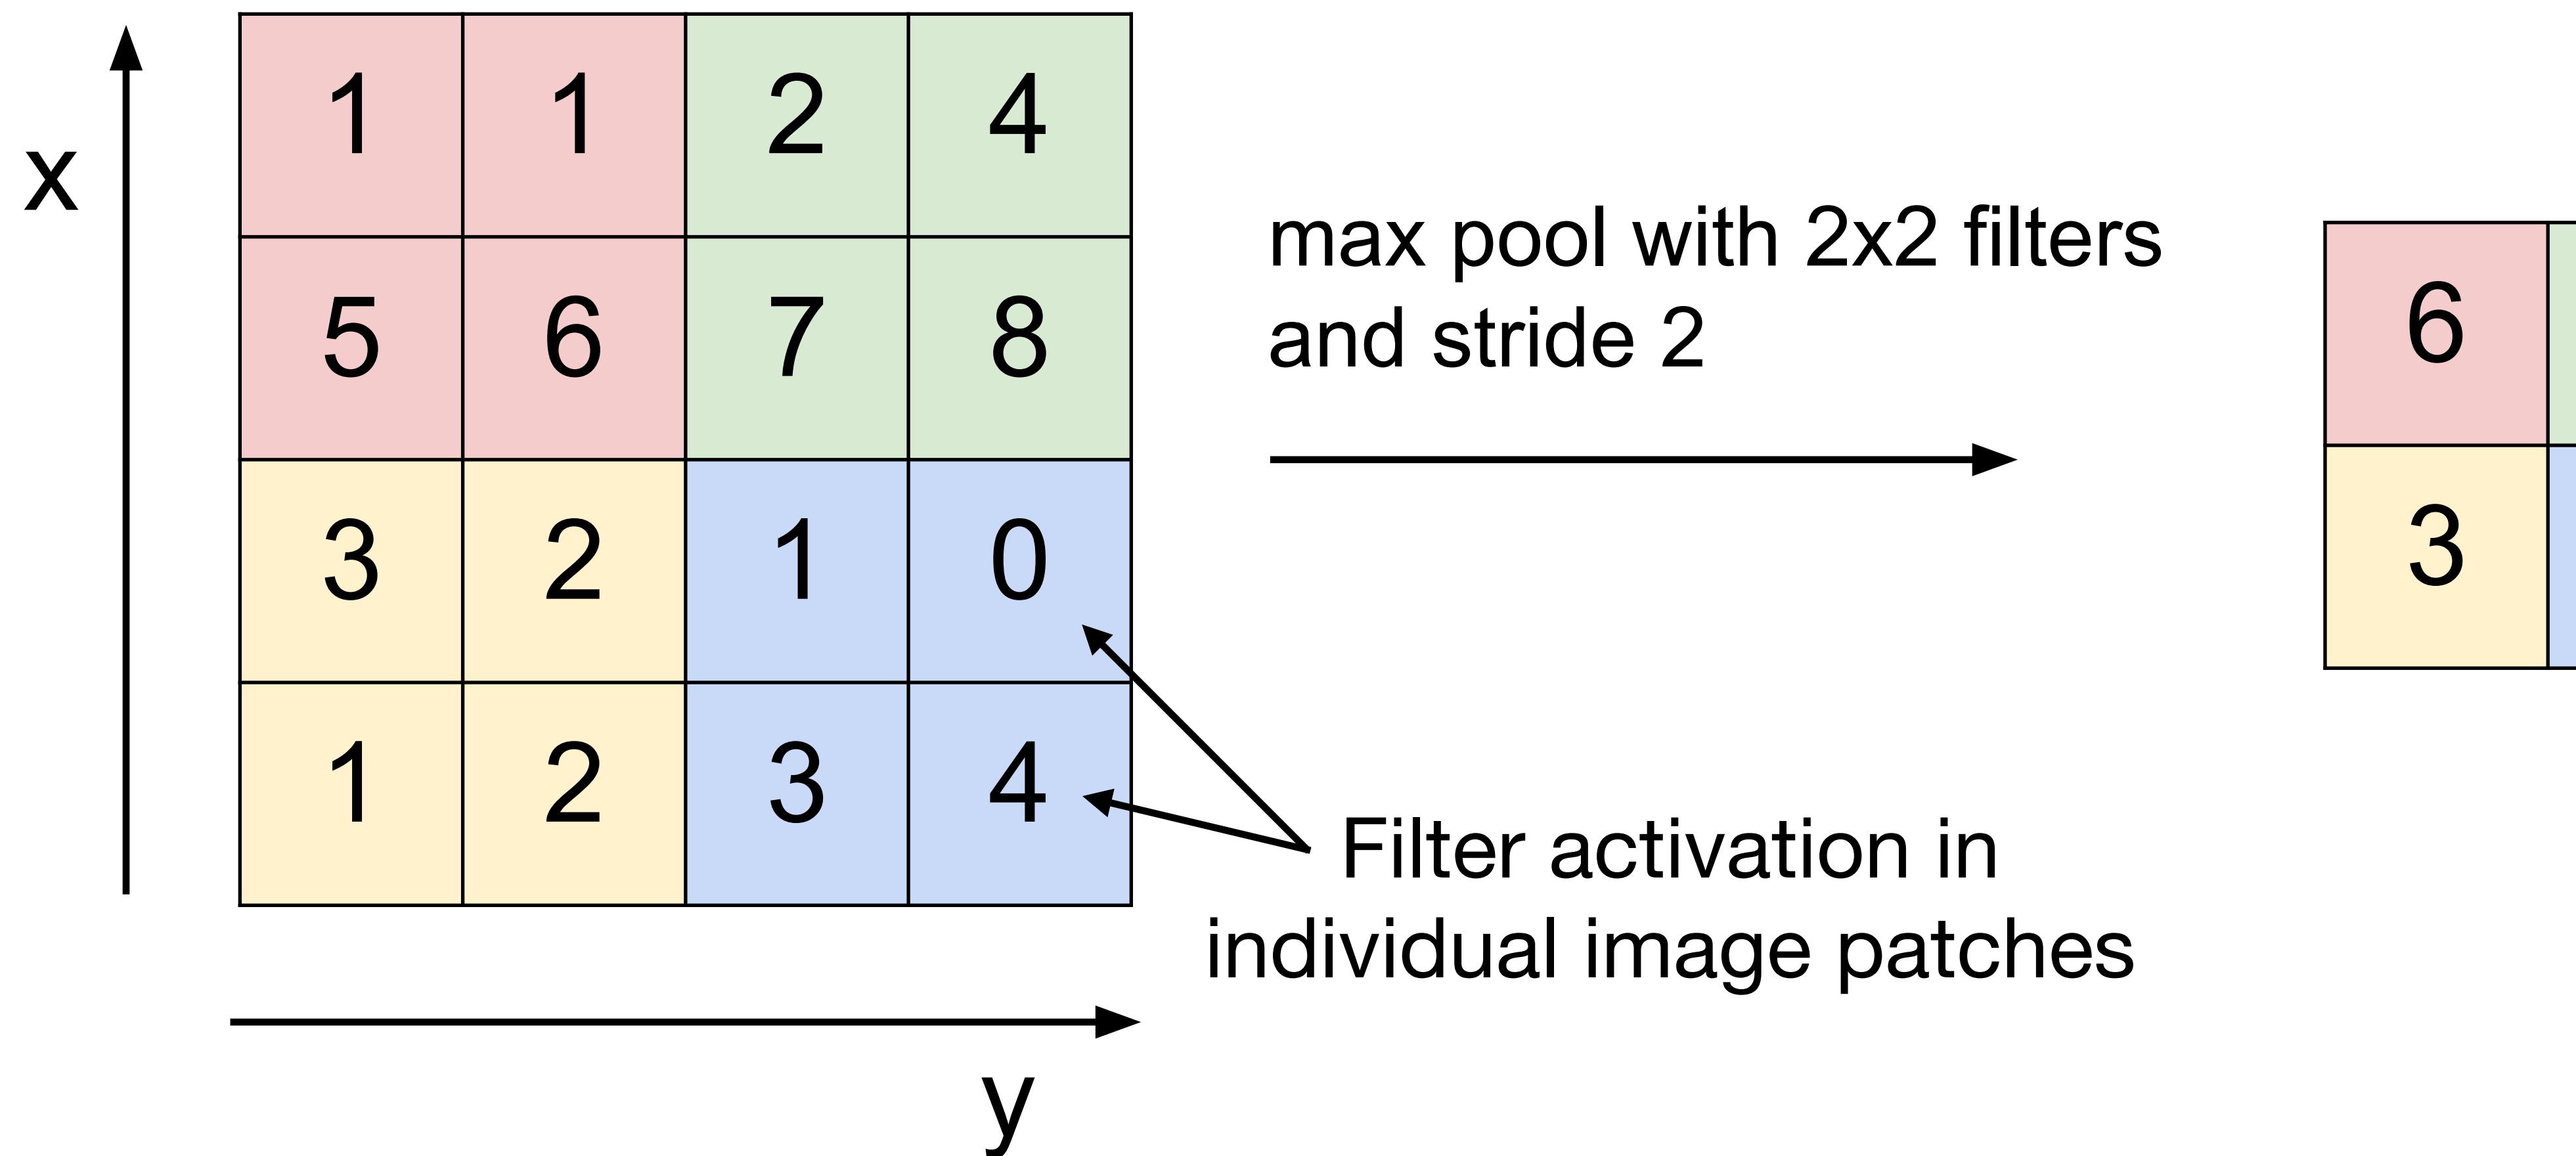


max pool with 2x2 filters  
and stride 2



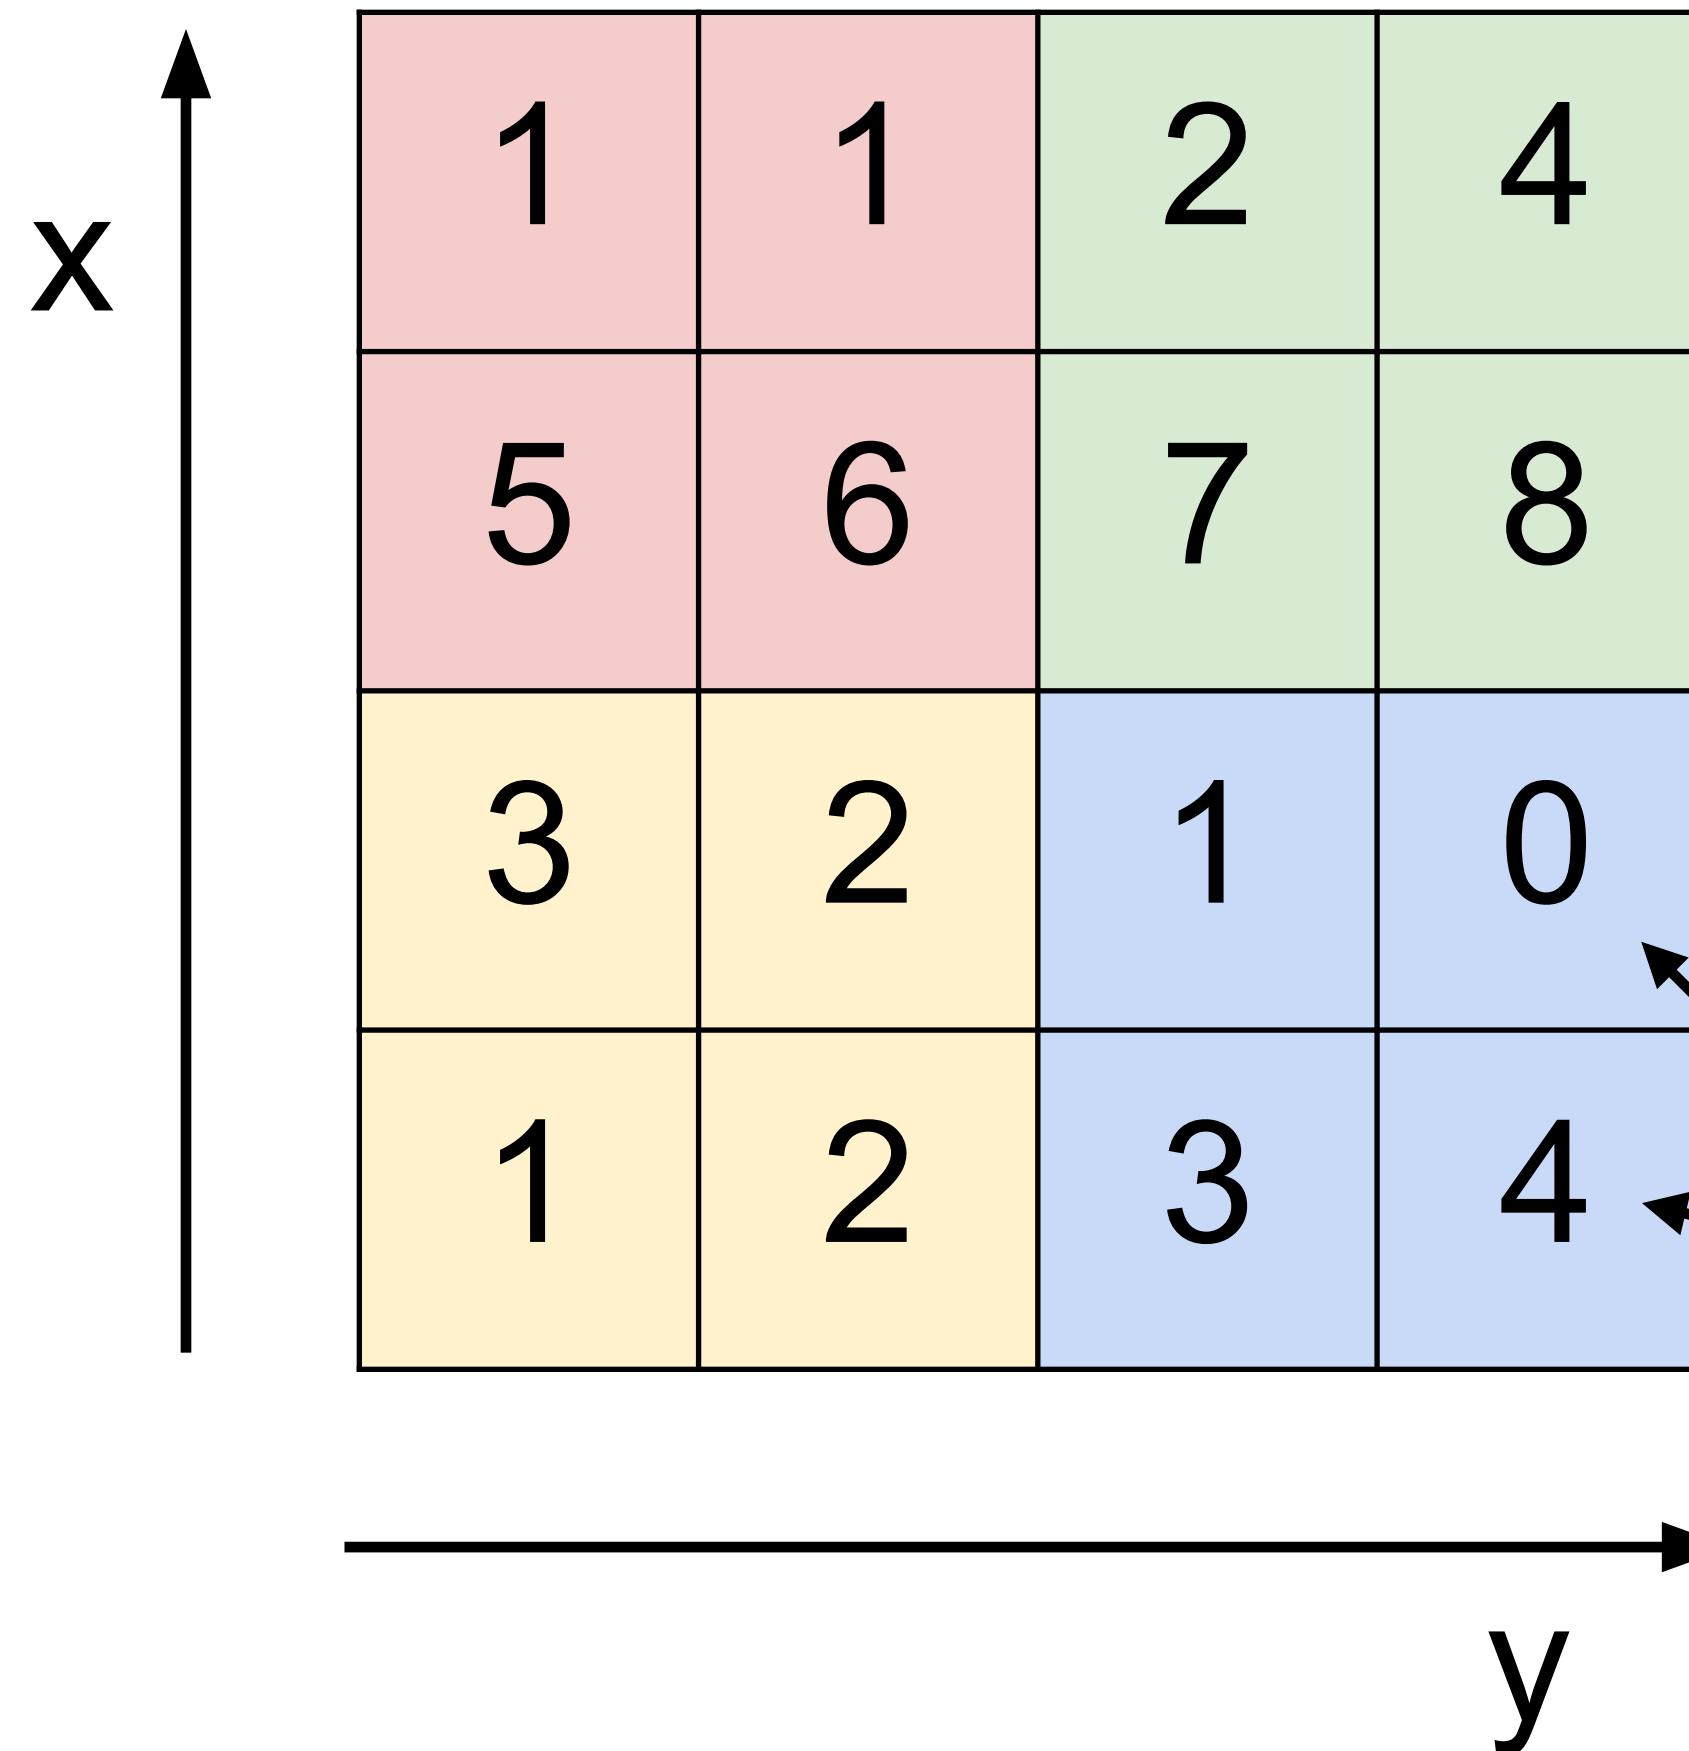
# MAX POOLING

Single depth slice



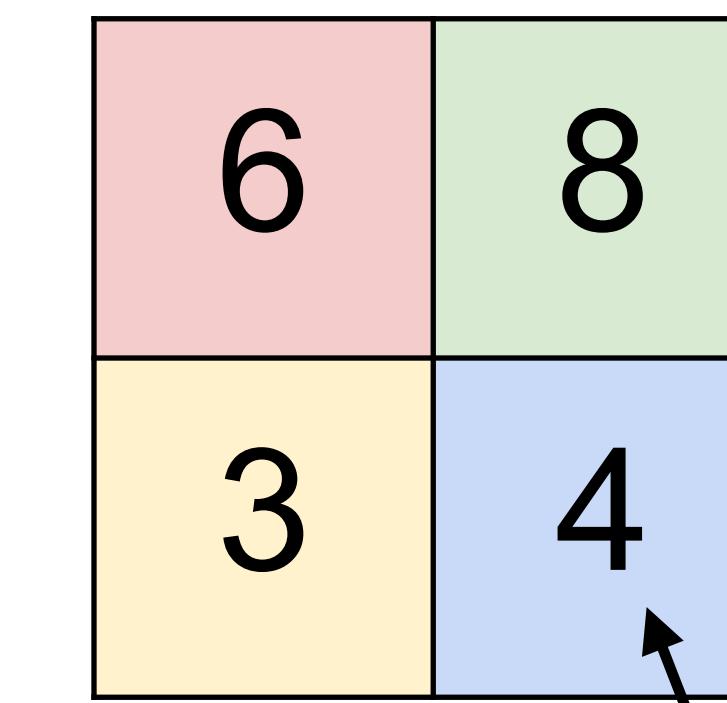
# MAX POOLING

Single depth slice



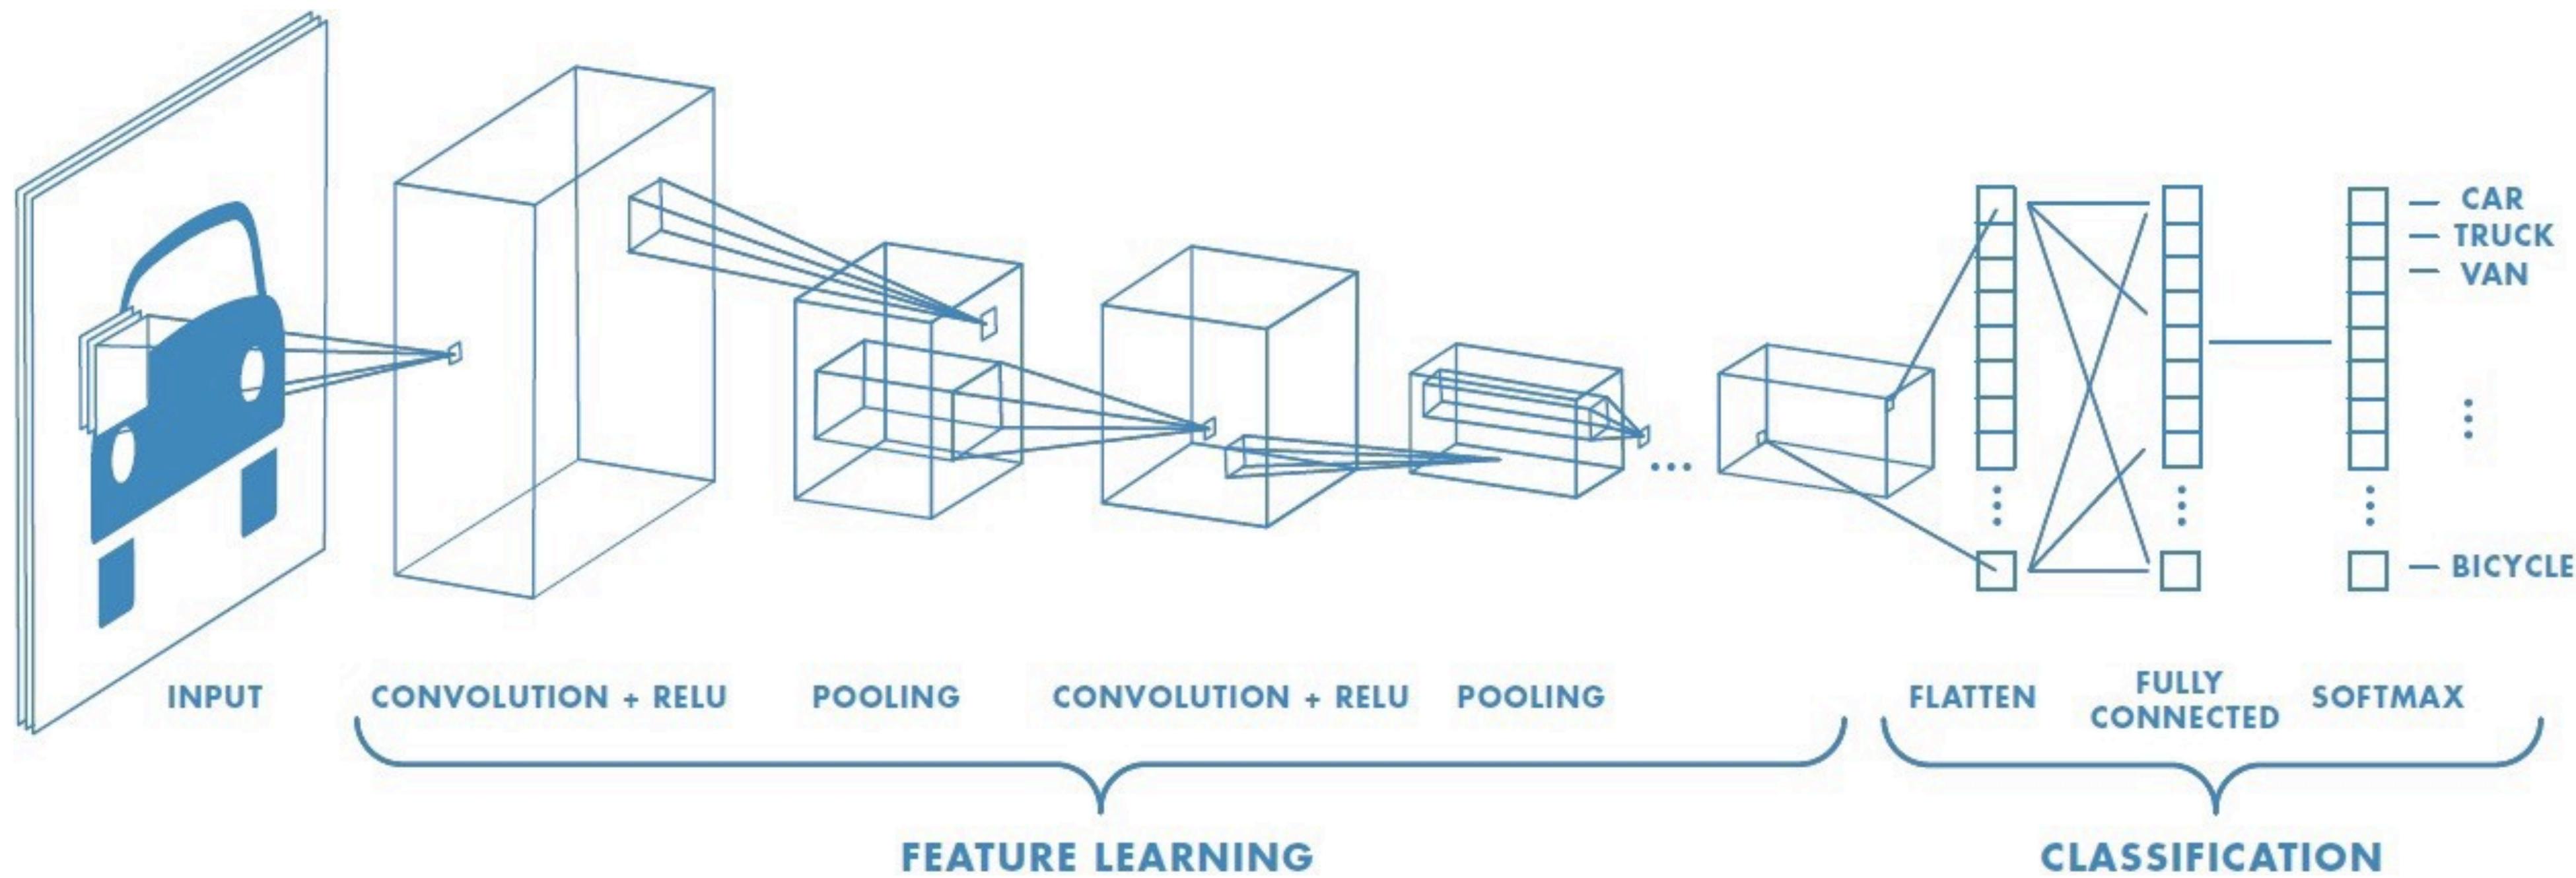
max pool with 2x2 filters  
and stride 2

Filter activation in  
individual image patches

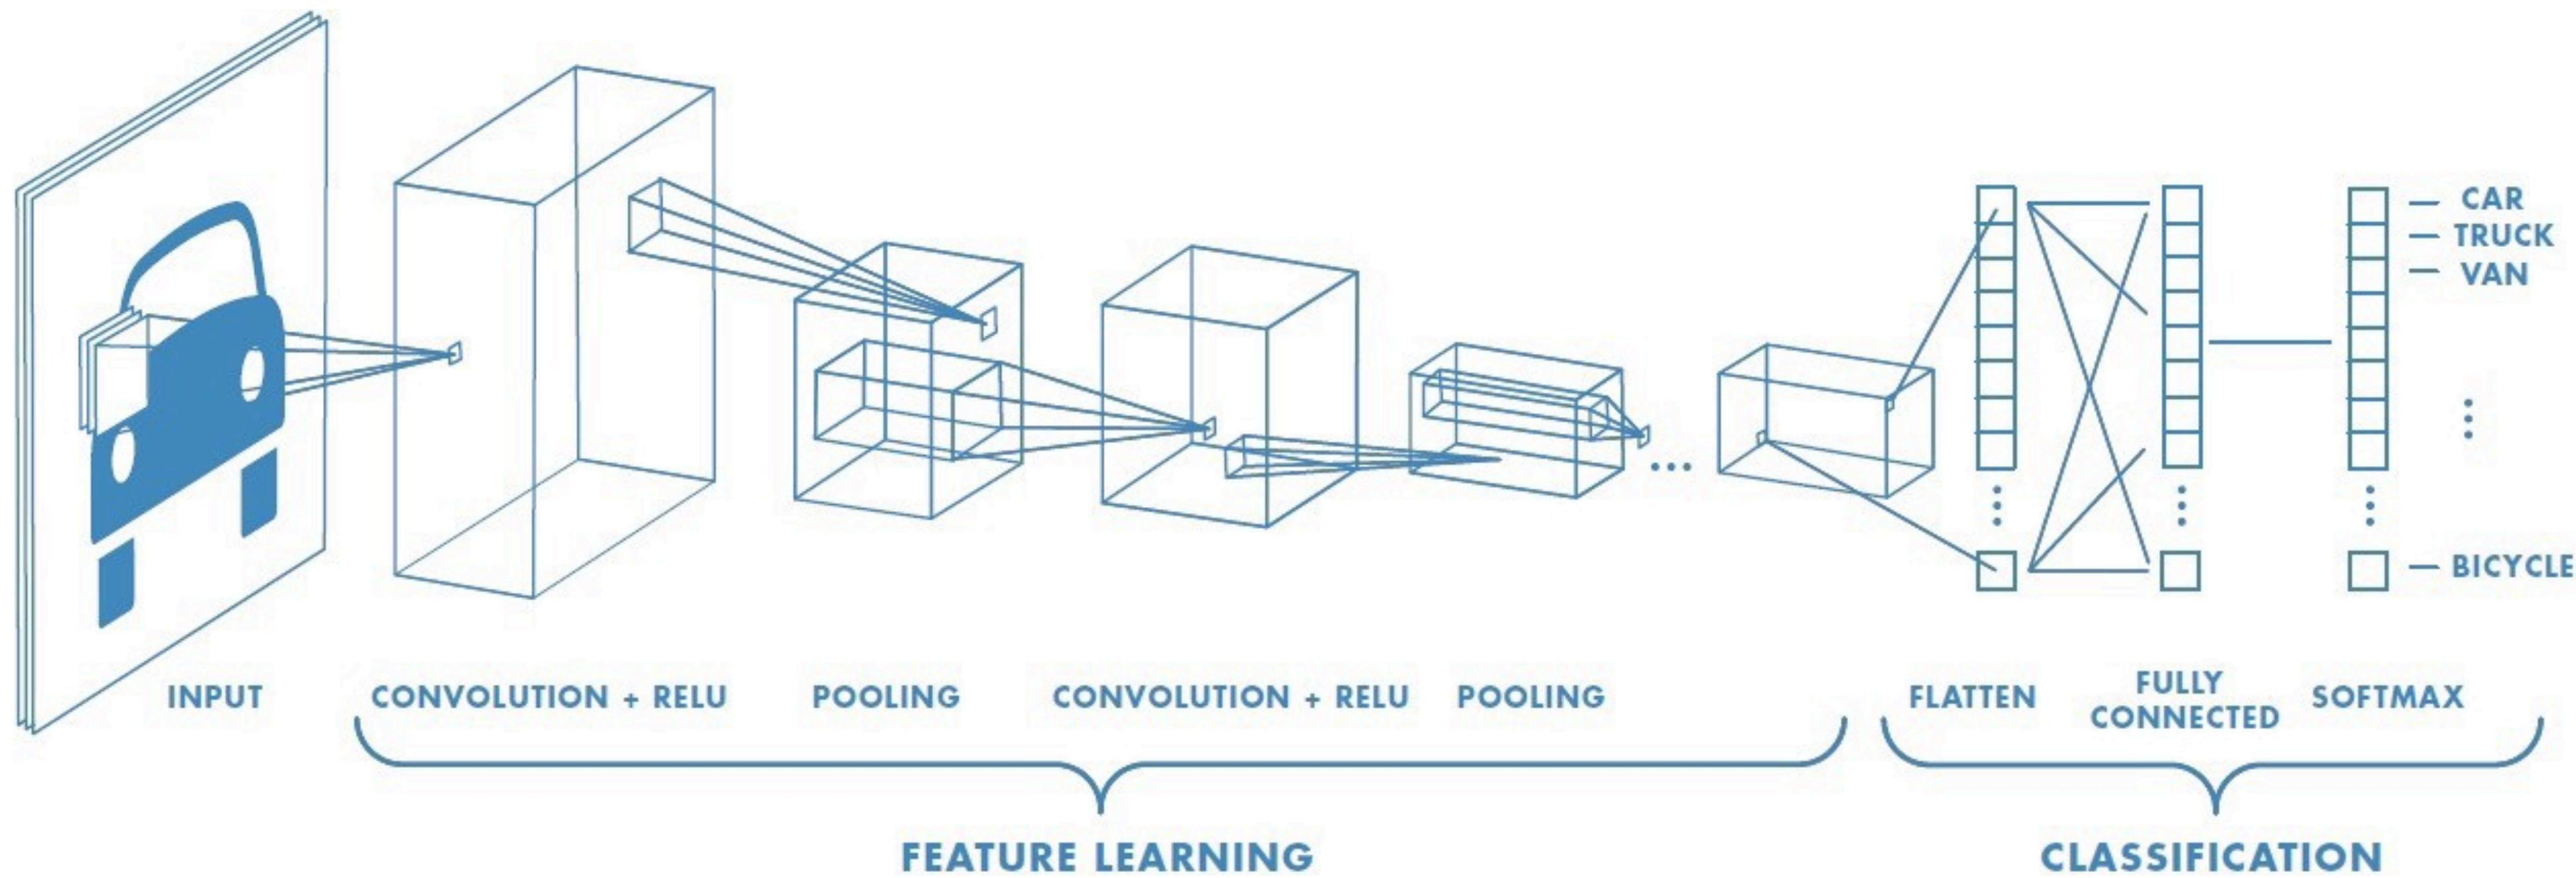


Maximum filter activation  
across adjacent patches

# Convolutional neural networks



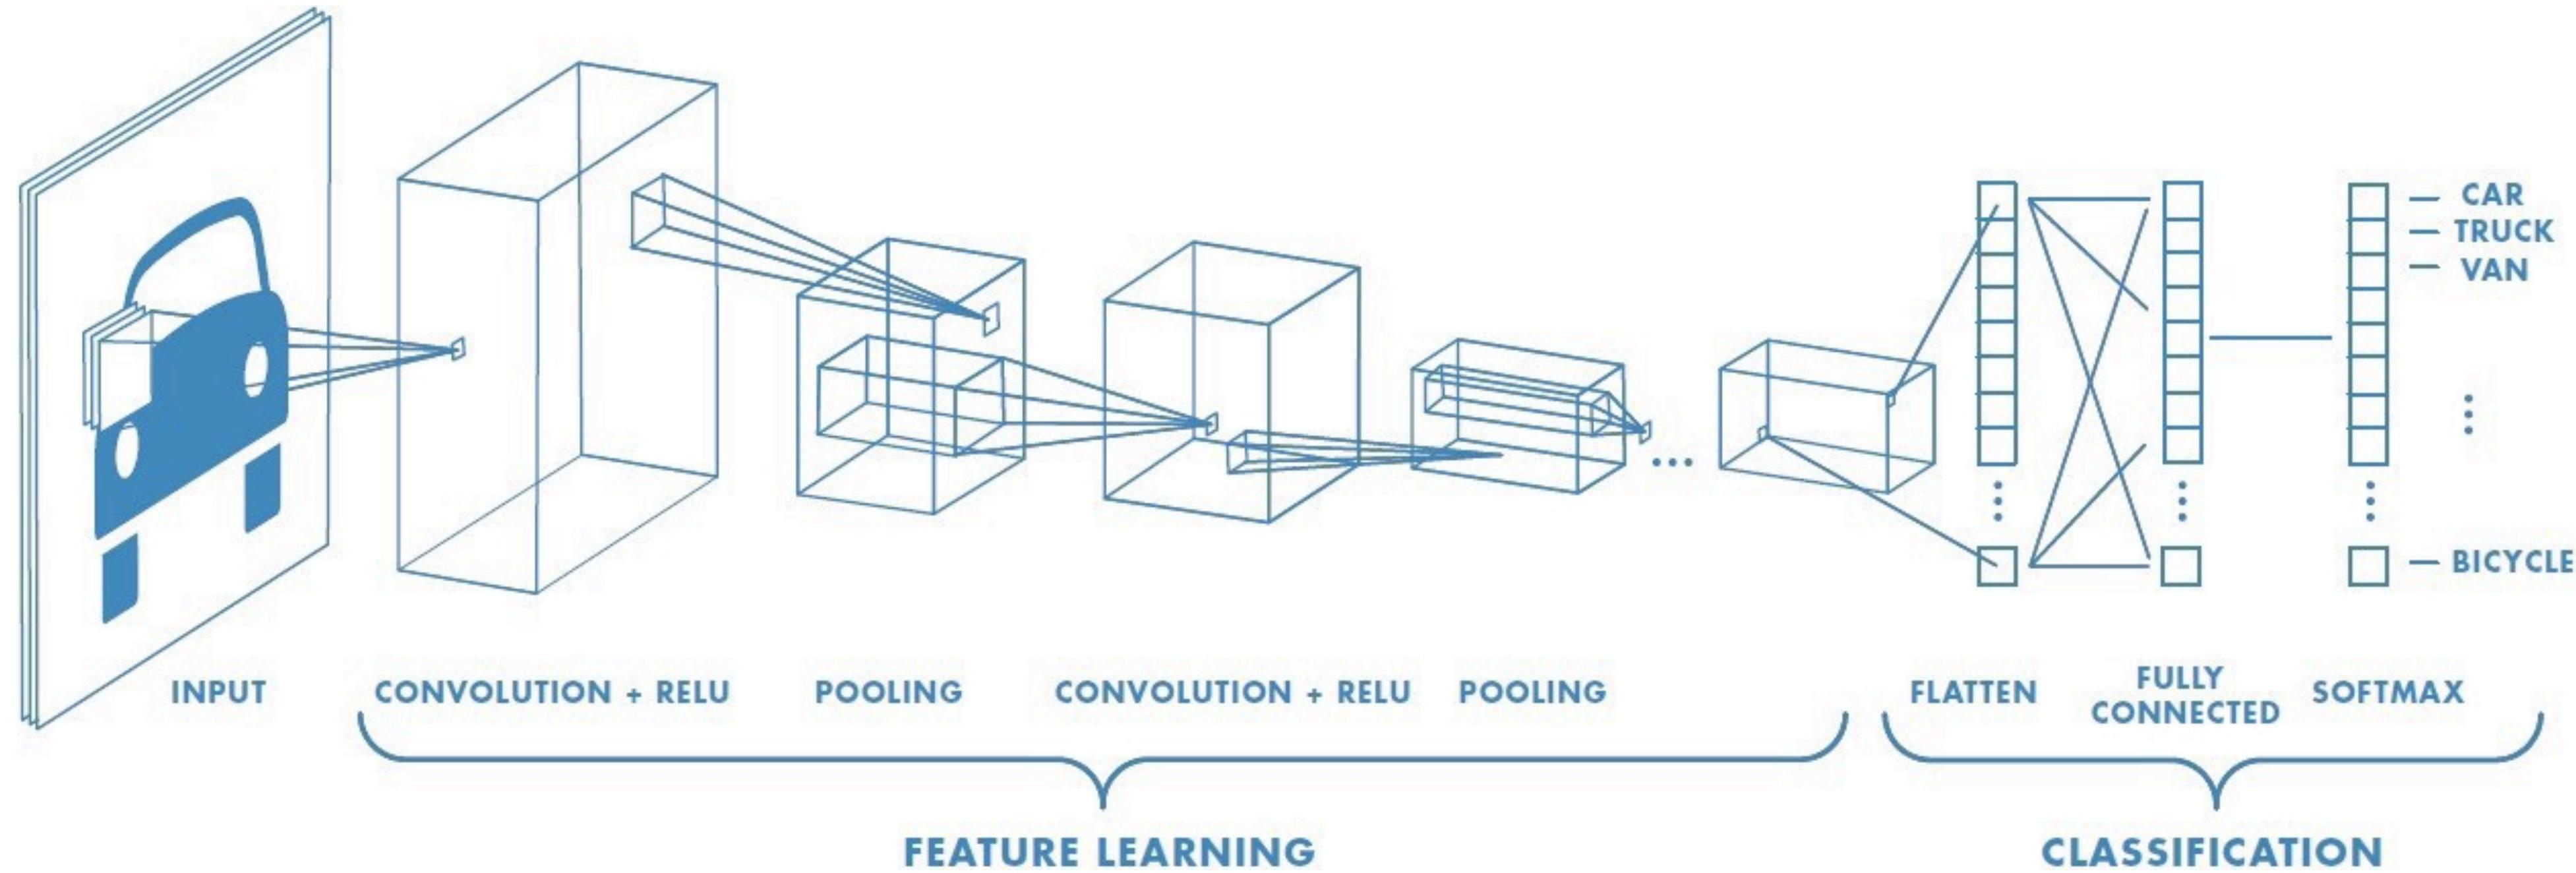
# Convolutional neural networks



<https://towardsdatascience.com/a-comprehensive-guide-to-convolutional-neural-networks-the-eli5-way-3bd2b1164a53>

A CNN stacks together several alternating convolution and pooling layers, followed by a fully connected layer and a softmax output.

# Convolutional neural networks



<https://towardsdatascience.com/a-comprehensive-guide-to-convolutional-neural-networks-the-eli5-way-3bd2b1164a53>

A CNN stacks together several alternating convolution and pooling layers, followed by a fully connected layer and a softmax output.

Filters, weights in fully connected layer, and biases learned by optimizing cross-entropy loss via stochastic gradient descent.

# Interpreting the filters learned by a CNN

# Interpreting the filters learned by a CNN

Use neural network for binary classification, e.g. faces versus not faces, cars versus not cars, etc.

# Interpreting the filters learned by a CNN

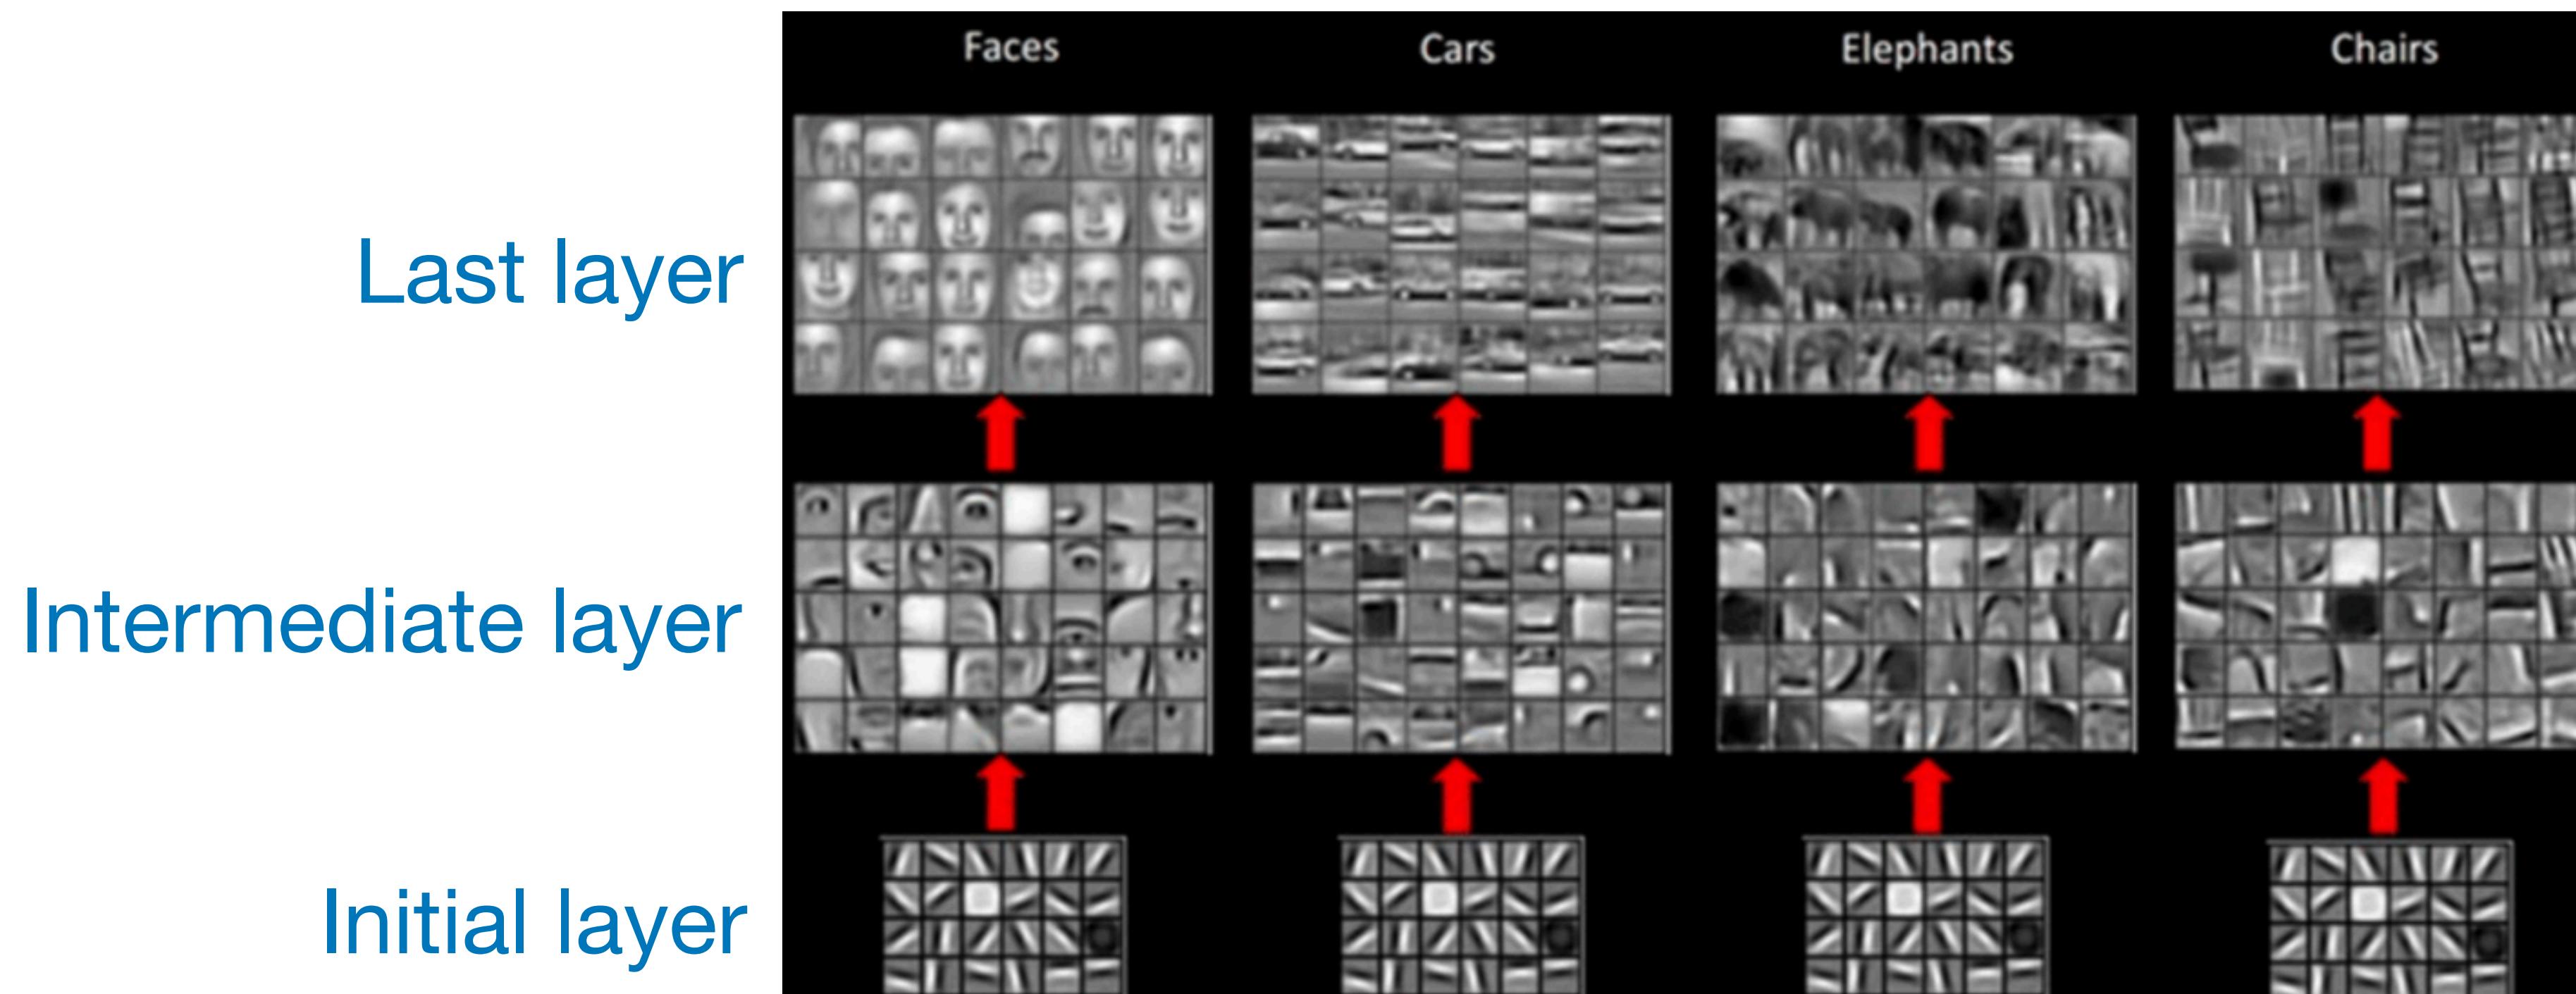
Use neural network for binary classification, e.g. faces versus not faces, cars versus not cars, etc.

For each neuron at each layer, find input image that activates it most strongly.

# Interpreting the filters learned by a CNN

Use neural network for binary classification, e.g. faces versus not faces, cars versus not cars, etc.

For each neuron at each layer, find input image that activates it most strongly.



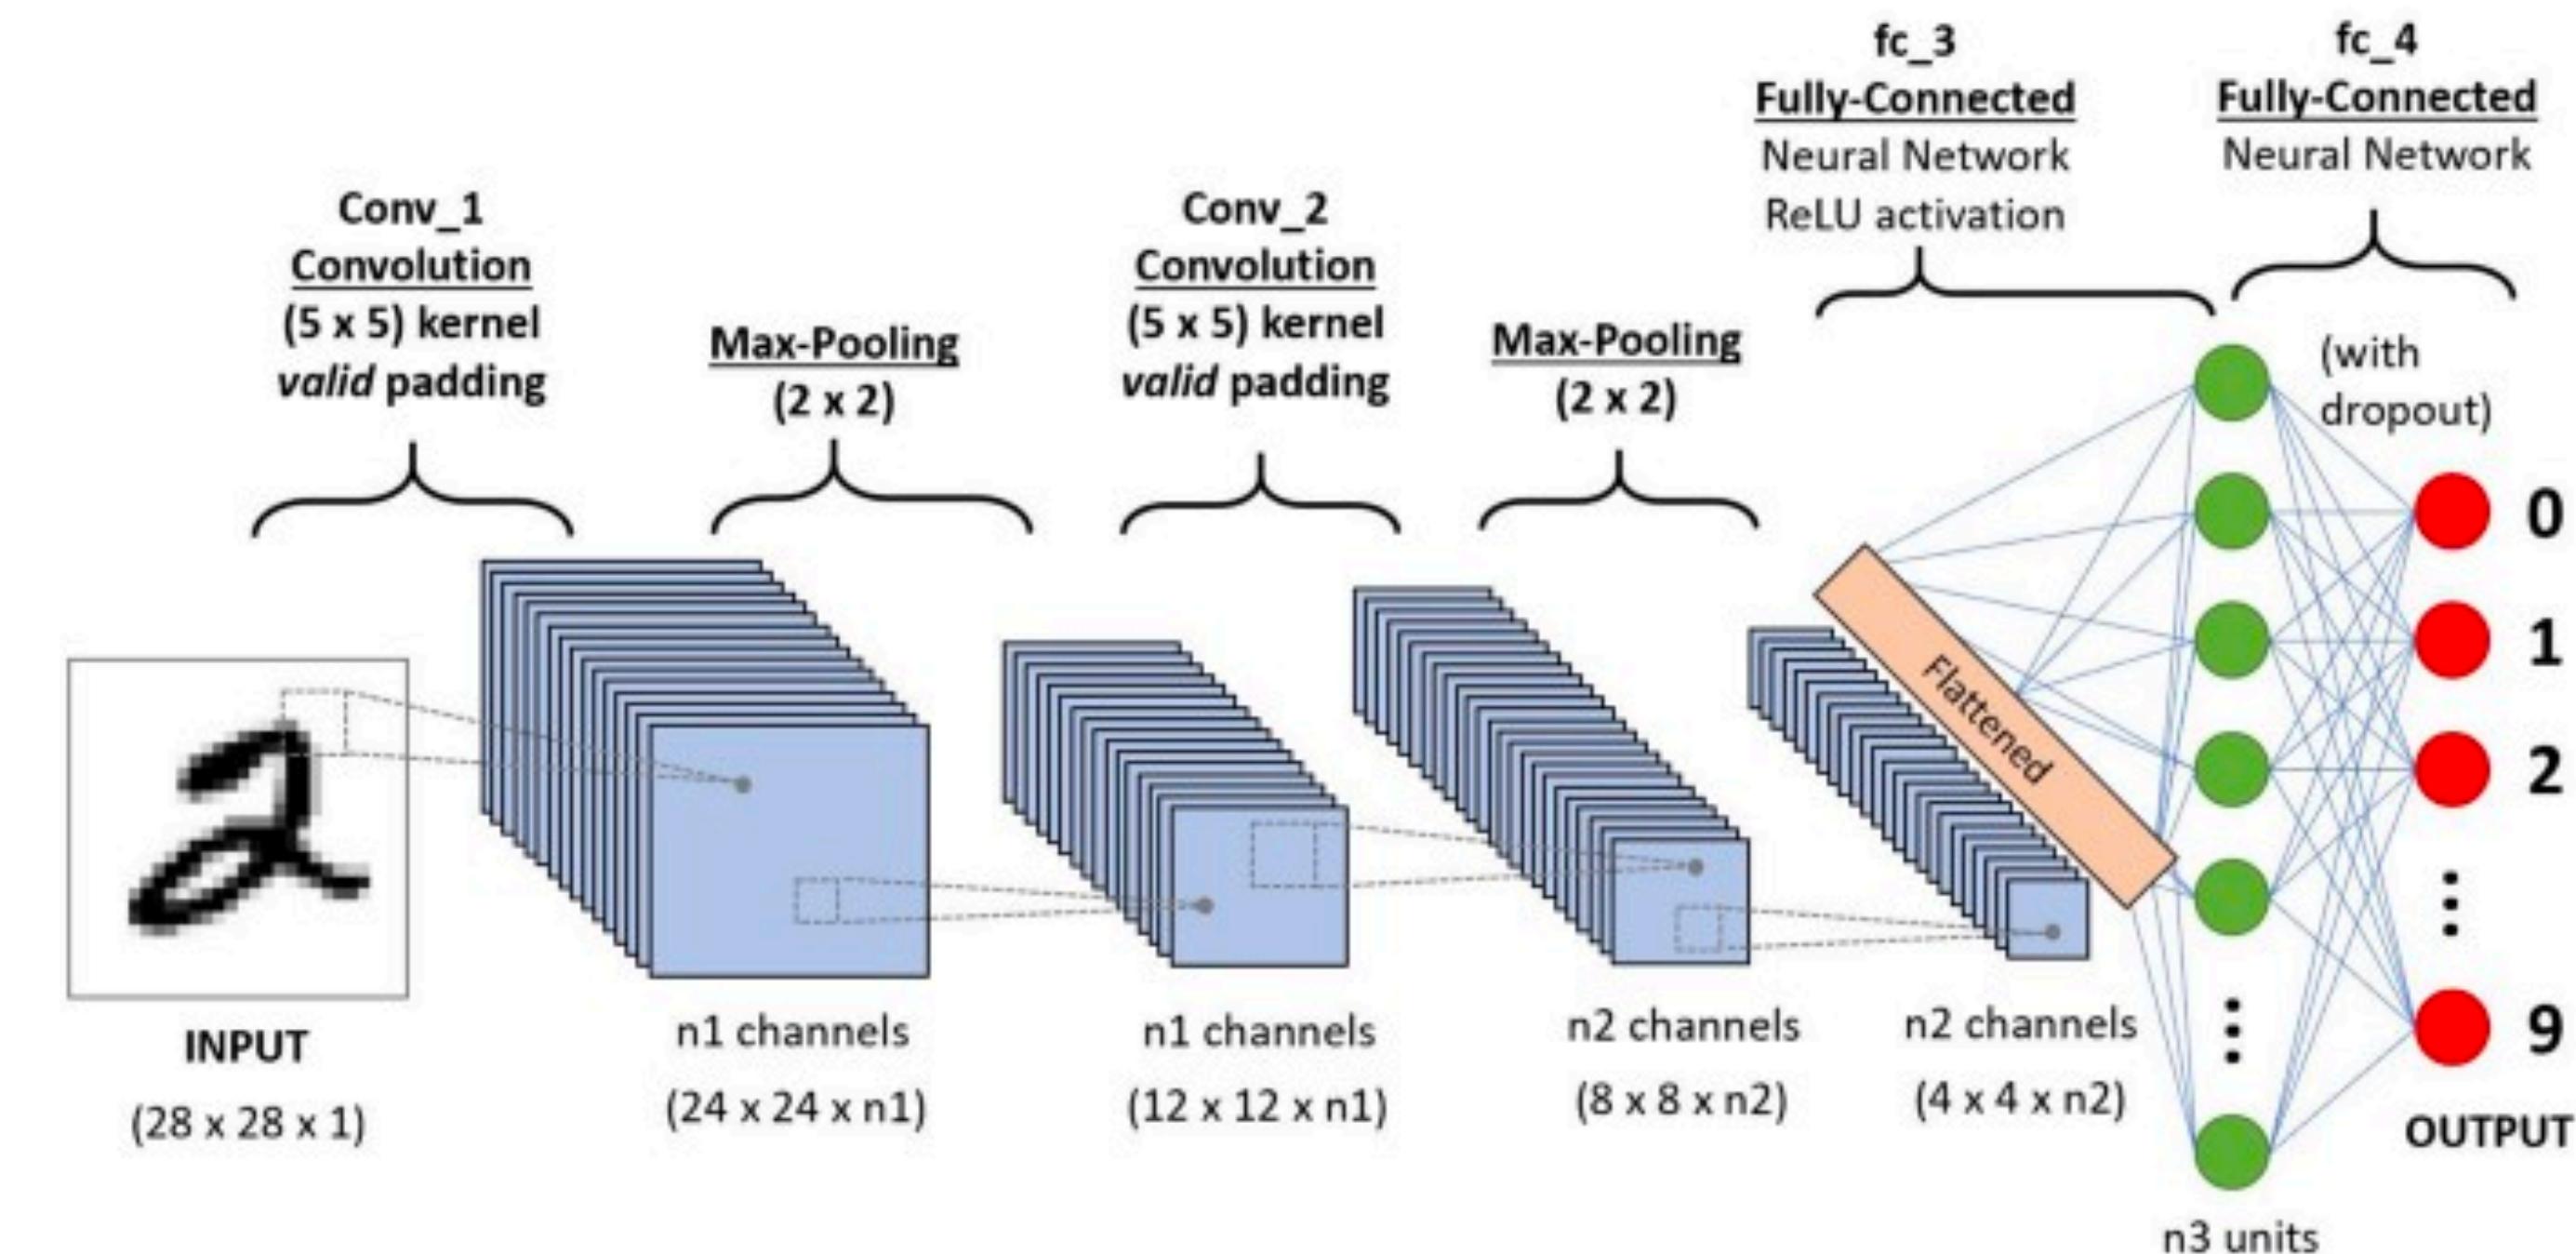
# The original CNN architecture

# The original CNN architecture

LeNet architecture for hand-written digit recognition (1989).

# The original CNN architecture

LeNet architecture for hand-written digit recognition (1989).



<https://towardsdatascience.com/a-comprehensive-guide-to-convolutional-neural-networks-the-eli5-way-3bd2b1164a53>

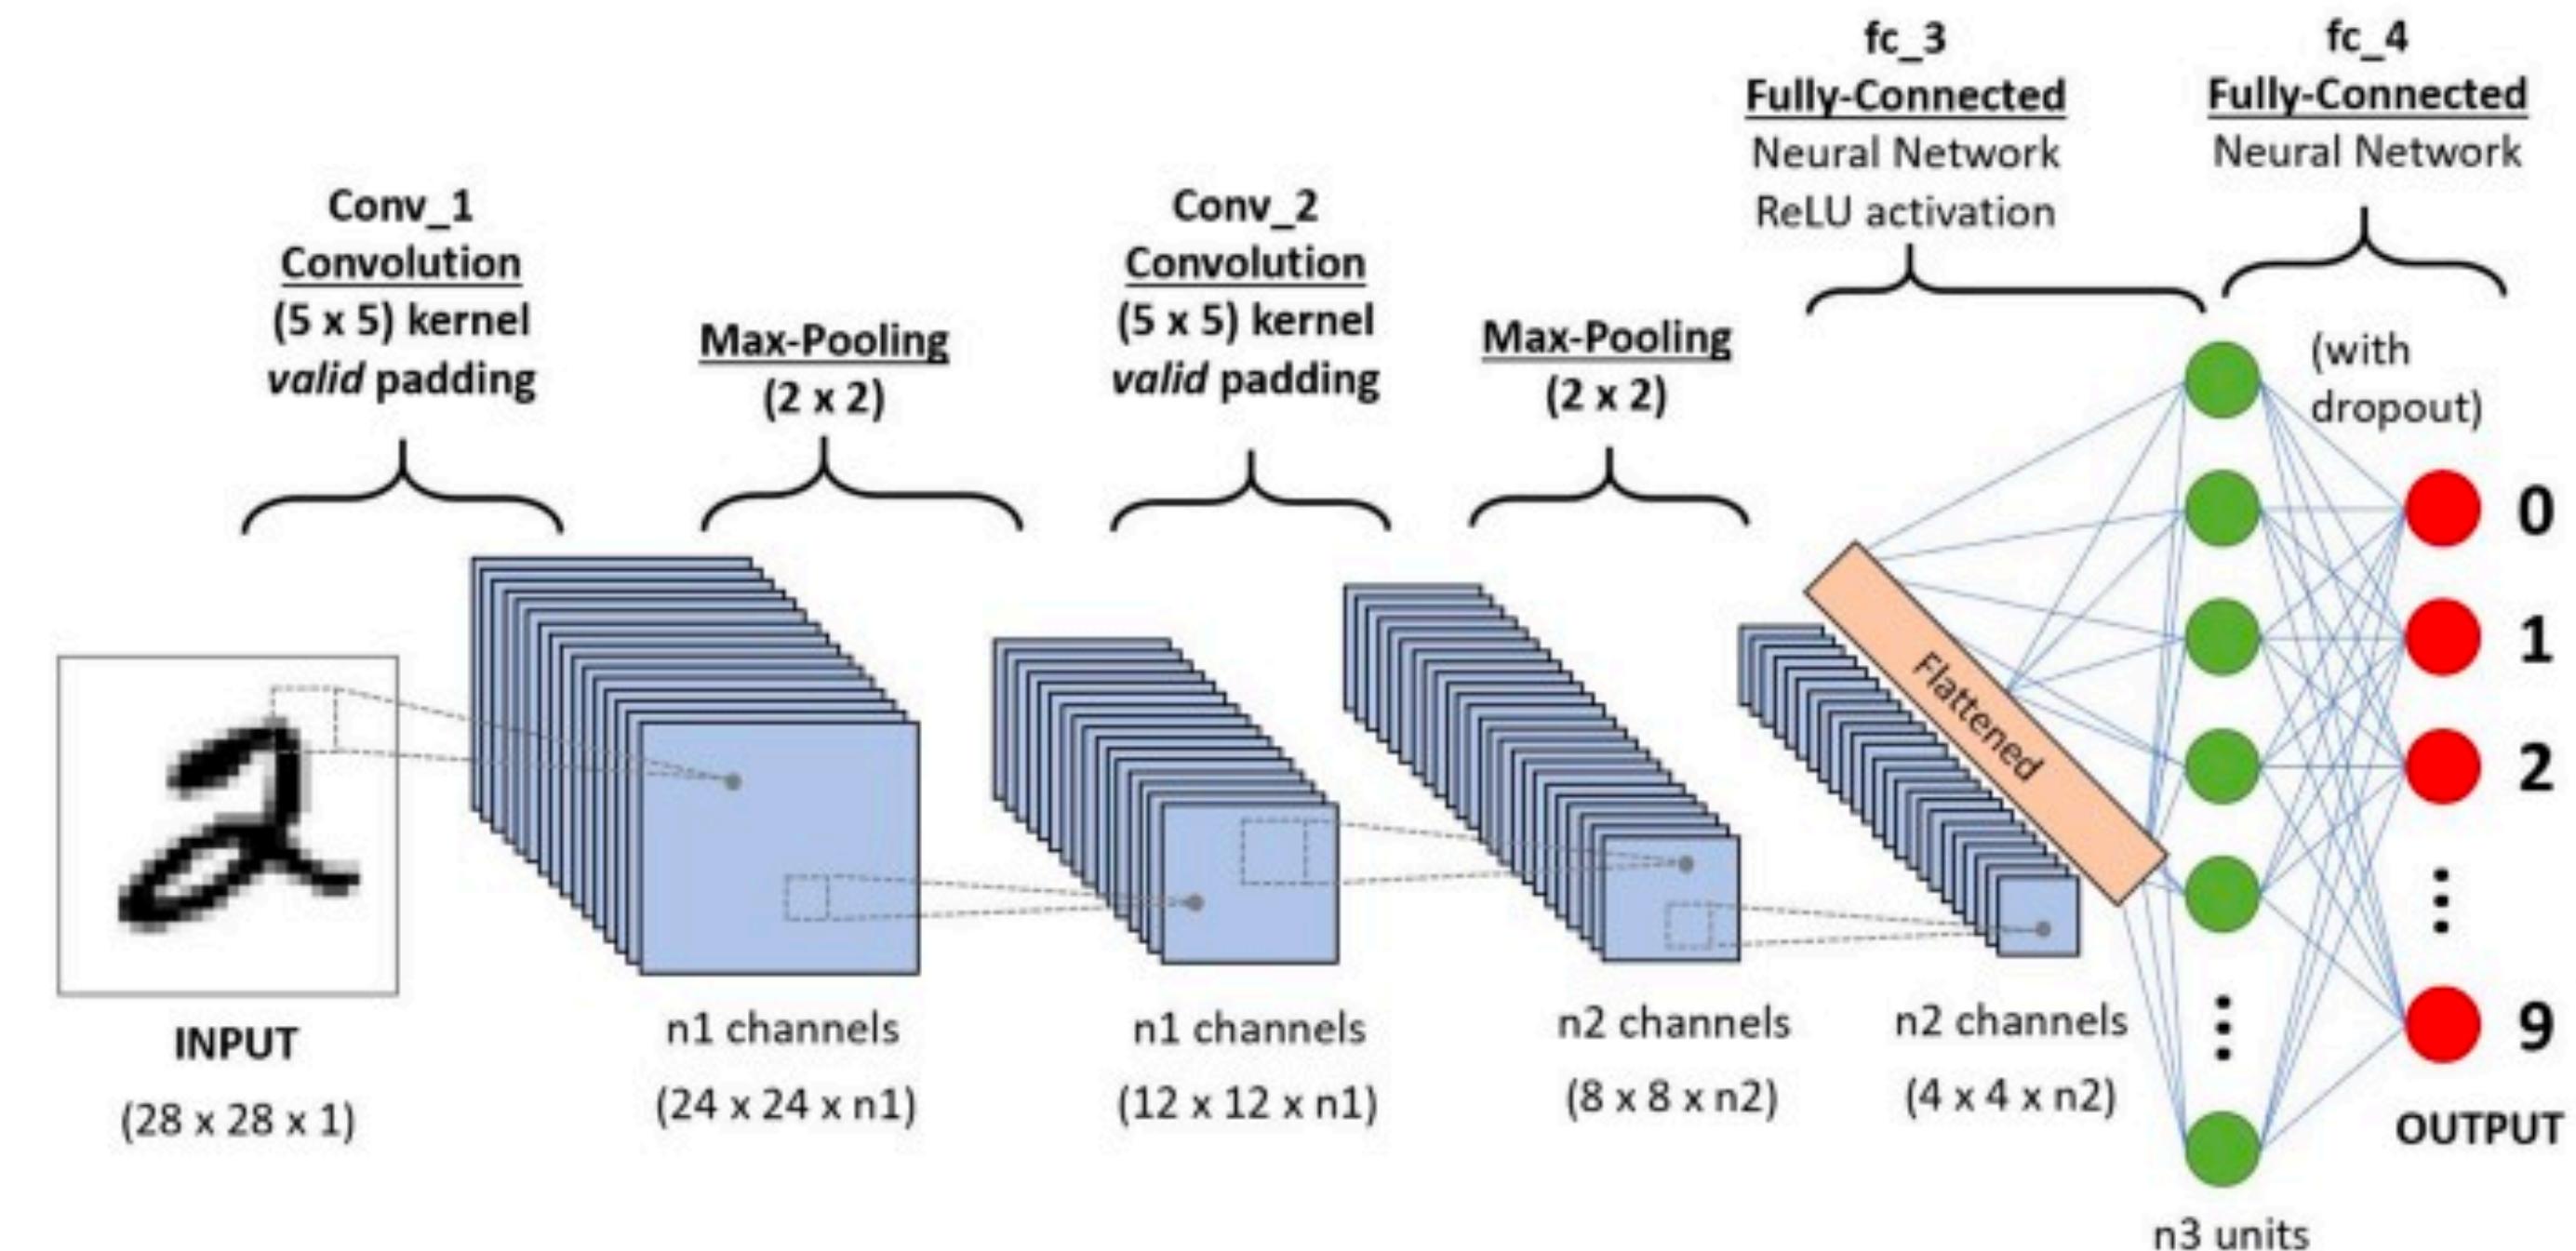


Yann LeCun

# The original CNN architecture

LeNet architecture for hand-written digit recognition (1989).

Idea existed decades ago but data and computing power only became available in 2010s.



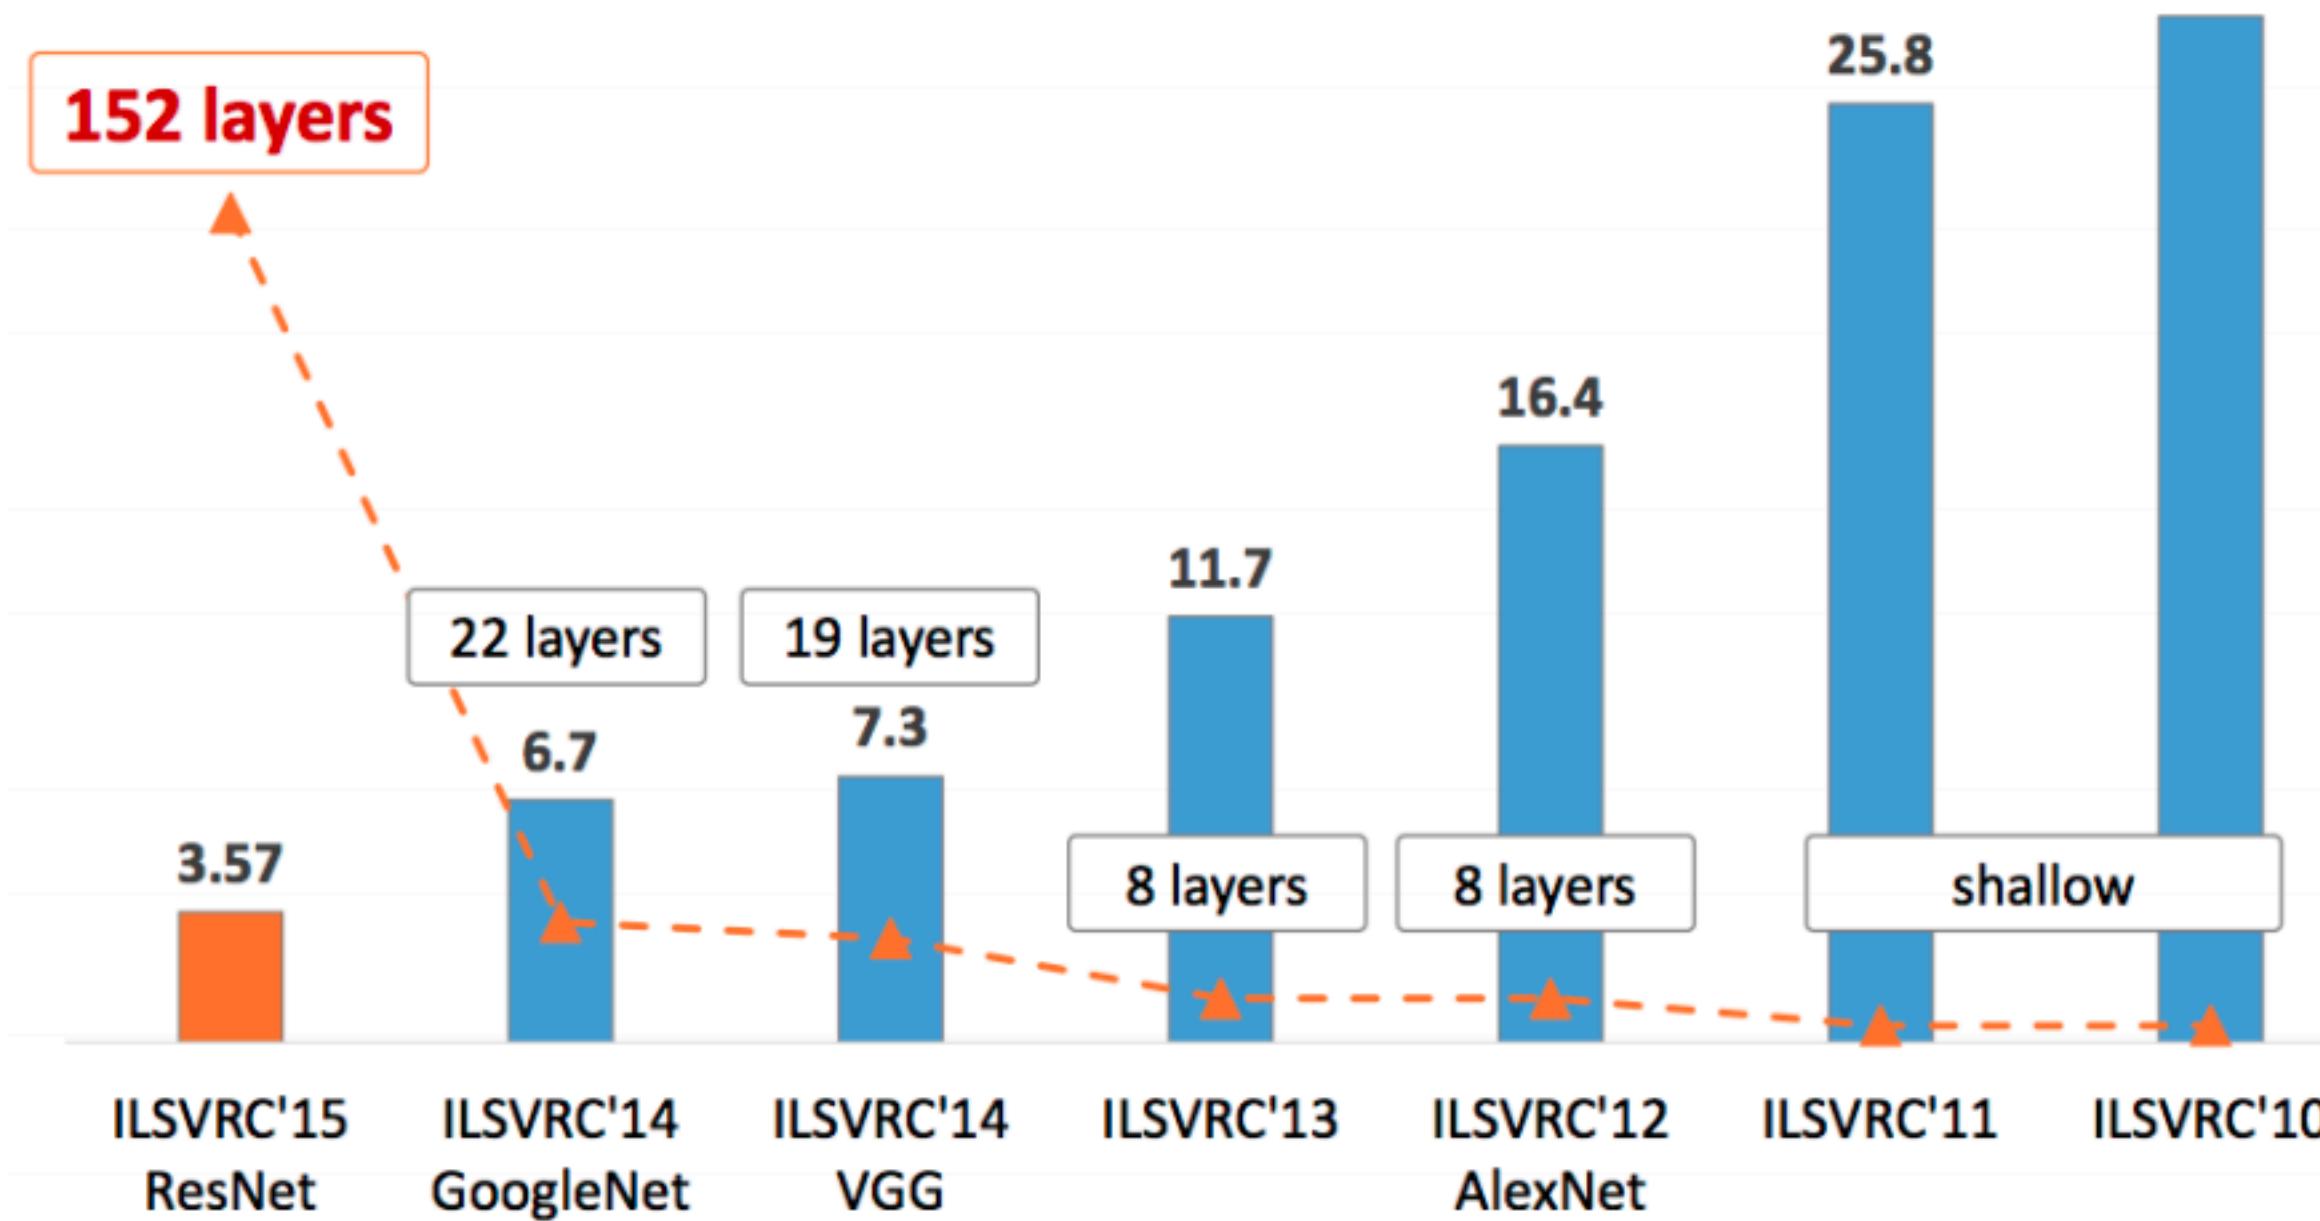
<https://towardsdatascience.com/a-comprehensive-guide-to-convolutional-neural-networks-the-eli5-way-3bd2b1164a53>



Yann LeCun

# Modern CNN architectures

**Classification:** ImageNet Challenge top-5 error



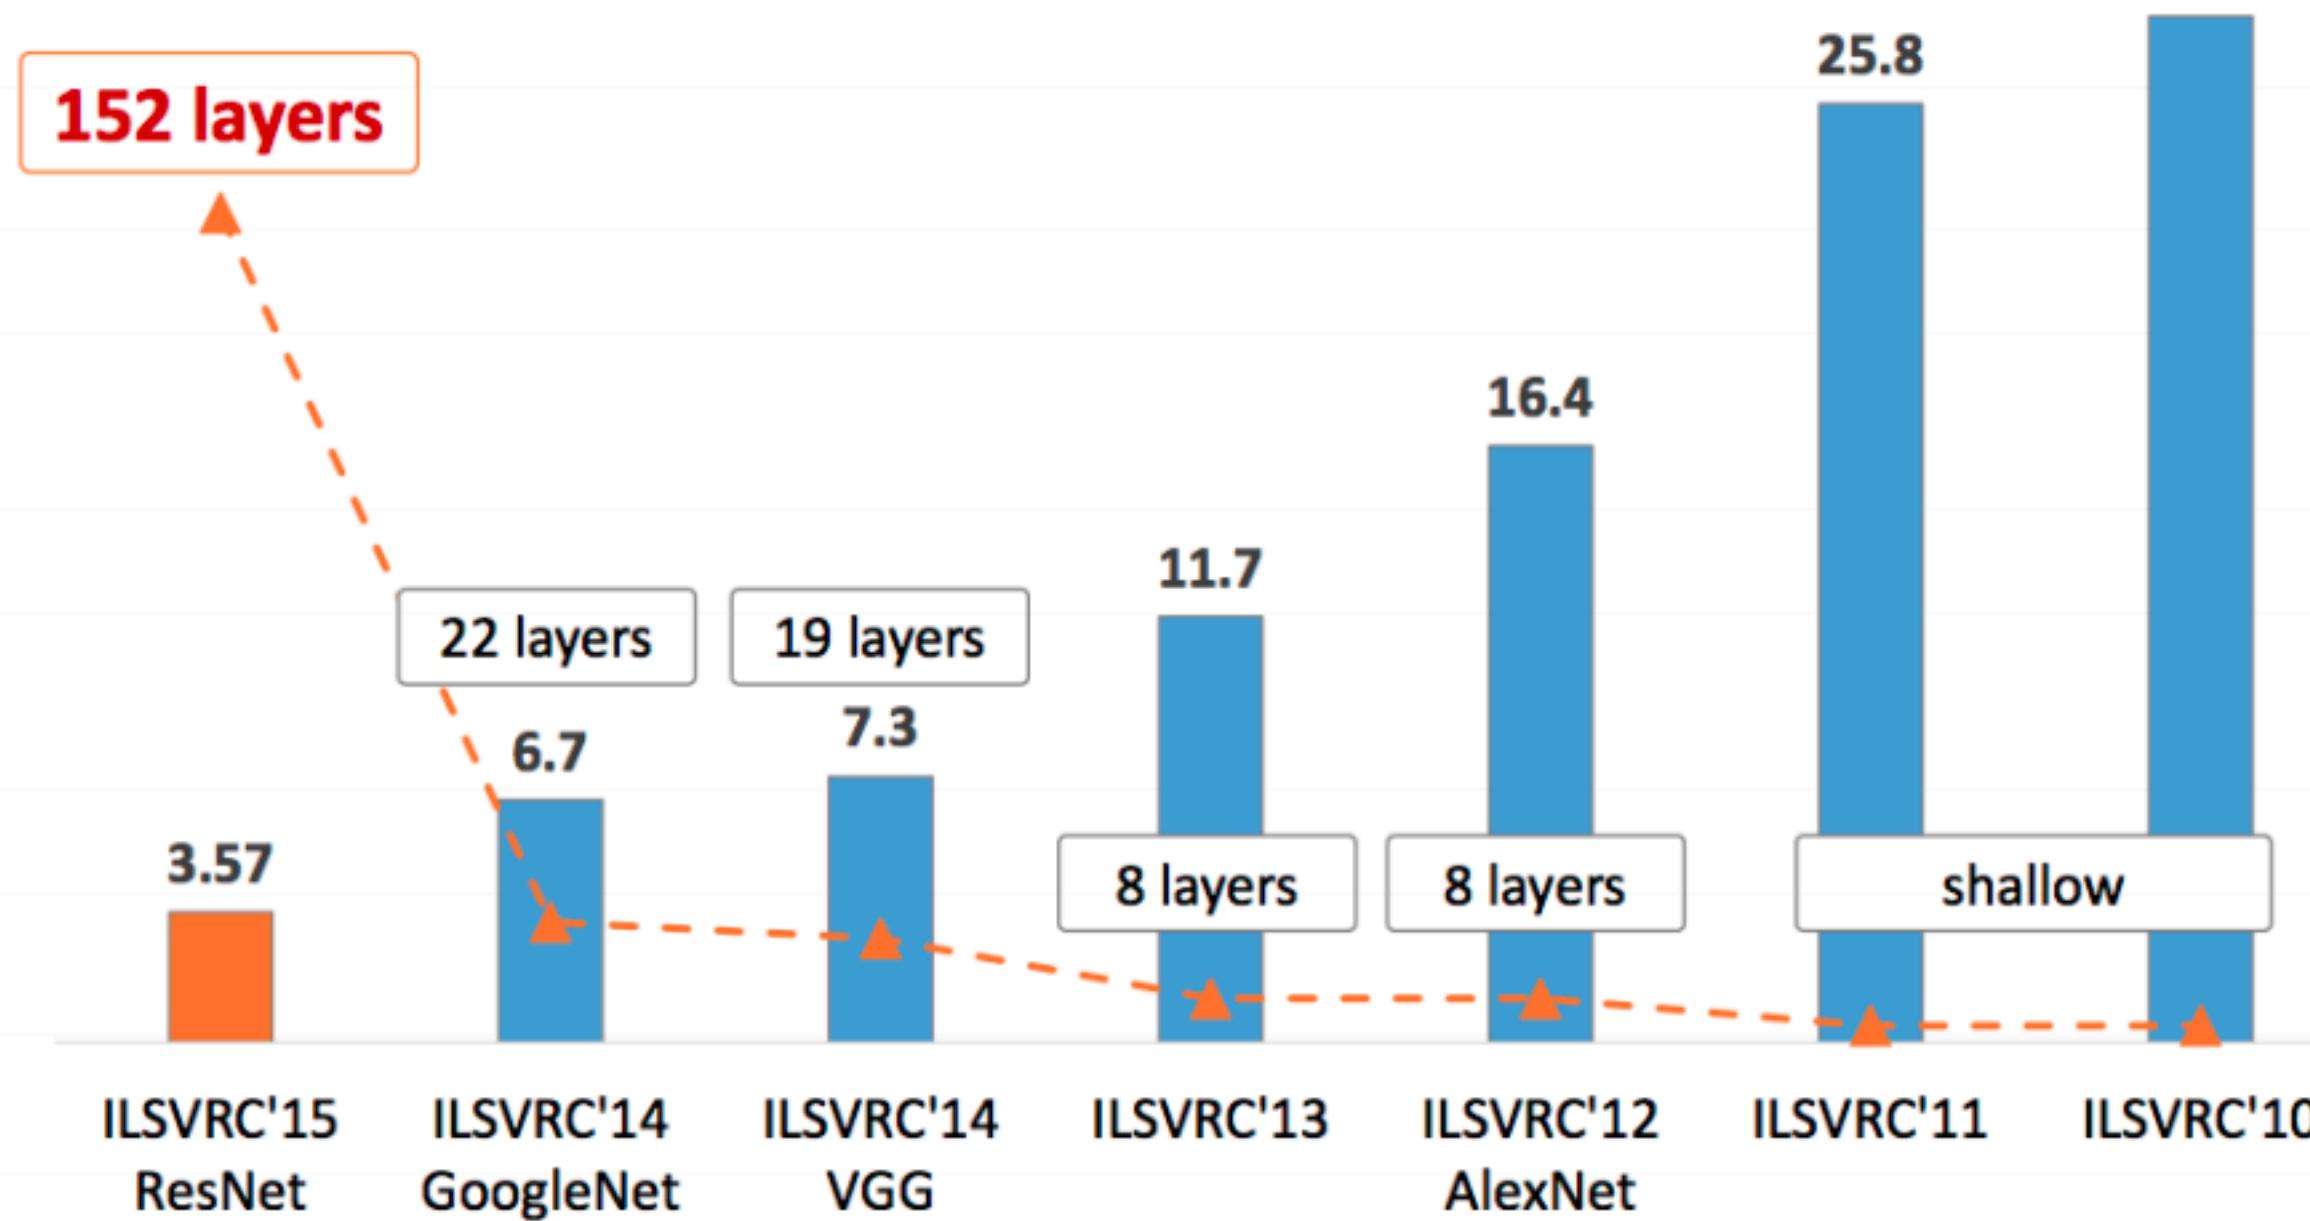
Model	Size (M)	Top-1/top-5 error (%)	# layers	Model description
AlexNet	238	41.00/18.00	8	5 conv + 3 fc layers
VGG-16	540	28.07/9.33	16	13 conv + 3 fc layers
VGG-19	560	27.30/9.00	19	16 conv + 3 fc layers
GoogleNet	40	29.81/10.04	22	21 conv + 1 fc layers
ResNet-50	100	22.85/6.71	50	49 conv + 1 fc layers
ResNet-152	235	21.43/3.57	152	151 conv + 1 fc layers

[https://www.researchgate.net/figure/The-comparison-of-different-CNN-architectures-on-model-size-classification-error-rate\\_tbl1\\_320199404](https://www.researchgate.net/figure/The-comparison-of-different-CNN-architectures-on-model-size-classification-error-rate_tbl1_320199404)

<https://medium.com/@RaghavPrabhu/cnn-architectures-lenet-alexnet-vgg-googlenet-and-resnet-7c81c017b848>

# Modern CNN architectures

**Classification:** ImageNet Challenge top-5 error



Model	Size (M)	Top-1/top-5 error (%)	# layers	Model description
AlexNet	238	41.00/18.00	8	5 conv + 3 fc layers
VGG-16	540	28.07/9.33	16	13 conv + 3 fc layers
VGG-19	560	27.30/9.00	19	16 conv + 3 fc layers
GoogleNet	40	29.81/10.04	22	21 conv + 1 fc layers
ResNet-50	100	22.85/6.71	50	49 conv + 1 fc layers
ResNet-152	235	21.43/3.57	152	151 conv + 1 fc layers

[https://www.researchgate.net/figure/The-comparison-of-different-CNN-architectures-on-model-size-classification-error-rate\\_tbl1\\_320199404](https://www.researchgate.net/figure/The-comparison-of-different-CNN-architectures-on-model-size-classification-error-rate_tbl1_320199404)

<https://medium.com/@RaghavPrabhu/cnn-architectures-lenet-alexnet-vgg-googlenet-and-resnet-7c81c017b848>

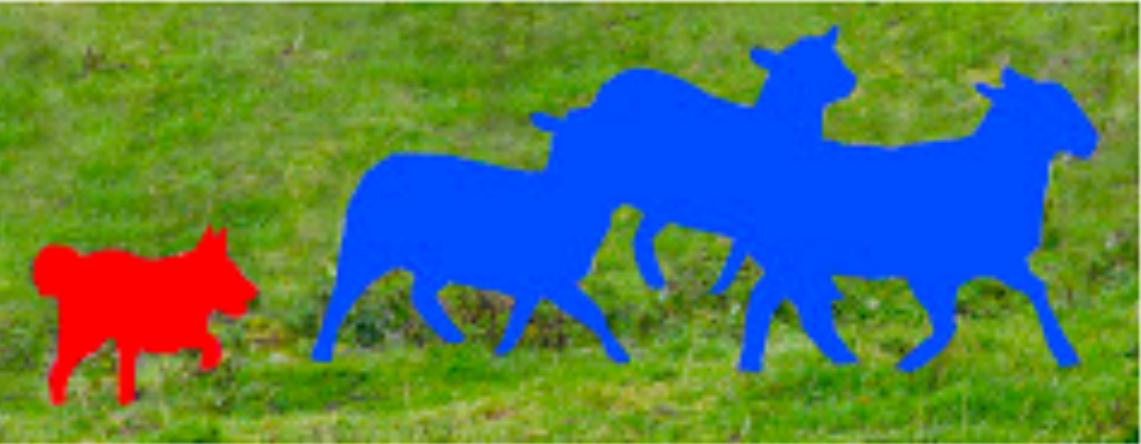
CNNs are getting progressively deeper with time.

# Other applications of CNNs

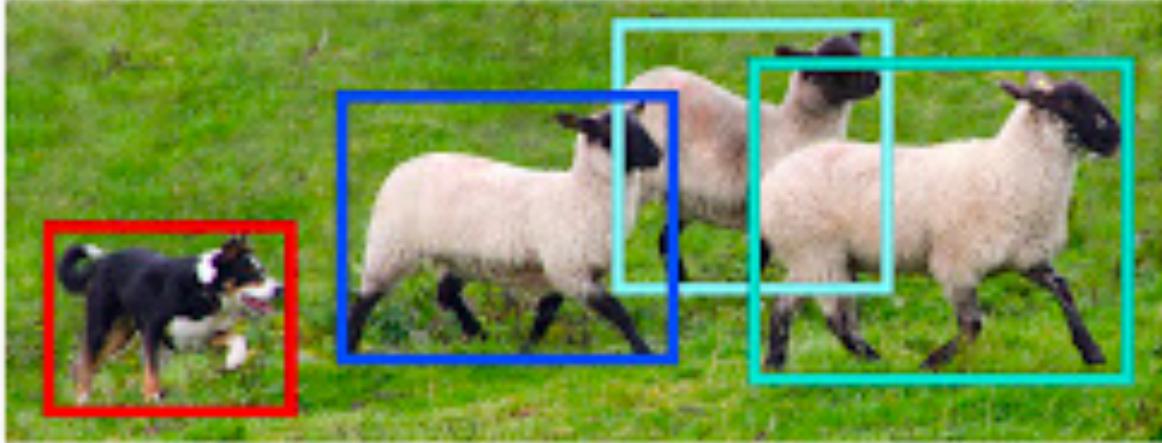
# Other applications of CNNs



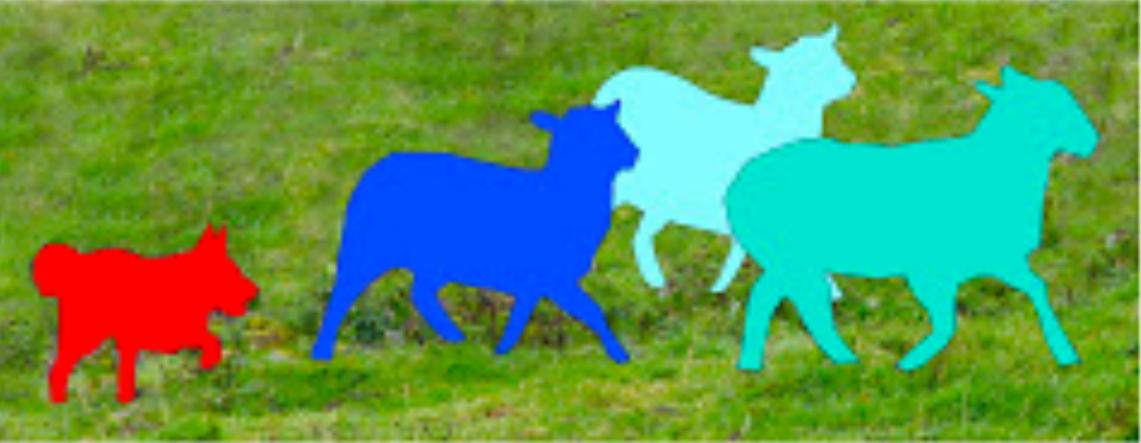
Image Recognition



Semantic Segmentation



Object Detection



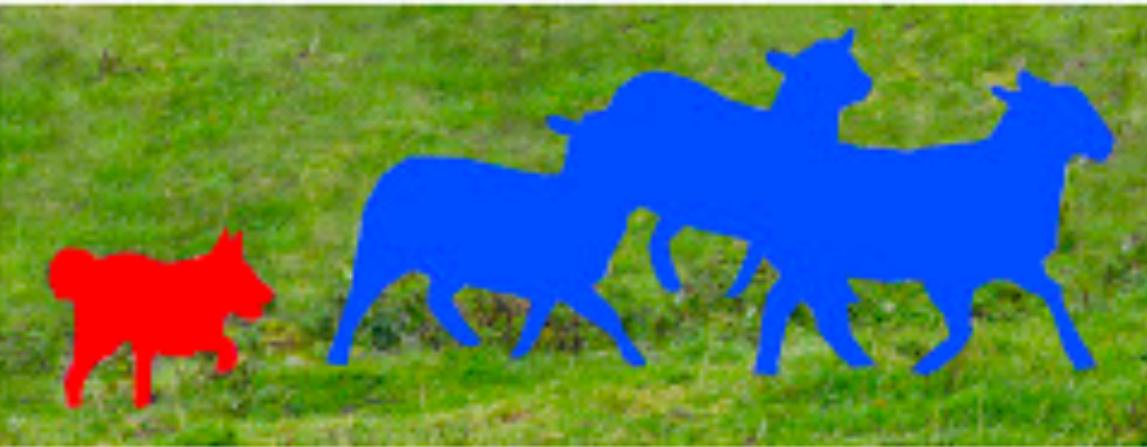
Instance Segmentation

<http://manipulation.csail.mit.edu/segmentation.html>

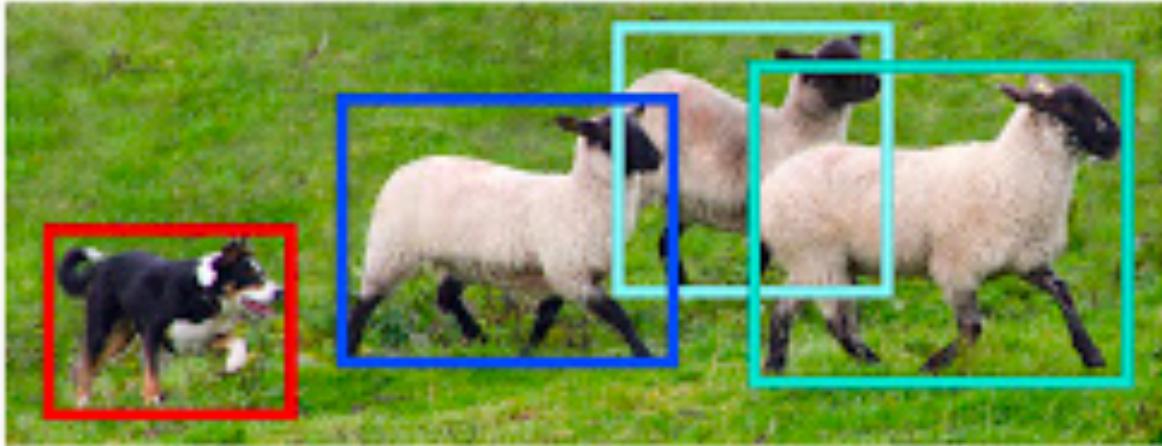
# Other applications of CNNs



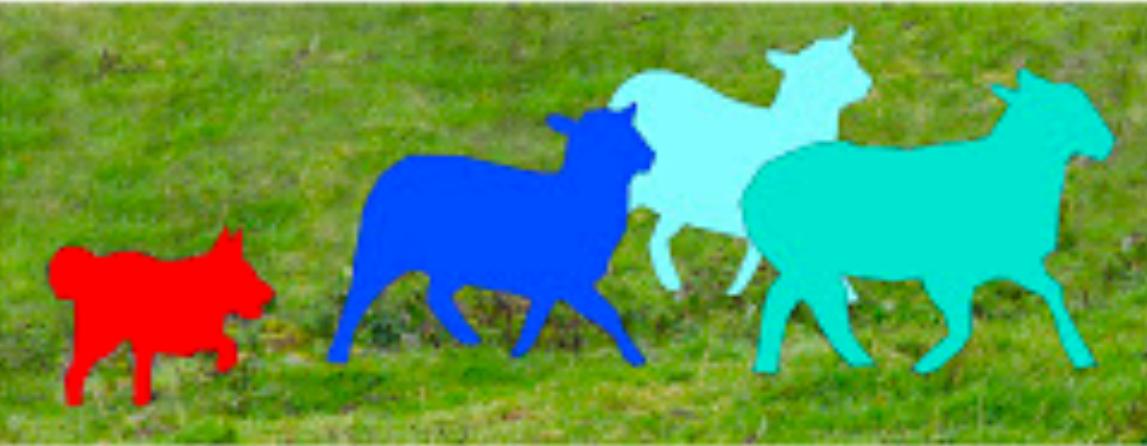
Image Recognition



Semantic Segmentation



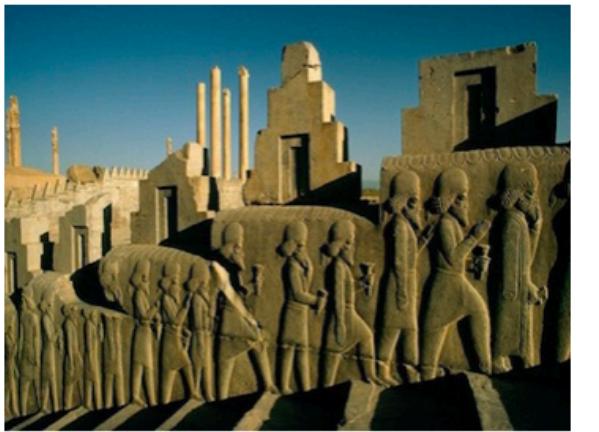
Object Detection



Instance Segmentation

<http://manipulation.csail.mit.edu/segmentation.html>

content image



style image



generated image



Ancient city of Persepolis

The Starry Night (Van Gogh)

Persepolis  
in Van Gogh style

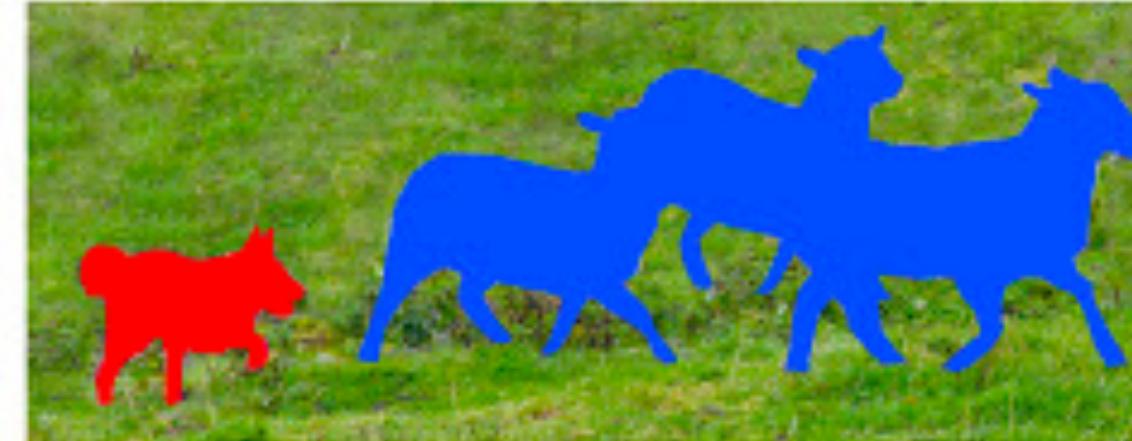
<https://towardsdatascience.com/neural-style-transfer-on-real-time-video-with-full-implementable-code-ac2dbc0e9822>

Style Transfer

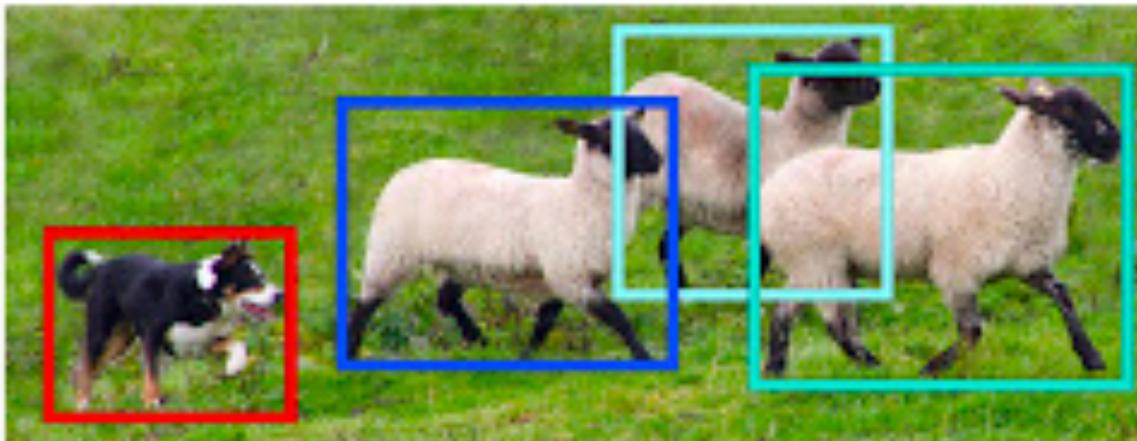
# Other applications of CNNs



Image Recognition



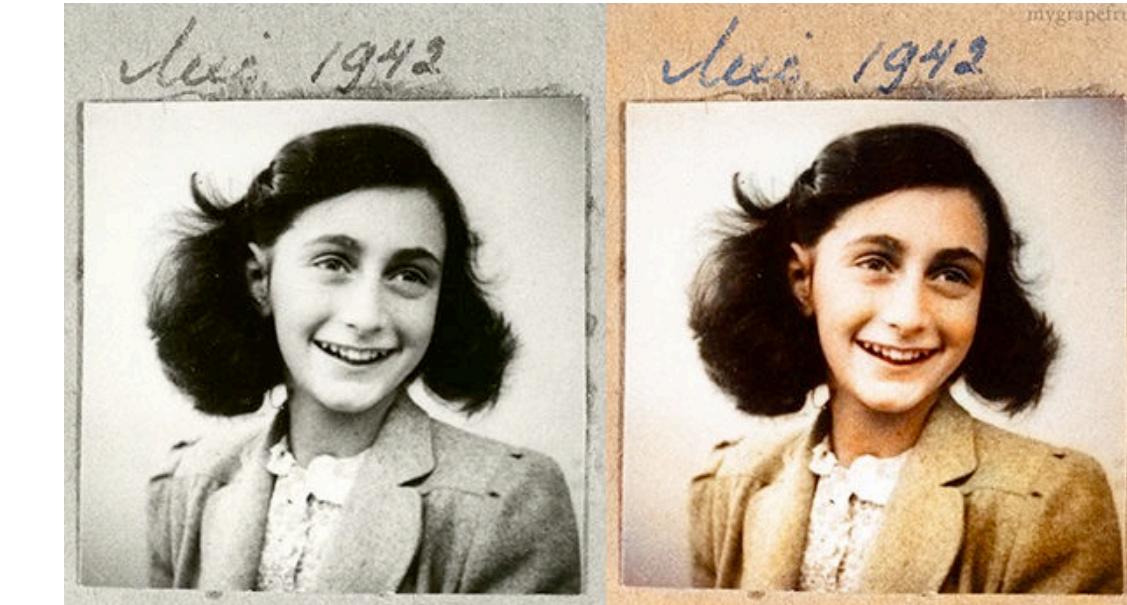
Semantic Segmentation



Object Detection



Instance Segmentation

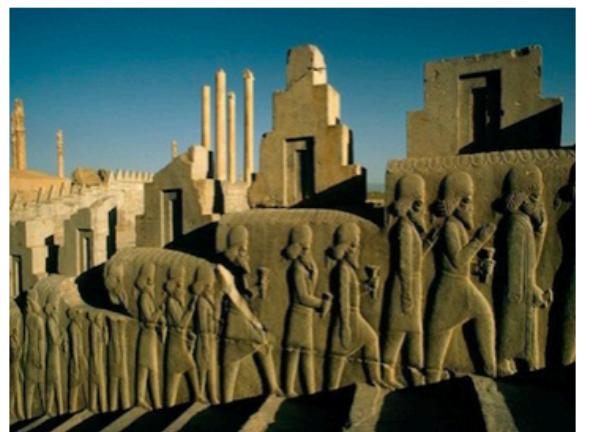


<https://medium.com/@nabilibin14/chain-colorization-using-cdcgans-and-cnn-from-learned-deep-prior-1405dab48df3>

Image colorization

<http://manipulation.csail.mit.edu/segmentation.html>

content image



style image



generated image



Ancient city of Persepolis

The Starry Night (Van Gogh)

Persepolis  
in Van Gogh style

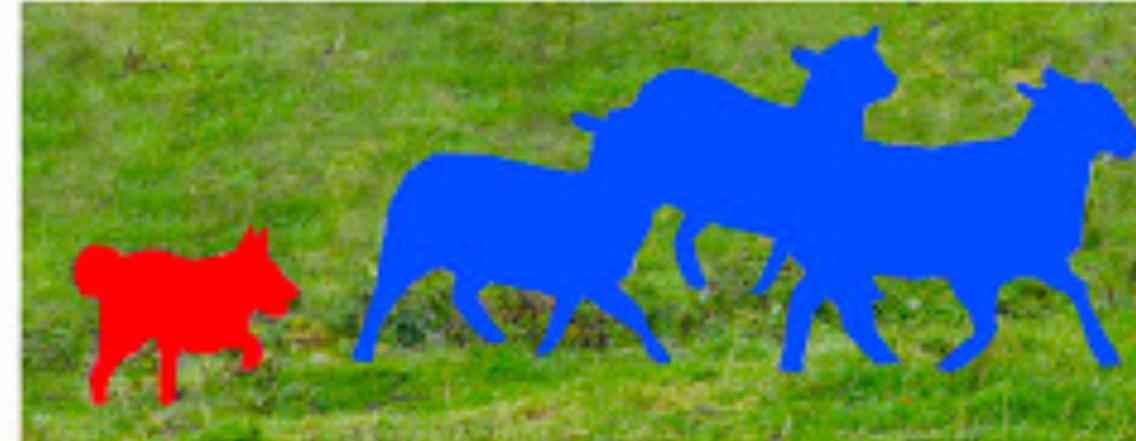
<https://towardsdatascience.com/neural-style-transfer-on-real-time-video-with-full-implementable-code-ac2dbc0e9822>

Style Transfer

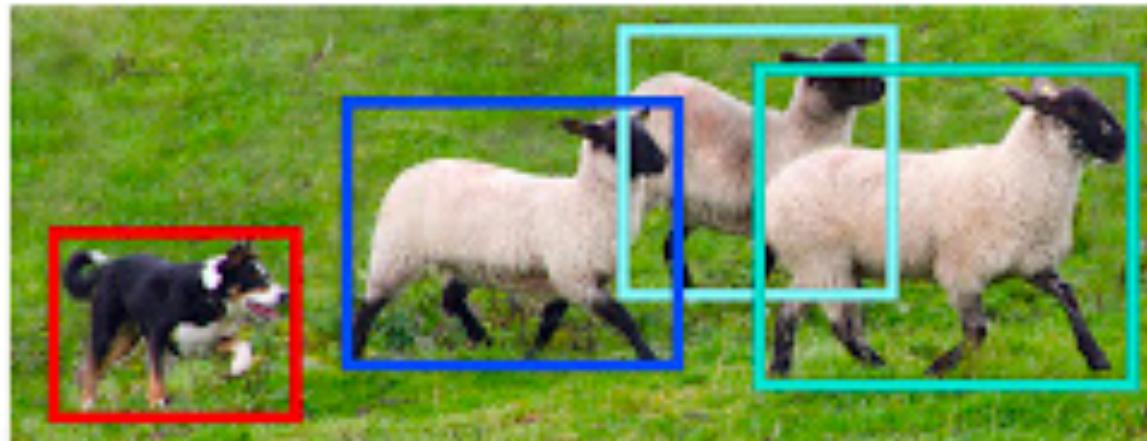
# Other applications of CNNs



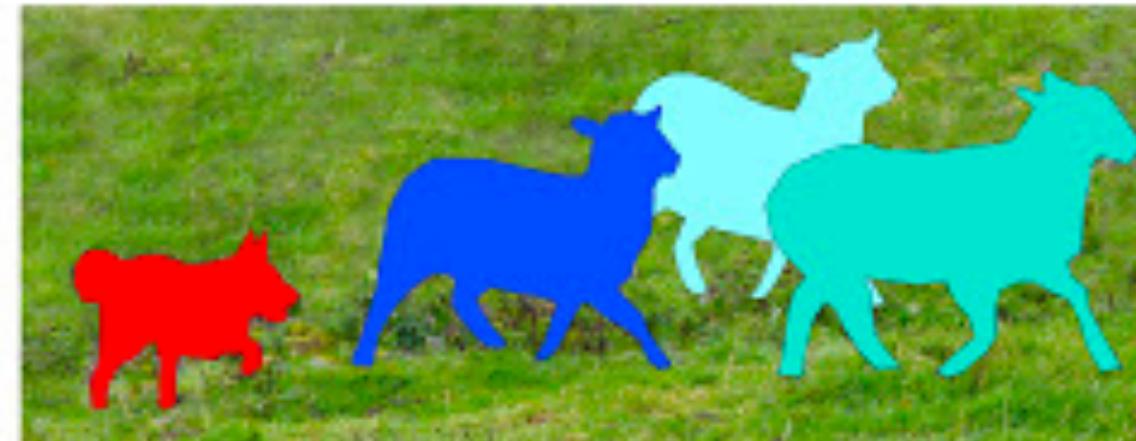
Image Recognition



Semantic Segmentation



Object Detection



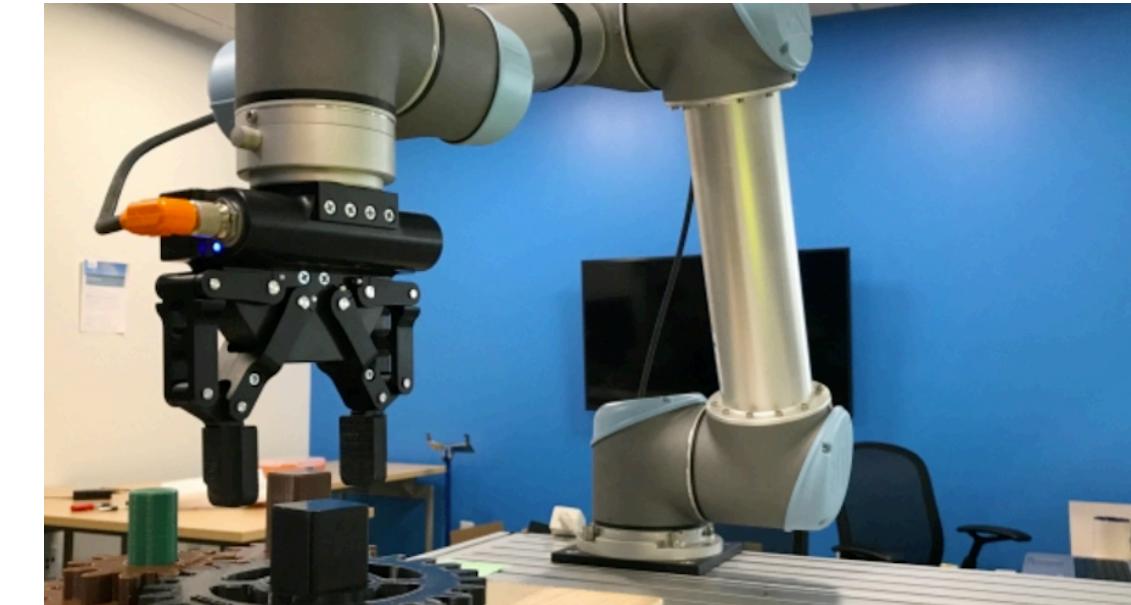
Instance Segmentation

<http://manipulation.csail.mit.edu/segmentation.html>



<https://medium.com/@nabilibin14/chain-colorization-using-cdcgans-and-cnn-from-learned-deep-prior-1405dab48df3>

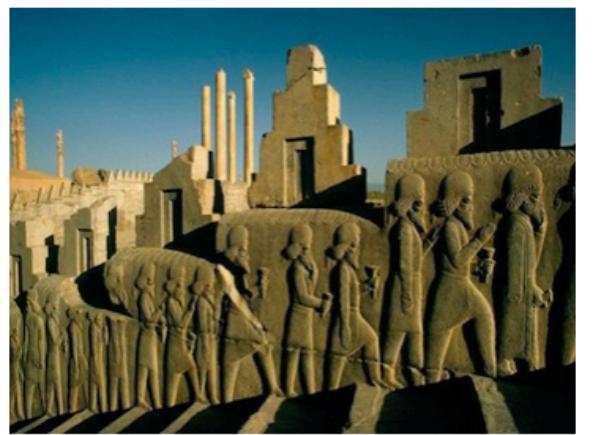
Image colorization



<https://www.therobotreport.com/reinforcement-learning-industrial-robotics/>

Reinforcement learning

content image



style image



generated image



Ancient city of Persepolis

The Starry Night (Van Gogh)

Persepolis  
in Van Gogh style

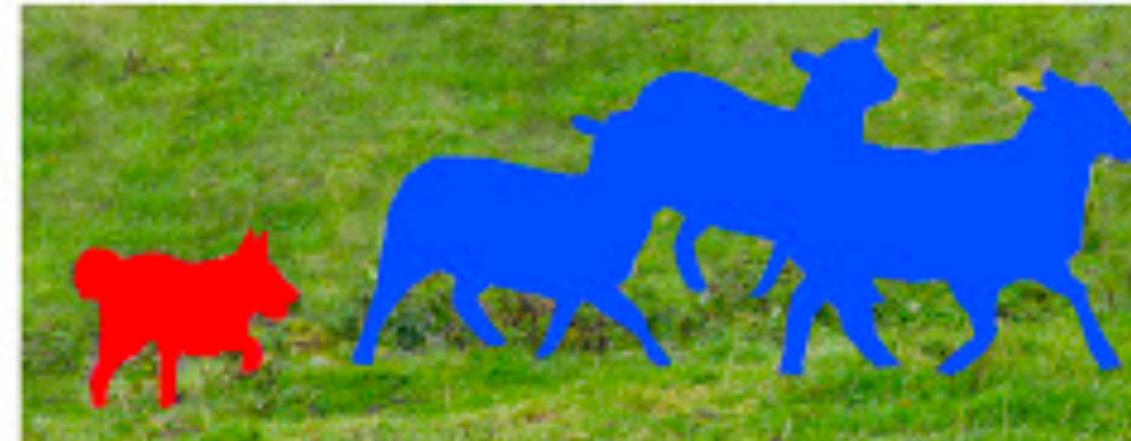
<https://towardsdatascience.com/neural-style-transfer-on-real-time-video-with-full-implementable-code-ac2dbc0e9822>

Style Transfer

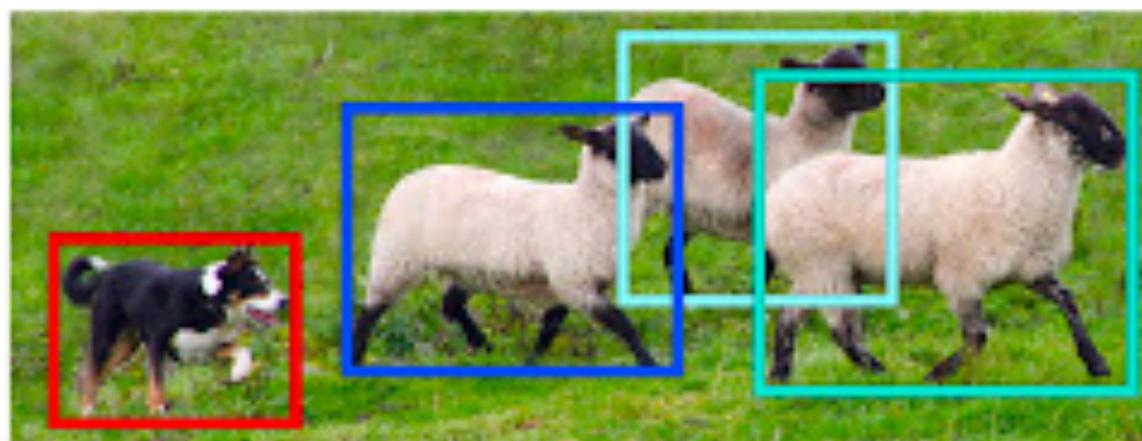
# Other applications of CNNs



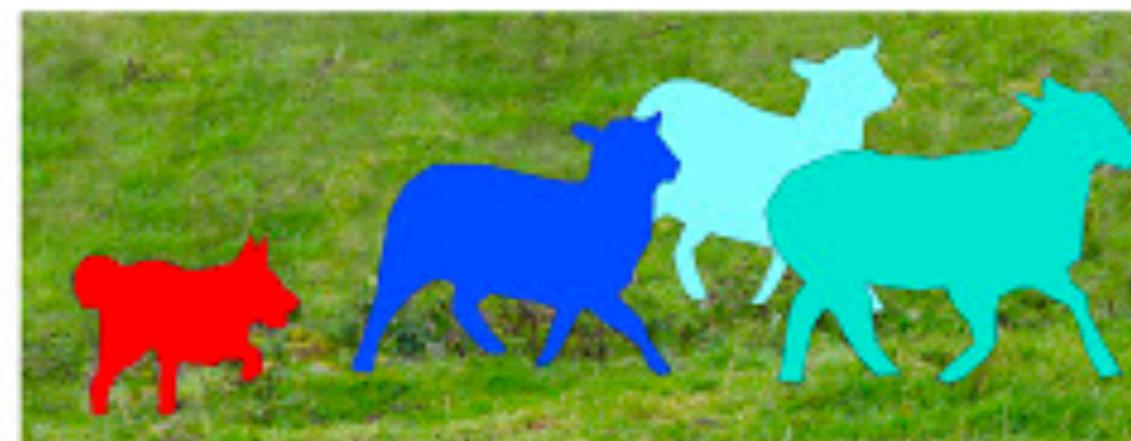
Image Recognition



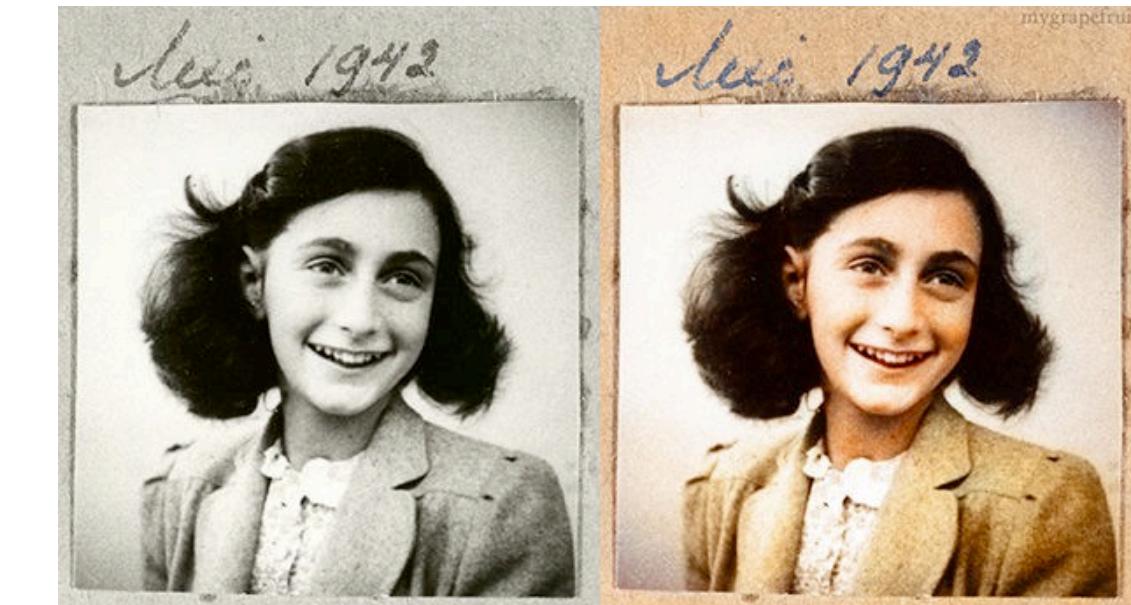
Semantic Segmentation



Object Detection



Instance Segmentation



<https://medium.com/@nabilibani/chain-colorization-using-cdcgans-and-cnn-from-learned-deep-prior-1405dab48df3>

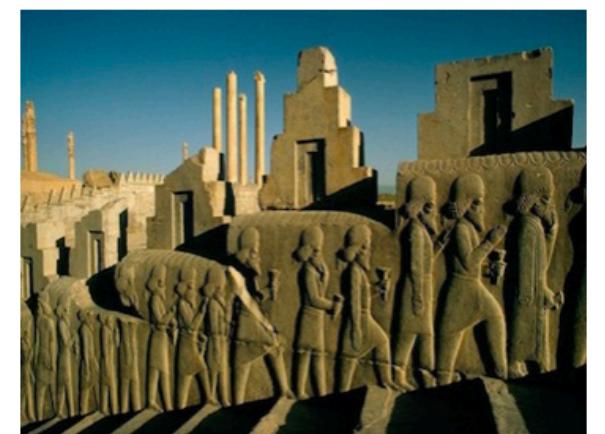
Image colorization



<https://www.therobotreport.com/reinforcement-learning-industrial-robotics/>

Reinforcement learning

content image



+



=



generated image

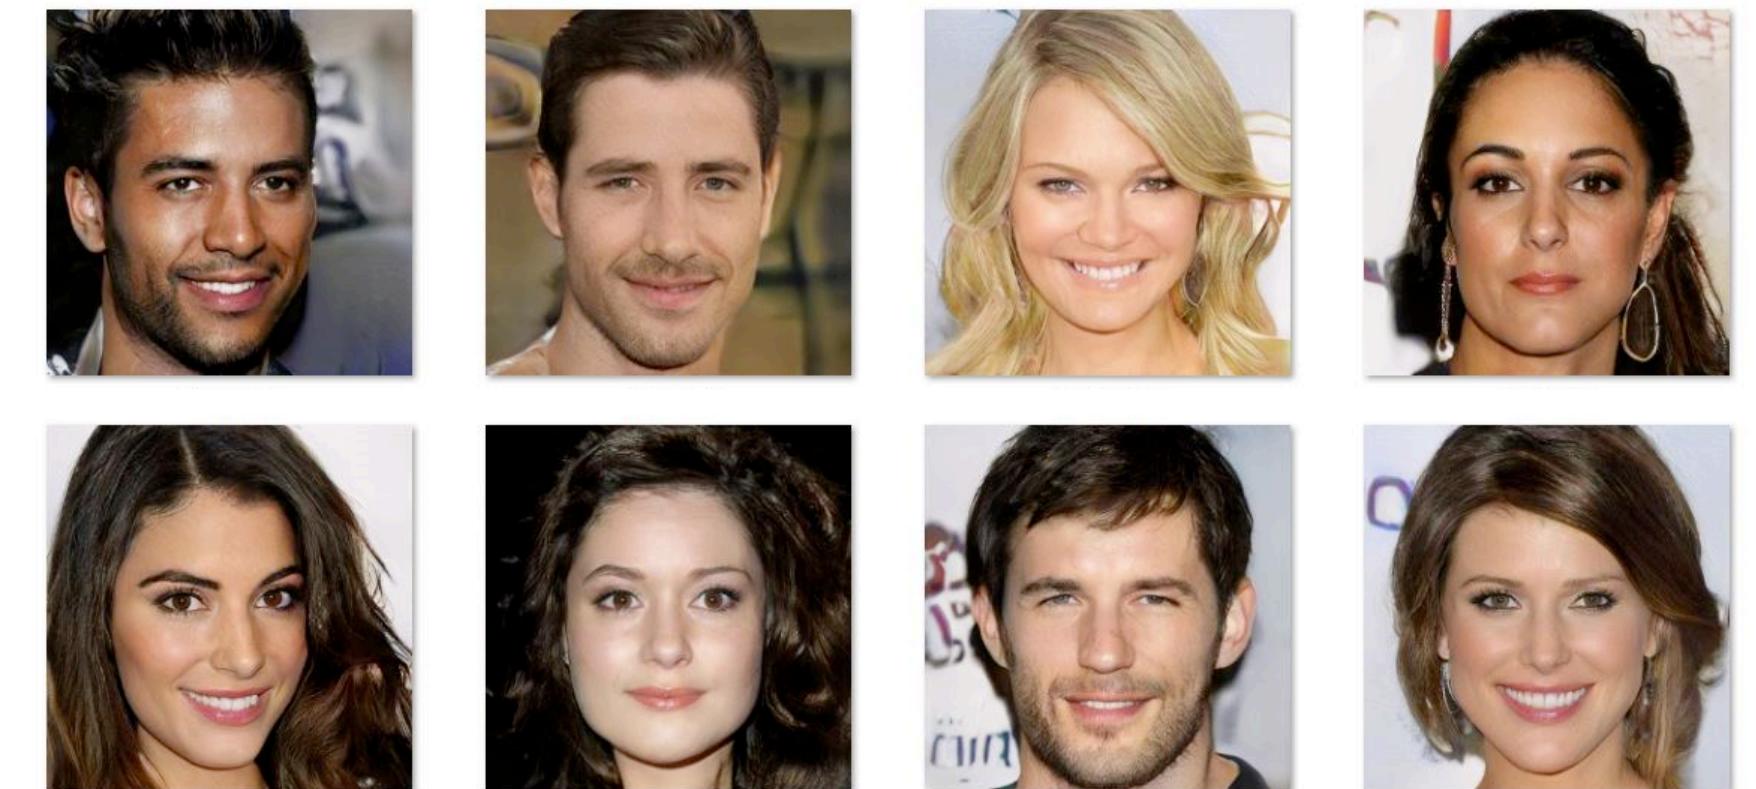
Ancient city of Persepolis

The Starry Night (Van Gogh)

Persepolis  
in Van Gogh style

<https://towardsdatascience.com/neural-style-transfer-on-real-time-video-with-full-implementable-code-ac2dbc0e9822>

Style Transfer



<https://medium.datadriveninvestor.com/artificial-intelligence-gans-can-create-fake-celebrity-faces-44fe80d419f7>

Generative models

# Summary

# Summary

- Image classification is a prototypical task in image processing.

# Summary

- Image classification is a prototypical task in image processing.
- A convolution sweeps over an image, looking for a specific pattern.

# Summary

- Image classification is a prototypical task in image processing.
- A convolution sweeps over an image, looking for a specific pattern.
- Convolutions are local (applied to image patches) and translation-invariant (applied equally to patches across the whole image).

# Summary

- Image classification is a prototypical task in image processing.
- A convolution sweeps over an image, looking for a specific pattern.
- Convolutions are local (applied to image patches) and translation-invariant (applied equally to patches across the whole image).
- Convolutional neural networks are a specialized architecture for image processing, consisting of alternating convolutional and pooling layers (feature learning), following by final fully connected layer (classification).

# Summary

- Image classification is a prototypical task in image processing.
- A convolution sweeps over an image, looking for a specific pattern.
- Convolutions are local (applied to image patches) and translation-invariant (applied equally to patches across the whole image).
- Convolutional neural networks are a specialized architecture for image processing, consisting of alternating convolutional and pooling layers (feature learning), following by final fully connected layer (classification).
- People have built increasingly deep CNNs, which have performed increasingly well. Image classification problem is essentially solved.

# Summary

- Image classification is a prototypical task in image processing.
- A convolution sweeps over an image, looking for a specific pattern.
- Convolutions are local (applied to image patches) and translation-invariant (applied equally to patches across the whole image).
- Convolutional neural networks are a specialized architecture for image processing, consisting of alternating convolutional and pooling layers (feature learning), following by final fully connected layer (classification).
- People have built increasingly deep CNNs, which have performed increasingly well. Image classification problem is essentially solved.
- Many other image processing tasks can be addressed with CNNs.

# Summary

- Image classification is a prototypical task in image processing.
- A convolution sweeps over an image, looking for a specific pattern.
- Convolutions are local (applied to image patches) and translation-invariant (applied equally to patches across the whole image).
- Convolutional neural networks are a specialized architecture for image processing, consisting of alternating convolutional and pooling layers (feature learning), following by final fully connected layer (classification).
- People have built increasingly deep CNNs, which have performed increasingly well. Image classification problem is essentially solved.
- Many other image processing tasks can be addressed with CNNs.

[Quiz Practice](#)